

Kubernetes in 30 minutes

Mar 10, 2017

Daisuke Maki @lestrrat





Netscape Navigator



~~Netscape Navigator~~



Kubernetes



Κυβερνήτης




k8s

**“Something to do with
containers...?”**

Assumptions


- You know about containers
- You know about orchestration
- You know about the ~~painful~~ art of deploying stuff




Docker




“Random containers doing their own sh*t”




Compose



“Foreman on Acid”



Swarm




“Clustering and Discovery”

(hey, we’re getting somewhere!)



=



“Batteries Included”

K8s on GCP



GKE



GCE



Networking


...and others

Basic Concepts


Node

**Hosts running
k8s daemons**

Master




Node



Node



Node




Pod

基本

**Basic unit of
deployment in k8s**


Group of Containers




Container configurations



Shared storage



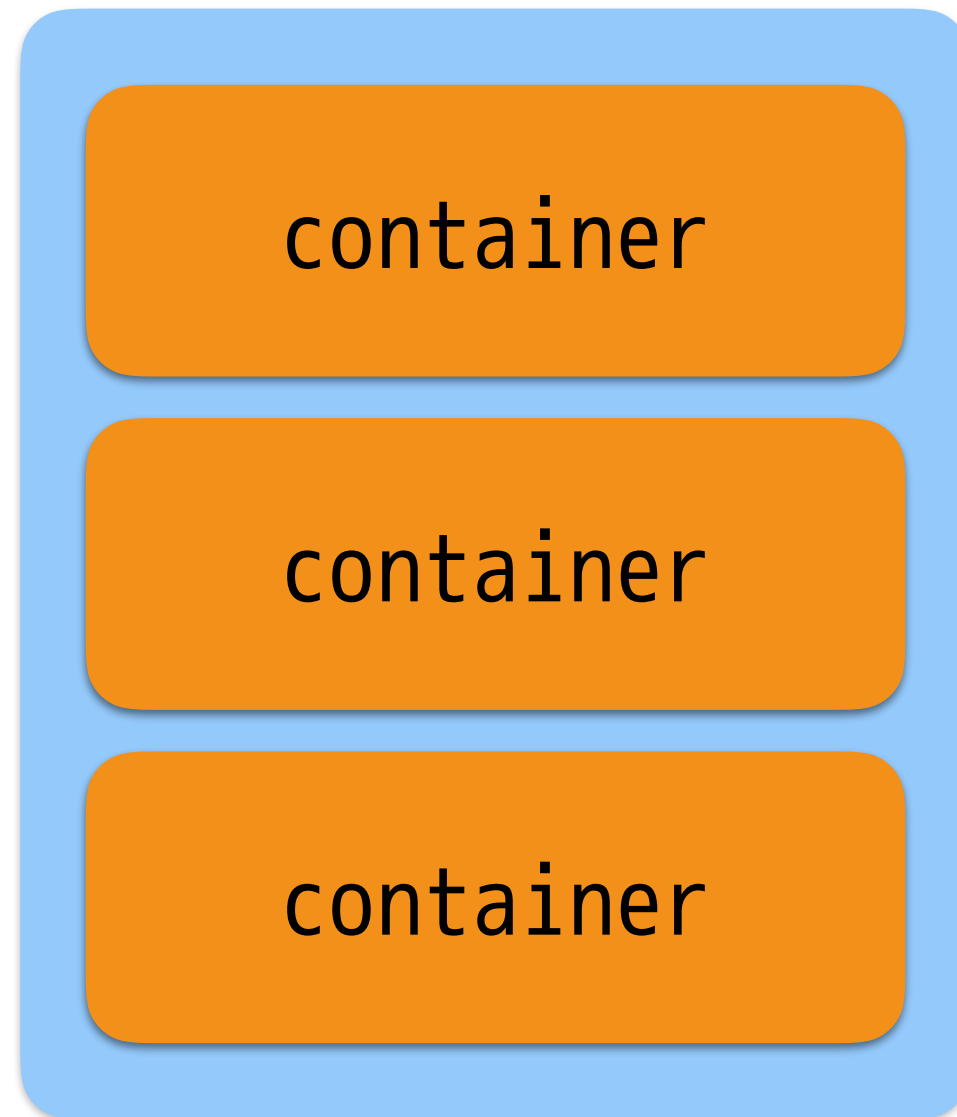
Pod



Pod



Pod



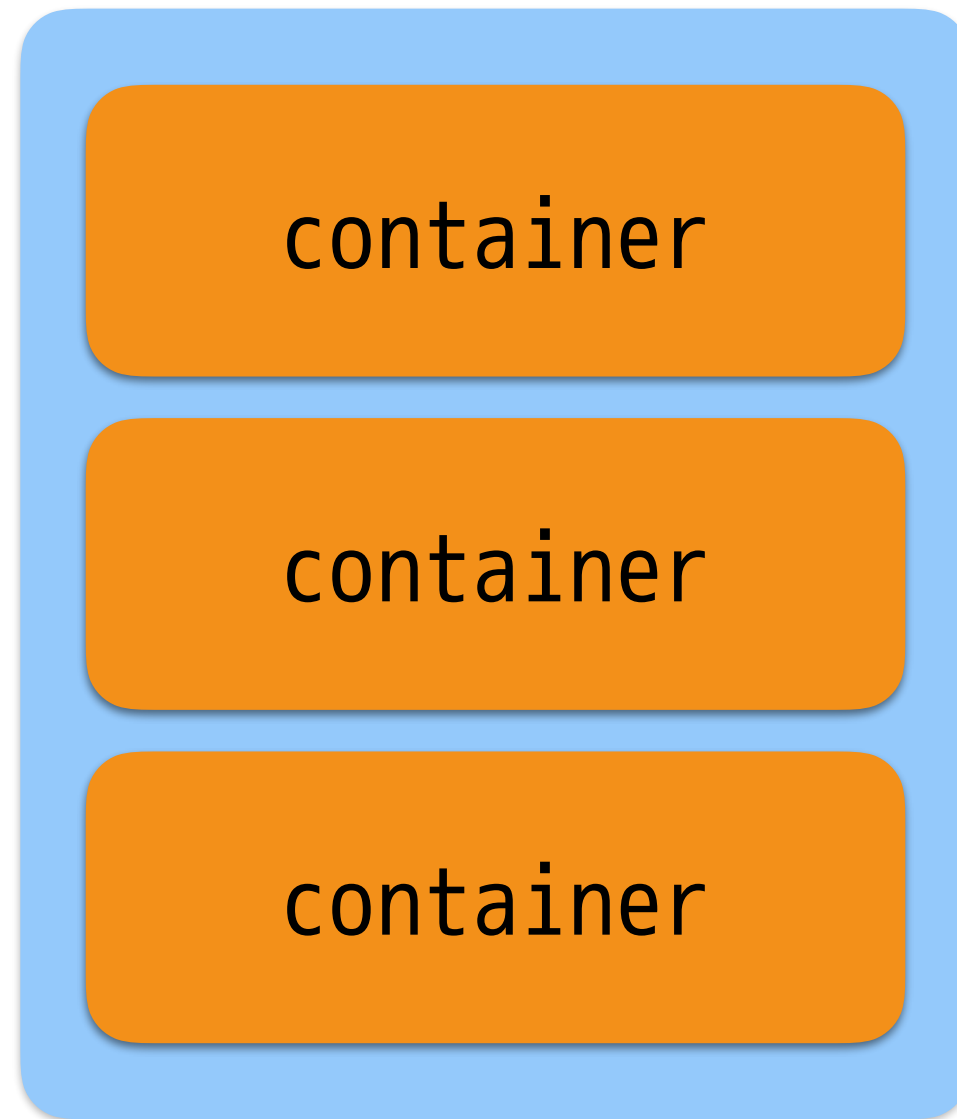
Pod

e.g.

nginx

app

redis/cache

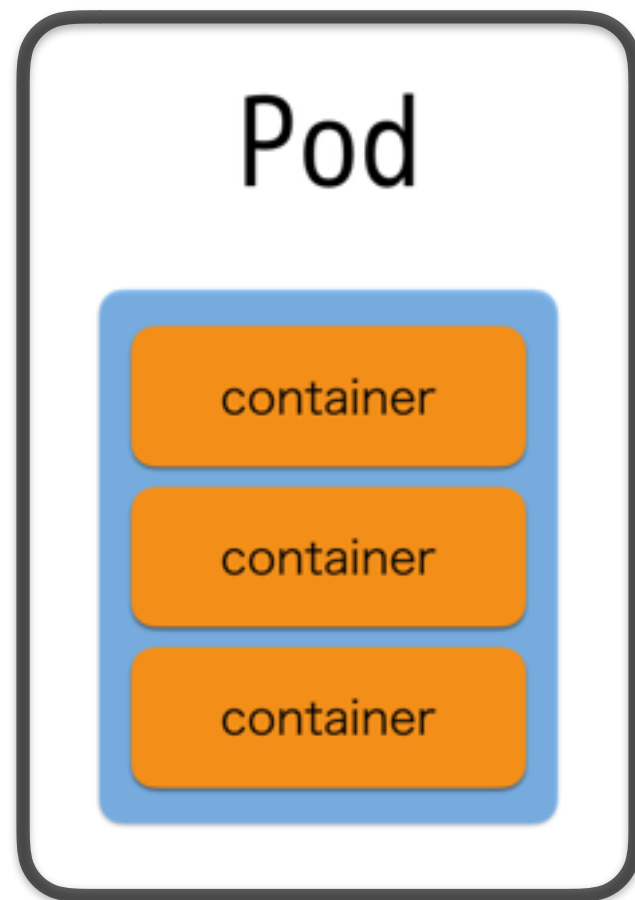


Scheduled together

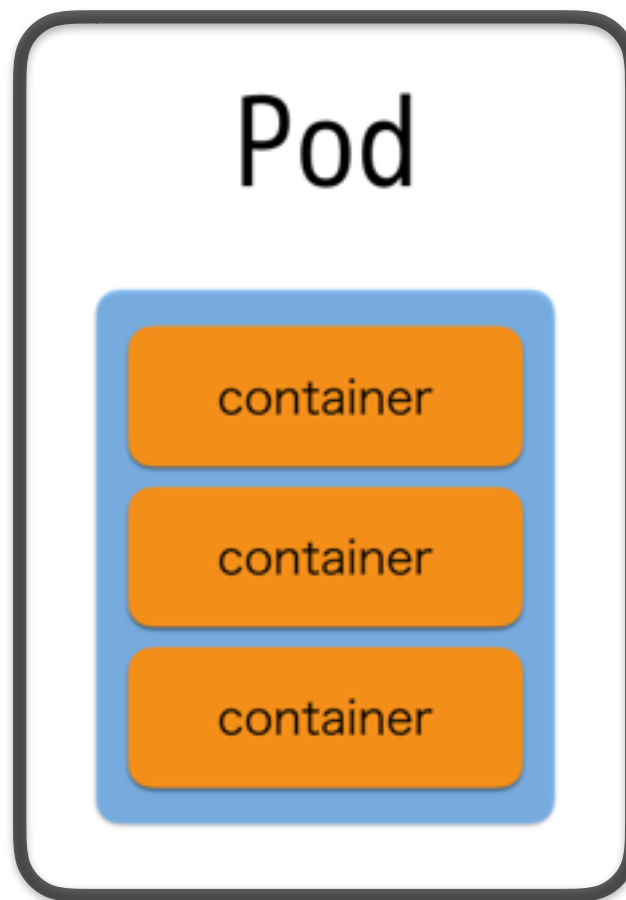
(“co-scheduled”)

**Guaranteed to be
on the same node
("co-located")**

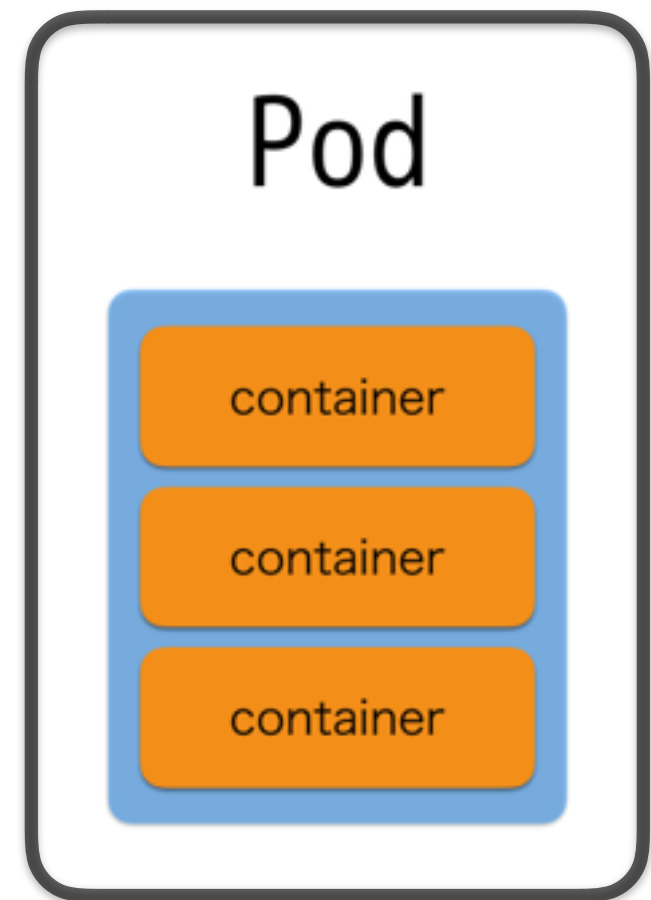
Node



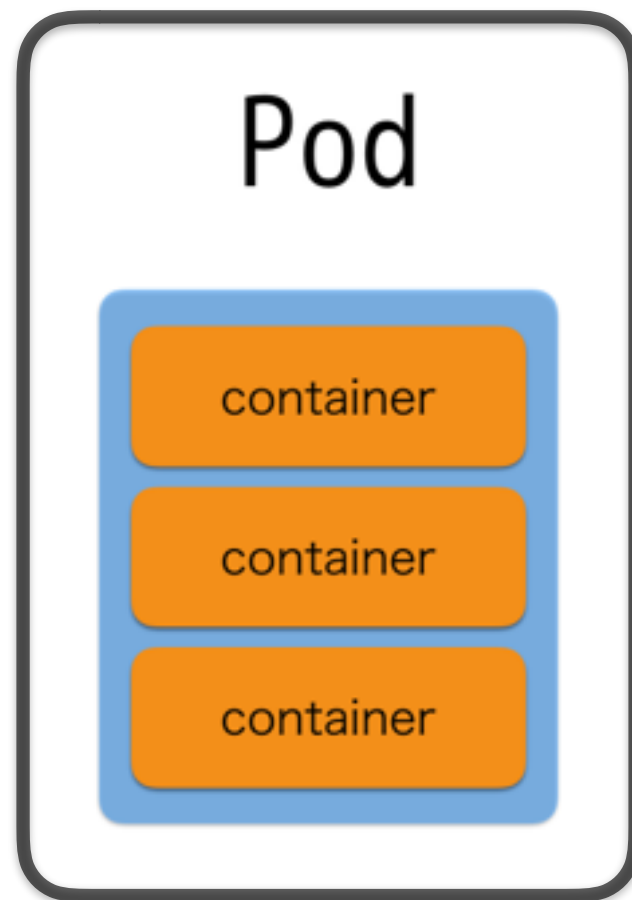
Node



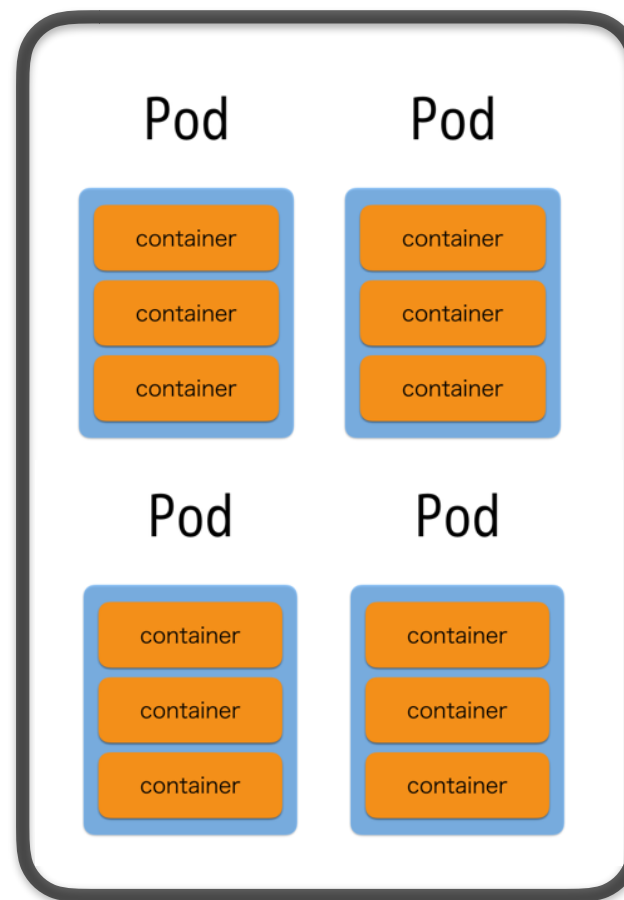
Node



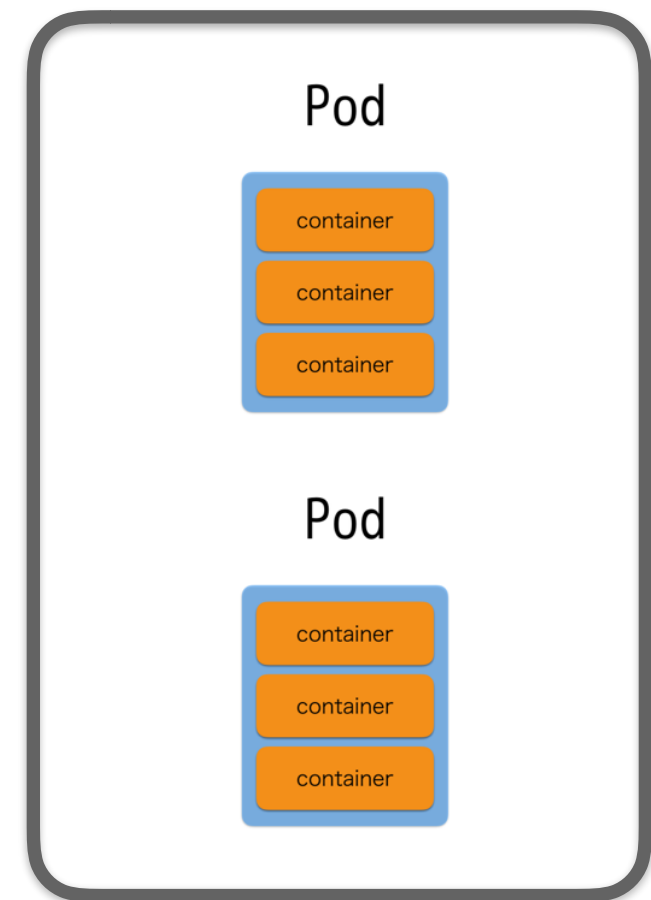
Node



Node




Node




**Depends on each node's resource availability
and each pod's resource requirements**

Node



Node



Node

Node

This will NEVER happen

container

container

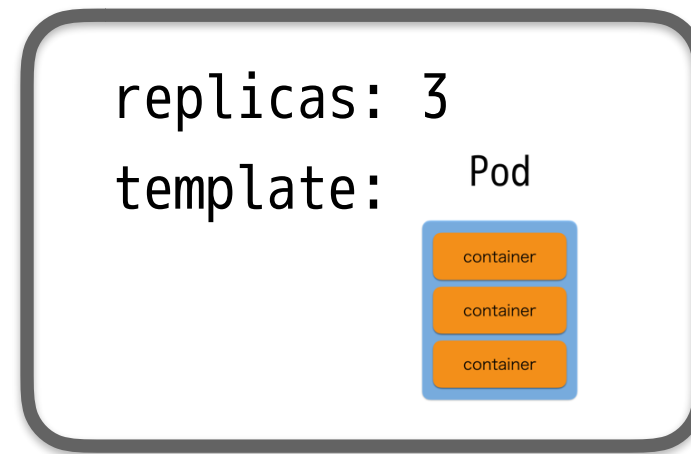
Miscellaneous

- **Each pod has its own IP address**
- **Pods are expected to be stateless**

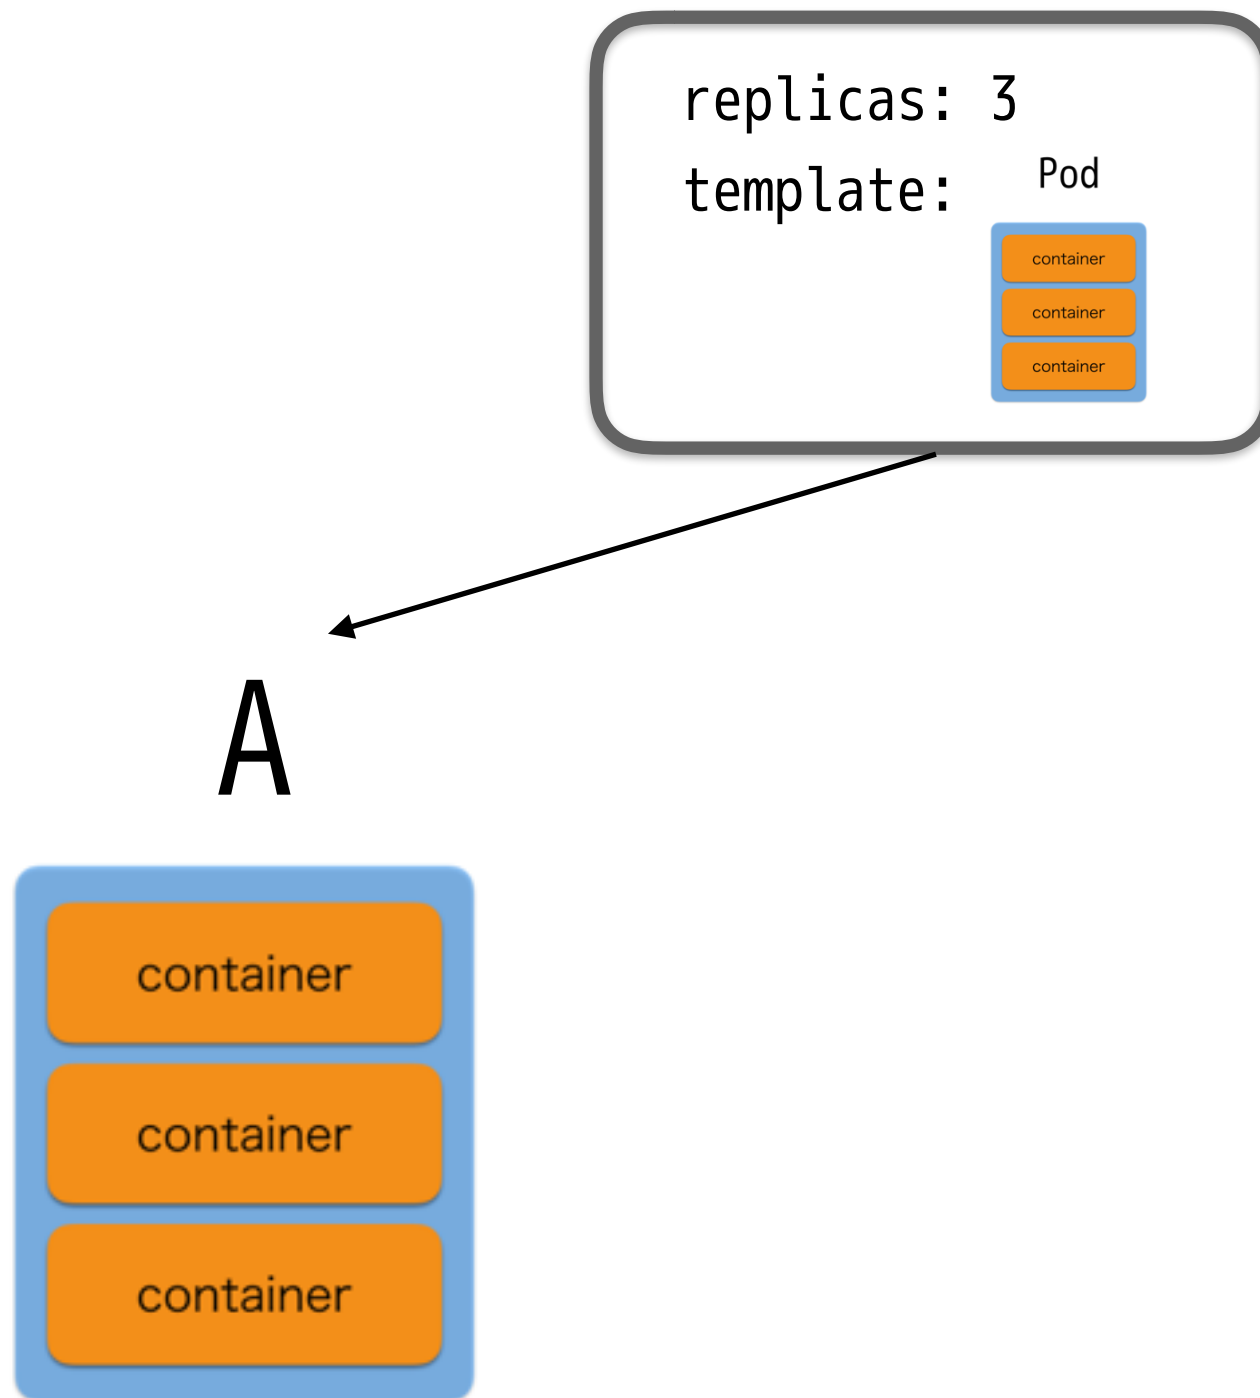
Replica Set

**Keeps track of
Pod replicas**


Replica Set



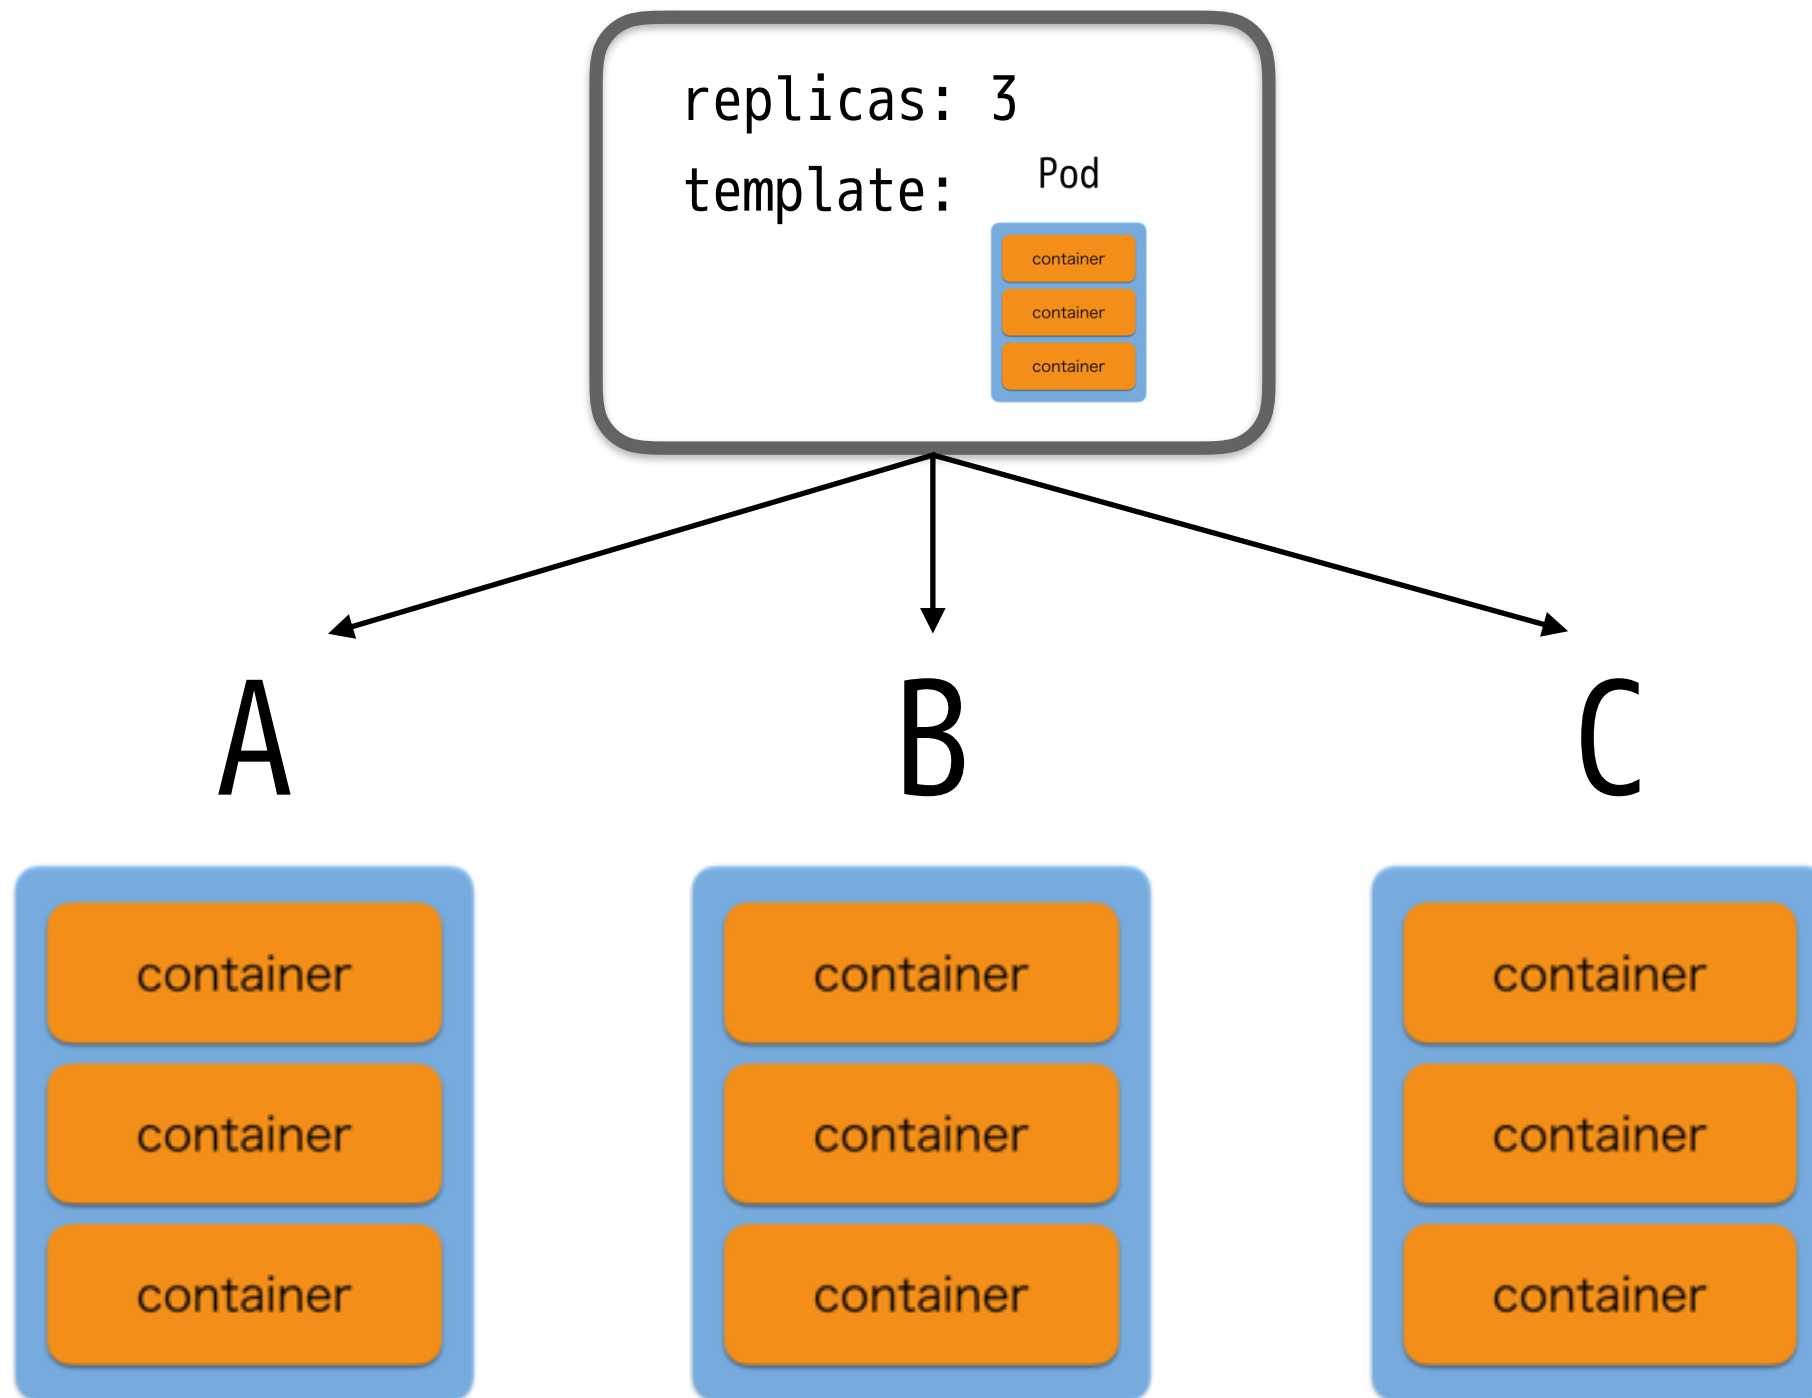
Replica Set



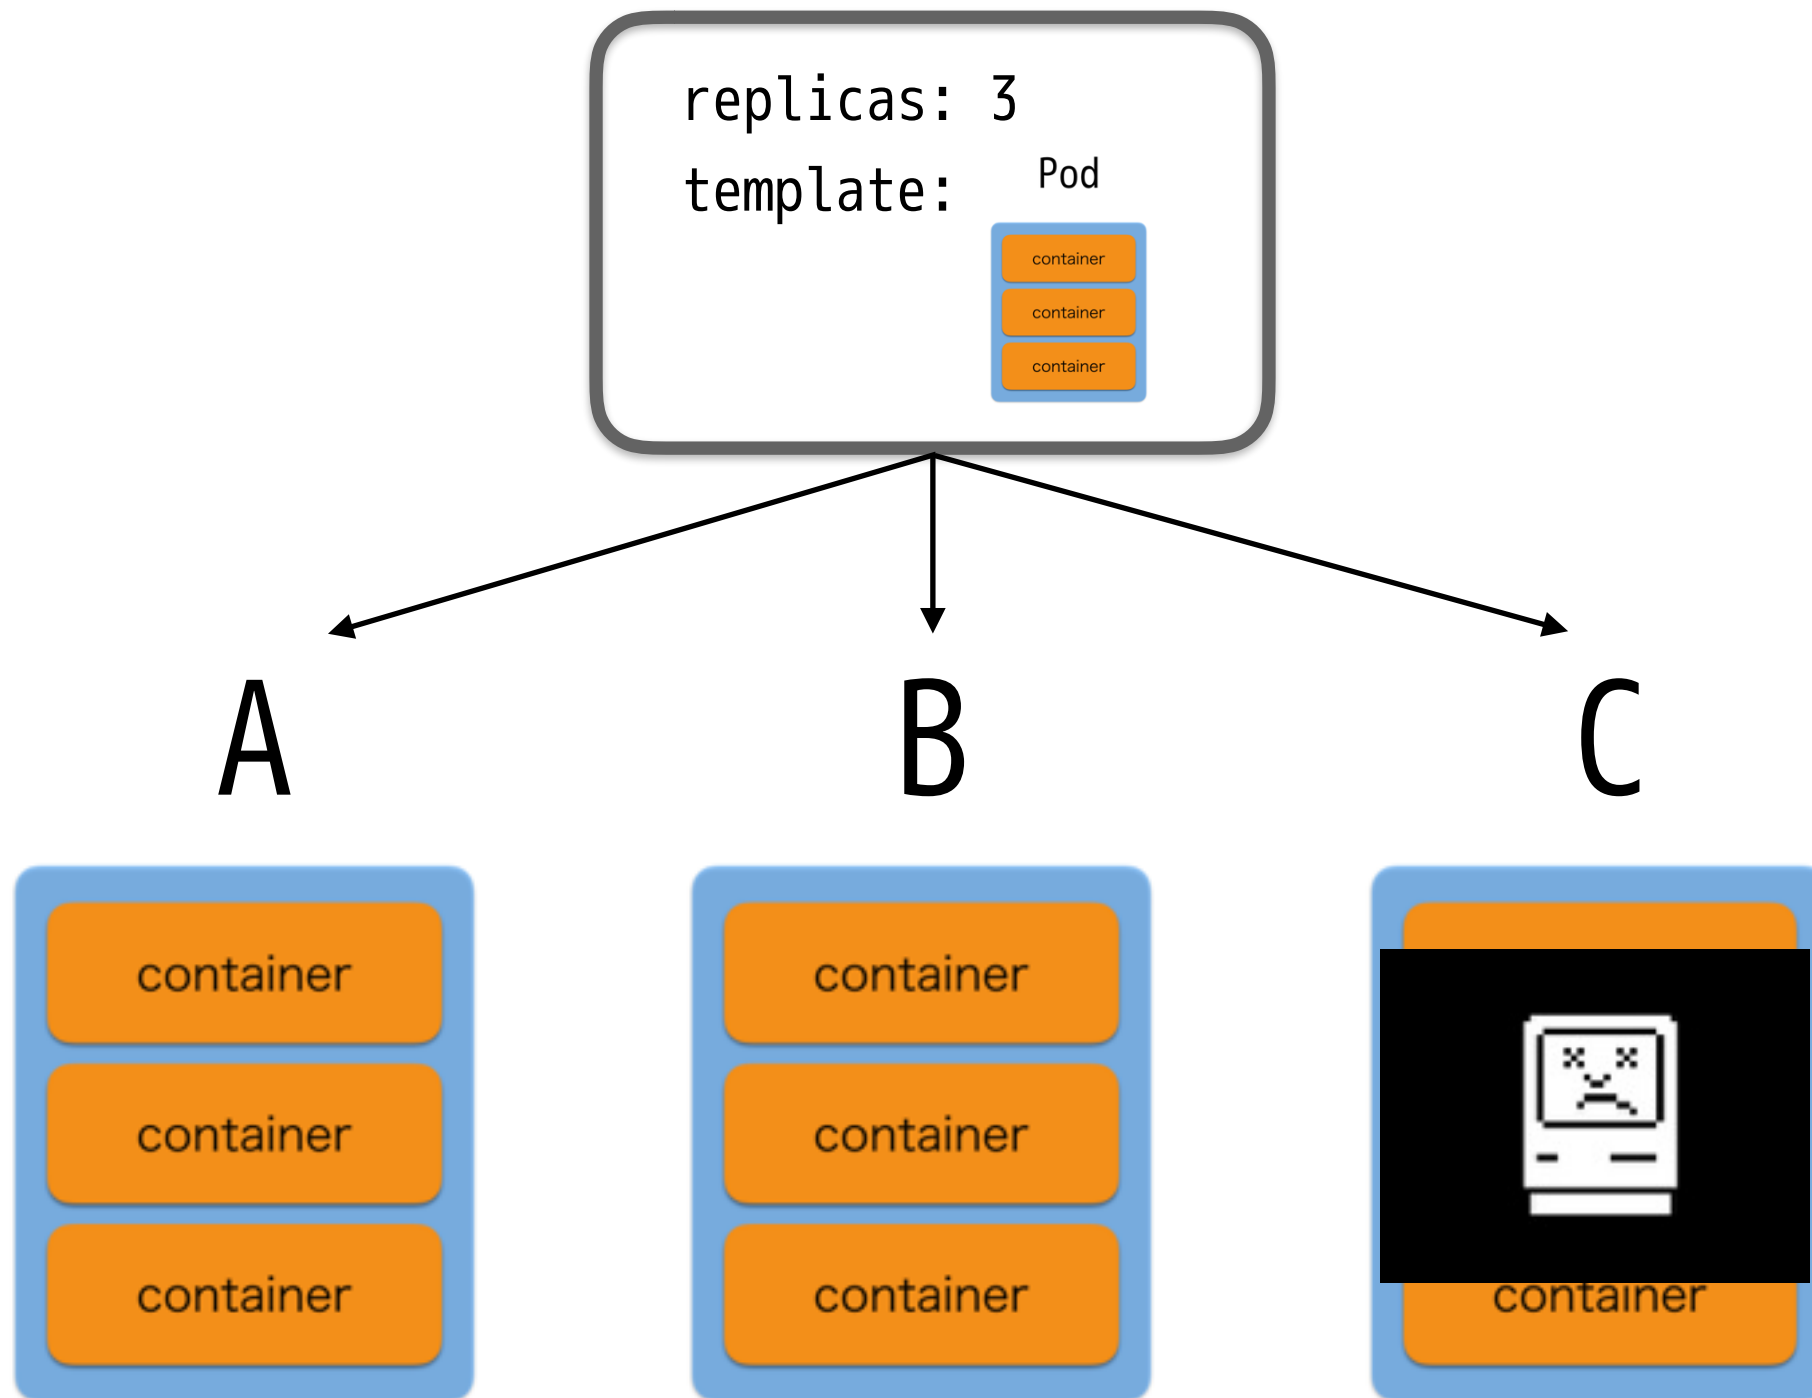
Replica Set




Replica Set



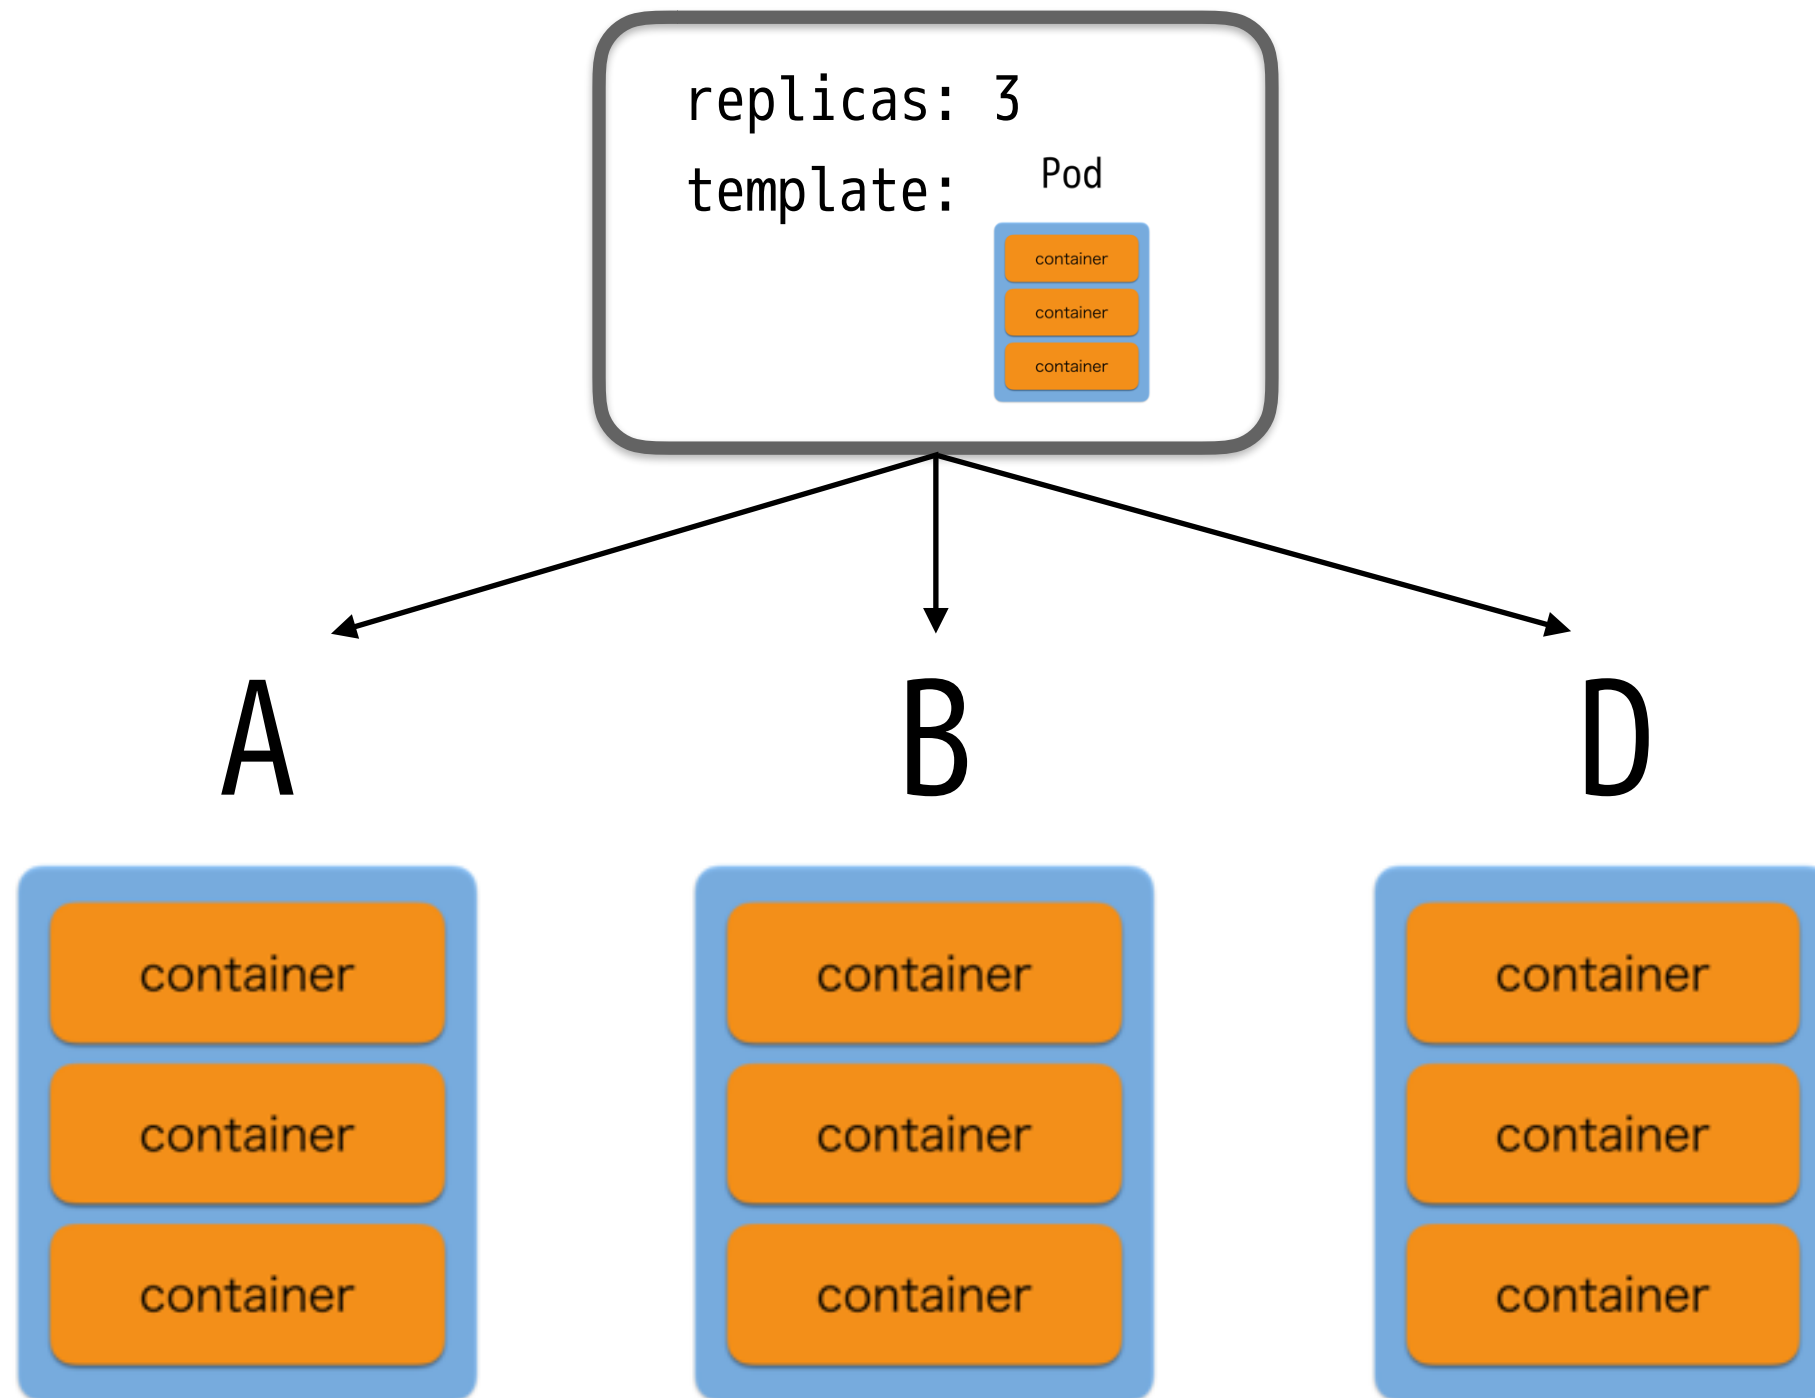
Replica Set



Replica Set



Replica Set



Deployment

Manages Replica Set state transitions

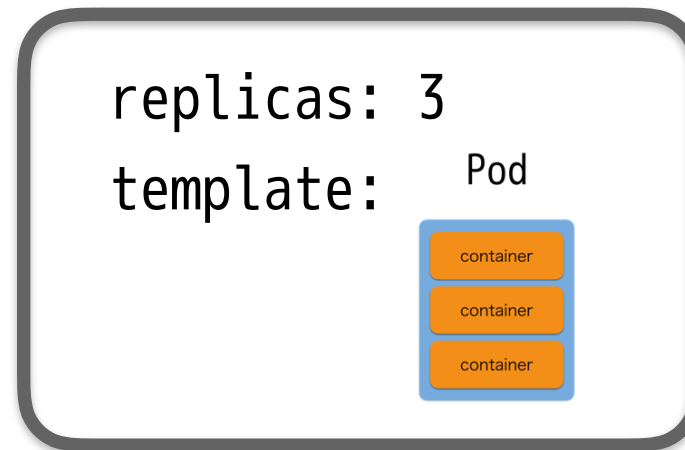
Deployment

replicas: 3

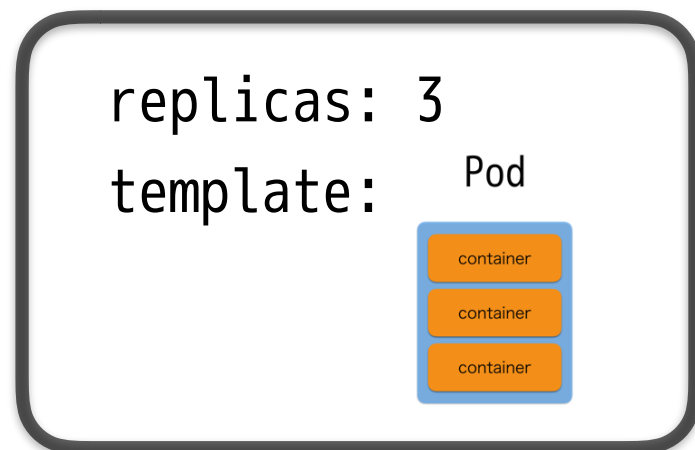
template: Pod




Deployment



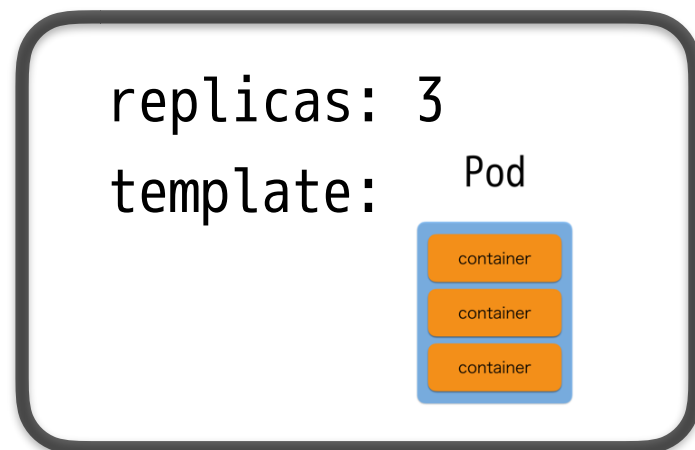
Replica Set A



Deployment



Replica Set A



Deployment

Change some settings,
please!

replicas: **2**
template: Pod

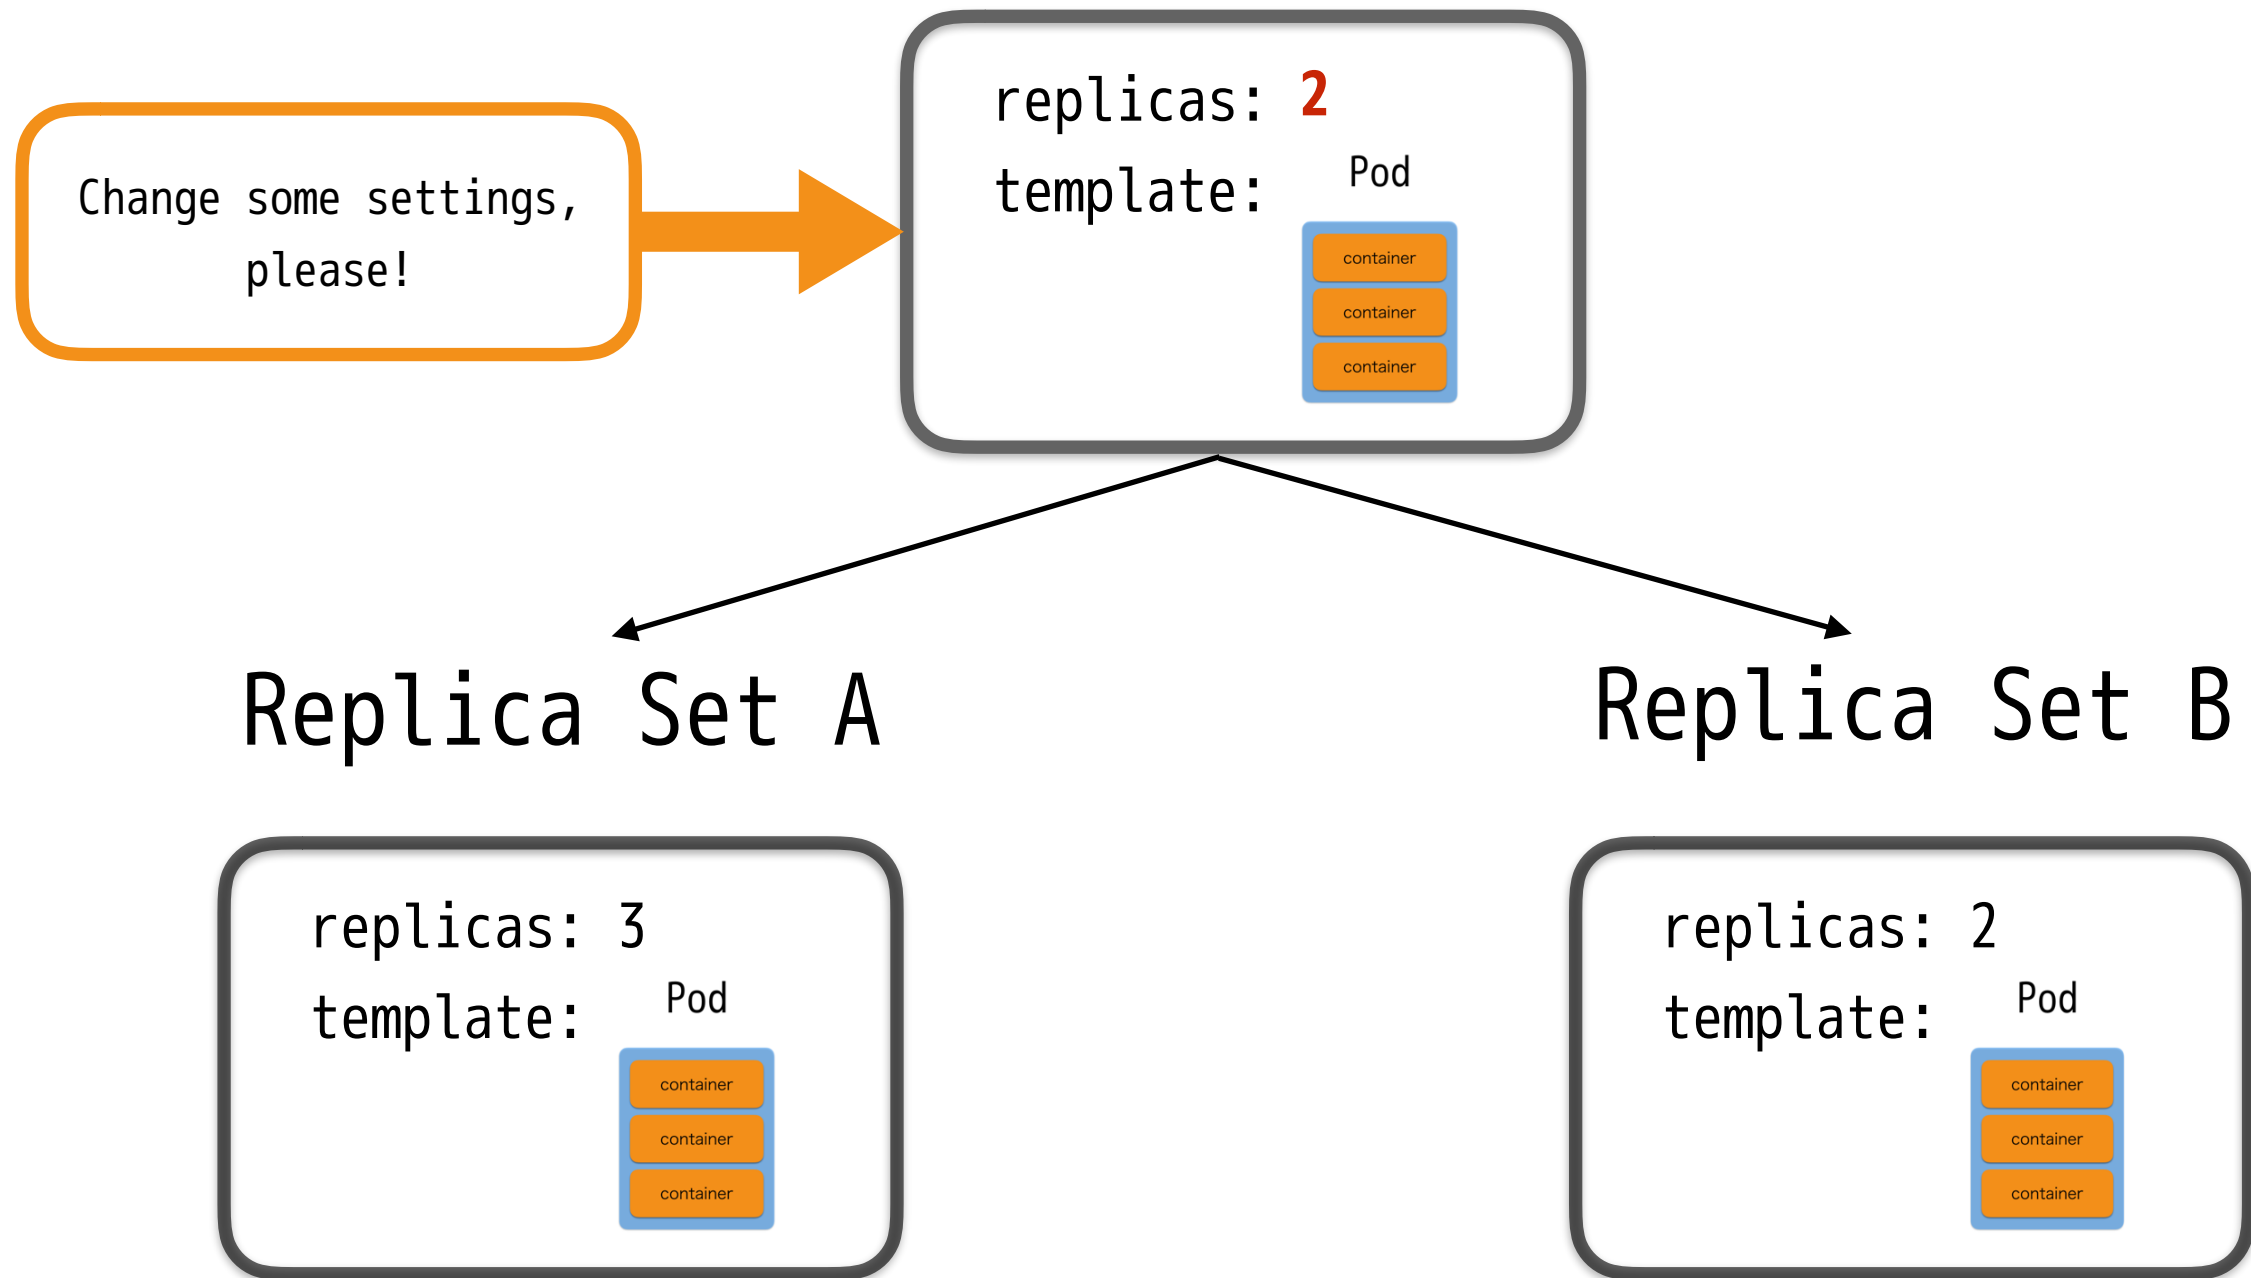
container
container
container

Replica Set A

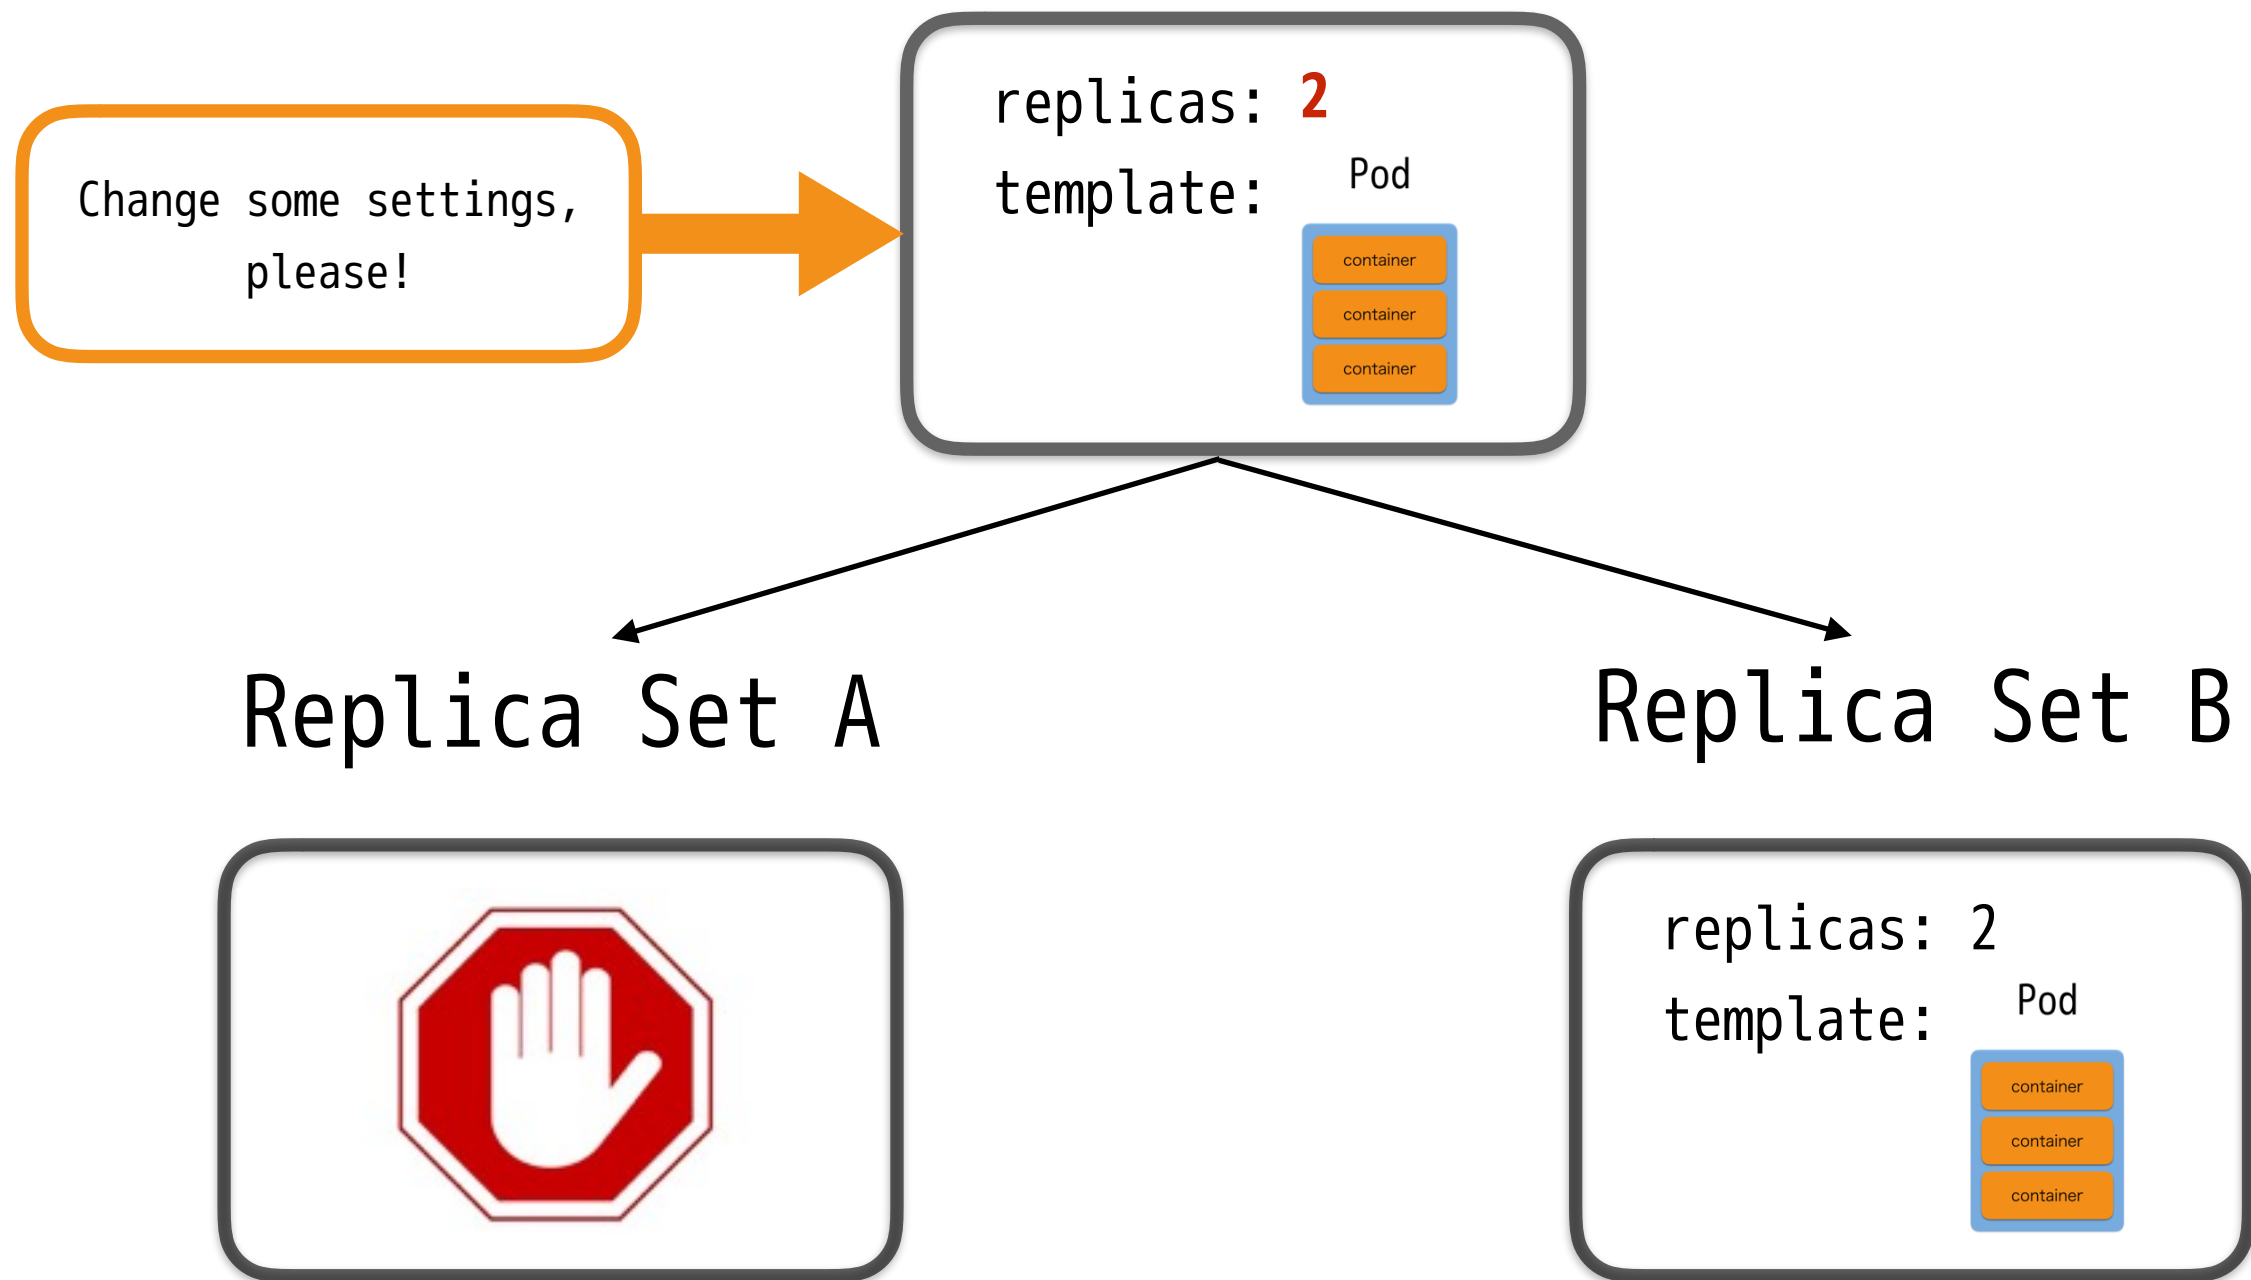
replicas: 3
template: Pod

container
container
container

Deployment



Deployment



Deployment

Change some settings,
please!

replicas: **2**
template: Pod



Replica Set B

replicas: 2
template: Pod



**Keeps track of
state change history**

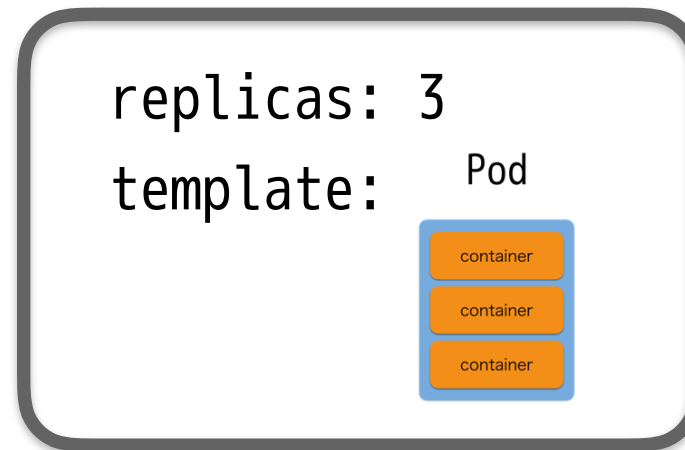
Deployment

replicas: 3

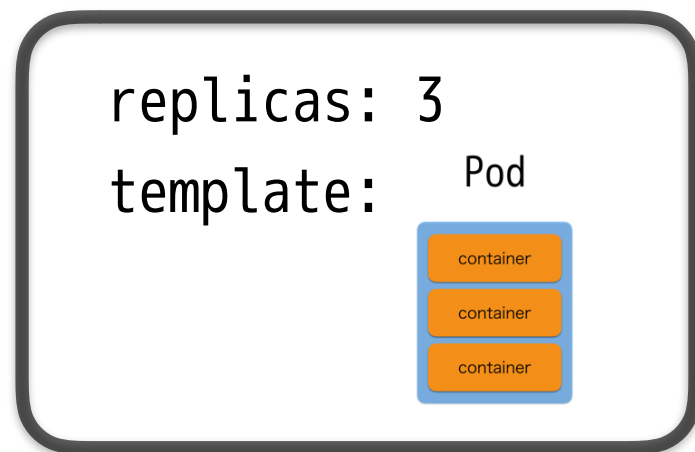
template: Pod



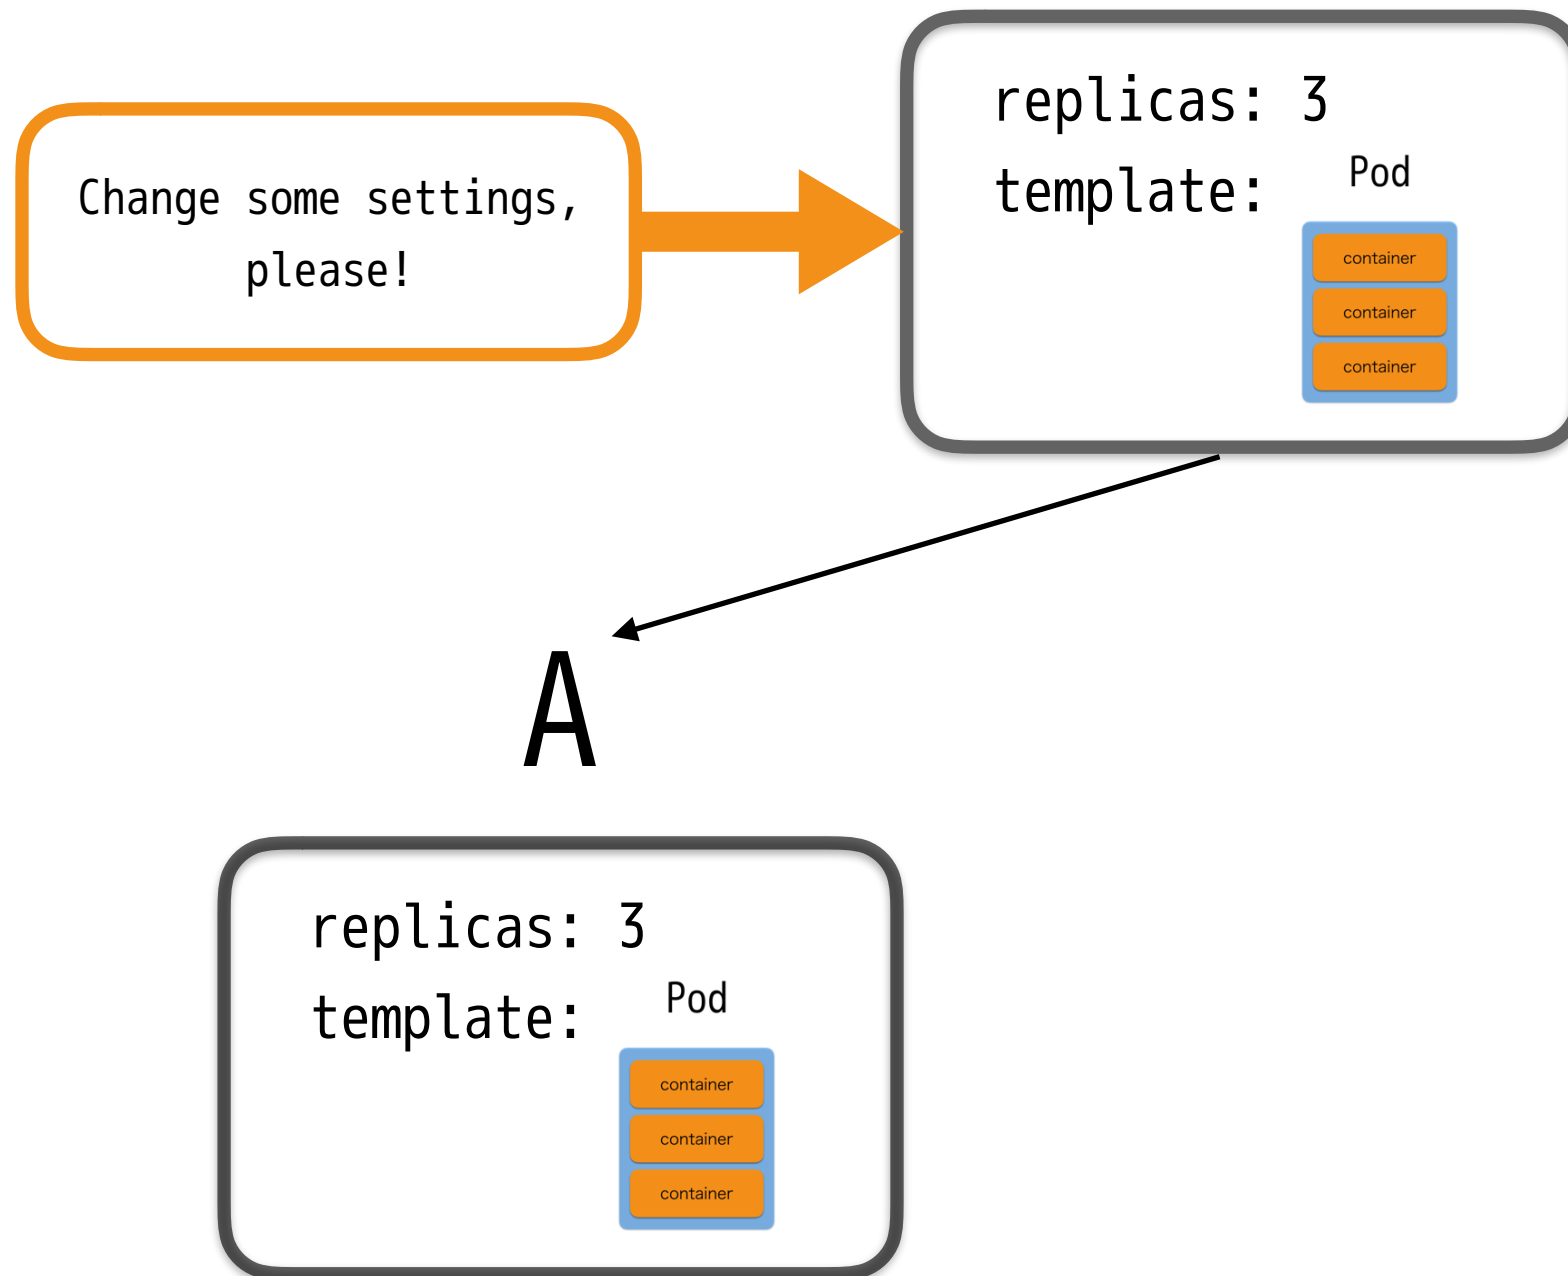
Deployment



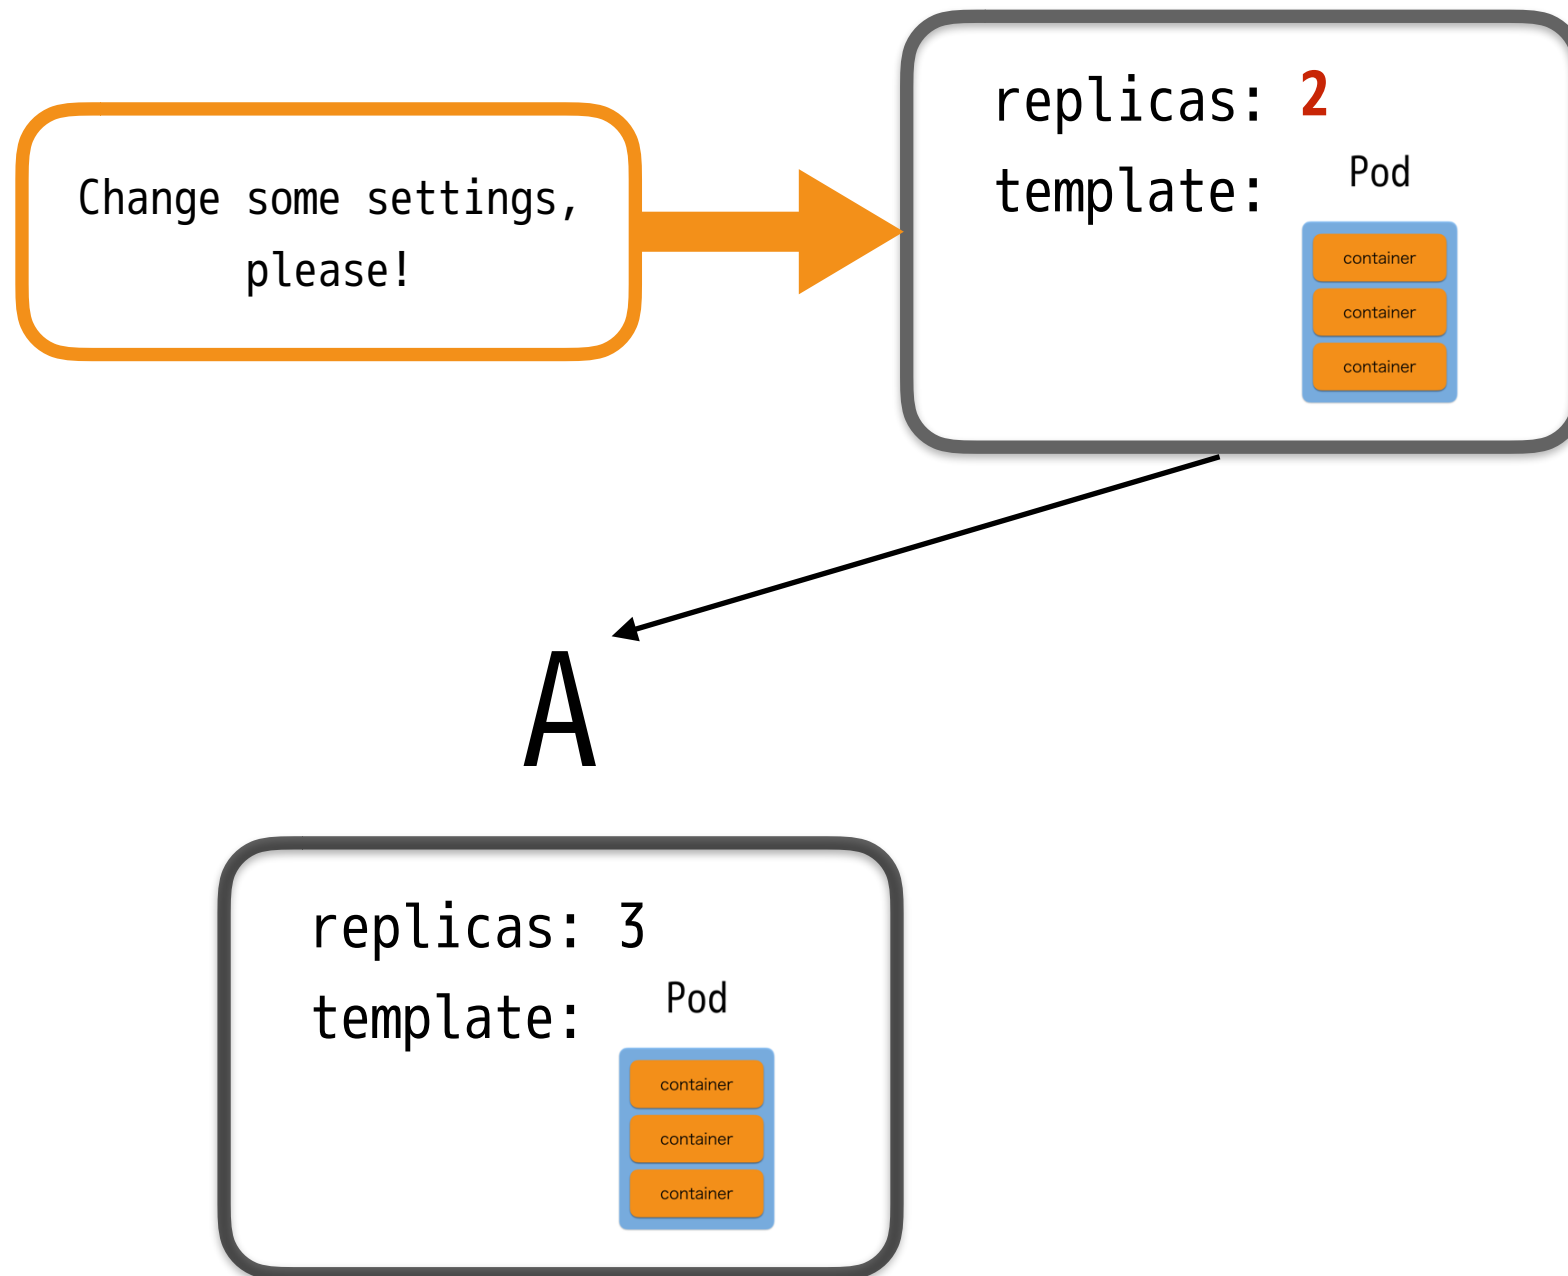
A



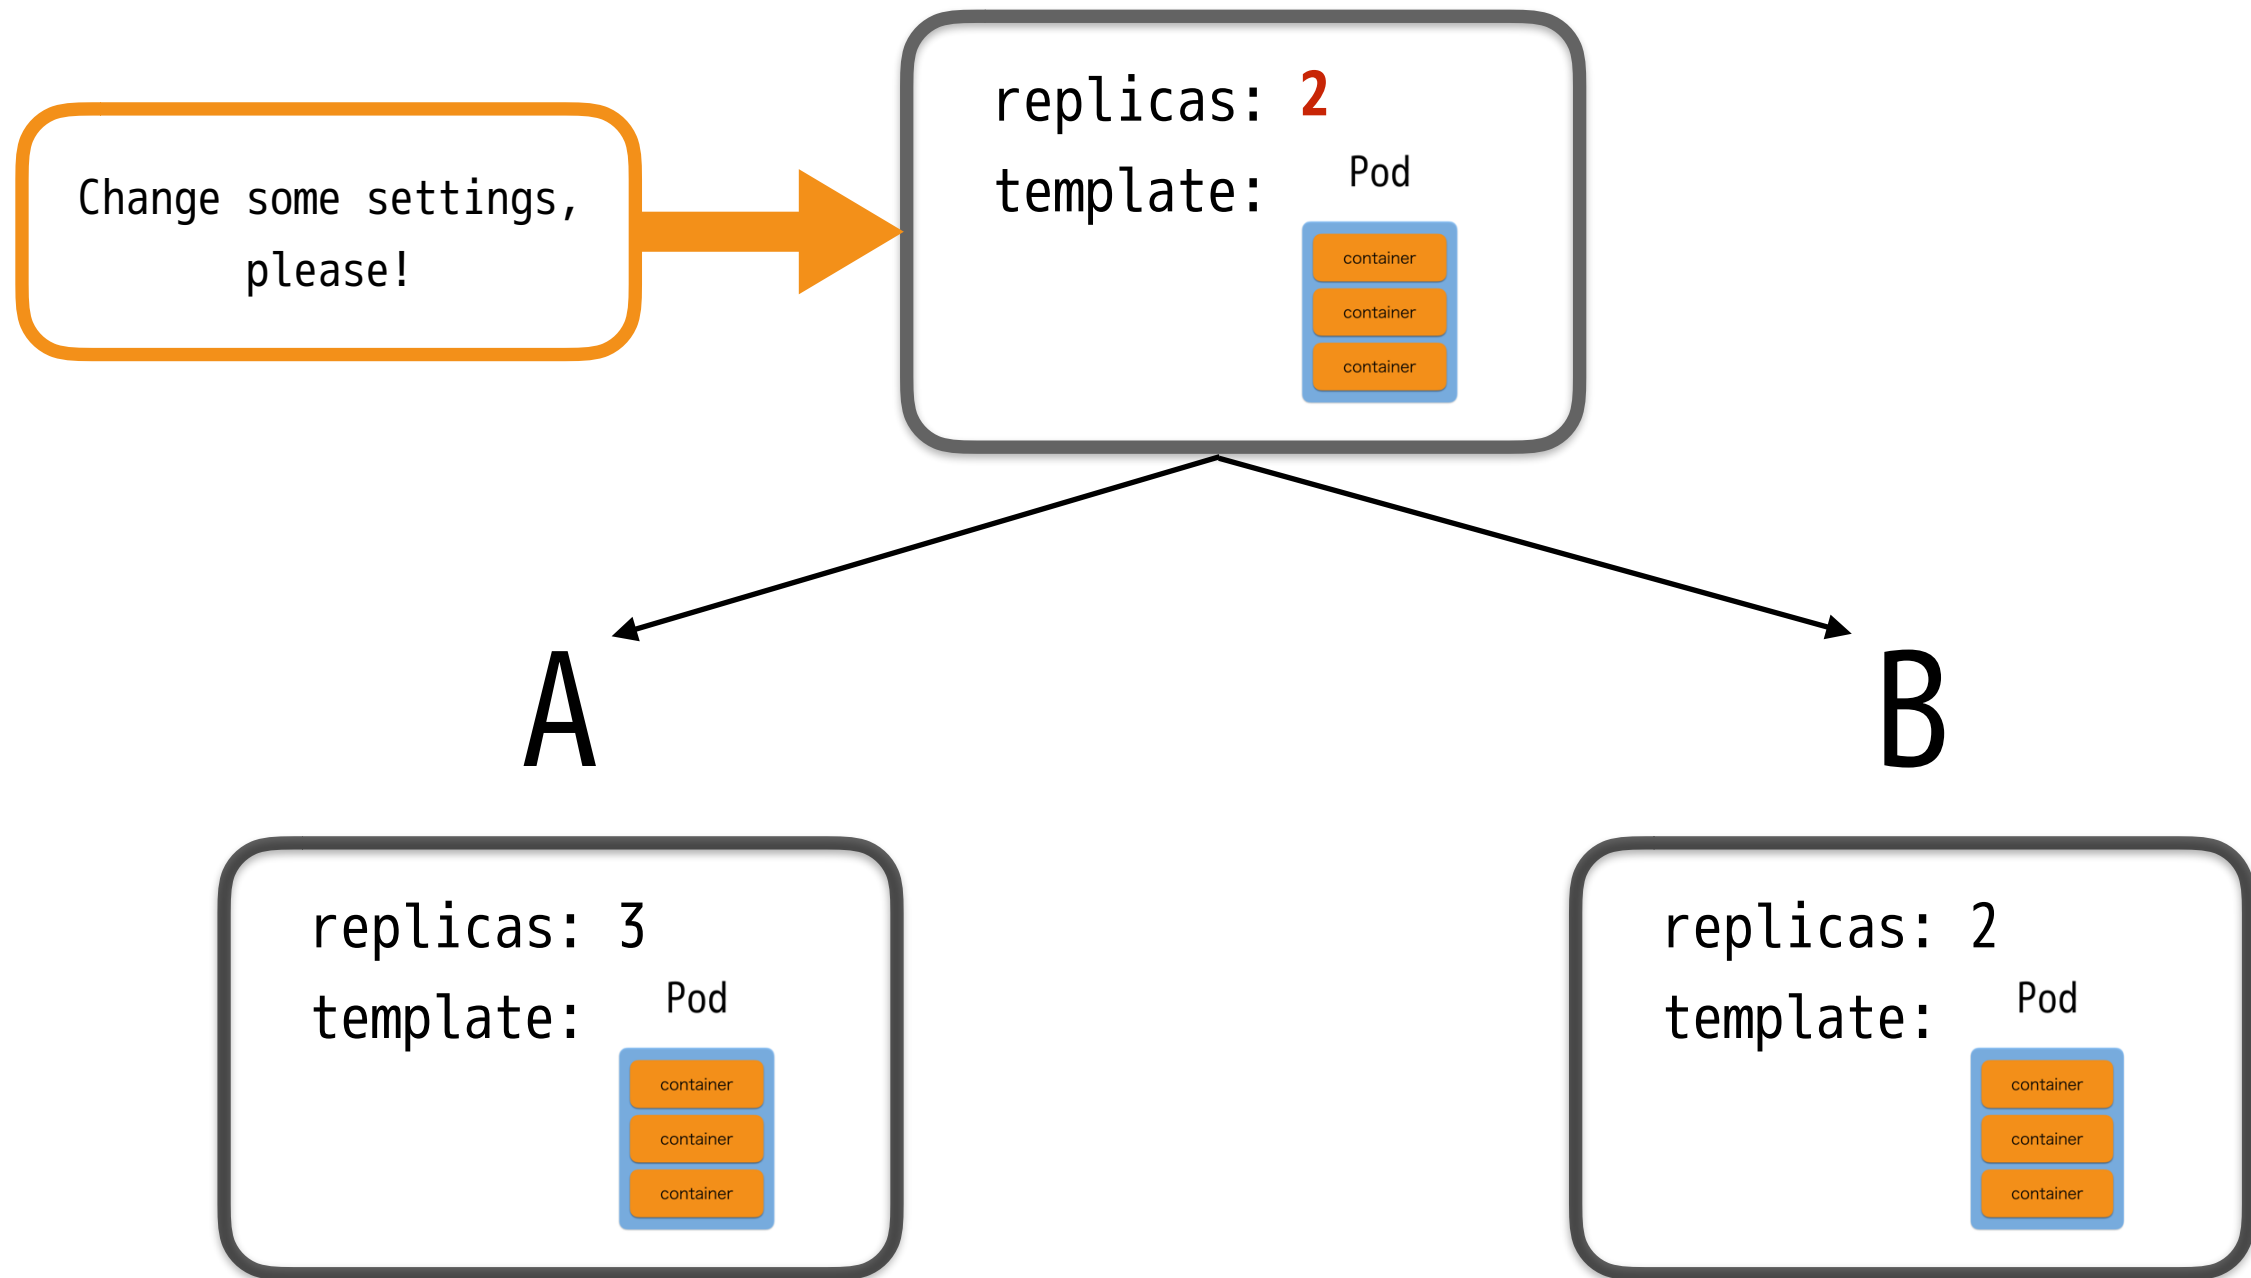
Deployment



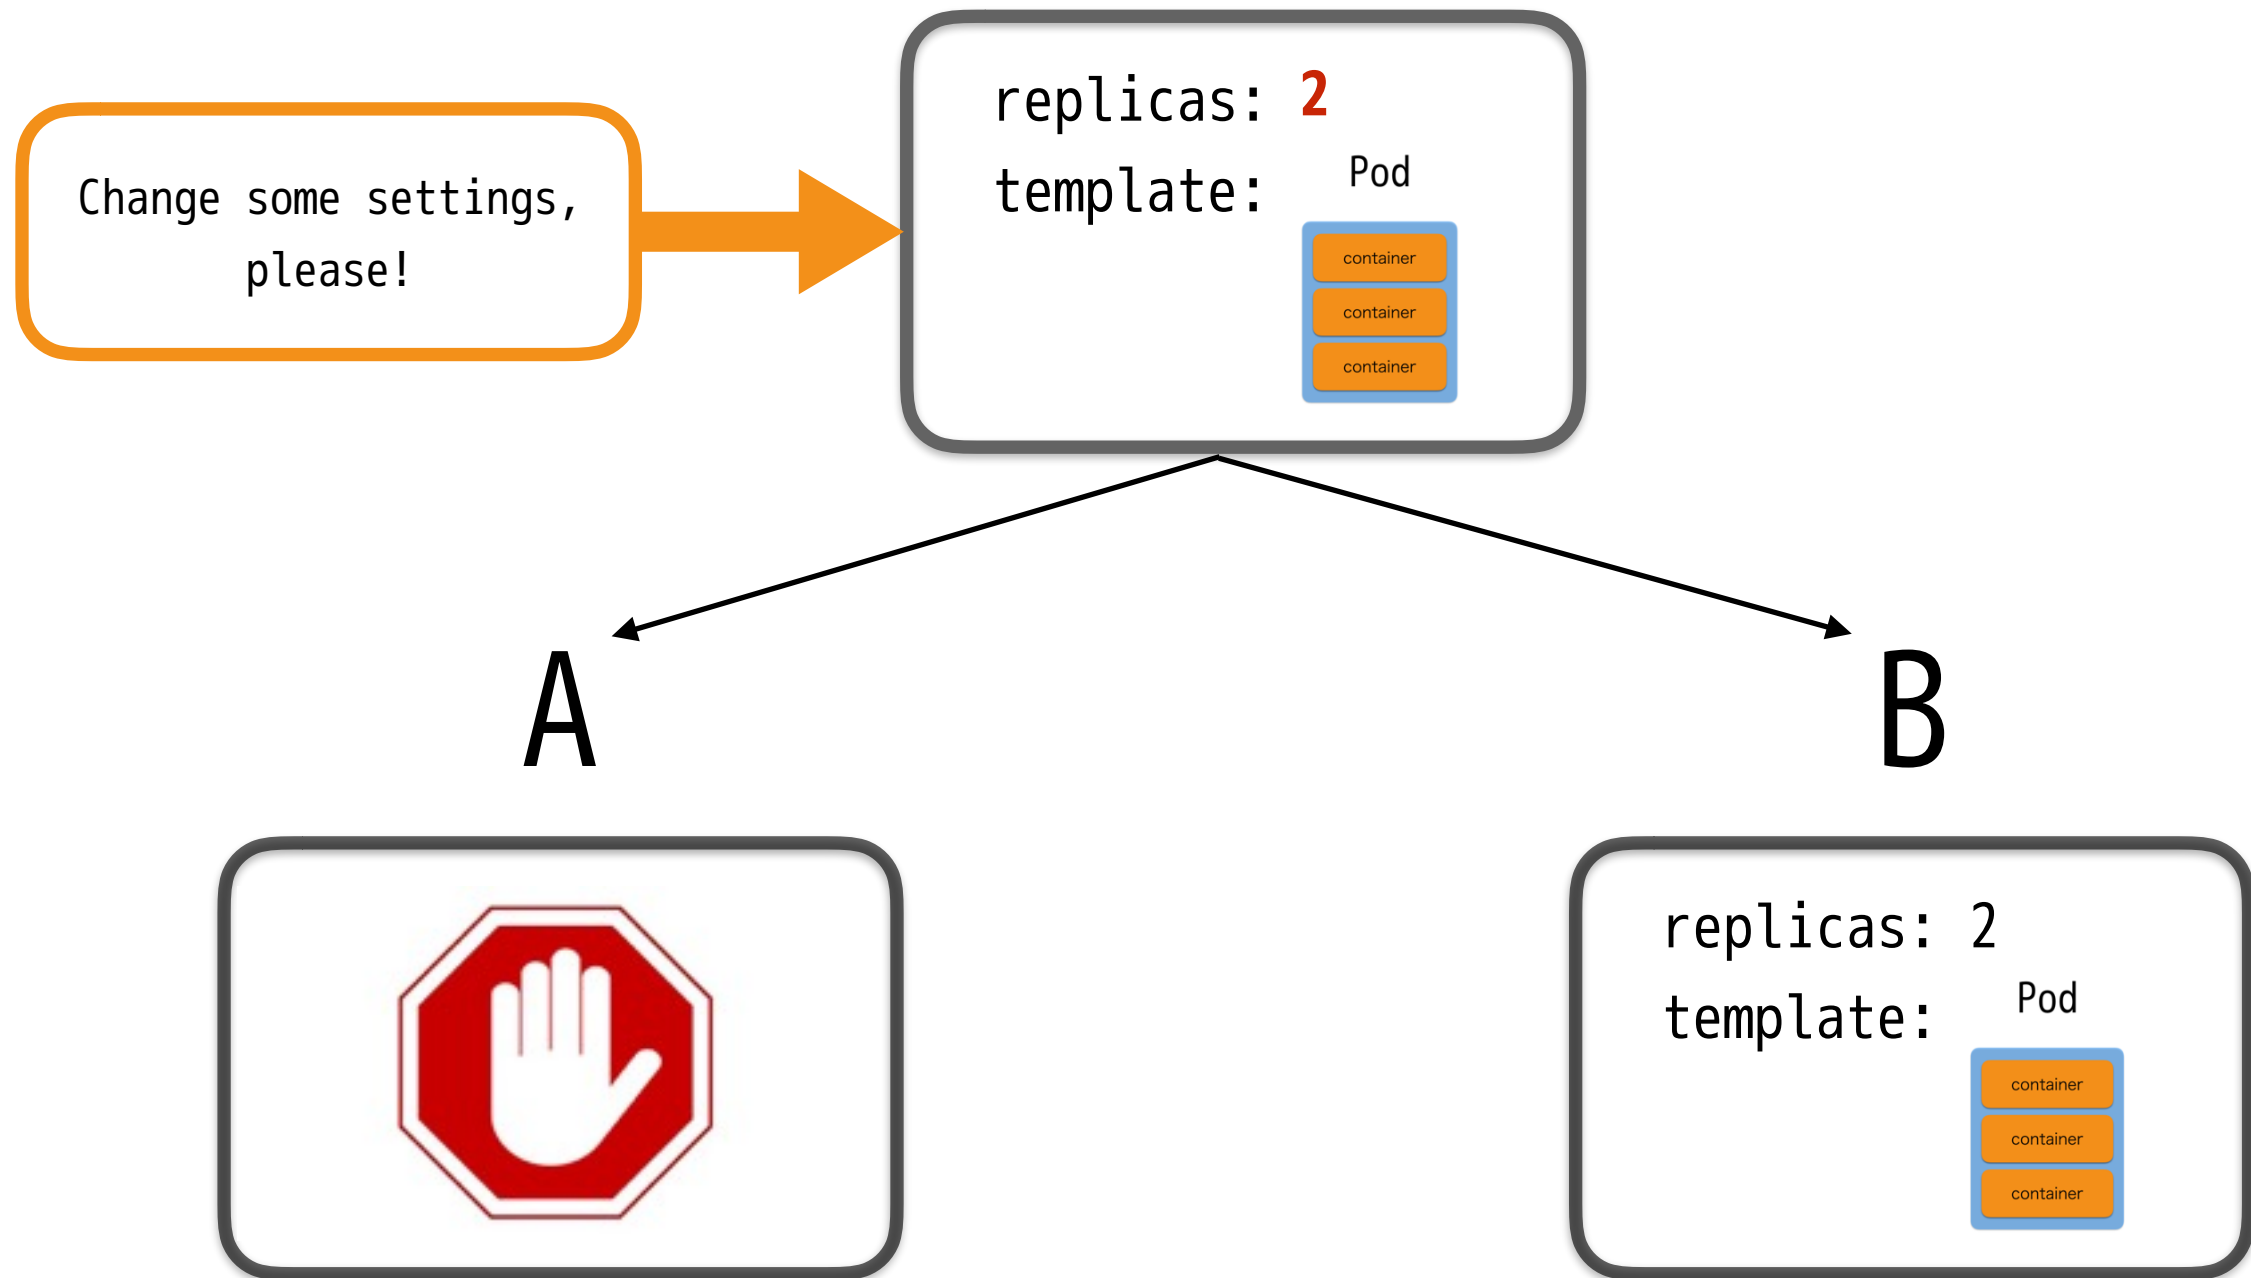
Deployment



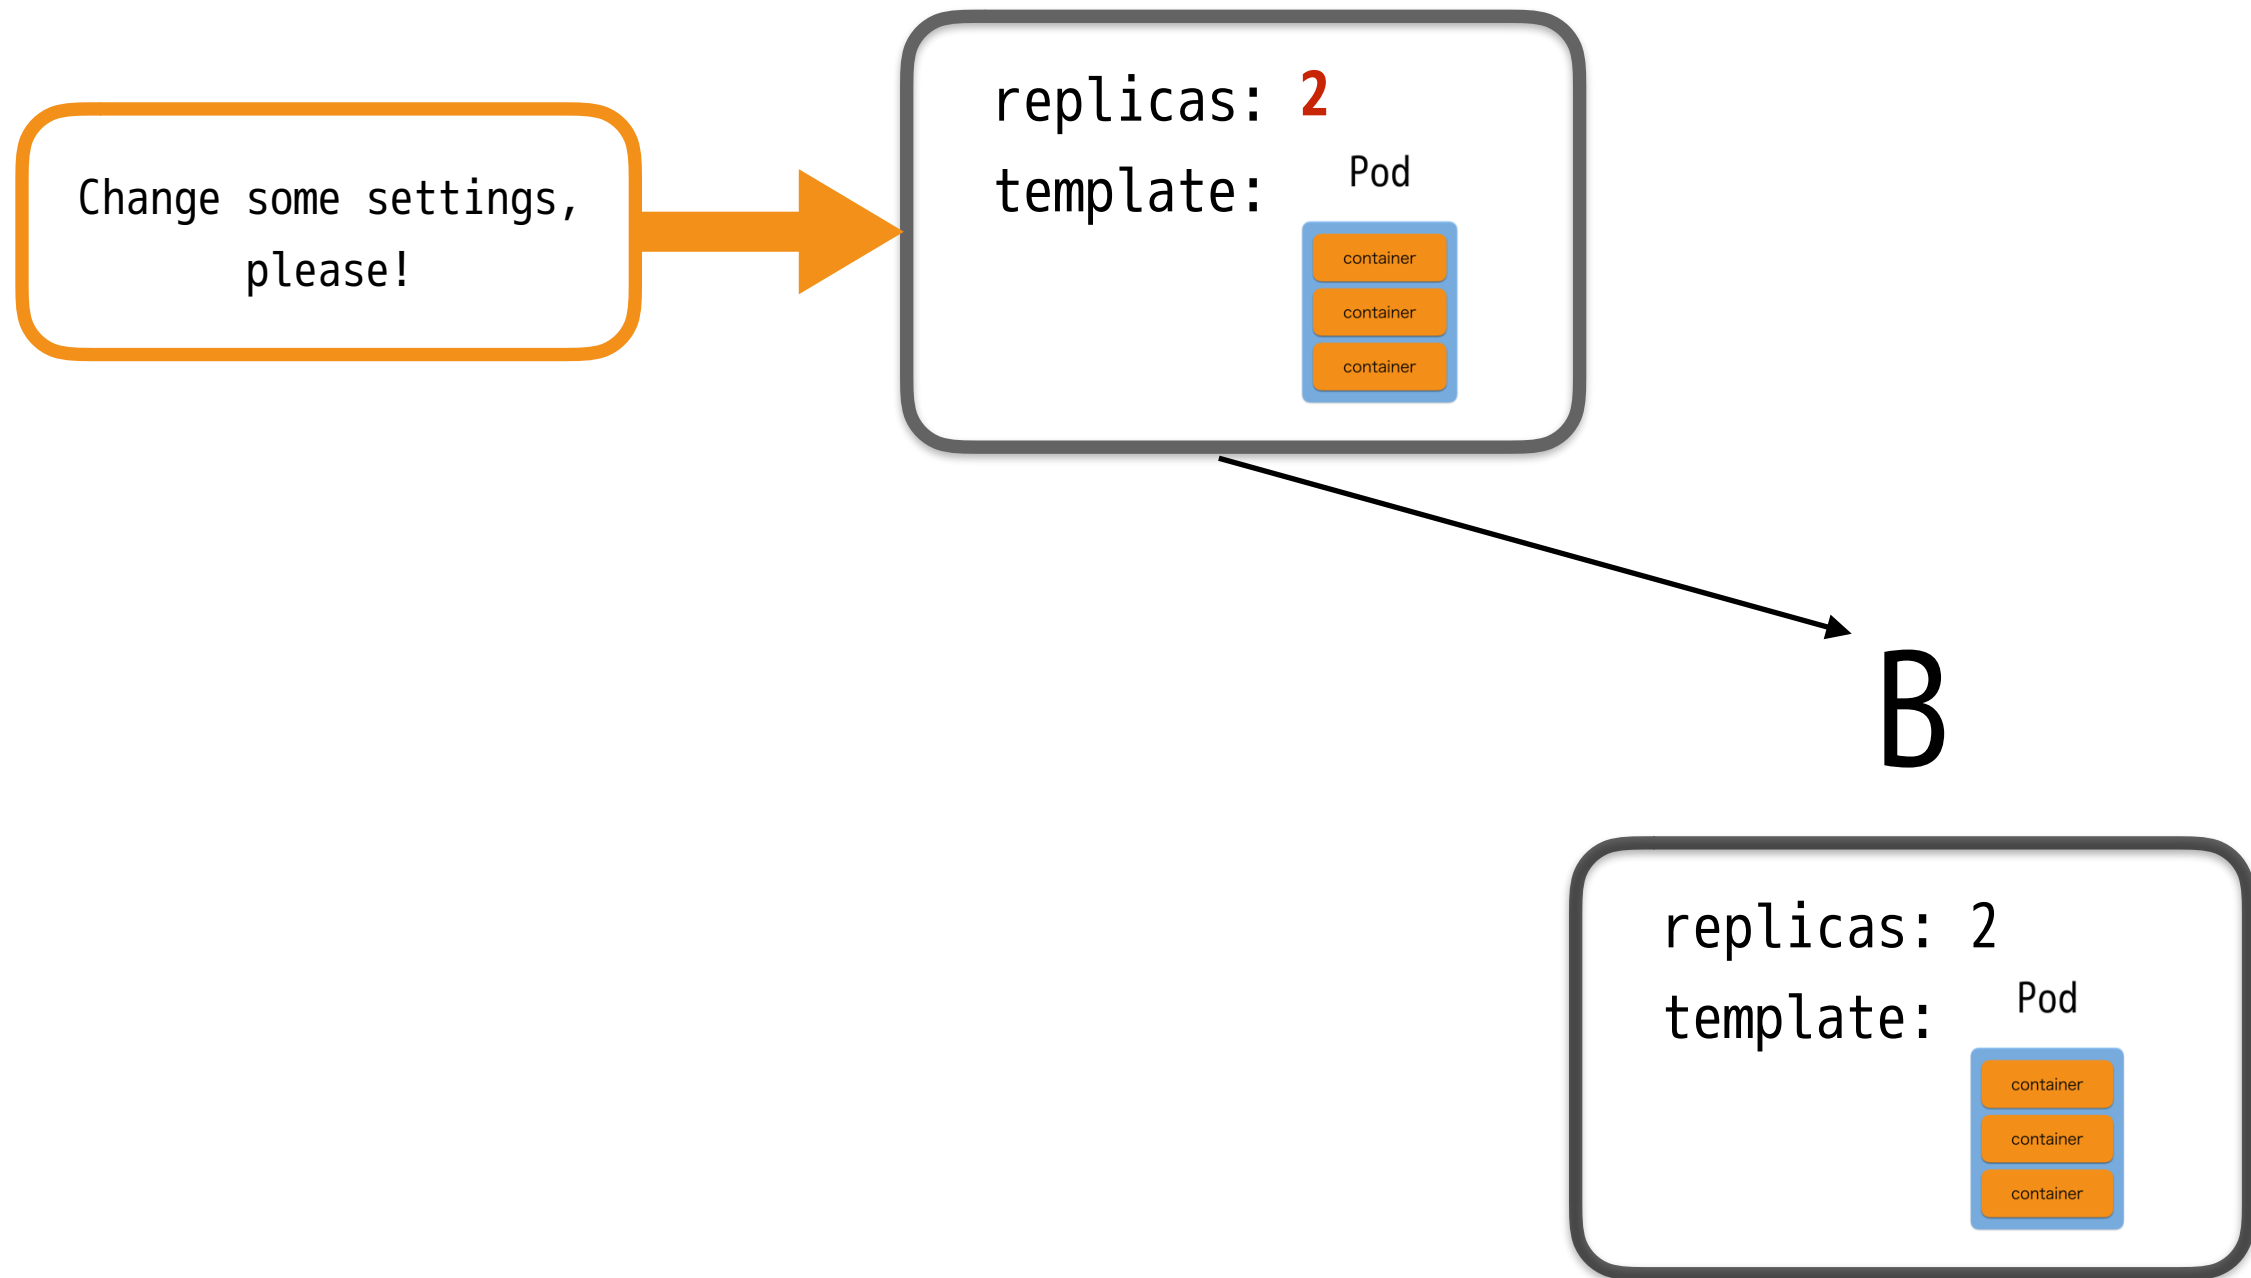
Deployment



Deployment



Deployment



Deployment

Oh wait, I actually
didn't want to do that...

replicas: **2**
template: Pod



B

replicas: 2
template: Pod



Deployment

Oh wait, I actually
didn't want to do that...

replicas: **3**
template: Pod



B

replicas: 2
template: Pod



Deployment

Oh wait, I actually
didn't want to do that...

replicas: **3**
template: Pod



B

rollback

replicas: 2
template: Pod



Deployment

Oh wait, I actually
didn't want to do that...

replicas: **3**
template: Pod



A

B

replicas: 3
template: Pod

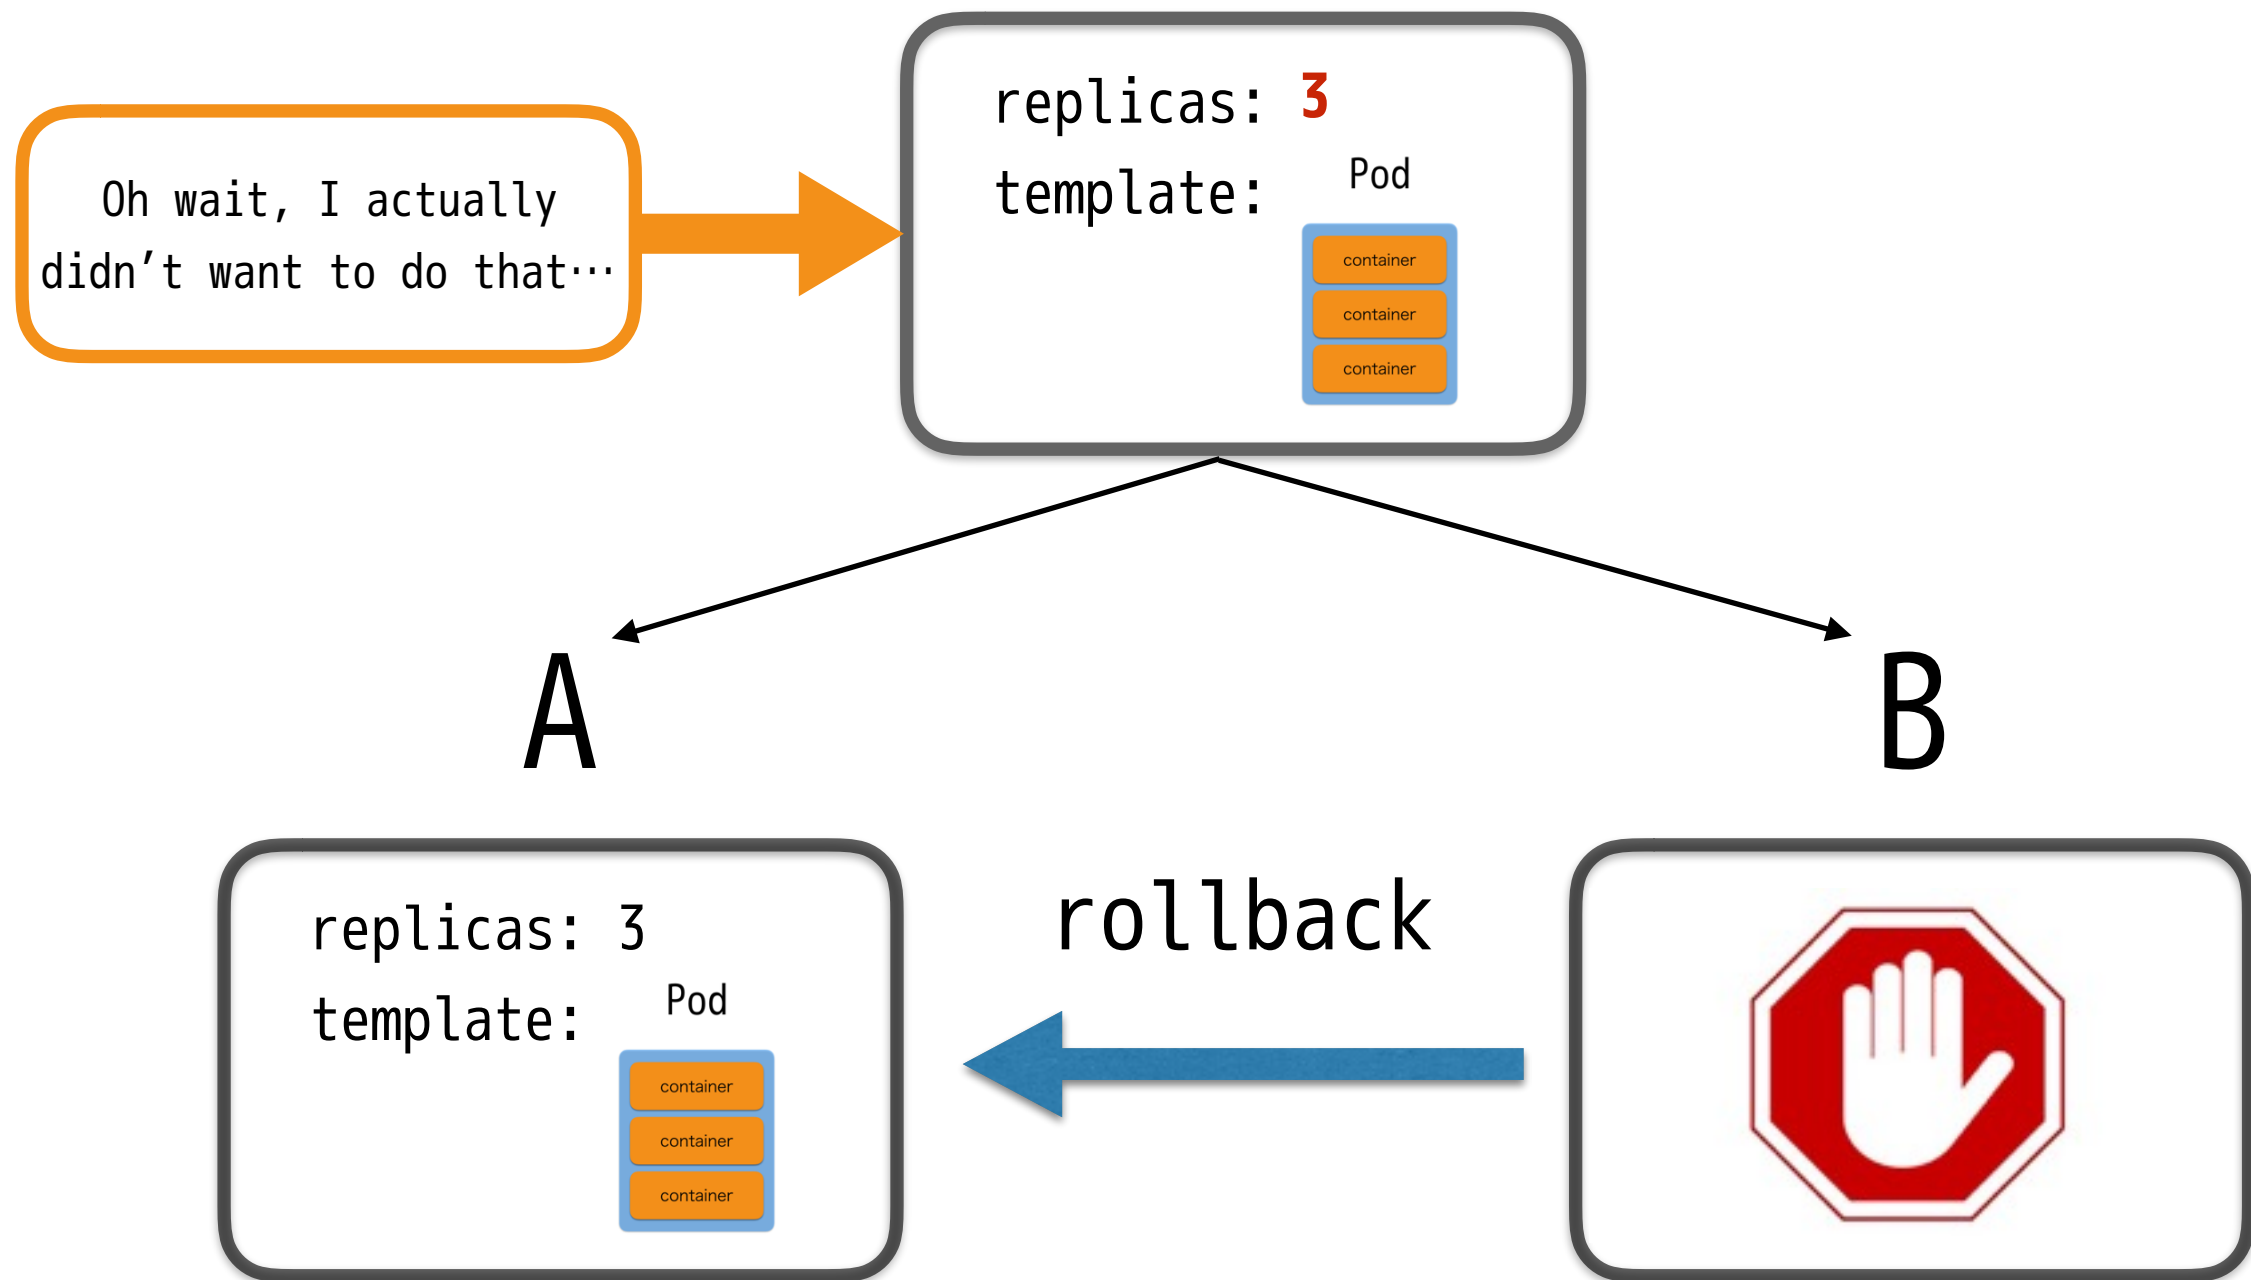


rollback

replicas: 2
template: Pod



Deployment



Deployment

Oh wait, I actually
didn't want to do that...

replicas: **3**
template: Pod

container
container
container

A

replicas: 3
template: Pod

container
container
container

Services

Logical set of Pods

(and ways to access them)

Raw Pod Access




Pod 10.0.96.2


Pod 10.0.96.3

Pod 10.0.96.3


Raw Pod Access




Raw Pod Access



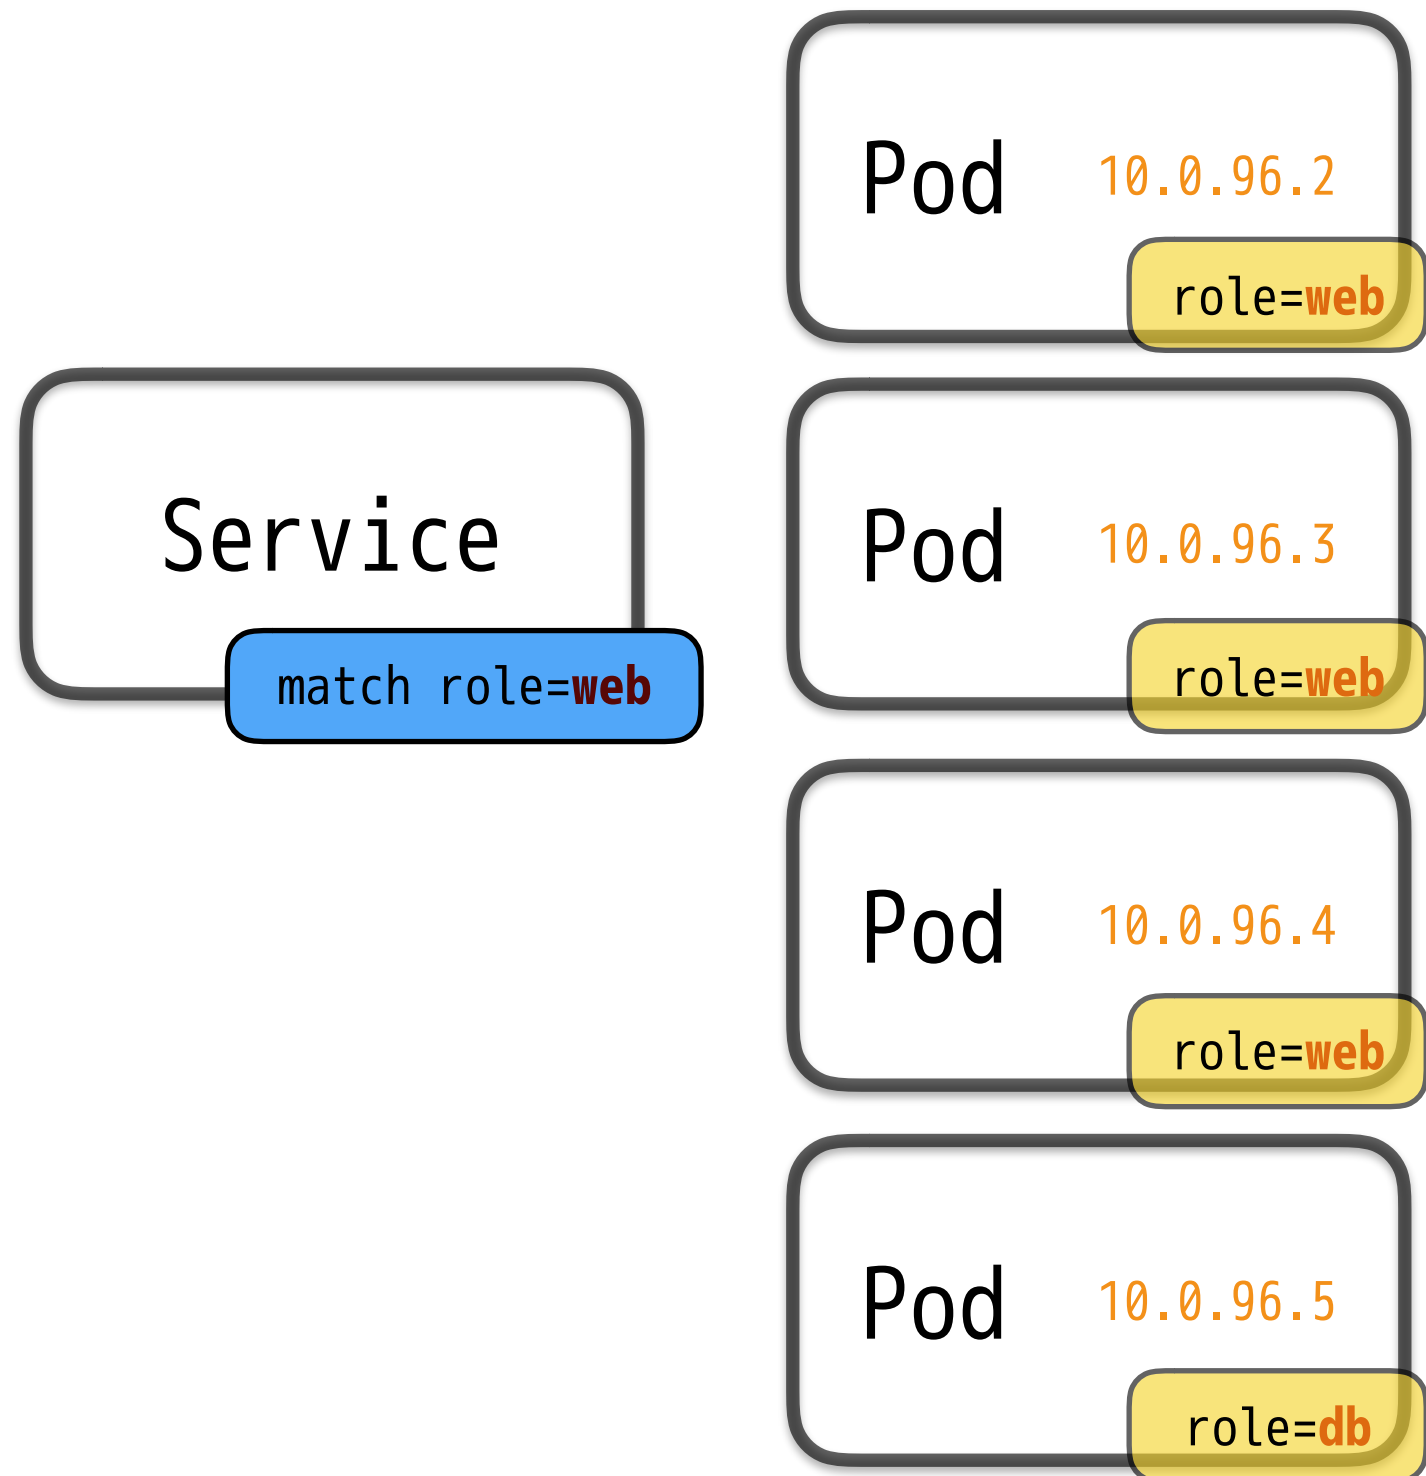
Raw Pod Access



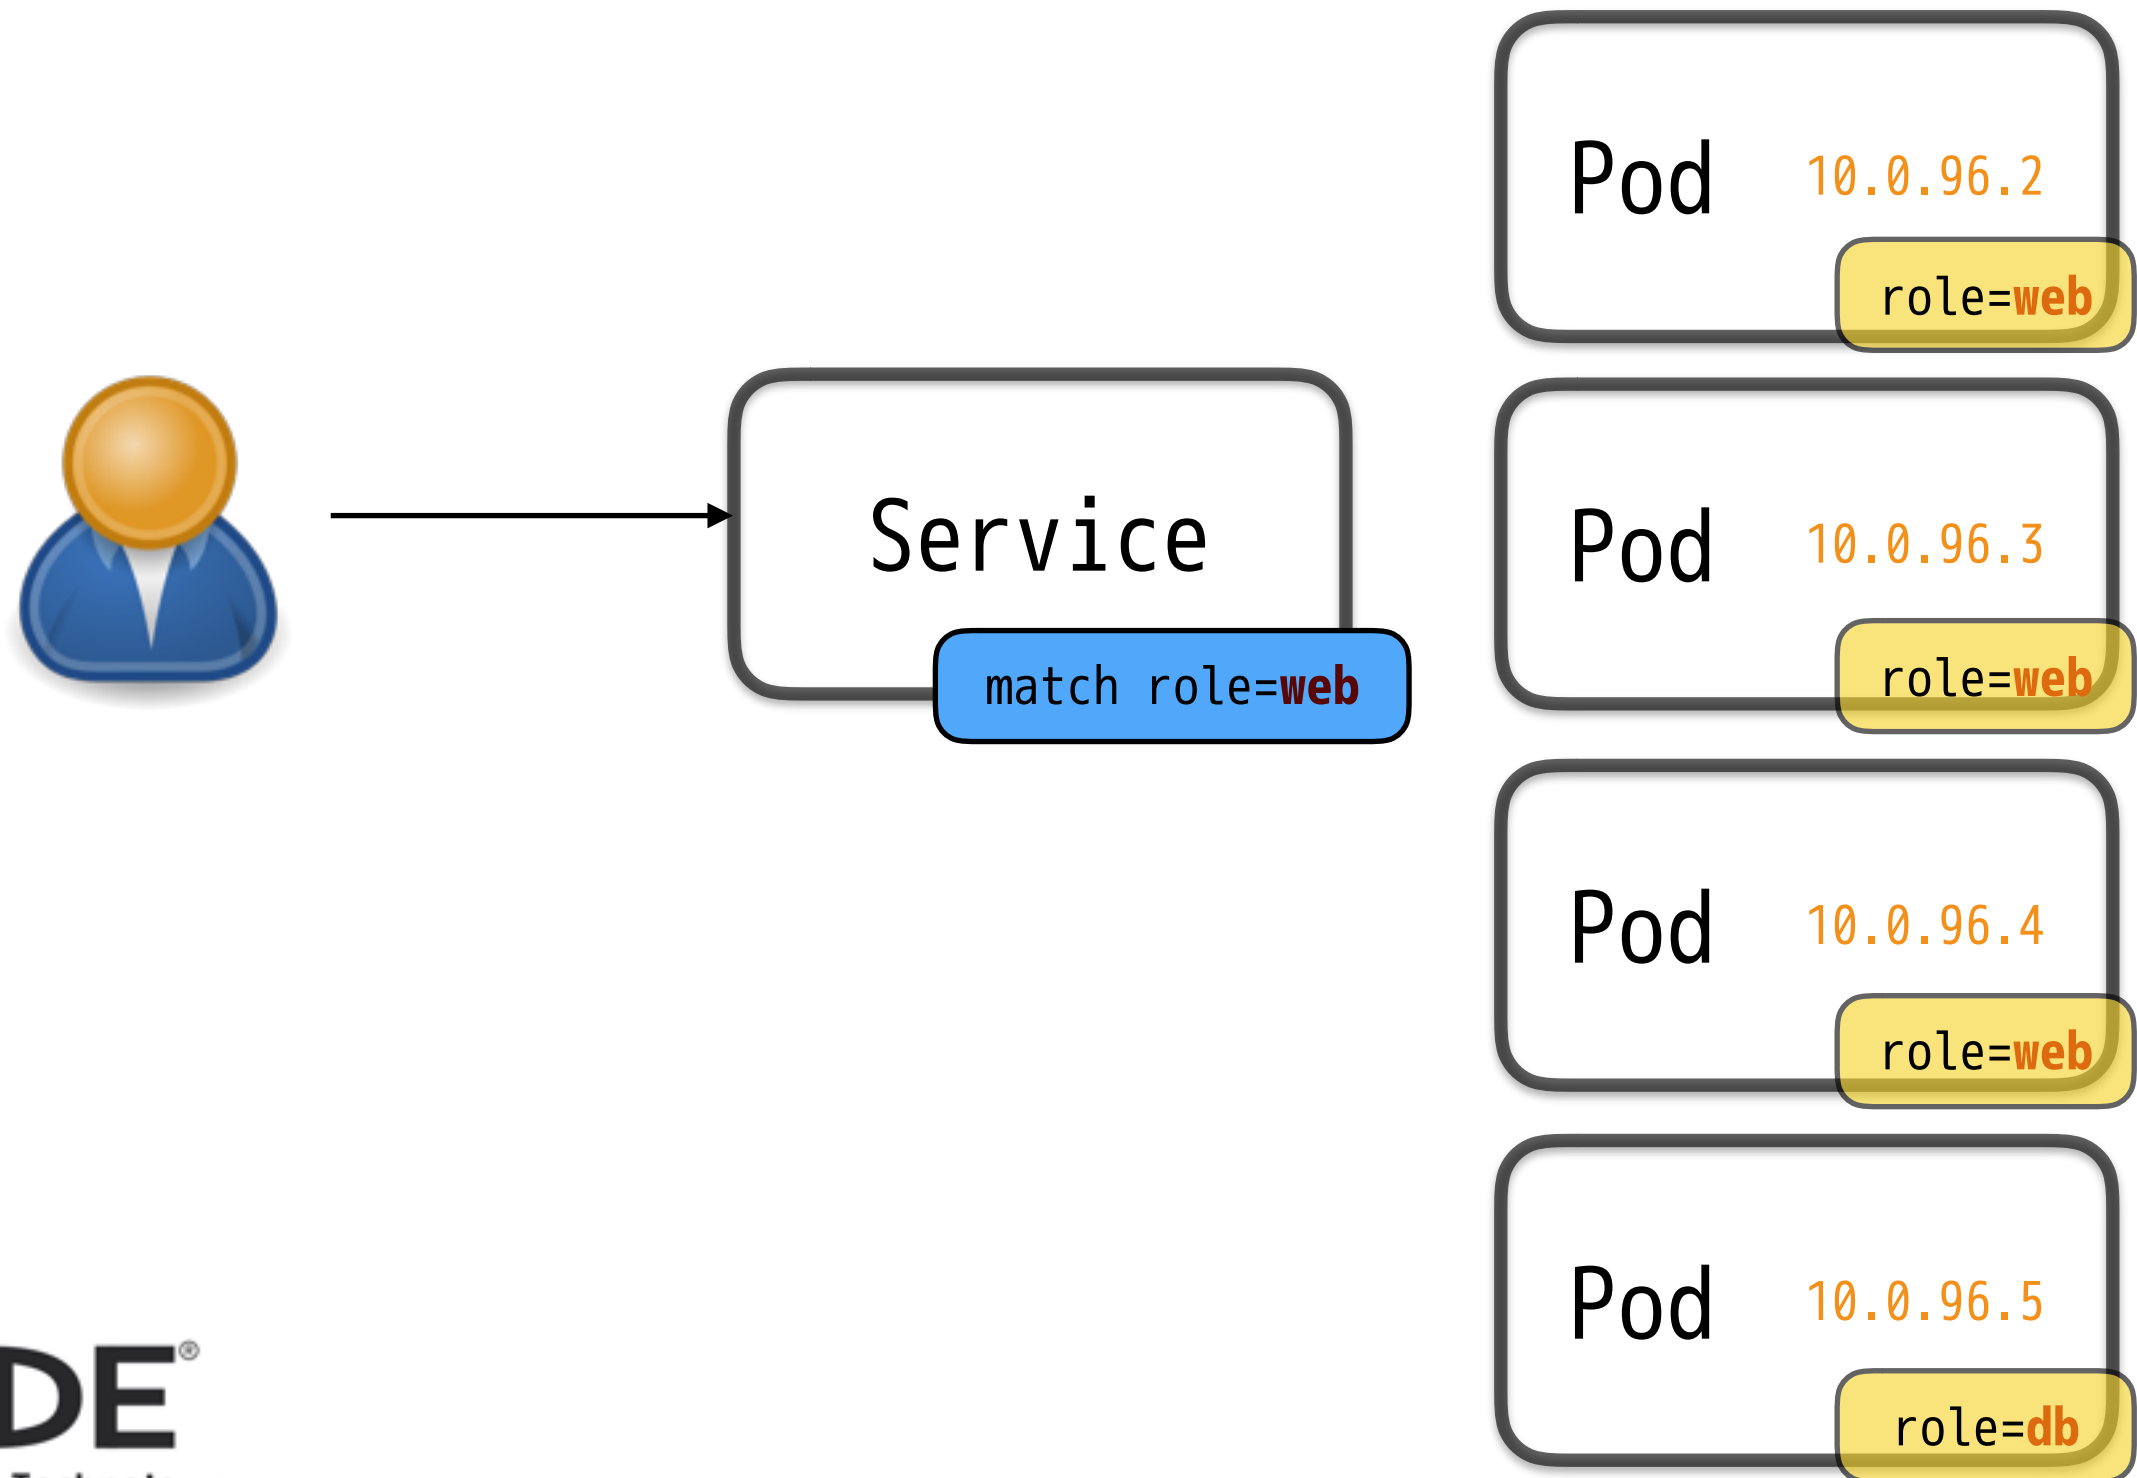
Raw Pod Access



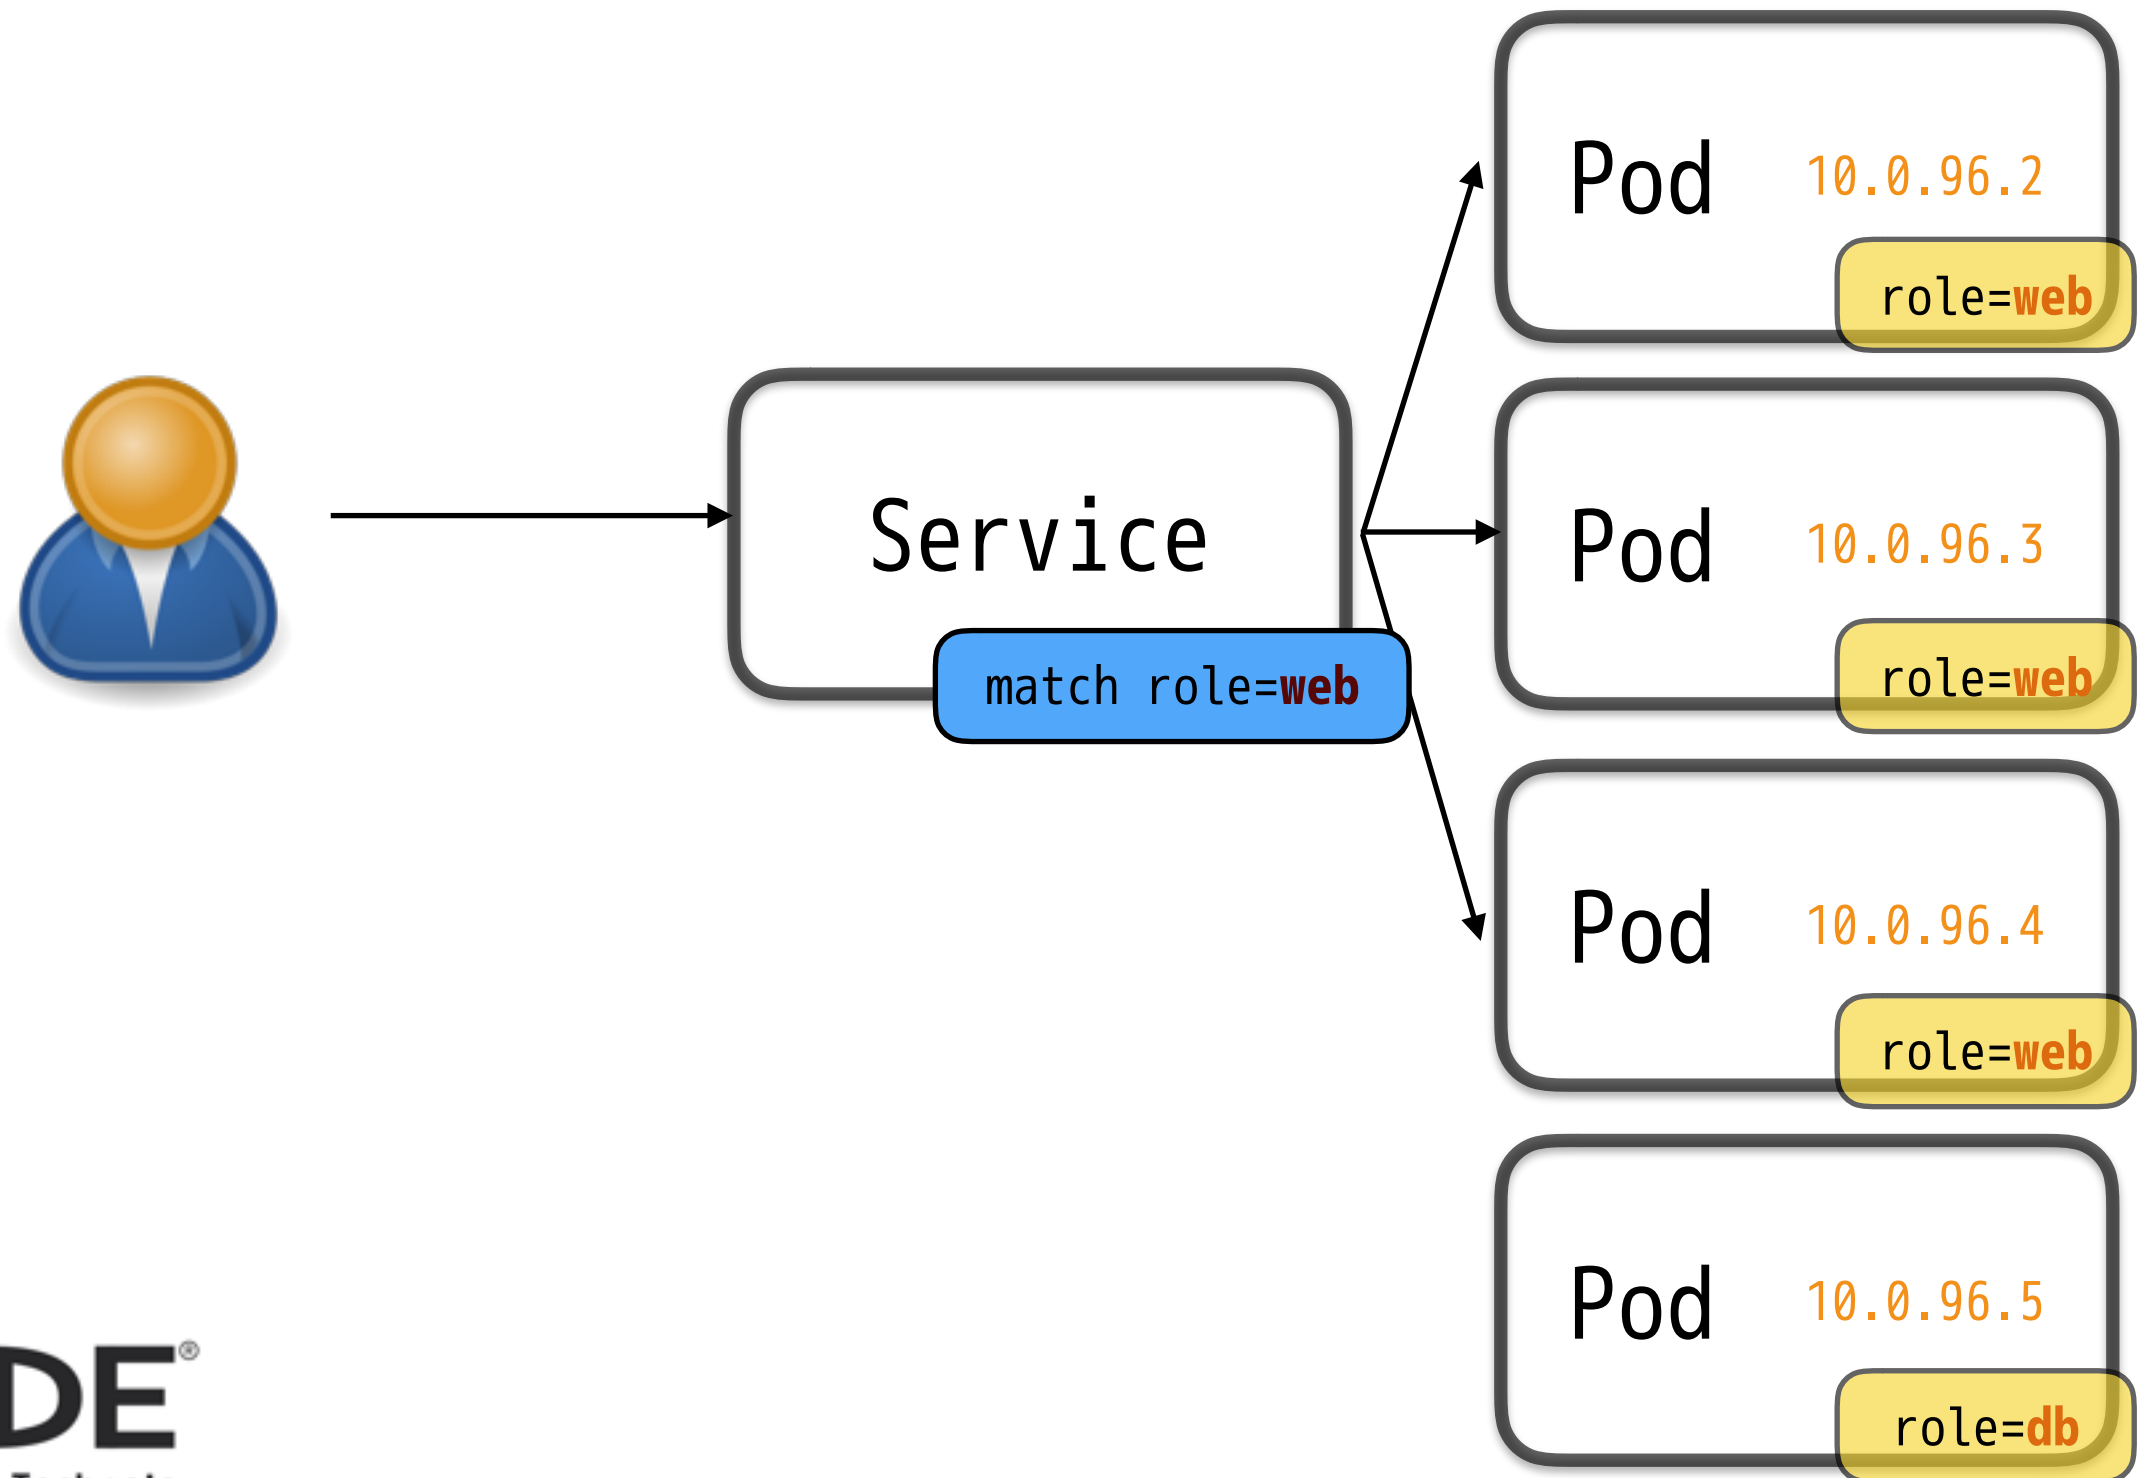
Access Via Service




Access Via Service




Access Via Service



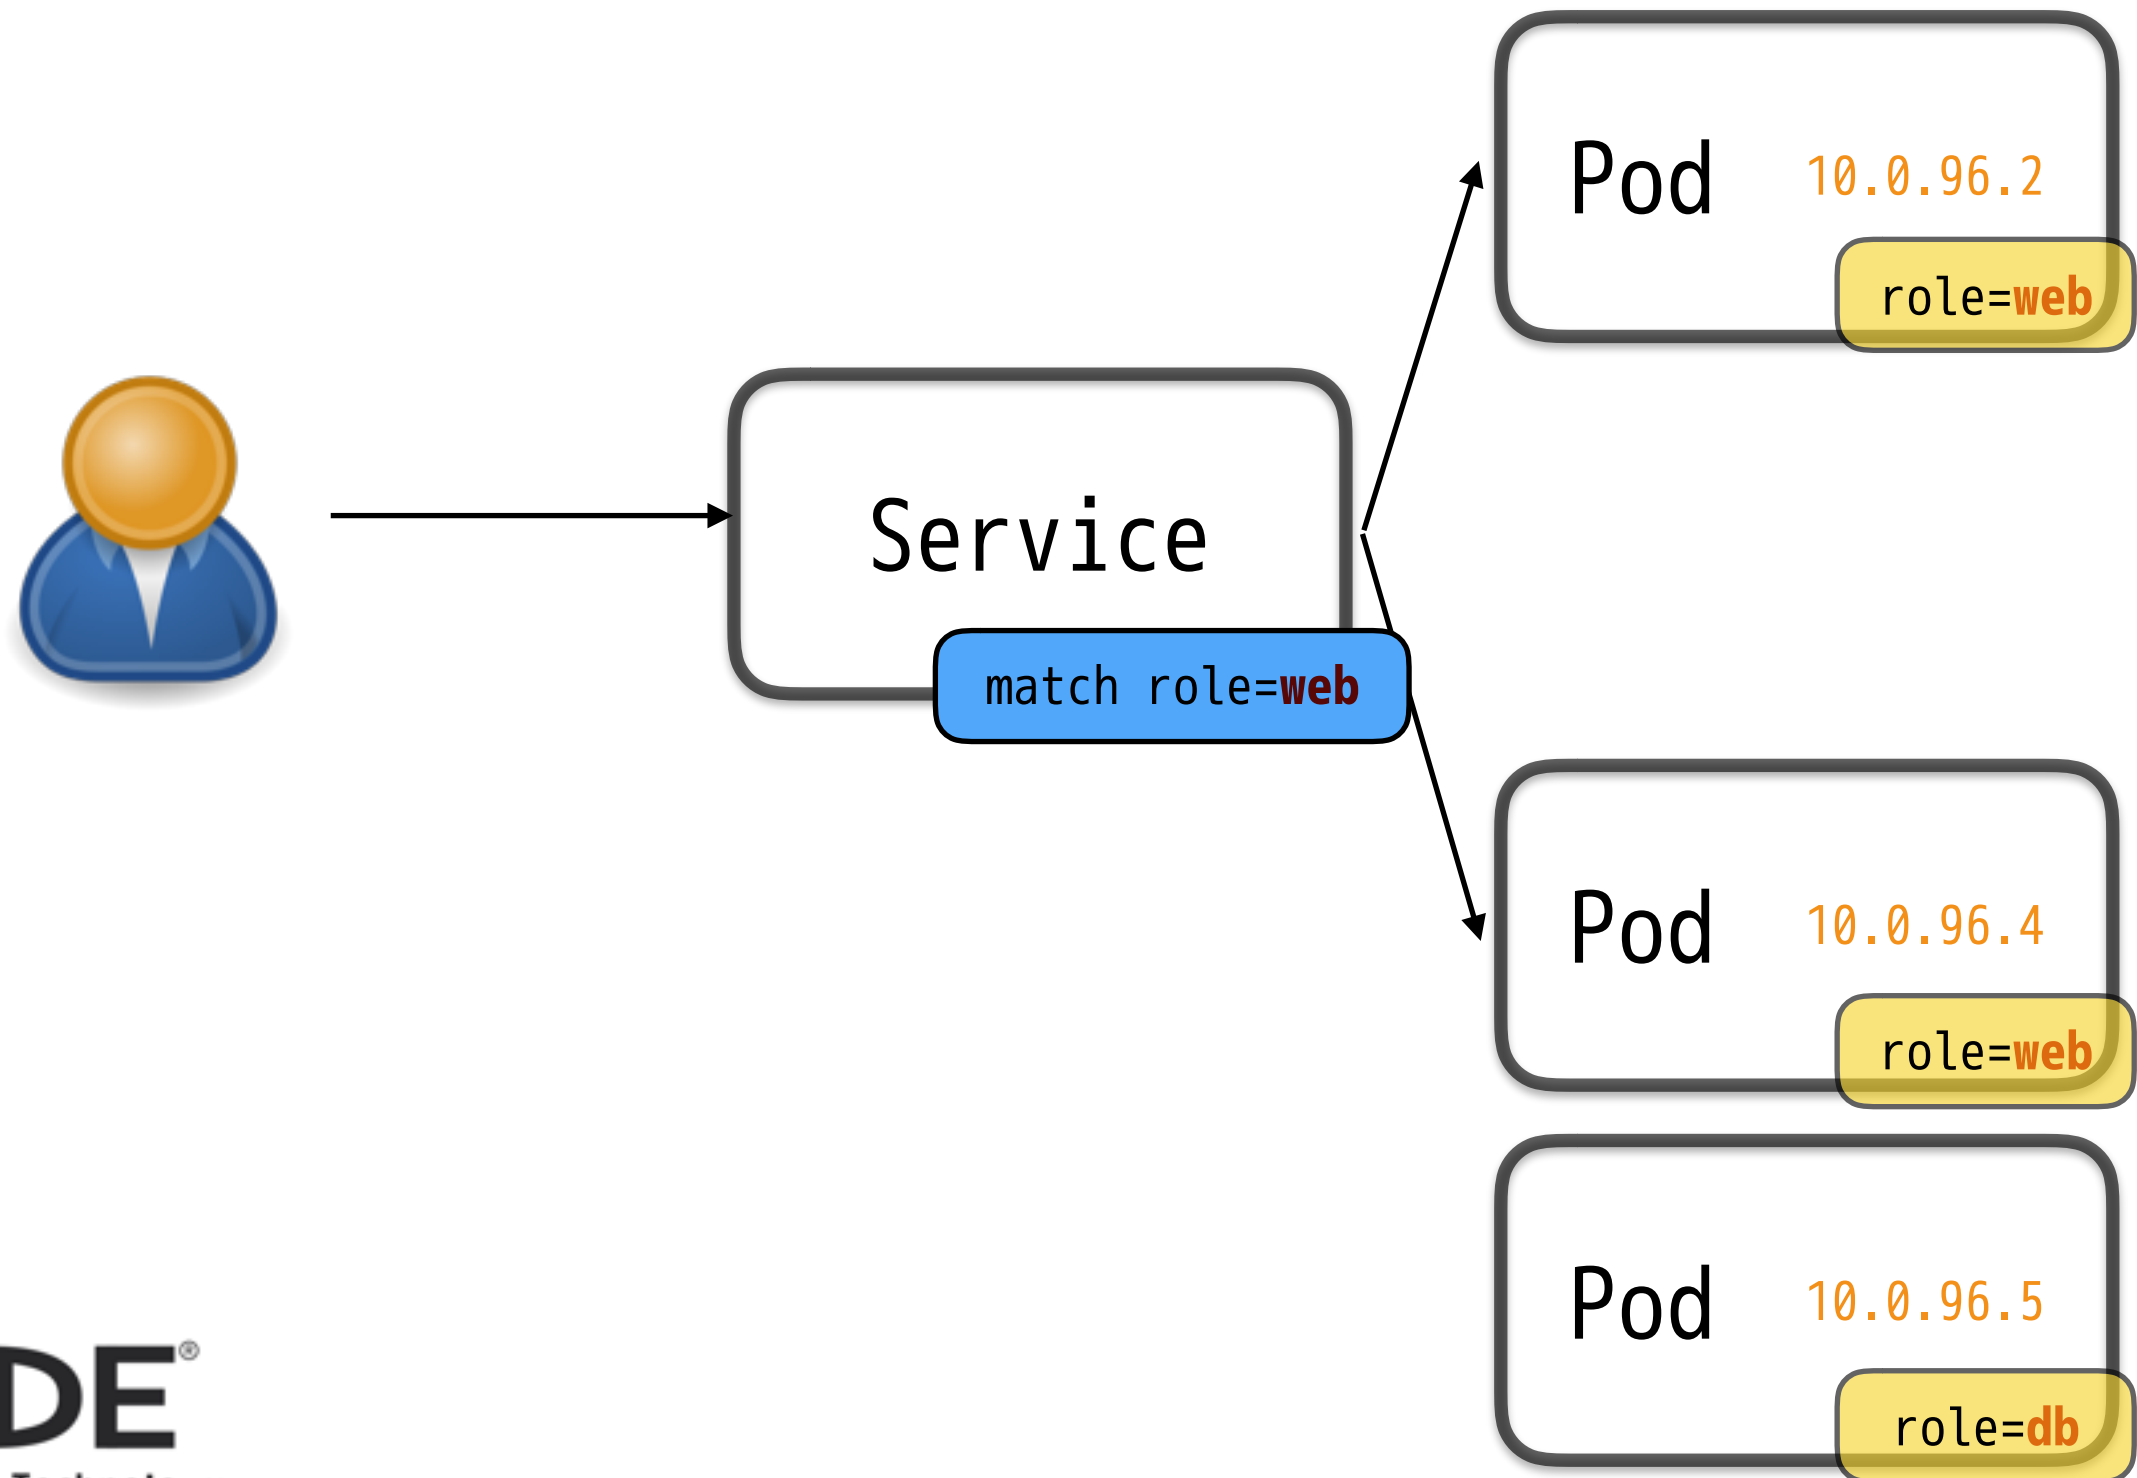
Access Via Service




Access Via Service



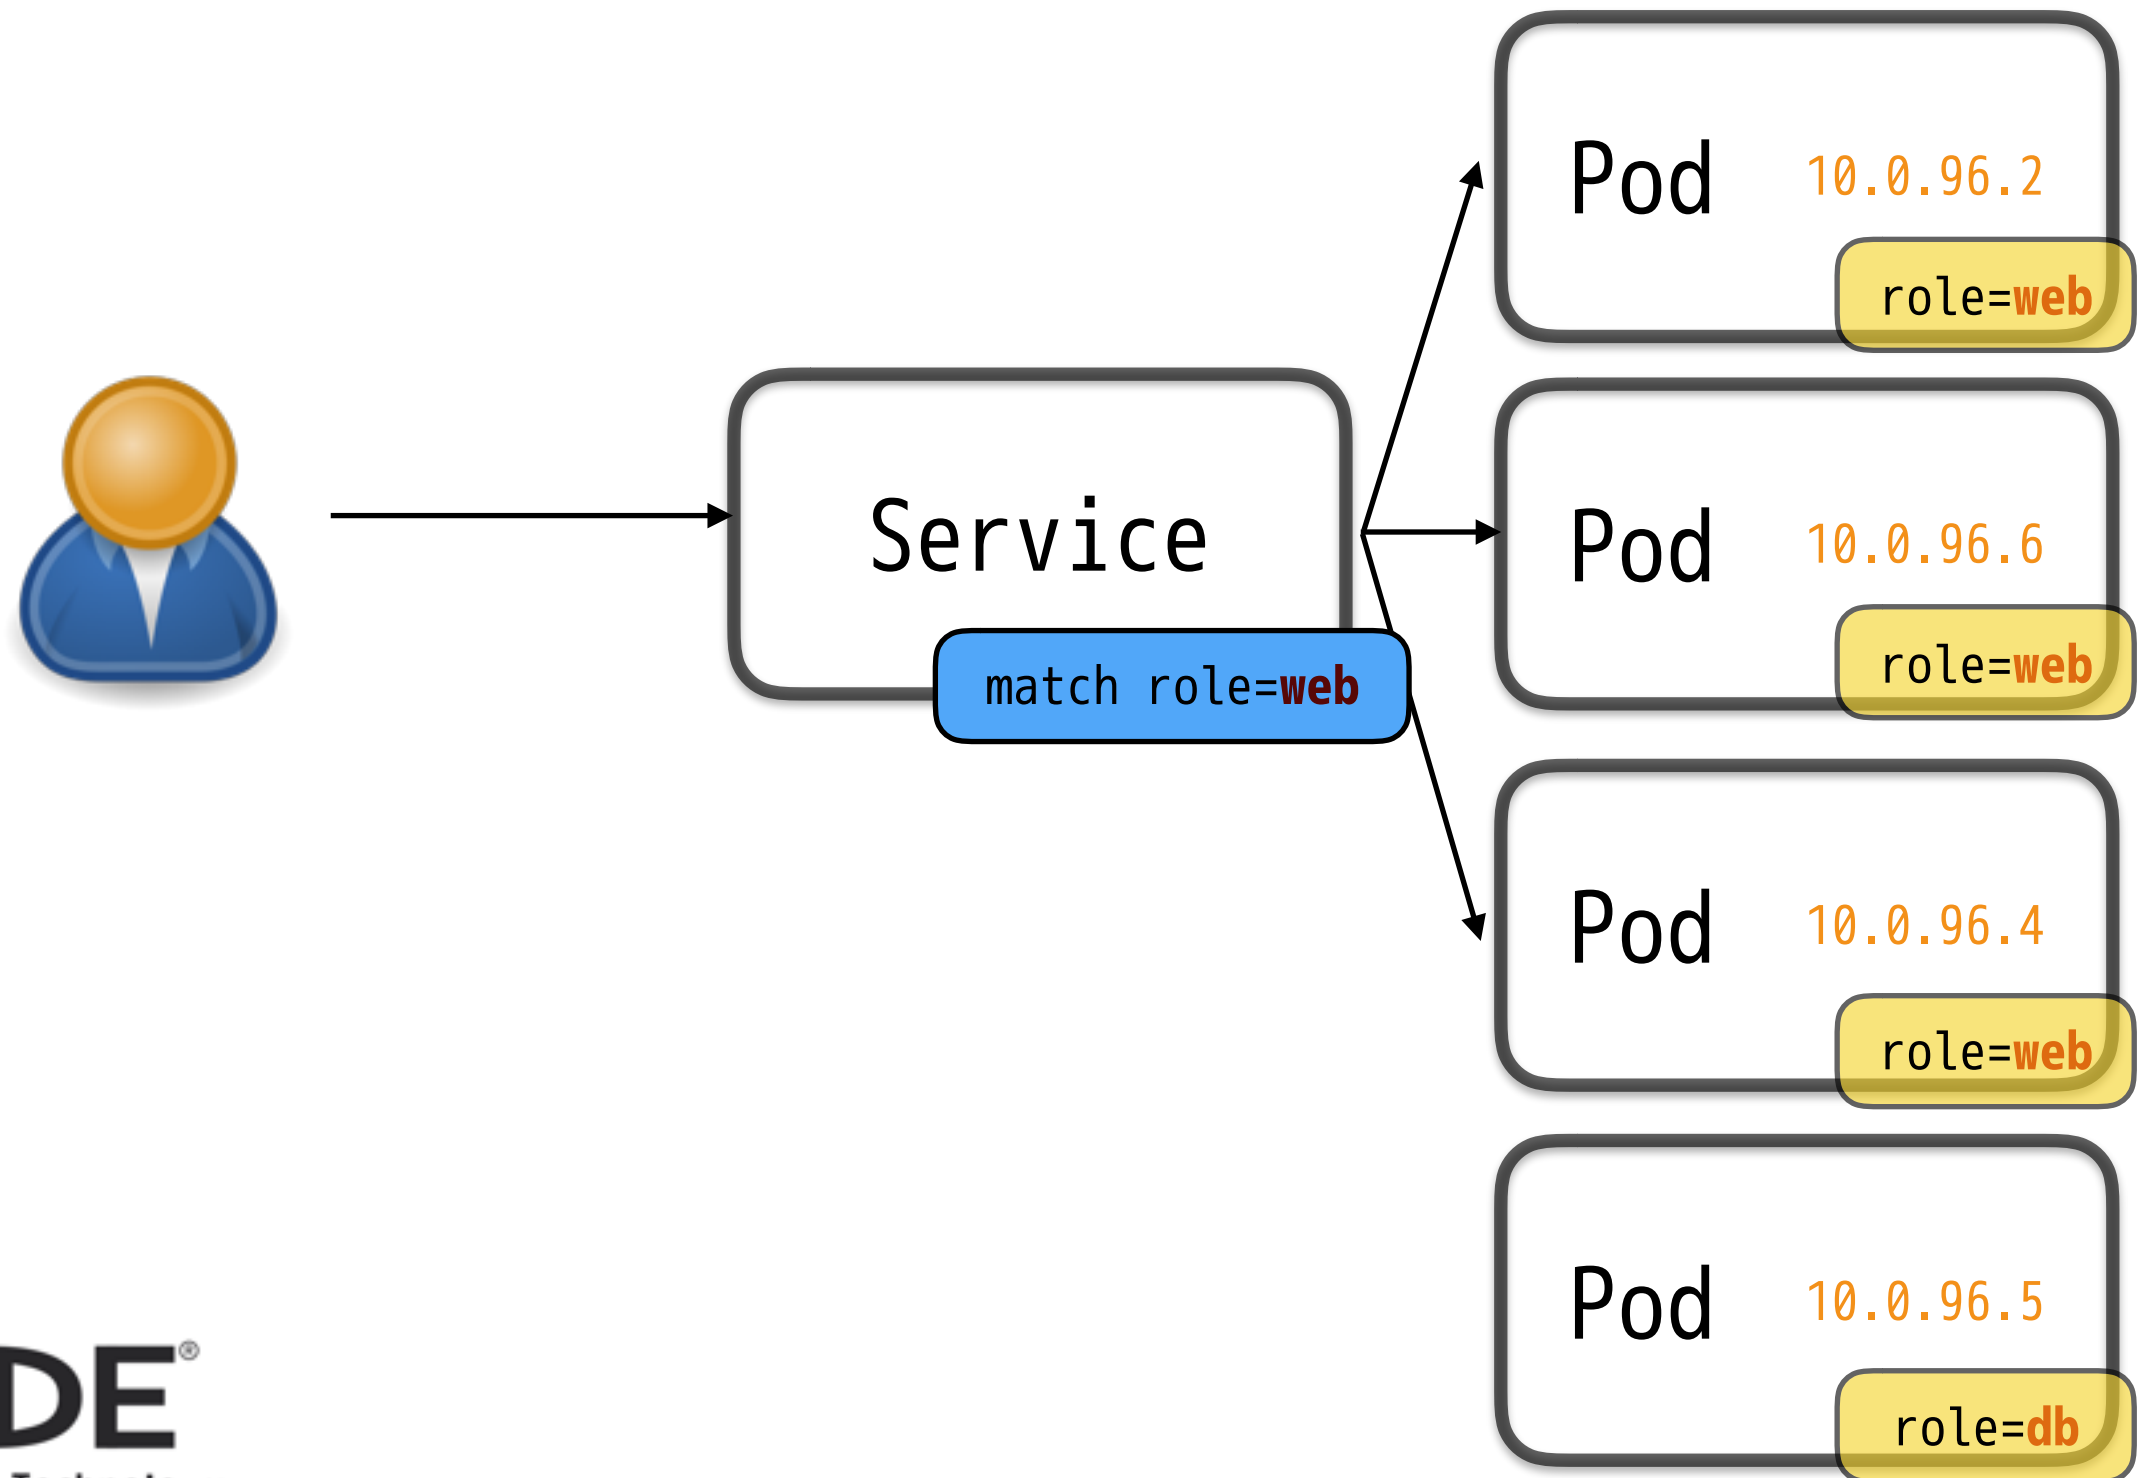
Access Via Service



Access Via Service



Access Via Service



Secrets

Store pieces of data in k8s

e.g. Identity Information

(securely)

(...in the future)


Secret

key1 → base64 value

key2 → base64 value

key3 → base64 value


Pod



```
container:  
  volumeMounts:  
    - name: certificates  
      mountPath: /etc/ssl/certs
```

```
volumes:  
  - name: certificates  
    secret:  
      secretName: ca-certificates
```

Pod



env:

- name: **foo-secret**
valueFrom:
secretKeyRef:
name: foo
value: secret-value

ConfigMaps

**Same as Secrets
(Unprotected)**

Ingress

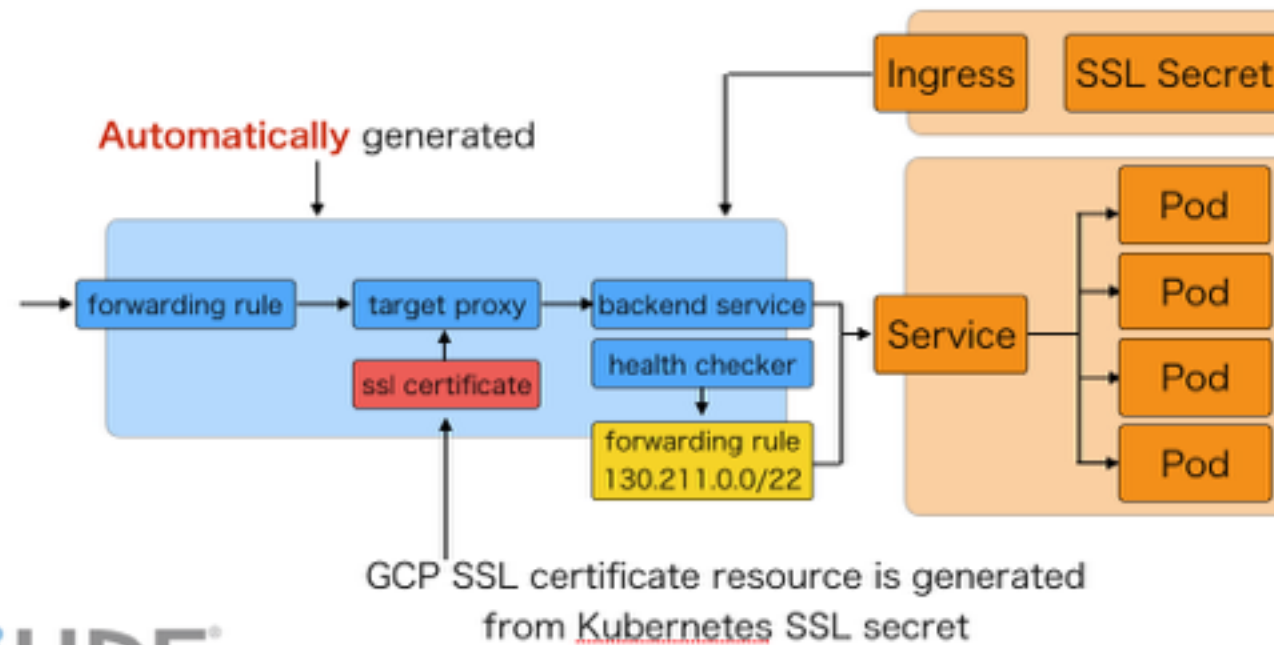
Services are for within the cluster only

**(external IPs allowed, but don' use that to serve
external requests)**

Inbound connections to internal cluster services

(New since 1.2)

GCP HTTPS LB + GKE (via Ingress)

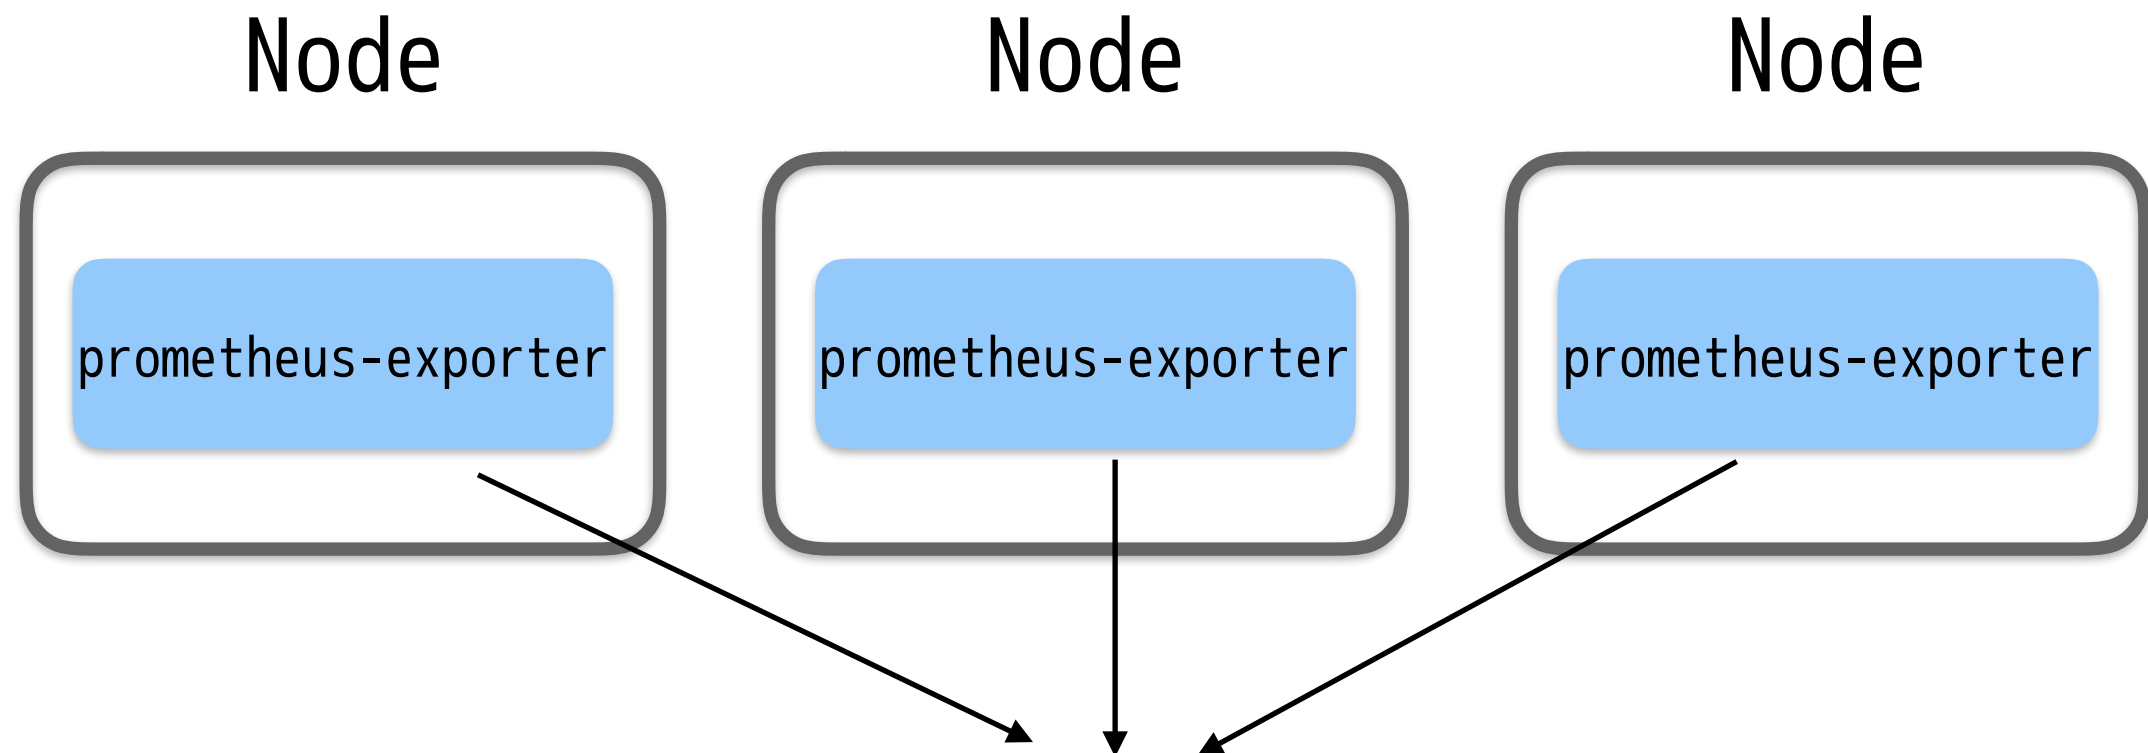


Please checkout my presentation

https://docs.google.com/presentation/d/11ZN6qgiuZZfVyhBK2hjp1vhp_5N0D0GnAmhPreS3L5A/pub?start=false&loop=false&delayms=3000

DaemonSets

Ensure nodes run a copy of a Pod



PetSets

~~ReplSets~~

StatefulSets

Use only when **really needed**

Keeps unique IDs in replicas
 $(x-0, x-1, x-2, \dots)$

Questions?