



APACHE  
AIRAVATA

# Introduction to First Half Course Topics

Marlon Pierce, Suresh Marru

Science Gateways Research Center, Indiana University

[marpierc@iu.edu](mailto:marpierc@iu.edu), [smarru@iu.edu](mailto:smarru@iu.edu)

# Monolithic Applications: Traditional Software Releases

- Software releases occur in discrete increments
- Software runs on clients' systems
- Releases may be frequent but they are still distinct
  - Firefox
  - OS system upgrades
- Traditional release cycles
  - Extensive testing
  - Alpha, beta, release candidates, and full releases
- Extensive recompiling and testing required after code changes
- Code changes require the entire release cycle to be repeated



# MicroServices: Software as a Service

- Does your software run as a service?
- Do you run this service yourself? Or does another part of your organization run the service that your team develops?
- Traditional release cycles don't work well
  - May make releases many times per day
  - Test-release-deploy takes too long
- You can be a little more tolerant of bugs discovered after release if you can fix quickly or roll back quickly.
- Get new features and improvements into production quickly.

Science gateways provide software as a service

The Netflix logo is displayed in a bold, red, sans-serif font. It is centered within a light gray rectangular box that has a subtle gradient.

# Science Gateways and Microservices

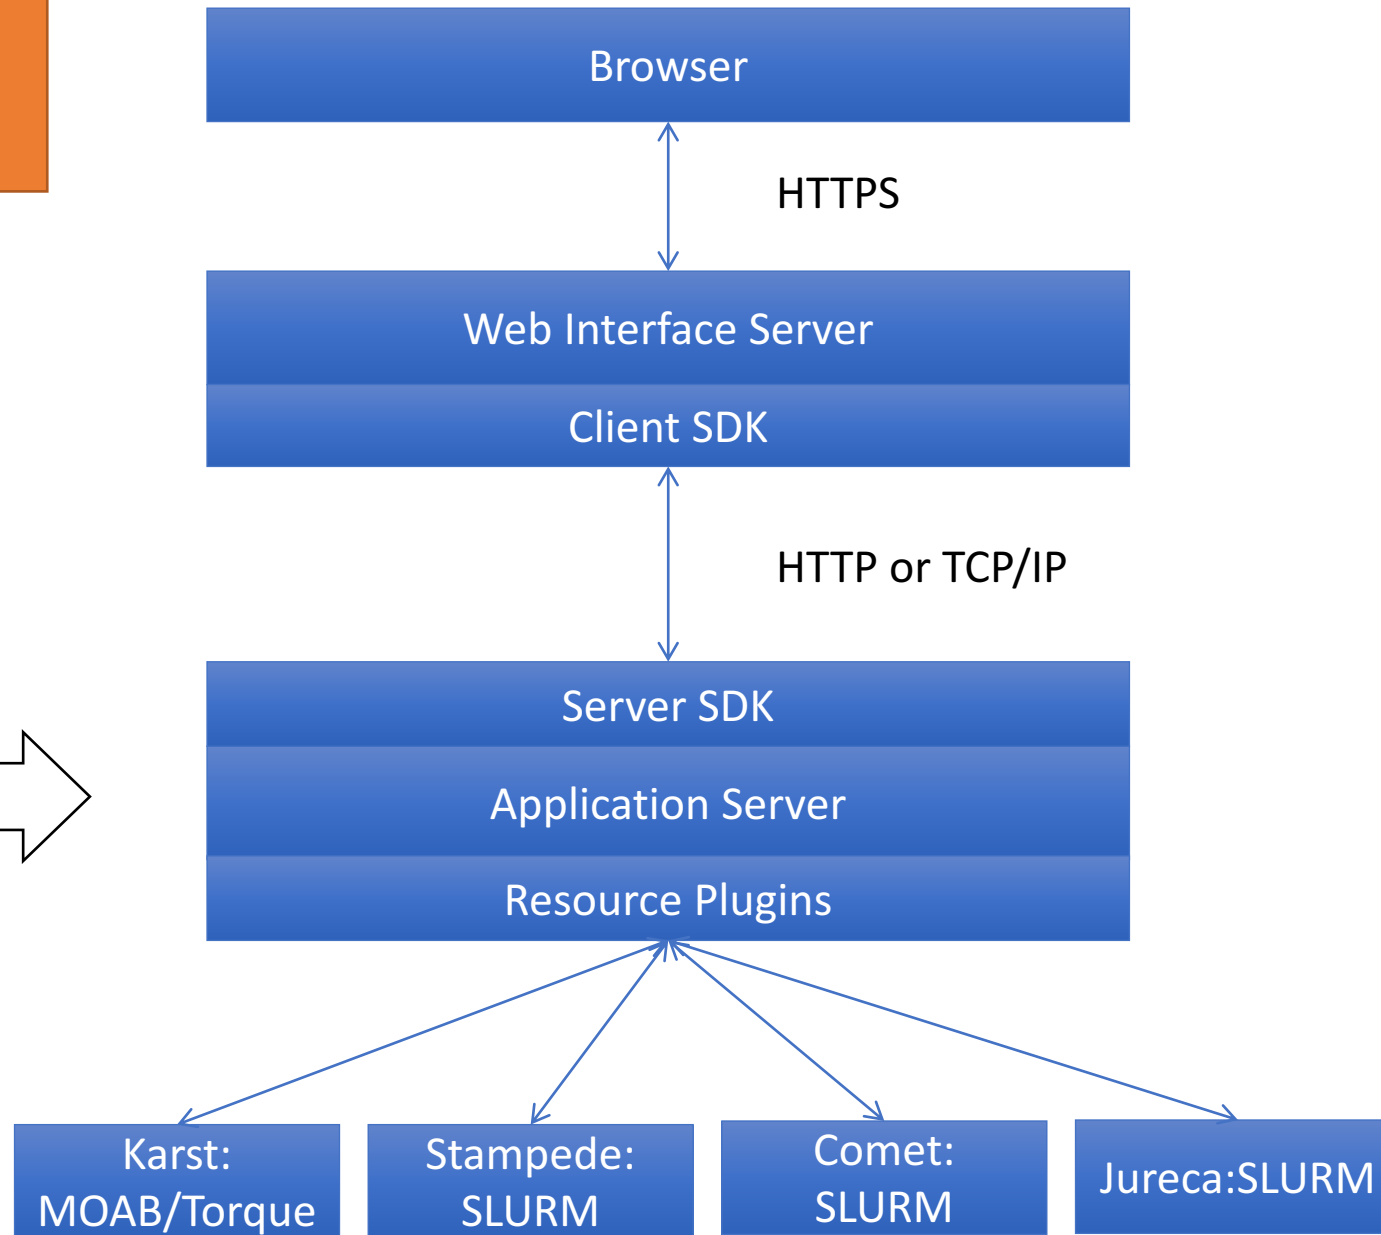


# What Is a Microservice?


- Develop a single application as a suite of small services
- Each service runs in its own process
- Services communicate with lightweight mechanisms
  - “Often an HTTP resource API”
  - But that has some problems
  - Messaging and hybrid approaches
- These services are built around business capabilities
- Independently deployable by fully automated deployment machinery.
- Minimum of centralized management of these services,
  - May be written in different programming languages
  - May use different data storage technologies.

## Recall the Gateway Octopus Diagram



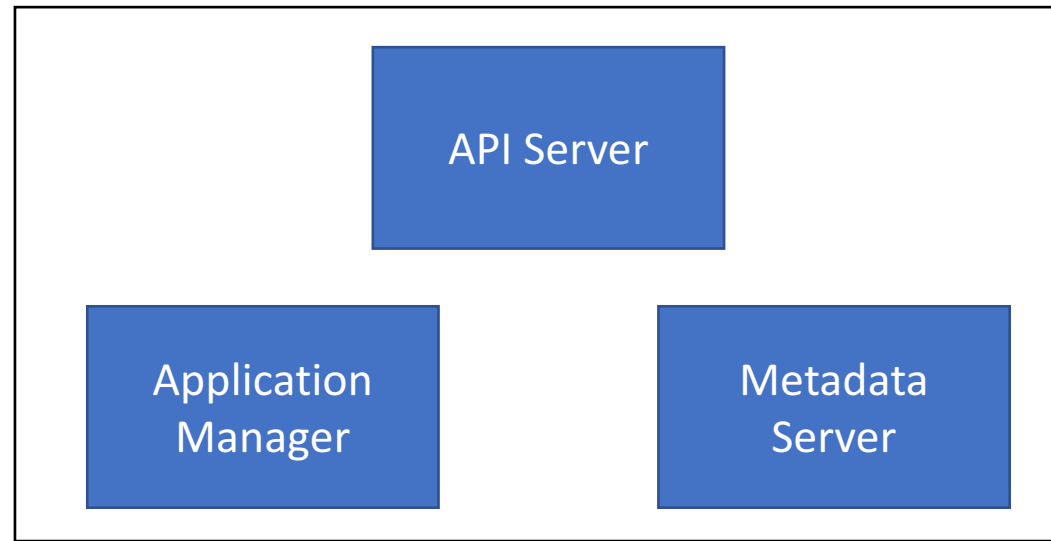
We will focus  
on this piece



# Basic Components of the Gateway App Server



# Decoupling the App Server






# How Do We Package and Where Do We Run All Those Microservices?

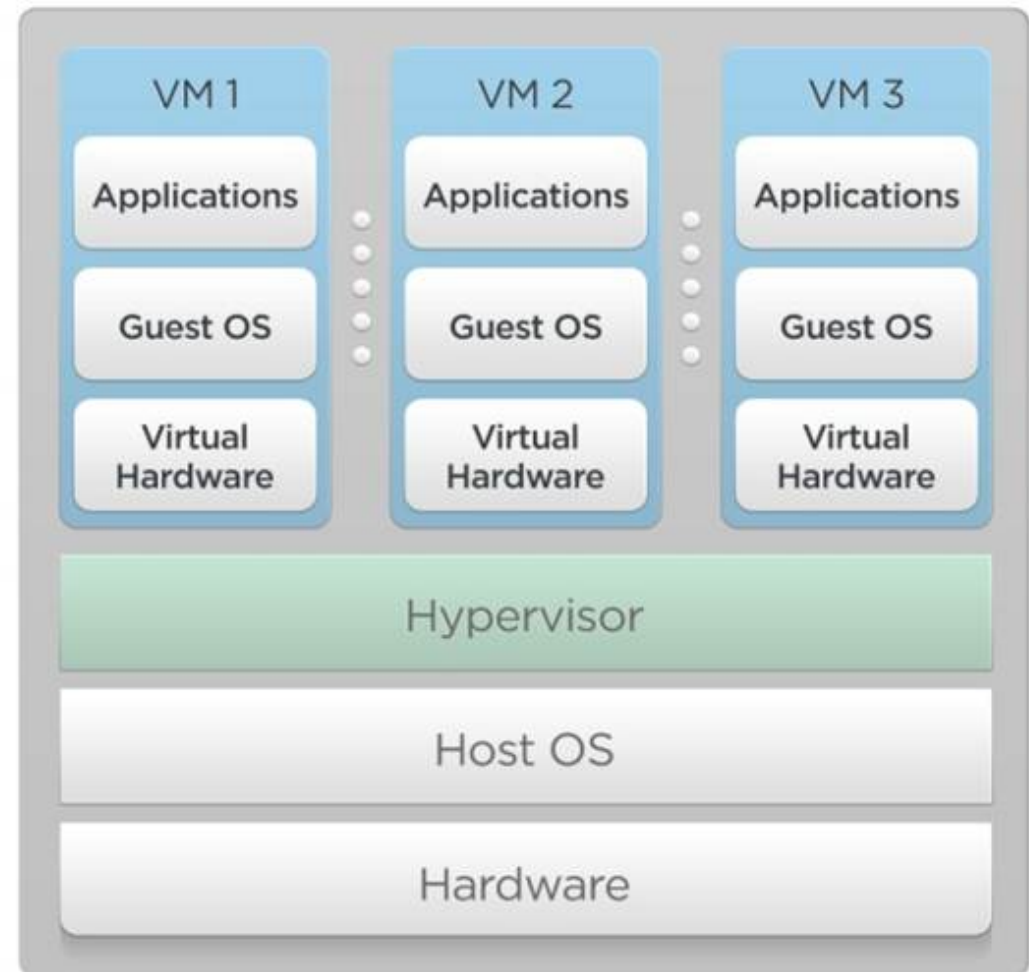
On the Cloud? In the Matrix?

# Virtualization, Containers, Docker



# Hypervisor Virtualization

- Hypervisors provide software emulated hardware and support multiple OS tenants.
- Type 1 Hypervisors run directly on the hardware
  - Kernel-based Virtual Machine (KVM)
  - Xen: Amazon uses (or used) this
- Type 2 Hypervisors run as another program in the HostOS
  - Virtual Box
  - VMWare



# Virtualization Benefits and Drawbacks


- Makes your development, testing, and deployment environments the same
- Supports Continuous Integration and Deployment

- Adds overhead
  - Relatively slow to create and destroy VMs
  - Virtualized hardware and networking are slower than real systems
  - Uses memory, disk, and CPU resources

You can run 1 VM per microservice, but there is a better way.

# Docker and Other Linux Containers

- Docker replaces the need for DIY containers
- Docker uses LXC containerization + AuFS
  - AuFS lets Docker containers share common files
  - AuFS also allows you to have efficient version control of container images



Because each container has its own thin writable container layer, and all changes are stored in this container layer, multiple containers can share access to the same underlying image and yet have their own data state.

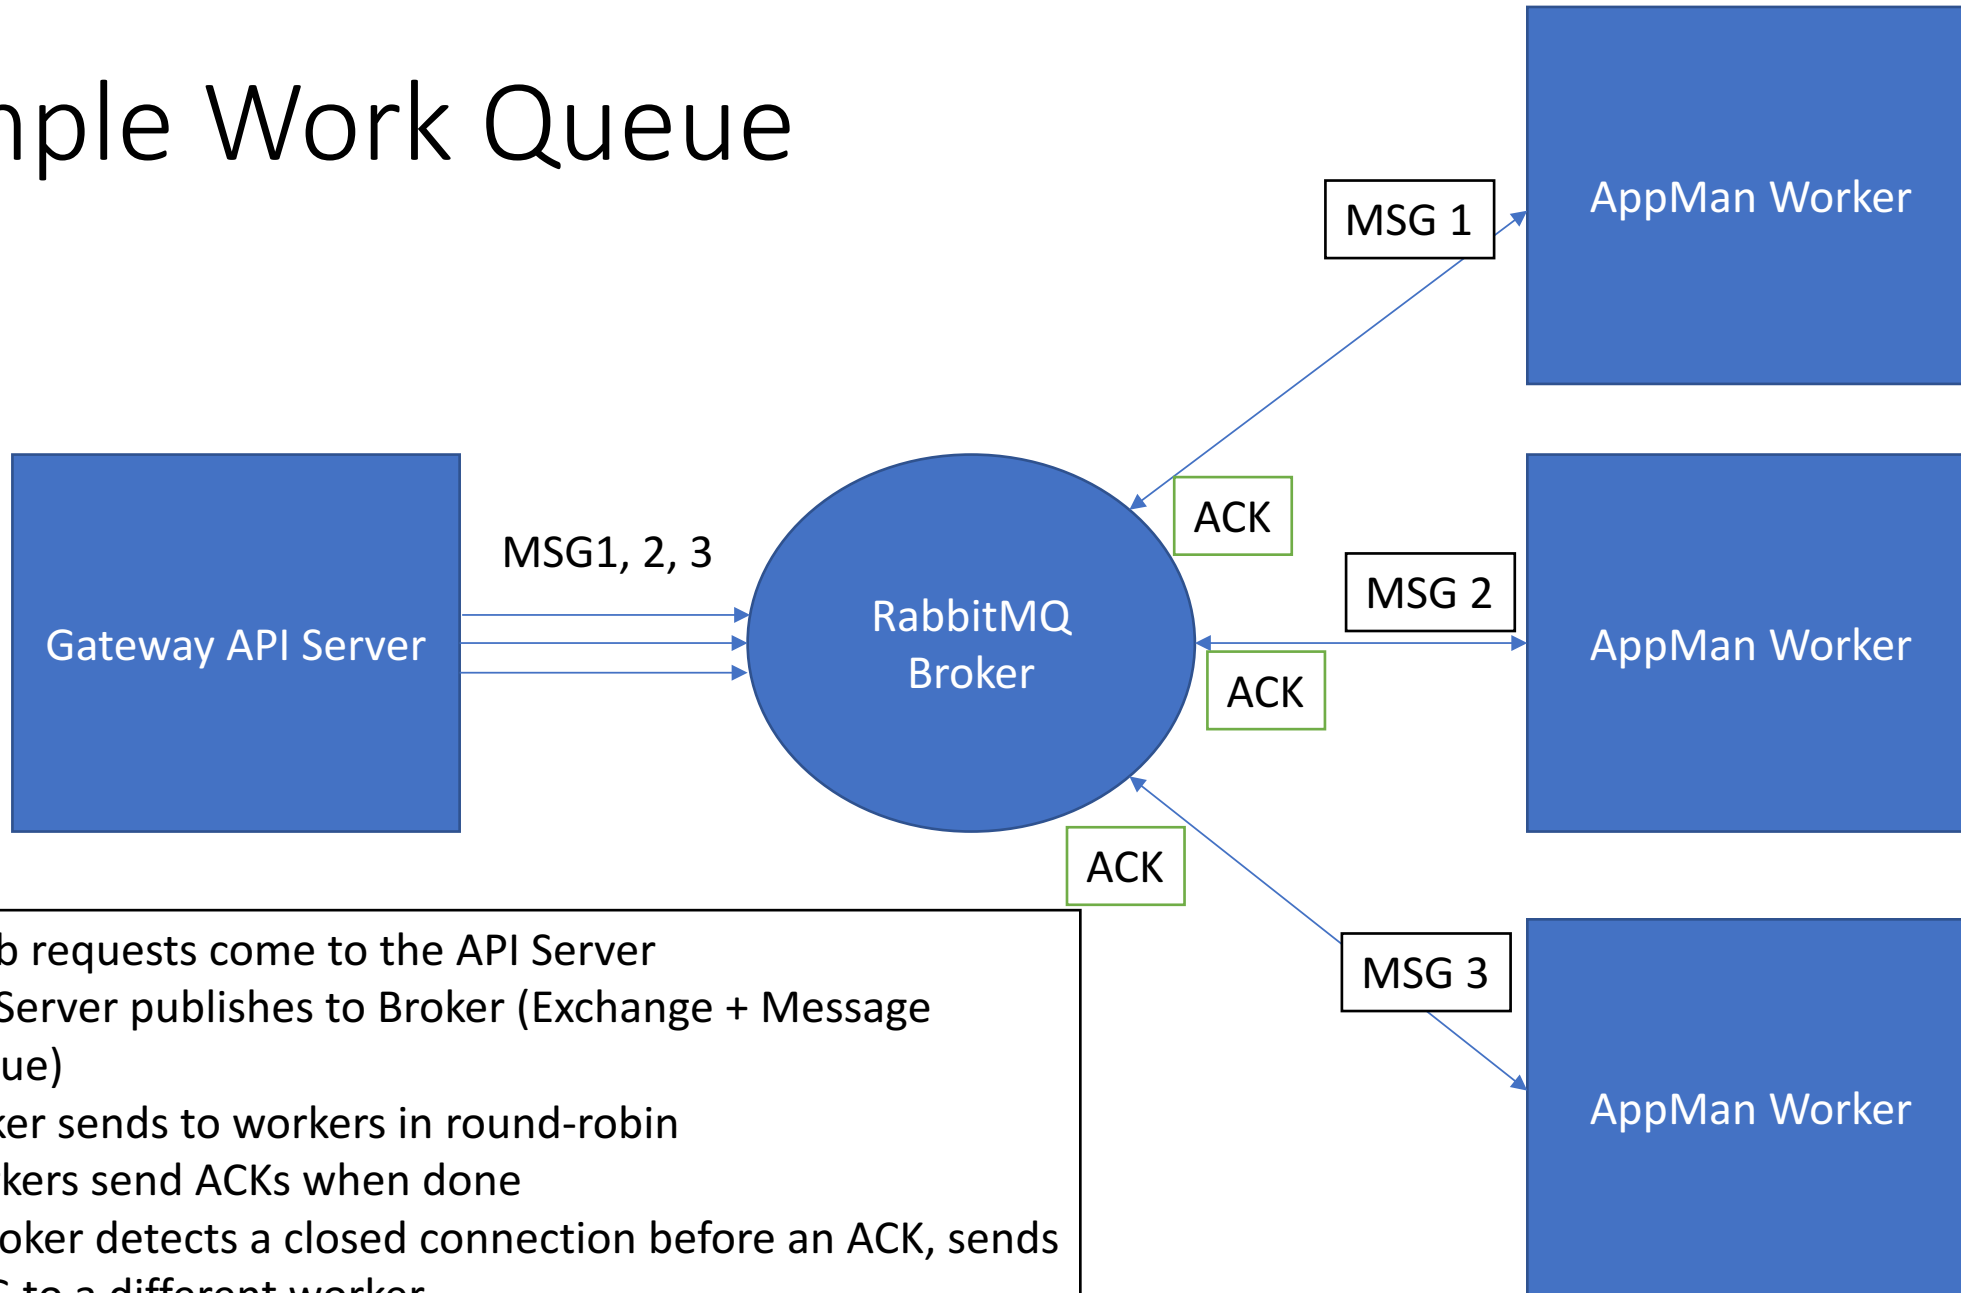
# How Do Microservices Communicate?

Push, pull, etc

# Messaging Systems: RabbitMQ, Apache Kafka



# Simple Work Queue





# How Can Components Expose their APIs and Data Models to Other Components?

And can we make this programming language independent?

# API and Metadata Model Design

Apache Thrift™


RESTful API

GET PUT POST DELETE

# How Can I Discover, Monitor, and Manage Services?

Can we learn some lessons from distributed systems  
research?

# Distributed State Management: Zookeeper, Consul



# How Do I Manage Logs from Microservices

And detect if there are problems



# How Can I Secure my API Server and Microservices?

How do I manage user identities, authentication and authorization?

# Security: OAuth2 and OpenIDConnect

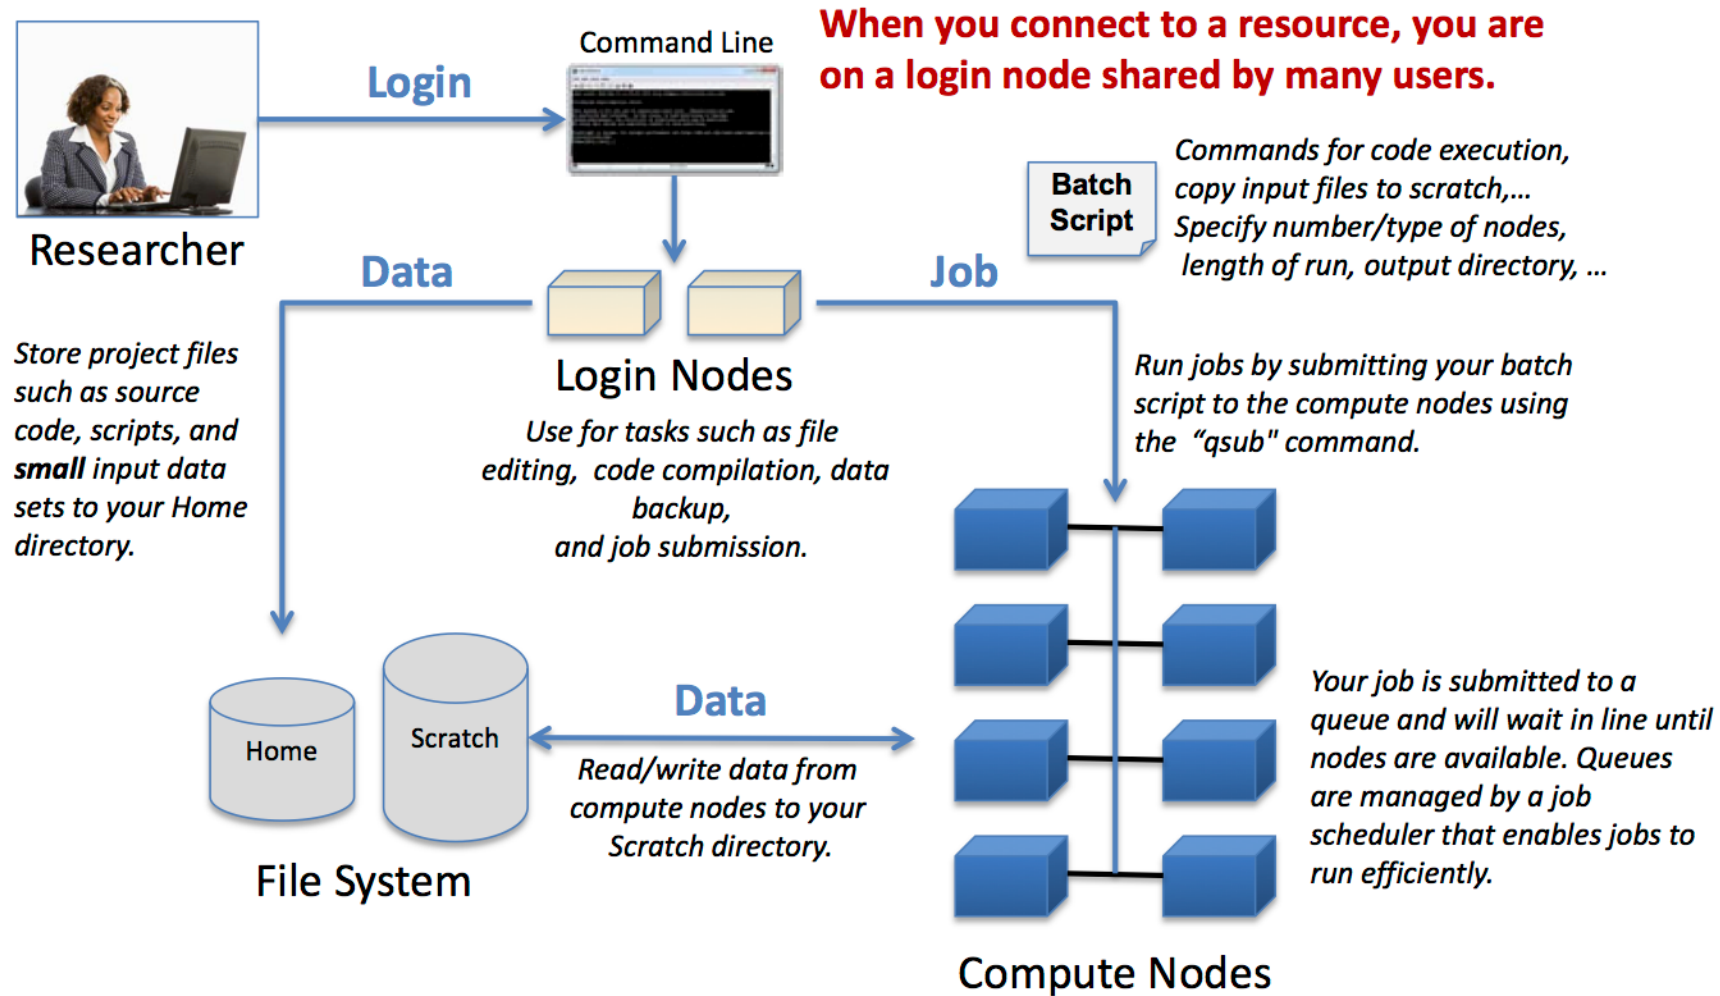




# How Does the Gateway Interact with Computing Clouds and Supercomputers?

How can we work reliably with unreliable resources?

# Running Jobs Overview

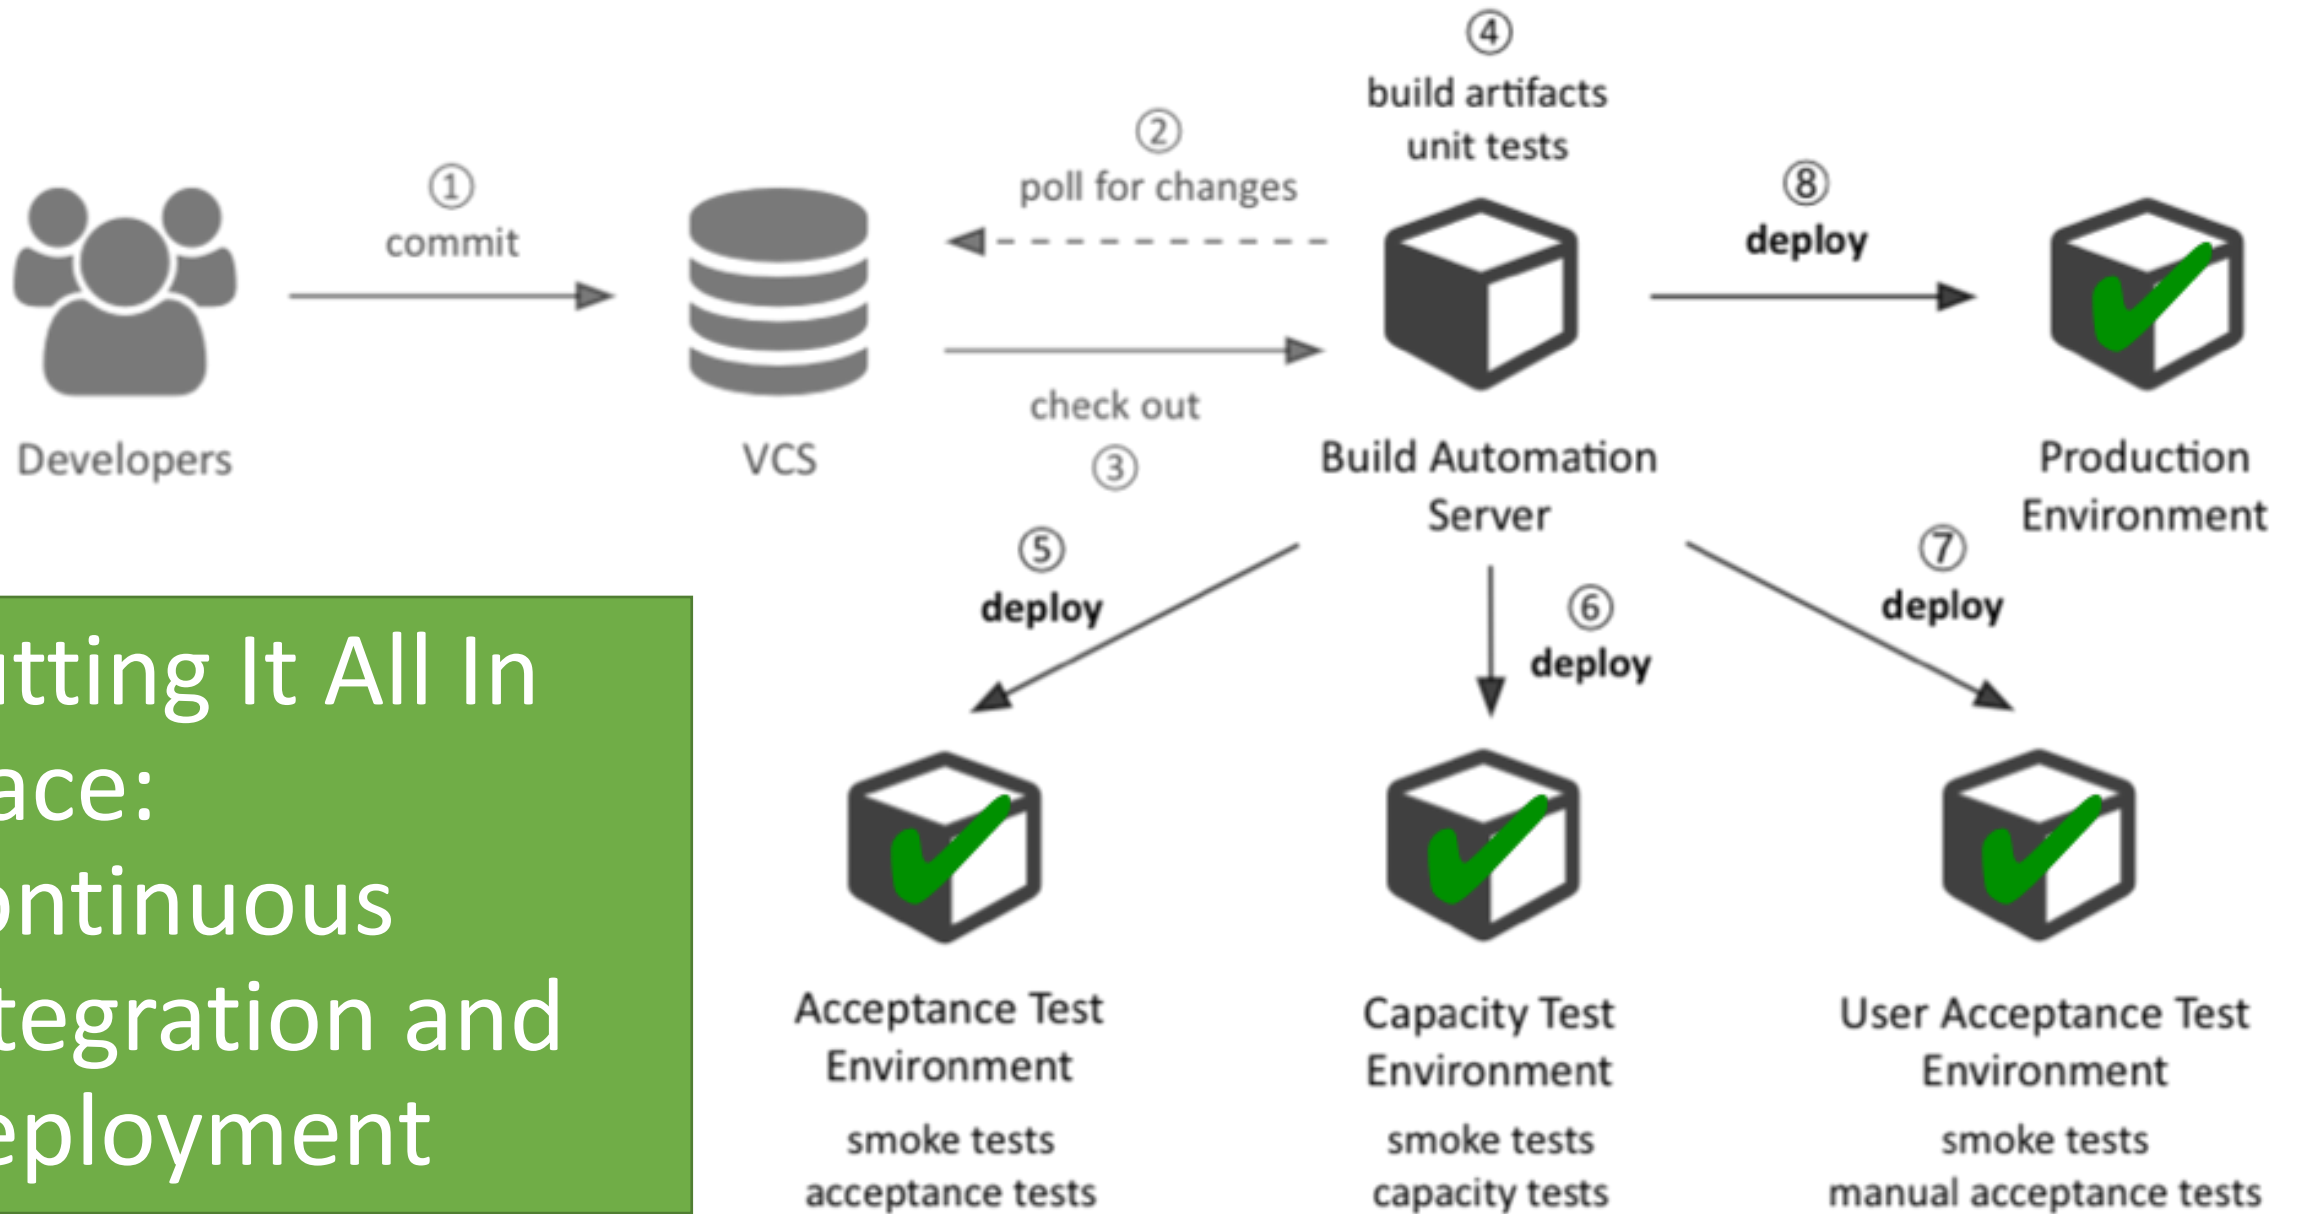


# Resource Management and Scheduling: Mesos, Aurora, and Torque



# How Can We Automate All of This?

How can we make our infrastructure reproducible?



Putting It All In  
Place:  
Continuous  
Integration and  
Deployment

# The Scientific Method as an Ongoing Process

Next  
Semester:  
Science!

