

# Introduction & Projects Overview

Marlon Pierce, Suresh Marru  
CSCI-B 649 Science Gateway Architectures  
Spring 2017 – Week 1 - January 13<sup>th</sup> 2017



# Todays Outline

- Apache Airavata Overview
- Microservice Architectures
- DevOps/NoOps Principles
- Overview of Project Themes
- Project Themes 1 & 2

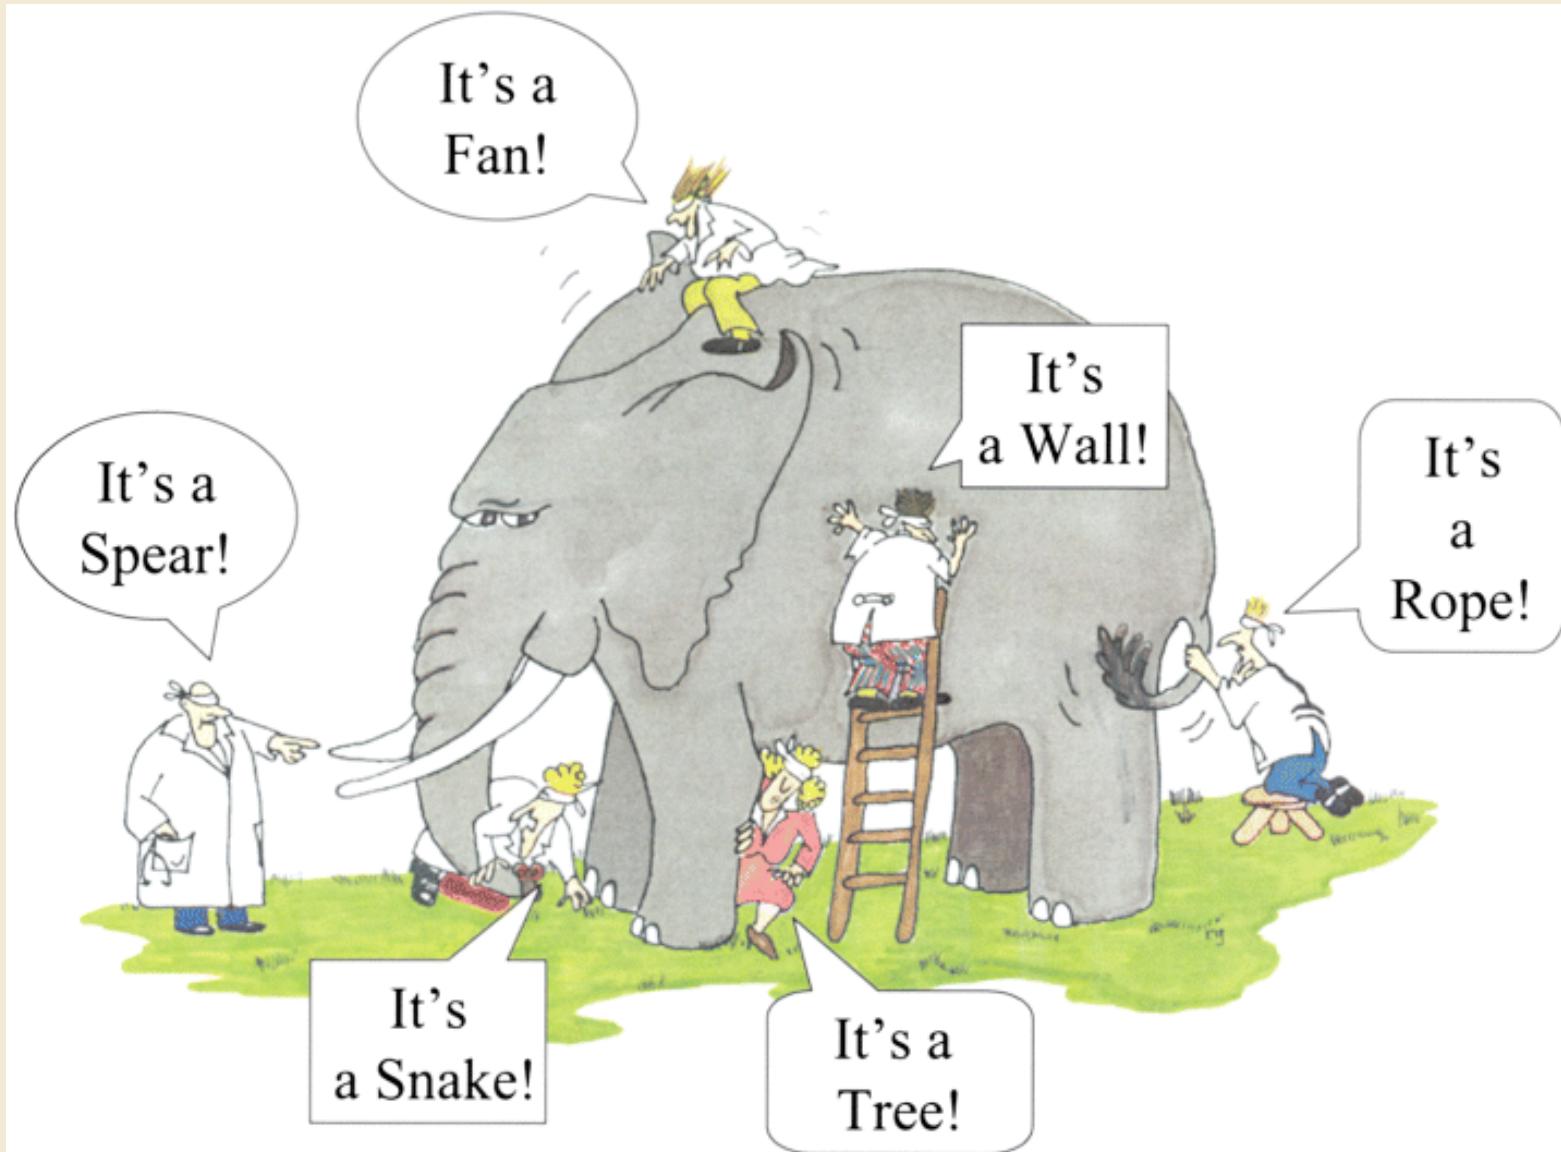


# What is a science gateway?

**science gateway** /sī' əns gāt' wā'/ *n.*

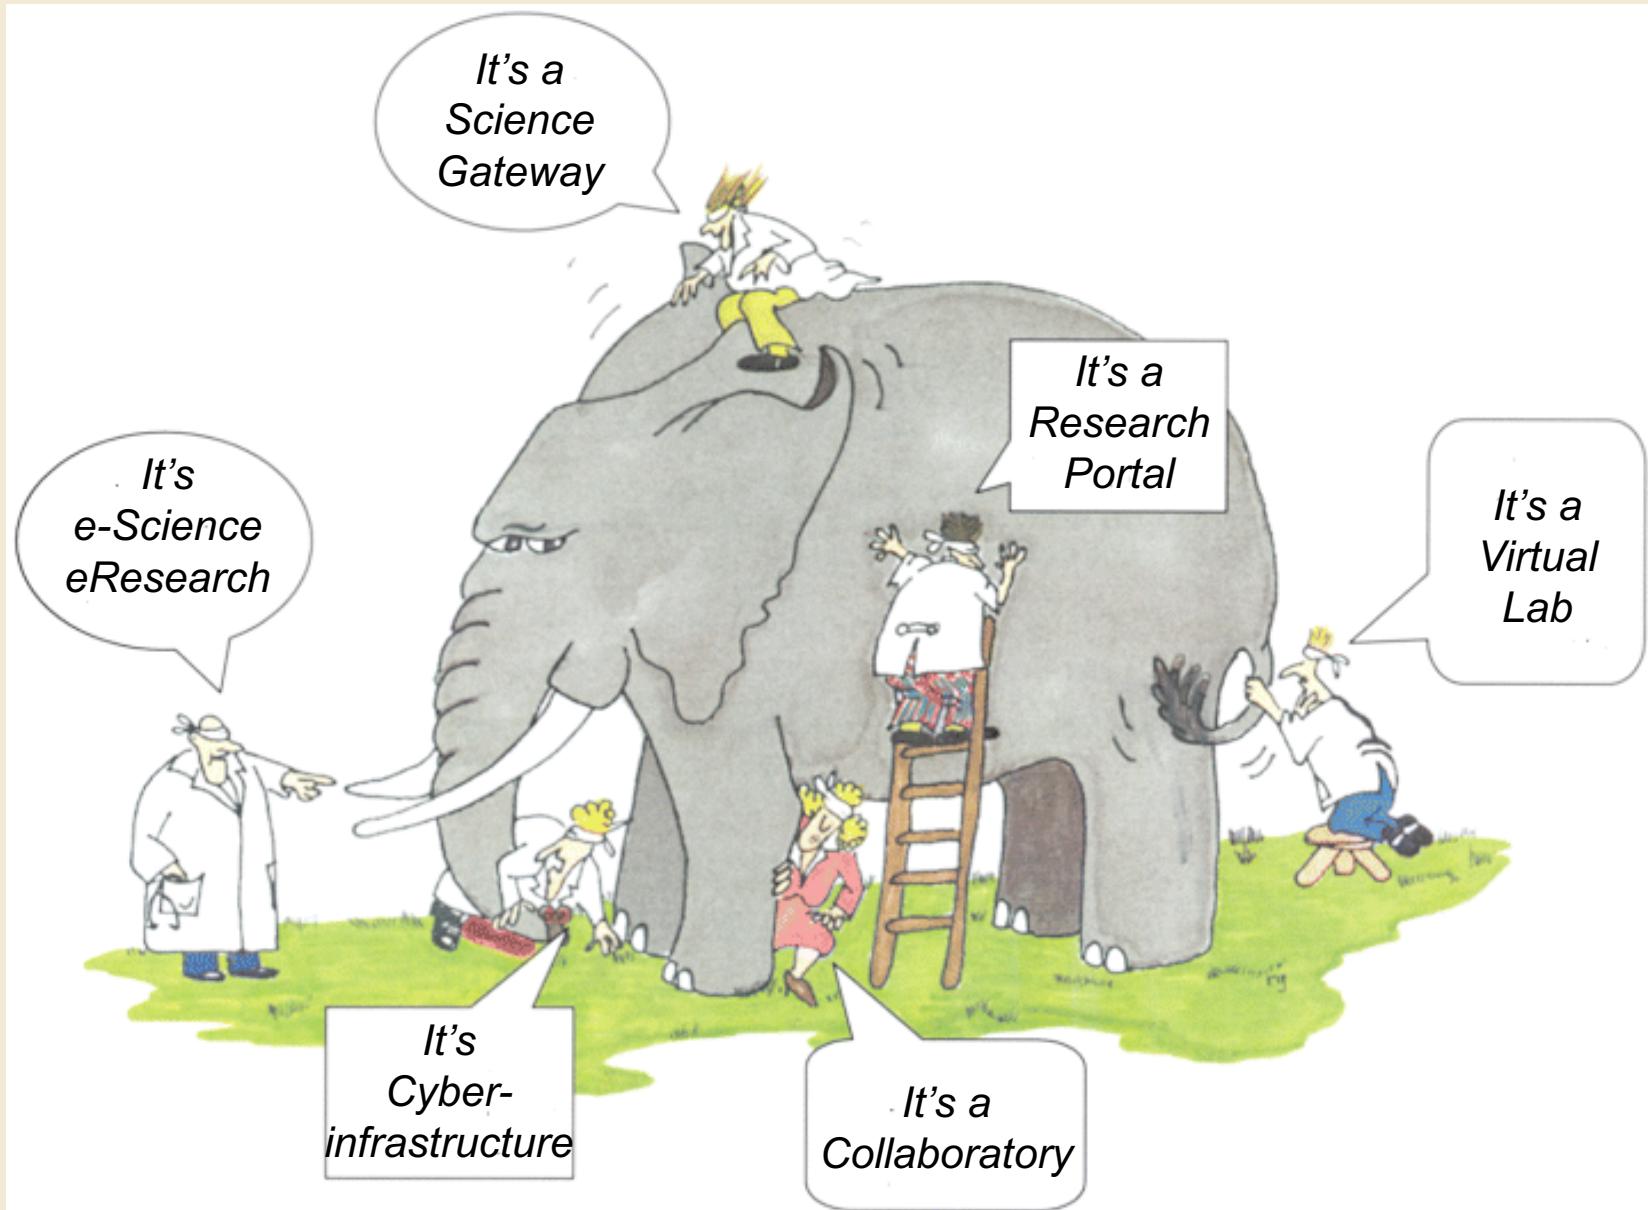
1. an online community space for science and engineering research and education.
2. a Web-based resource for accessing data, software, computing services, and equipment specific to the needs of a science or engineering discipline.





Source: Nancy Wilkins-Diehr





Source: Nancy Wilkins-Diehr



INDIANA UNIVERSITY BLOOMINGTON  
**SCHOOL OF INFORMATICS AND COMPUTING**

CSCI-B 649 Advanced Science Gateway Architectures

# For this class, Science Gateway is:

- Implement distributed systems to enable Science.
- Adopt architectural patterns used by web-scale companies.
- Usage and growth of gateways is elastic.
- Tie together diverse, heterogeneous Computational Resources (analogous to multiple cloud vendors).
- Have to be very secure (science is proprietary at least until published/patented/tenured)



# Anatomy of a Science Gateway

- Gateway User Interface
  - Web Portals
  - Desktop Clients
  - Social/ Collaboration Capabilities
- Security Infrastructure
- Execution Frameworks
  - Application Abstraction
  - Workflow construction & Enactment
  - Compute Resource Management
  - Scheduling
  - Messaging System
- Data Services
  - Replica Management
  - Provenance
  - Analyses
  - Reproducibility
  - Visualization



# What is Supercomputing?

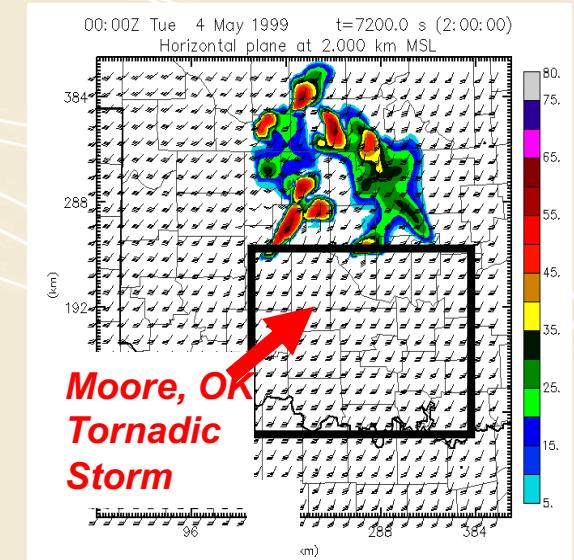
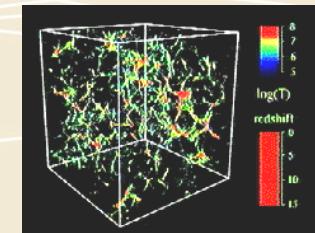
- ***Supercomputing*** is the **biggest, fastest computing right this minute.**
  - So, the definition of supercomputing is **constantly changing**.
  - GigaFLOPS, TeraFLOPS, PetaFLOPS, ExaFLOPS.....
- Supercomputing is also heard in context of :
  - *High Performance Computing (HPC)*,
  - *High End Computing (HEC)*,
  - Grid Computing
  - *Cyberinfrastructure (CI)*.

Source: *Supercomputing in Plain English*, Henry Neeman, University of Oklahoma



# What is HPC (hence Science Gateways) Used For?

- **Simulation** of physical phenomena, such as
  - Weather forecasting
  - Galaxy formation
  - Oil reservoir management
- **Data mining**: finding **needles** of information in a **haystack** of data, such as
  - Gene sequencing
  - Signal processing
  - Detecting storms that might produce tornados
- **Visualization**: turning a vast sea of **data** into **pictures** that a scientist can understand



Source: *Supercomputing in Plain English*, Henry Neeman, University of Oklahoma



# What a Cluster is ....

A cluster is a collection of small computers, called **nodes**, hooked together by an **interconnection network**.

It also **needs** software that allows the nodes to communicate over the interconnect.

A cluster **is** ... is all of these components working together as if they're one big computer ... a **super** computer.

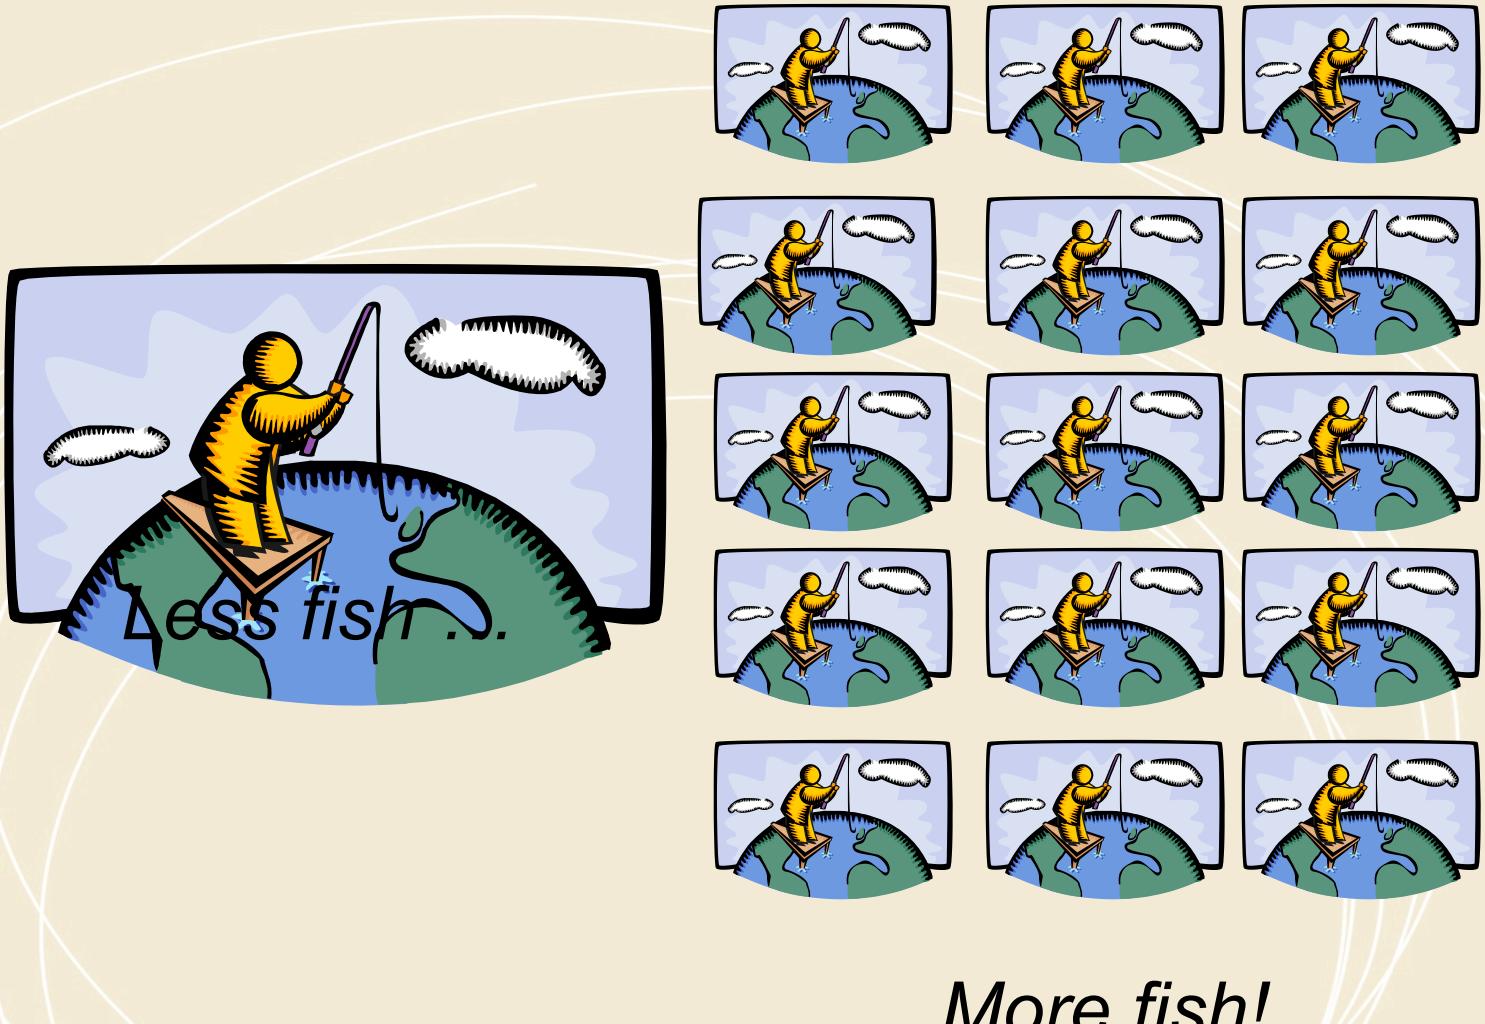


*Source: Supercomputing in Plain English, Henry Neeman, University of Oklahoma*



# Parallelism

**Parallelism** means  
doing multiple things at  
the same time: you can  
get more work done in  
the same time.



Source: *Supercomputing in Plain English*, Henry Neeman, University of Oklahoma



# What is Parallelism?

**Parallelism** is the use of multiple processing units – either processors or parts of an individual processor – to solve a problem, and in particular the use of multiple processing units operating concurrently on different parts of a problem. The different parts could be different tasks, or the same task on different pieces of the problem's data.

*Source: Supercomputing in Plain English, Henry Neeman, University of Oklahoma*



# Berekely Dwarfs

- <https://www2.eecs.berkeley.edu/Pubs/TechRpts/2006/EECS-2006-183.pdf>



Source: Prof. Geoffrey Fox



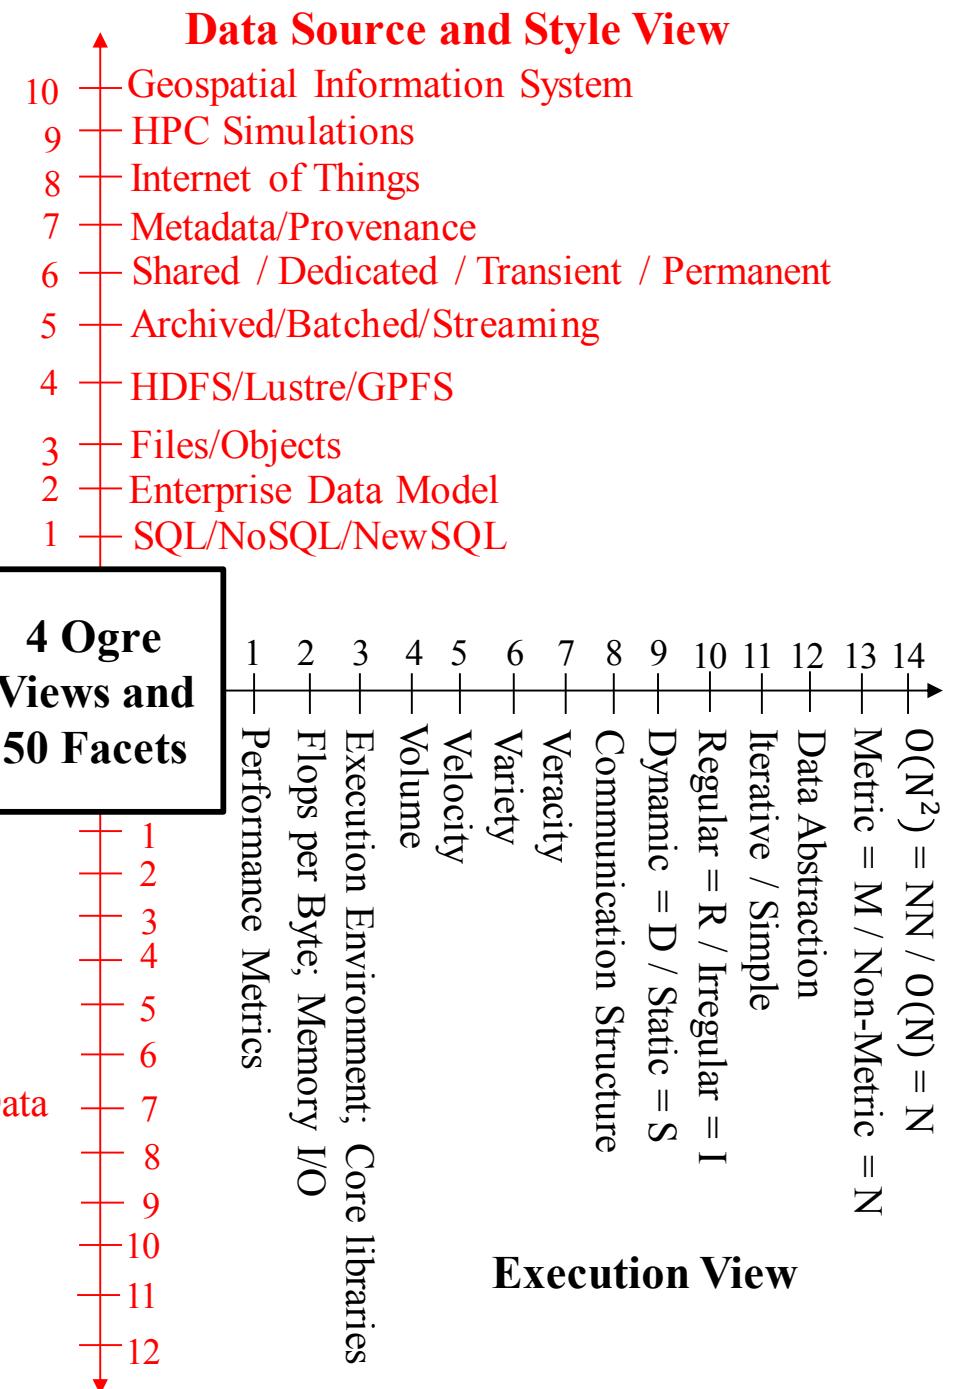
INDIANA UNIVERSITY  
SCHOOL OF

## Problem Architecture View

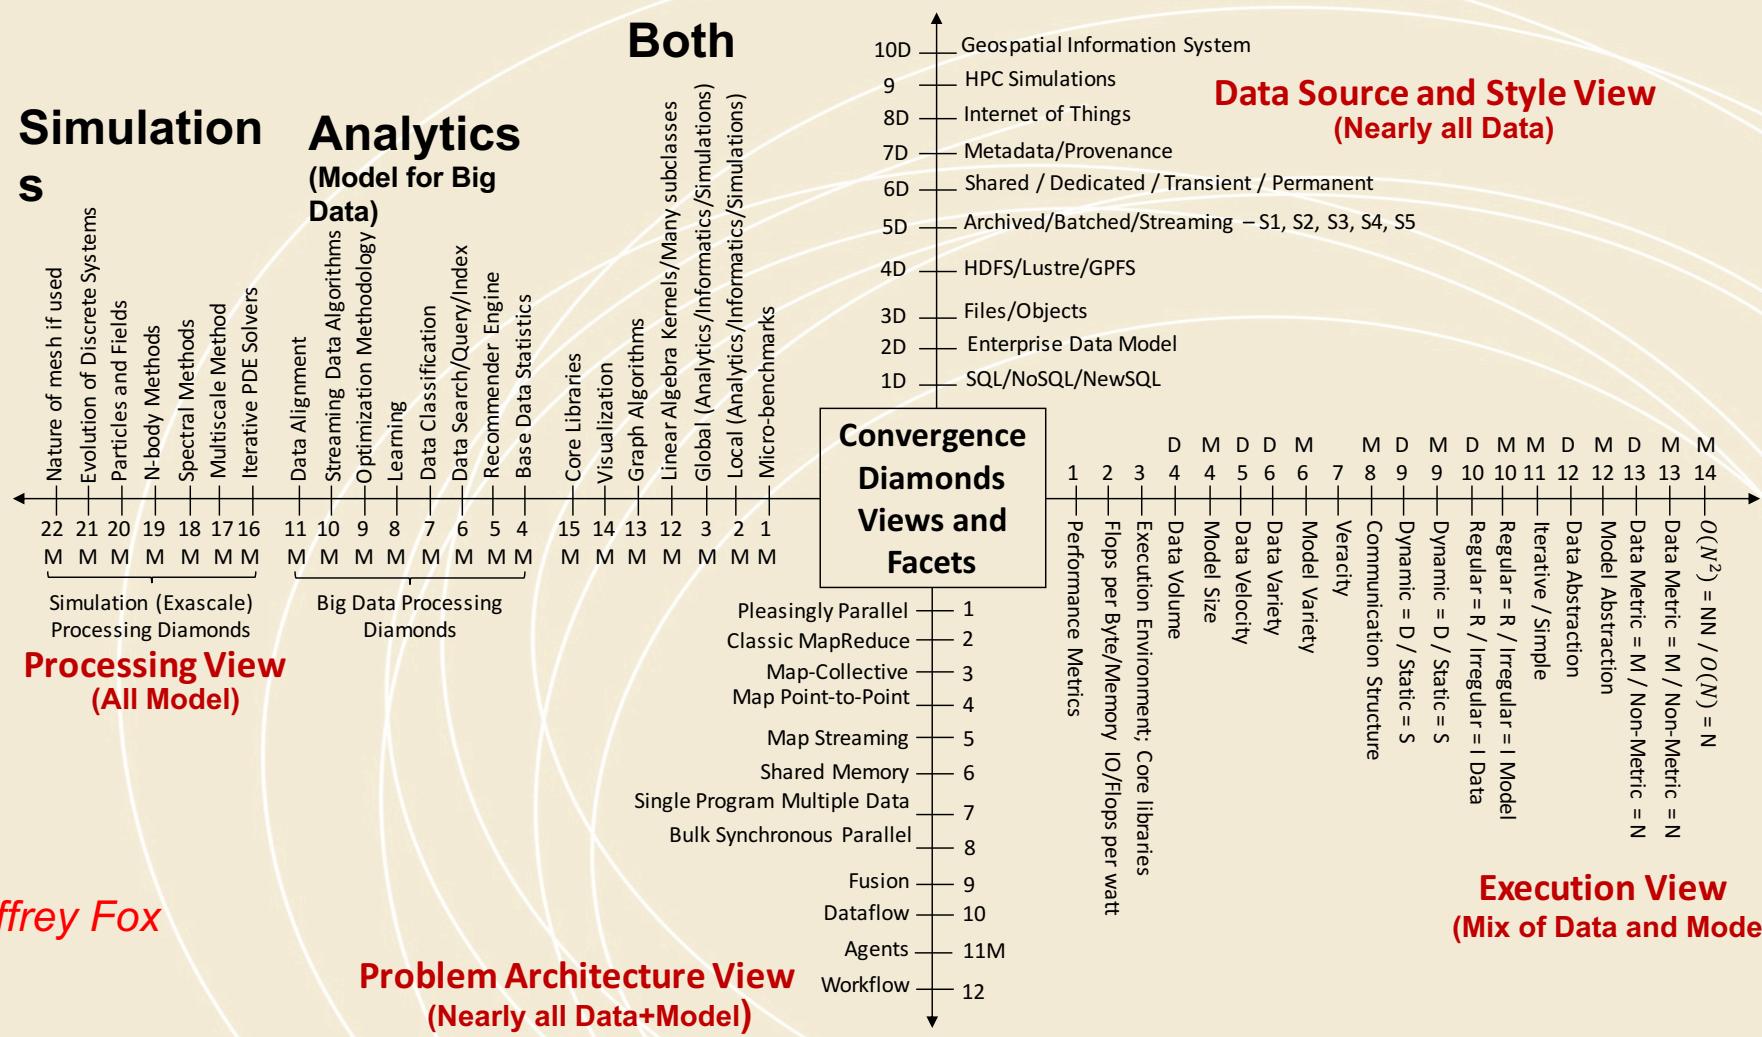
- Optimization Methodology
    - Streaming
    - Alignment
  - Linear Algebra Kernels
  - Graph Algorithms
  - Visualization
- Pleasingly Parallel
- Classic MapReduce
  - Map-Collective
  - Map Point-to-Point
  - Map Streaming
  - Shared Memory
  - Single Program Multiple Data
  - Bulk Synchronous Parallel
  - Fusion
  - Dataflow
  - Agents
  - Workflow

## Processing View

- Micro-benchmarks
- Local Analytics
- Global Analytics
- Base Statistics
- Recommendations
- Search / Query / Index
- Classification
- Learning



# 64 Features in 4 views for Unified Classification of Big Data and Simulation Applications



Source: Prof. Geoffrey Fox



# 6 Forms of MapReduce

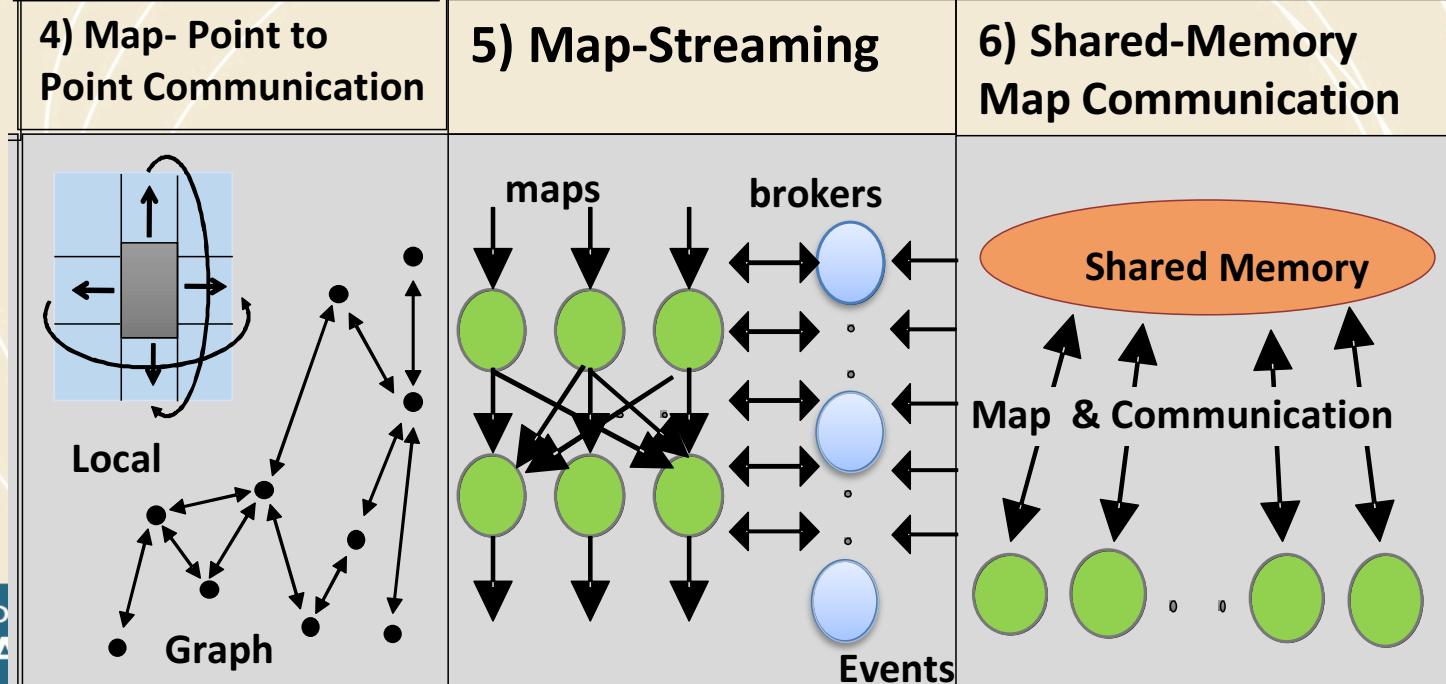
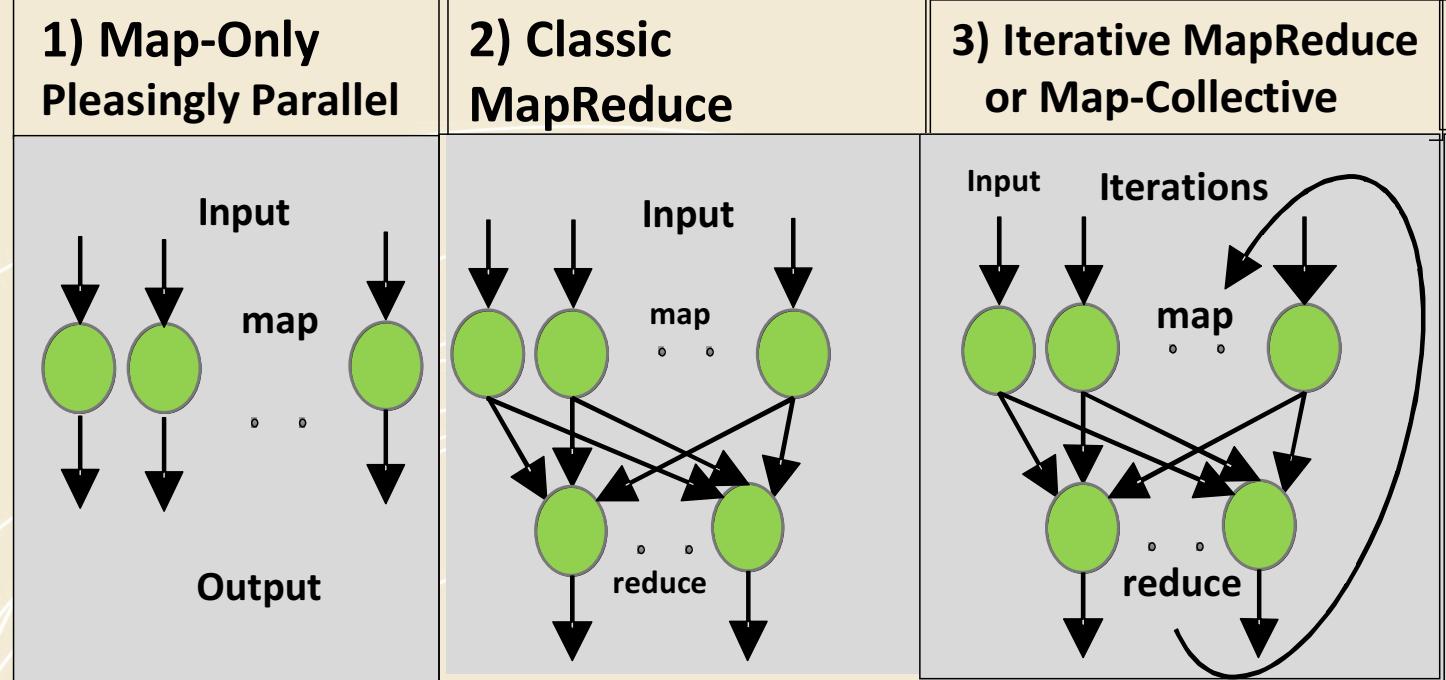
Describes Architecture of

- Problem (Model reflecting data)
- Machine
- Software

2 important variants (software) of Iterative MapReduce and Map-Streaming

- a) **"In-place"** HPC
- b) Flow for model and data

Source: Prof. Geoffrey Fox

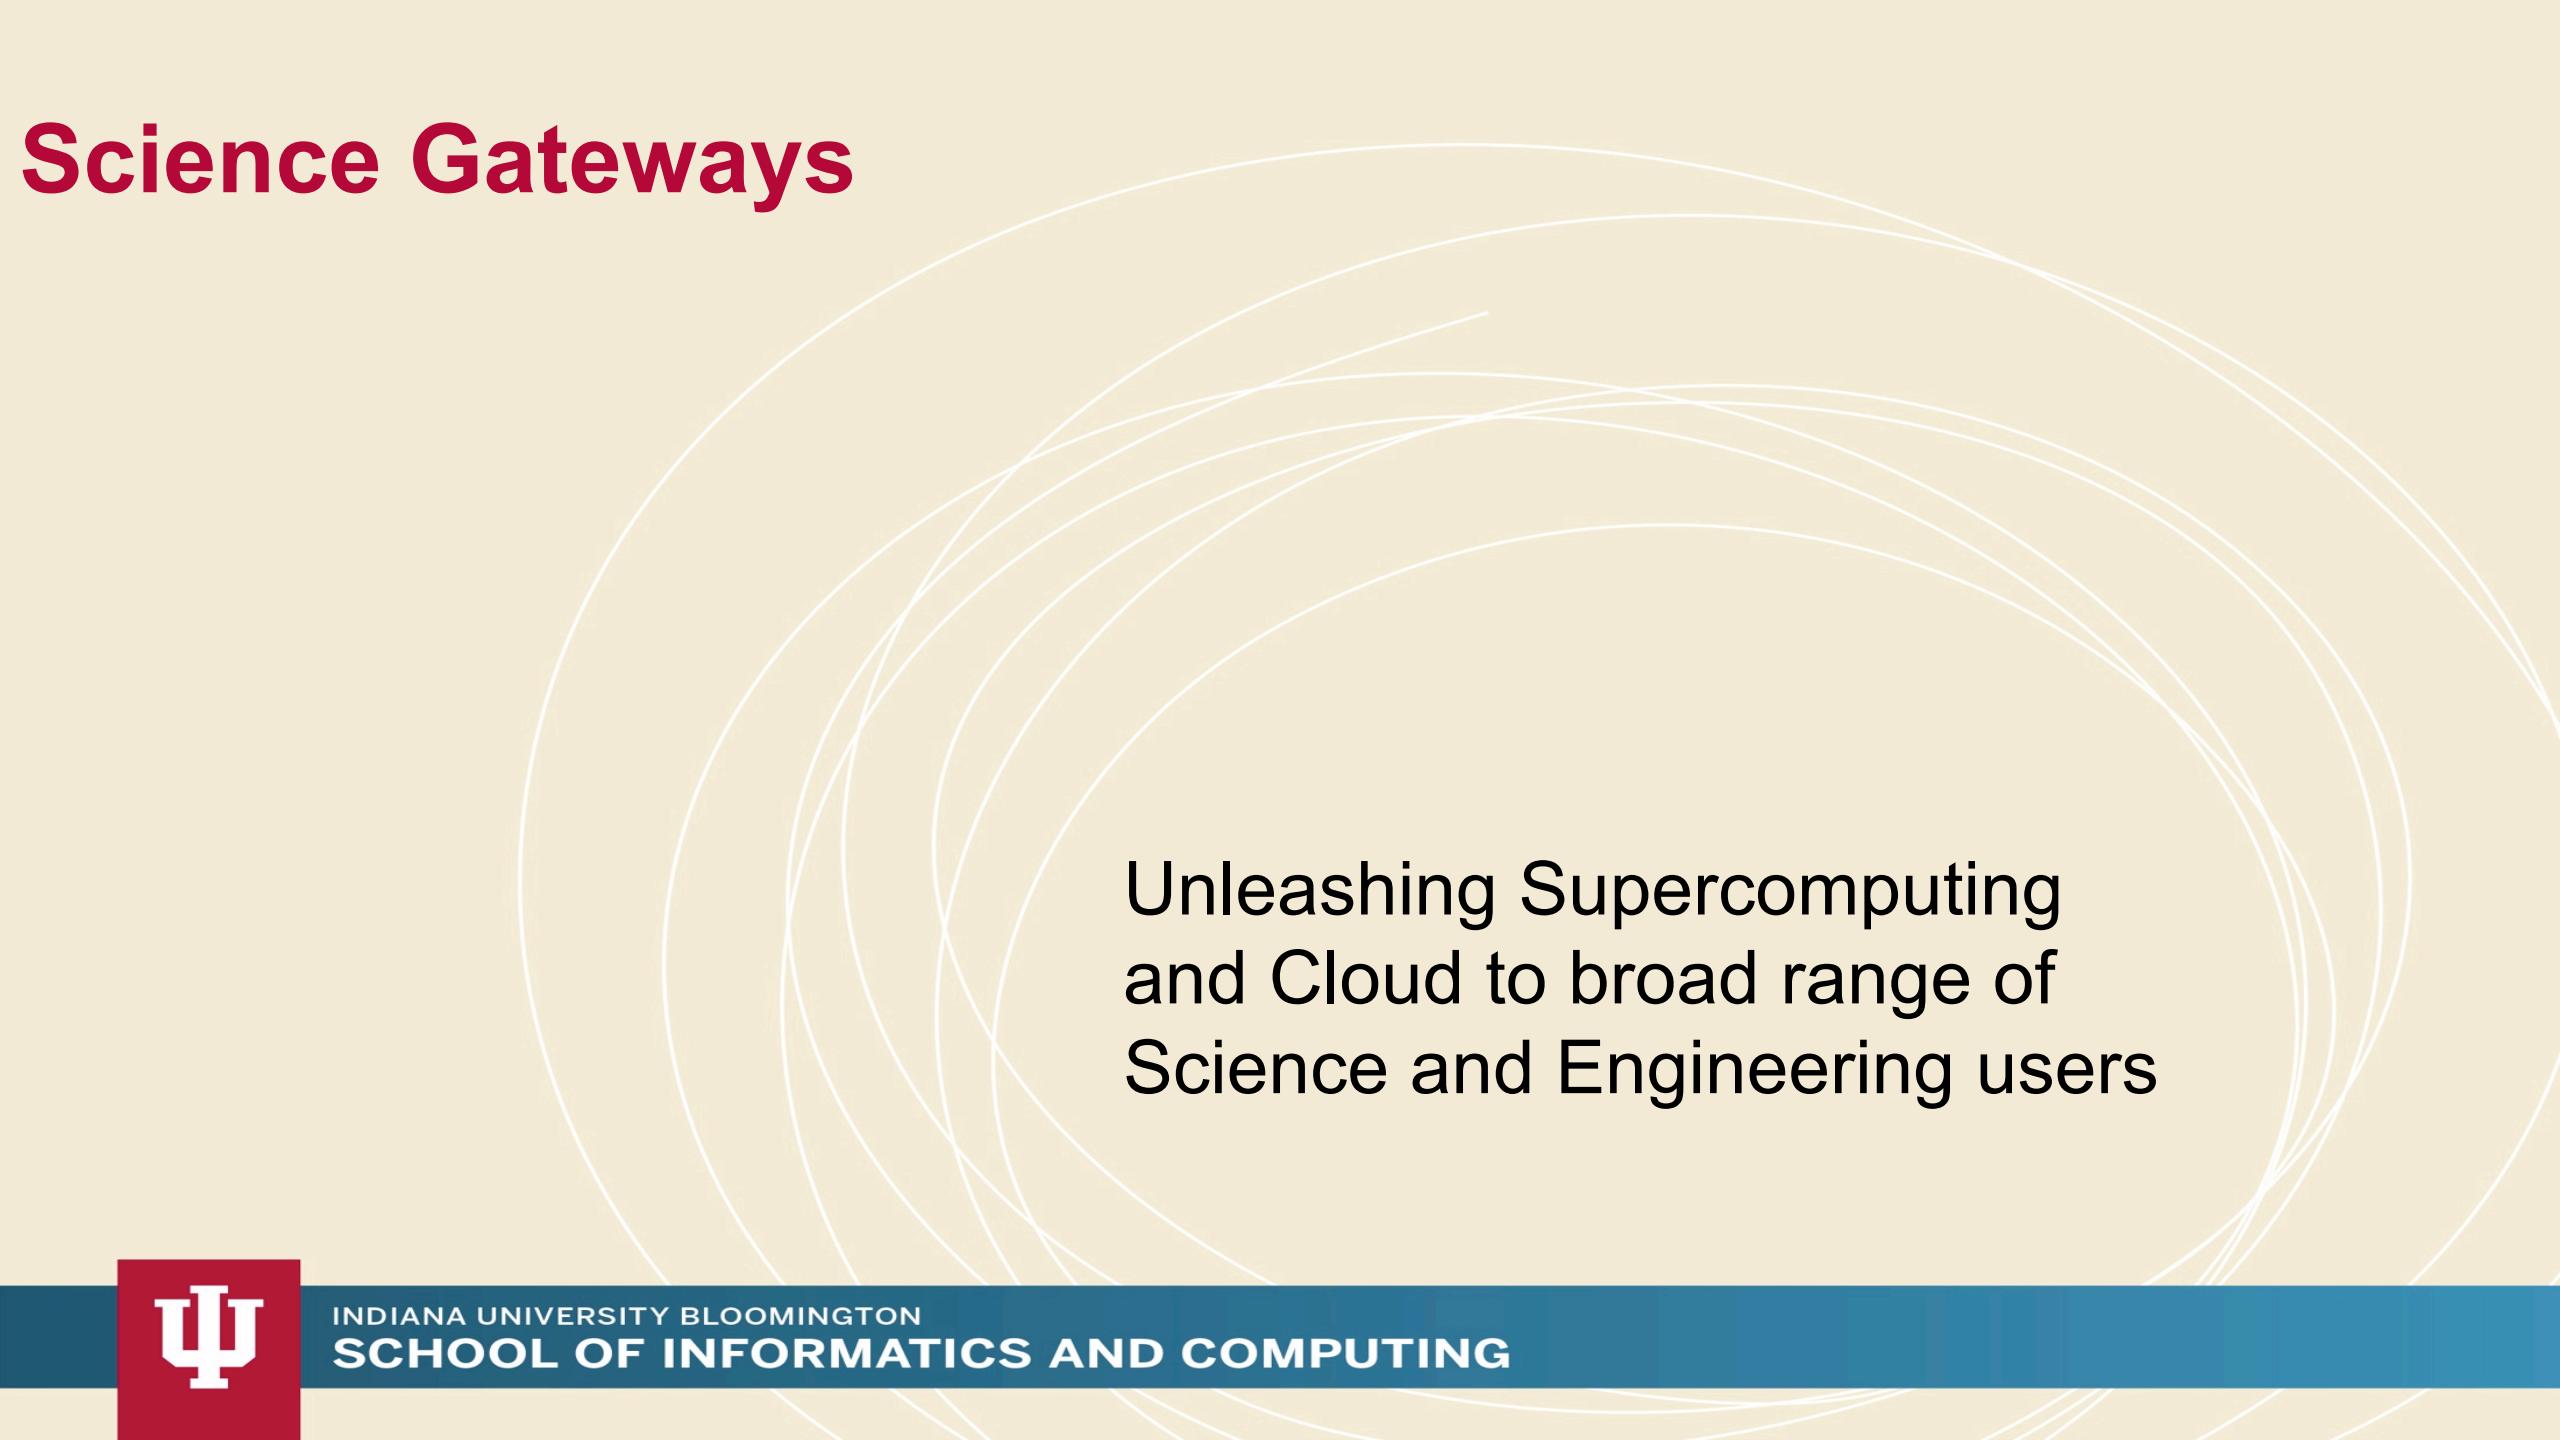


# **Gateways abstract these details from Users**

- Parallelism Types
  - Instruction Level Parallelism
  - Shared Memory Multithreading
  - Distributed Memory Multiprocessing
  - GPU Parallelism
  - Hybrid Parallelism (Shared + Distributed + GPU)
- Gateways should pick the right computations paradigm based on the problem.



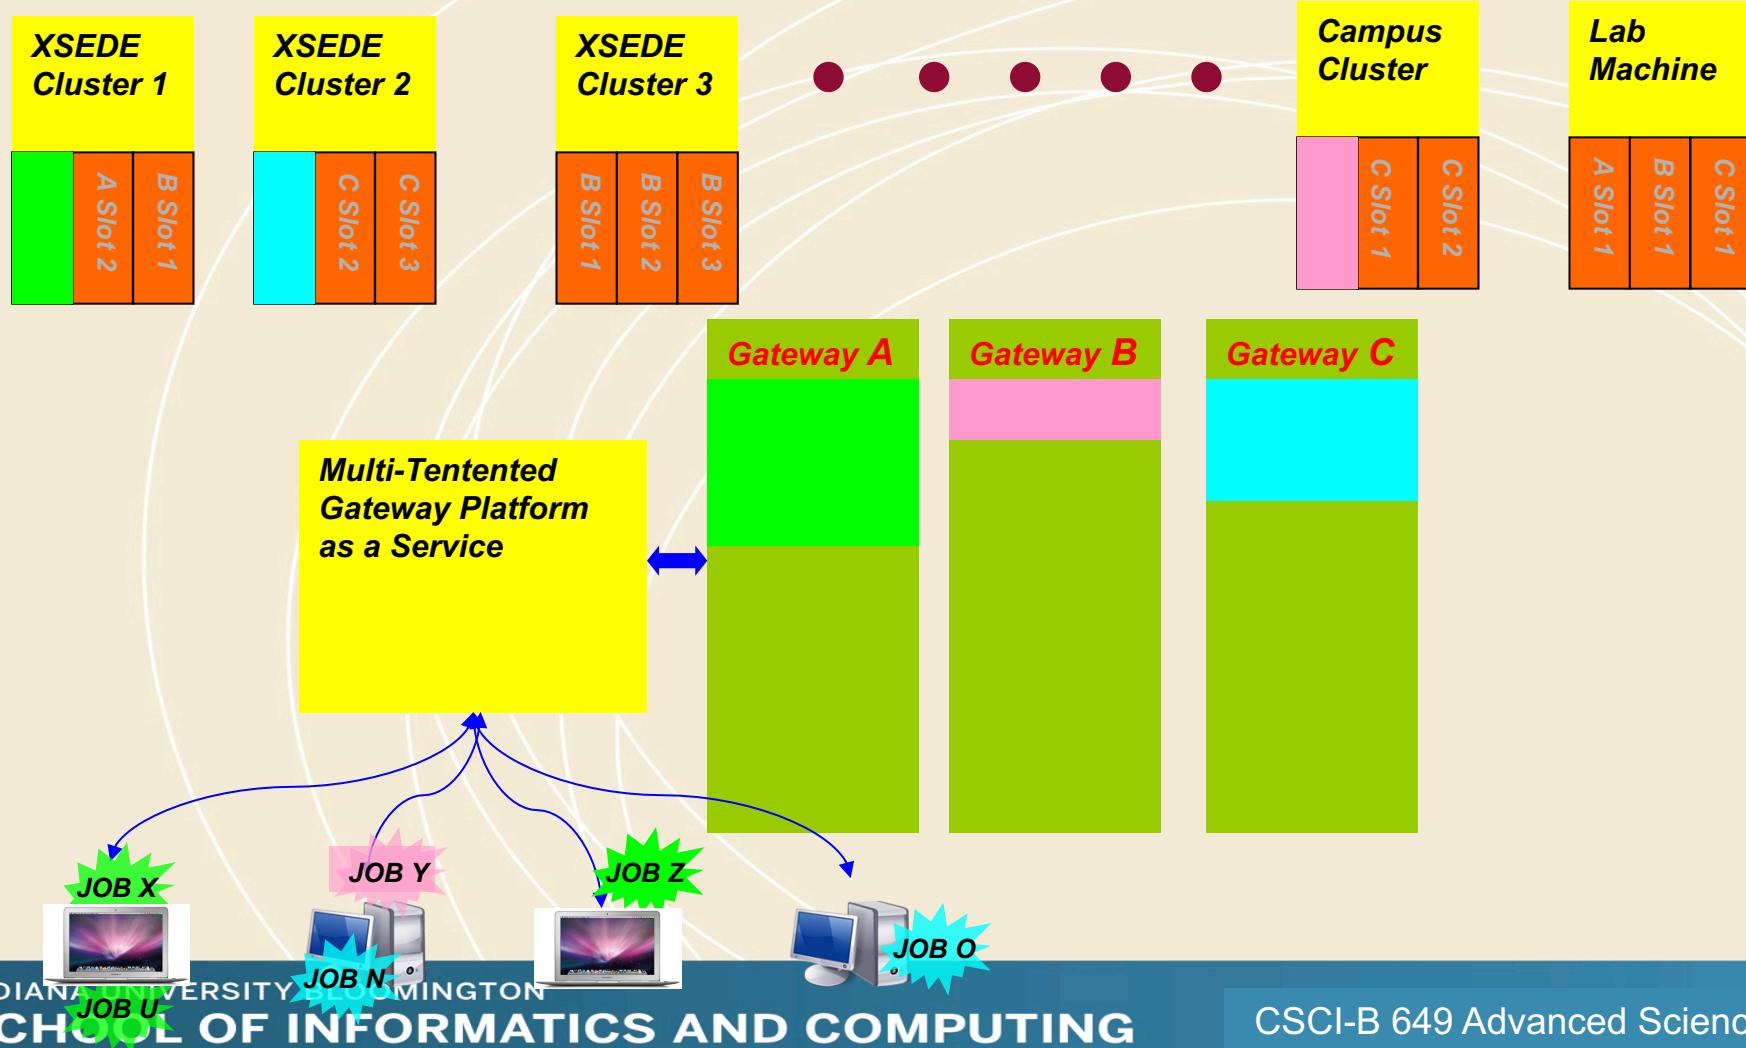
# Science Gateways



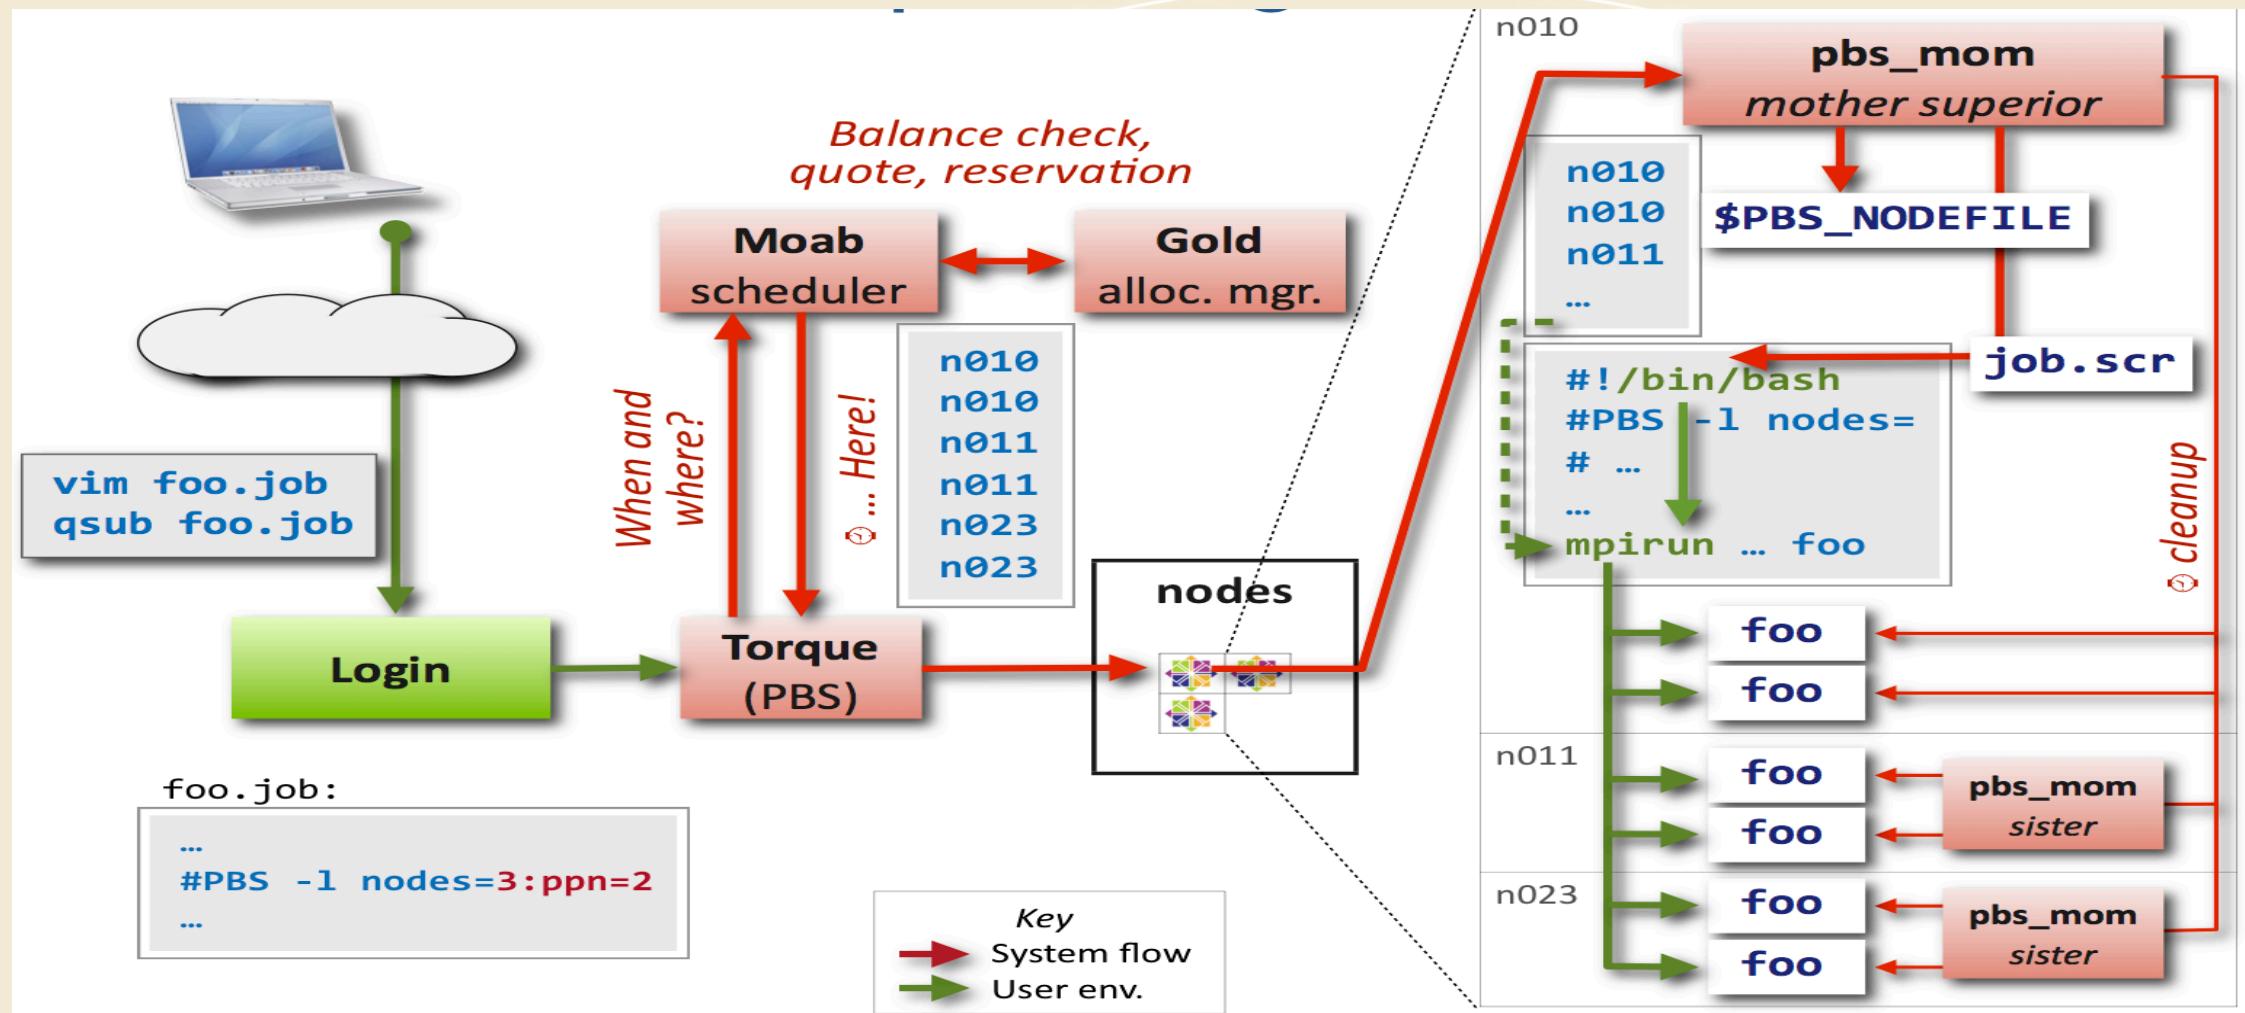
Unleashing Supercomputing  
and Cloud to broad range of  
Science and Engineering users



# Example: Science Gateways Federate Jobs across clusters



# Job Managers Flow



Source: [https://wiki.anl.gov/cnm/HPC/Submitting\\_and\\_Managing\\_Jobs](https://wiki.anl.gov/cnm/HPC/Submitting_and_Managing_Jobs)

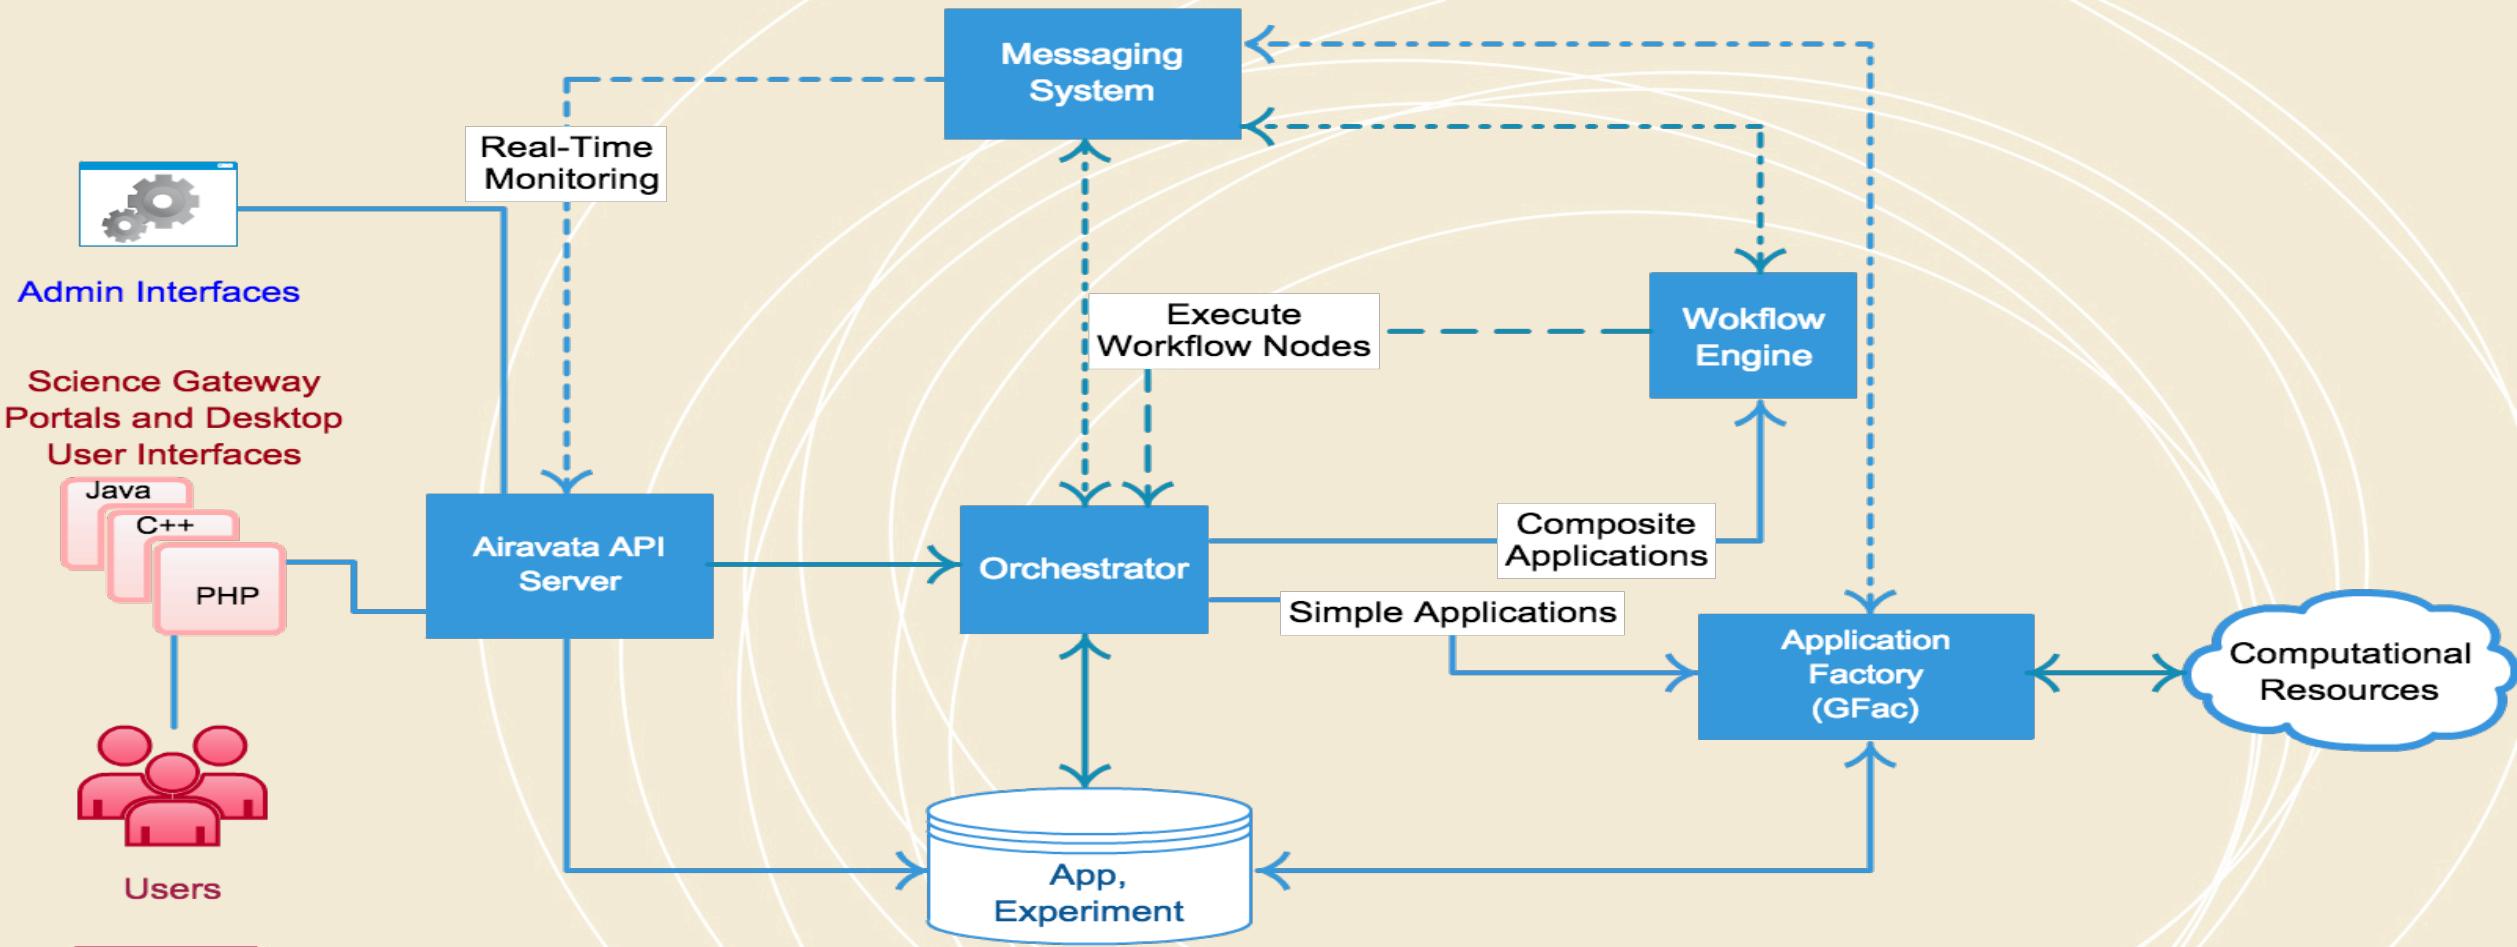


# Apache Airavata



INDIANA UNIVERSITY BLOOMINGTON  
**SCHOOL OF INFORMATICS AND COMPUTING**

# Airavata: Multi-Tentanted Gateway Middleware



# Airavata Overview

- Airavata is a general purpose distributed system software framework build on micro-service and component based architecture principles.
- Airavata provides capabilities to compose, manage, execute and monitor large scale applications and workflows on distributed computing resources.
- Airavata supports executions on local clusters, national grids, academic and commercial clouds.
- Airavata is inherently multi-tenanted.
- Airavata is increasingly evolving to support 4<sup>th</sup> Paradigm – “Data-Intensive Scientific Discovery”.

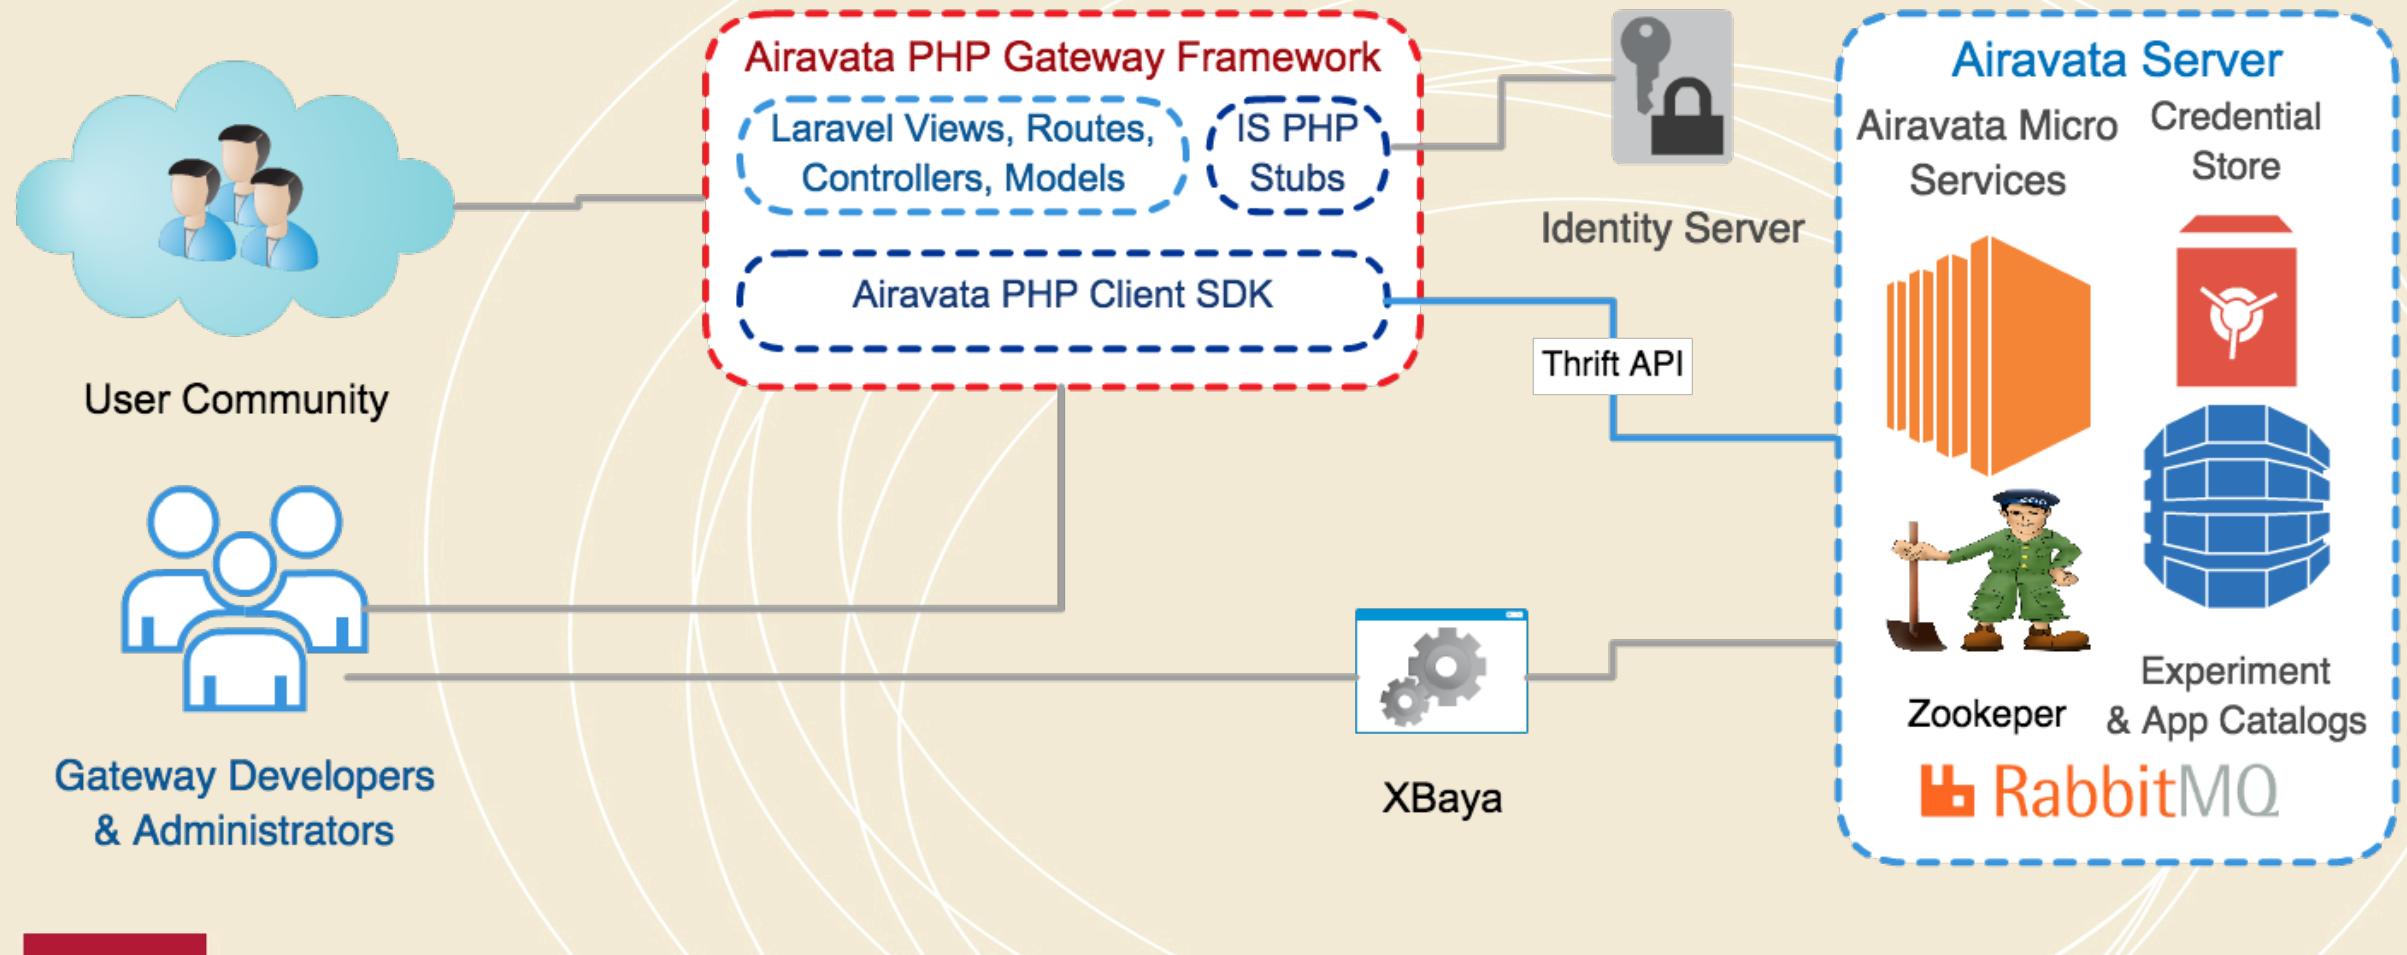


# Airavata as a Science Gateway Middleware

- Airavata is dominantly used to build science gateways.
- Airavata supports secured communications to HPC resources and empowers gateway operators to administer and monitor long running executions.
- A reference PHP based gateway is provided to illustrate the Airavata capabilities and can be used to customize science-centric gateways



# Reference Gateway <-> Airavata Services



# SciGaP – Powered by Airavata

- Science Gateway Platform as a Service (SciGaP) provides application programmer interfaces (APIs) to hosted generic infrastructure services that can be used by domain science communities to create Science Gateways.
- SciGaP hosted service platform is powered by Apache Airavata.
- SciGaP helps gateway developers to concentrate their efforts on building their scientific communities and not worry about operations.



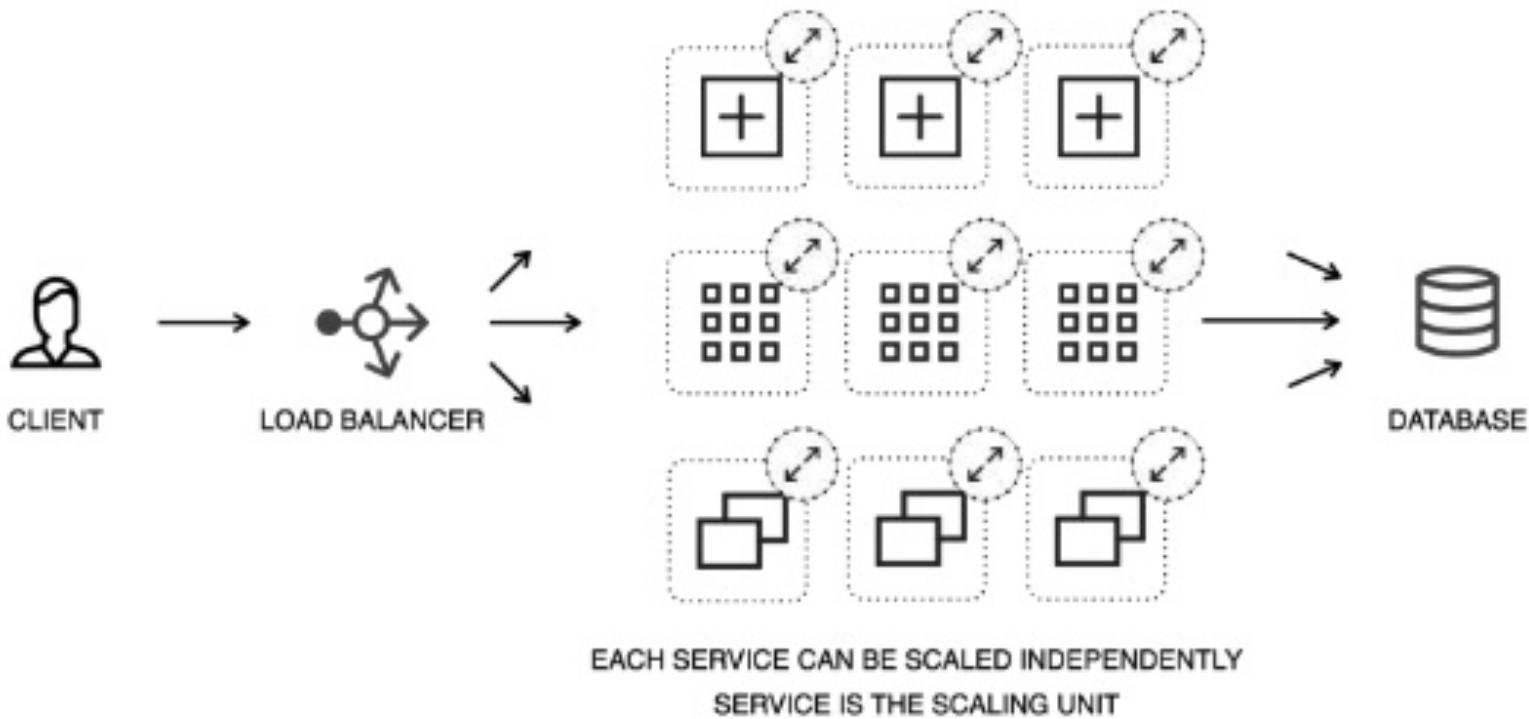
# Challenges for Gateways

- Providing a rich user experience
- Defining an API for the application server
- Defining the right sub-components for the application server.
- Implementing the components, wiring them together correctly.
- Supporting multiple gateway tenants
- Fault tolerance for components
- State management
- Continuous delivery
- Security management
- Supporting full scientific exploratory cycle

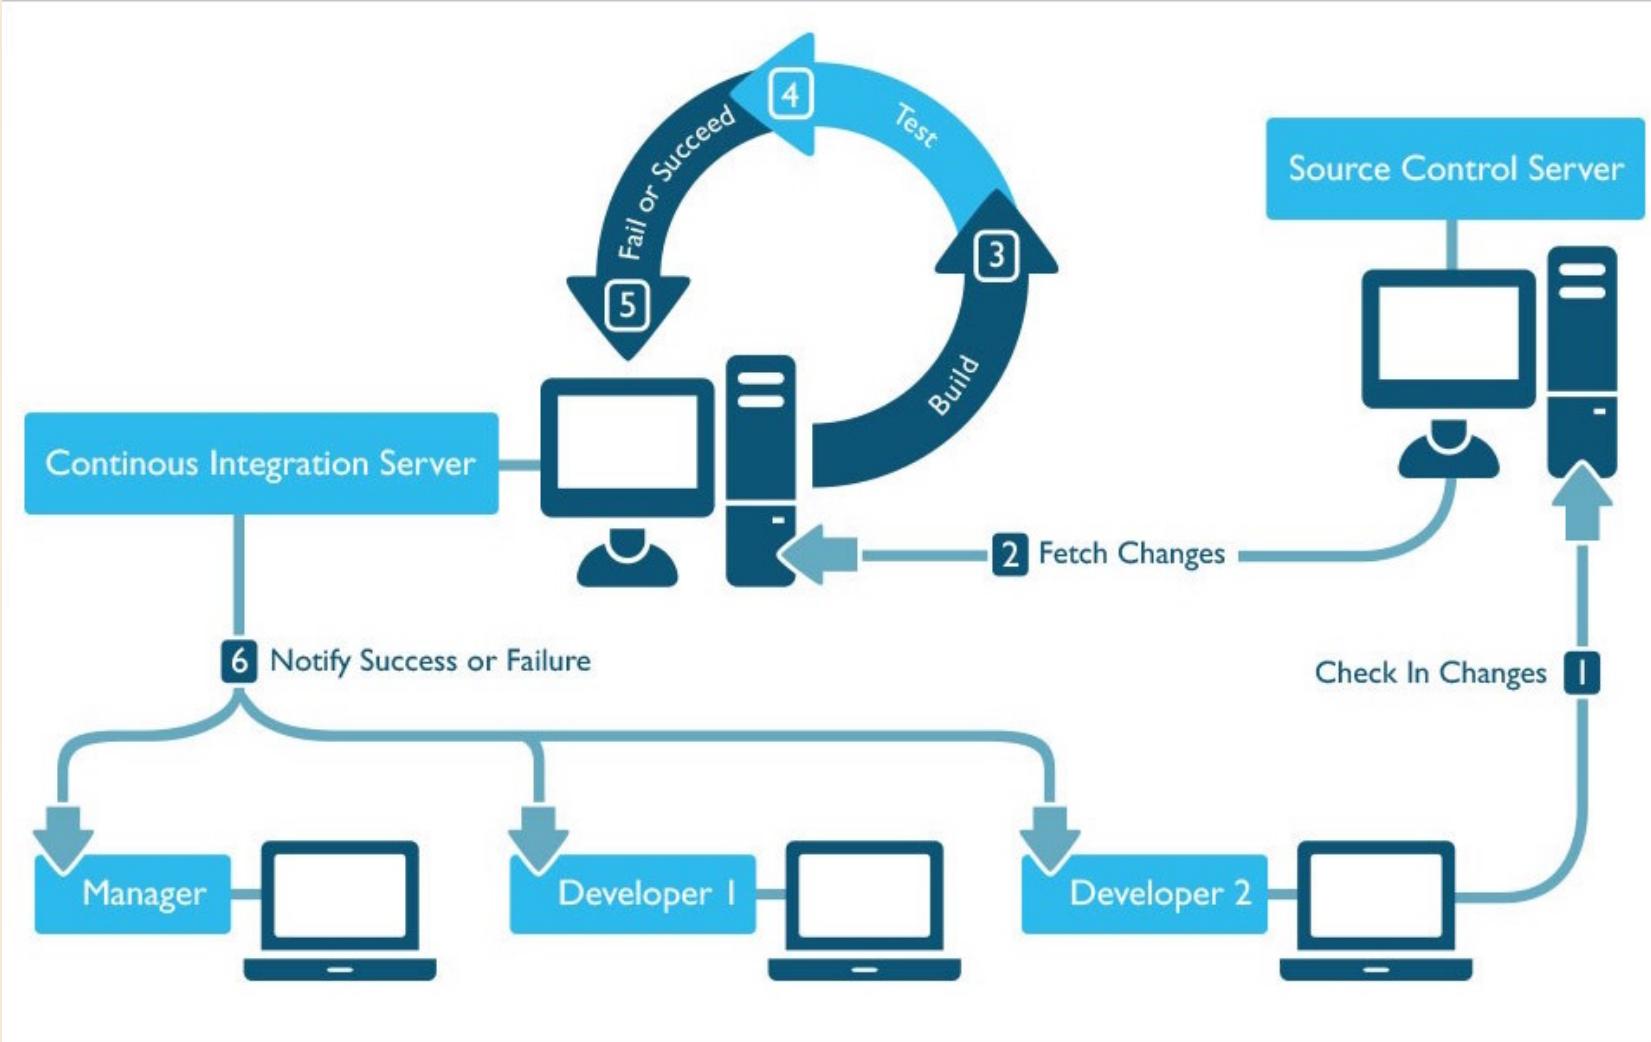


# Component Microservices

## Microservices



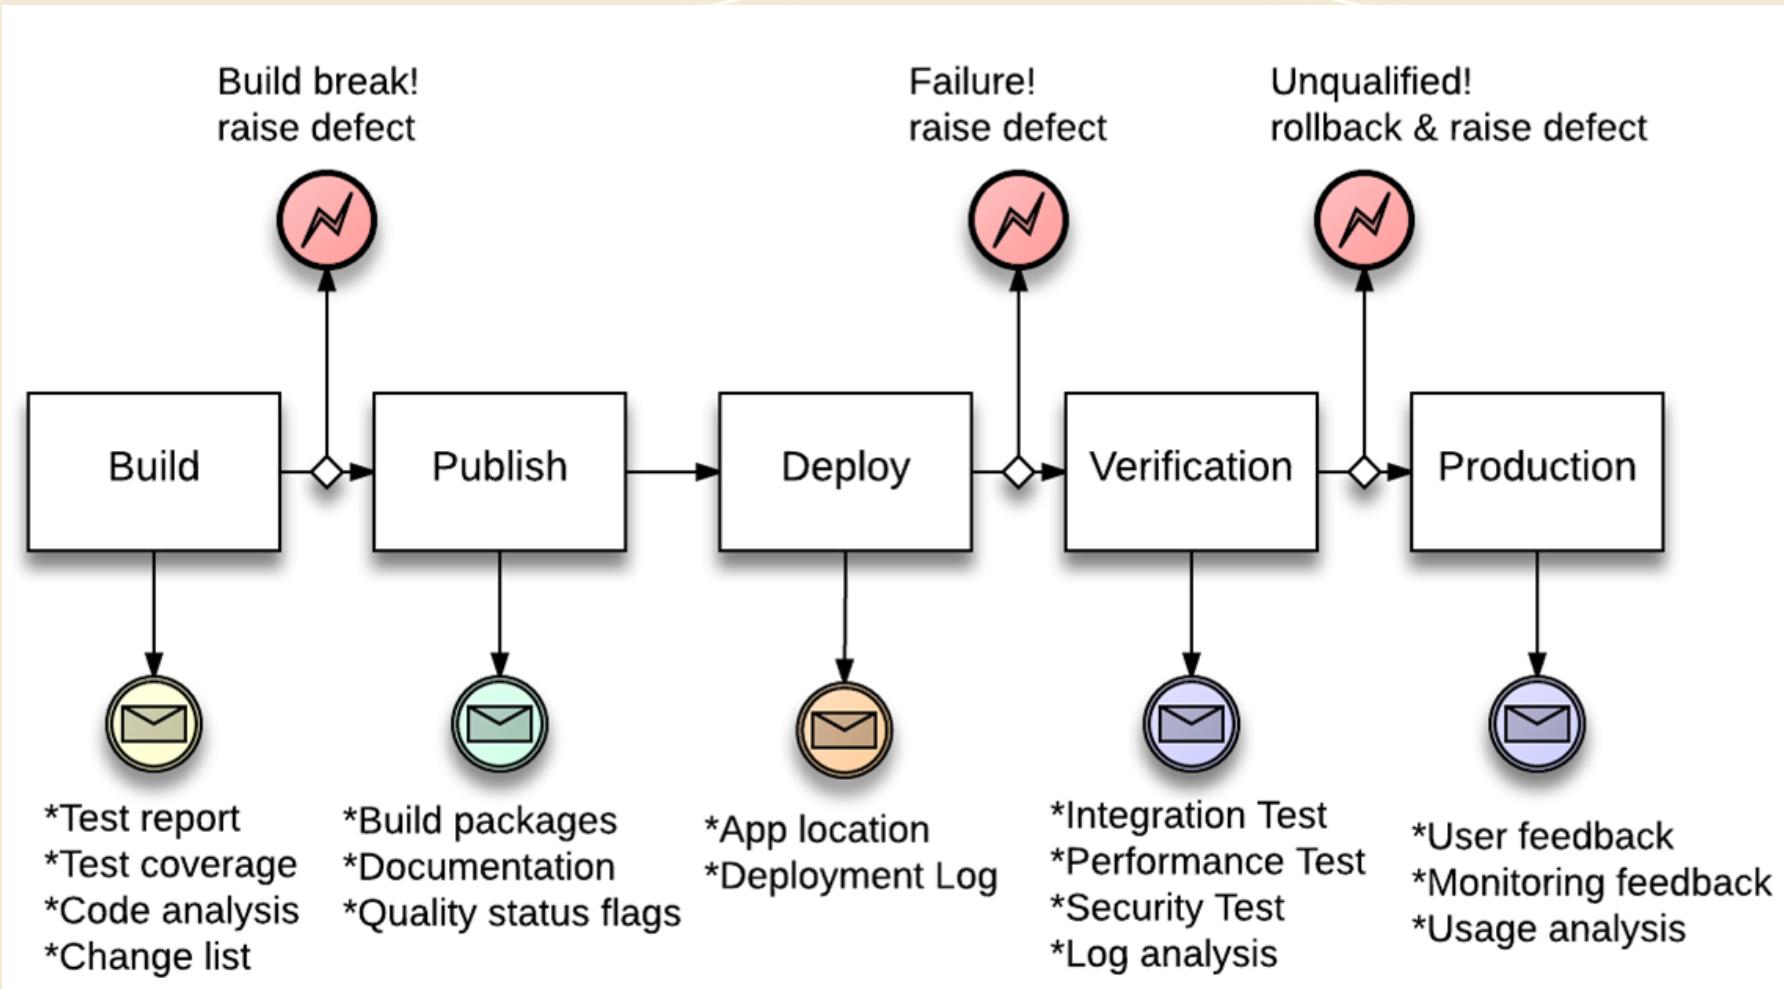
# Continuous Integration



Source: <https://insights.sei.cmu.edu/devops/2015/01/continuous-integration-in-devops-1.html>



# Continuous Deployments/Delivery



Source: [https://www.ibm.com/developerworks/community/blogs/c914709e-8097-4537-92ef-8982fc416138/entry/devops\\_in\\_practice\\_best\\_practices\\_for\\_adopting\\_continuous\\_delivery?lang=en](https://www.ibm.com/developerworks/community/blogs/c914709e-8097-4537-92ef-8982fc416138/entry/devops_in_practice_best_practices_for_adopting_continuous_delivery?lang=en)



# Projects Overview

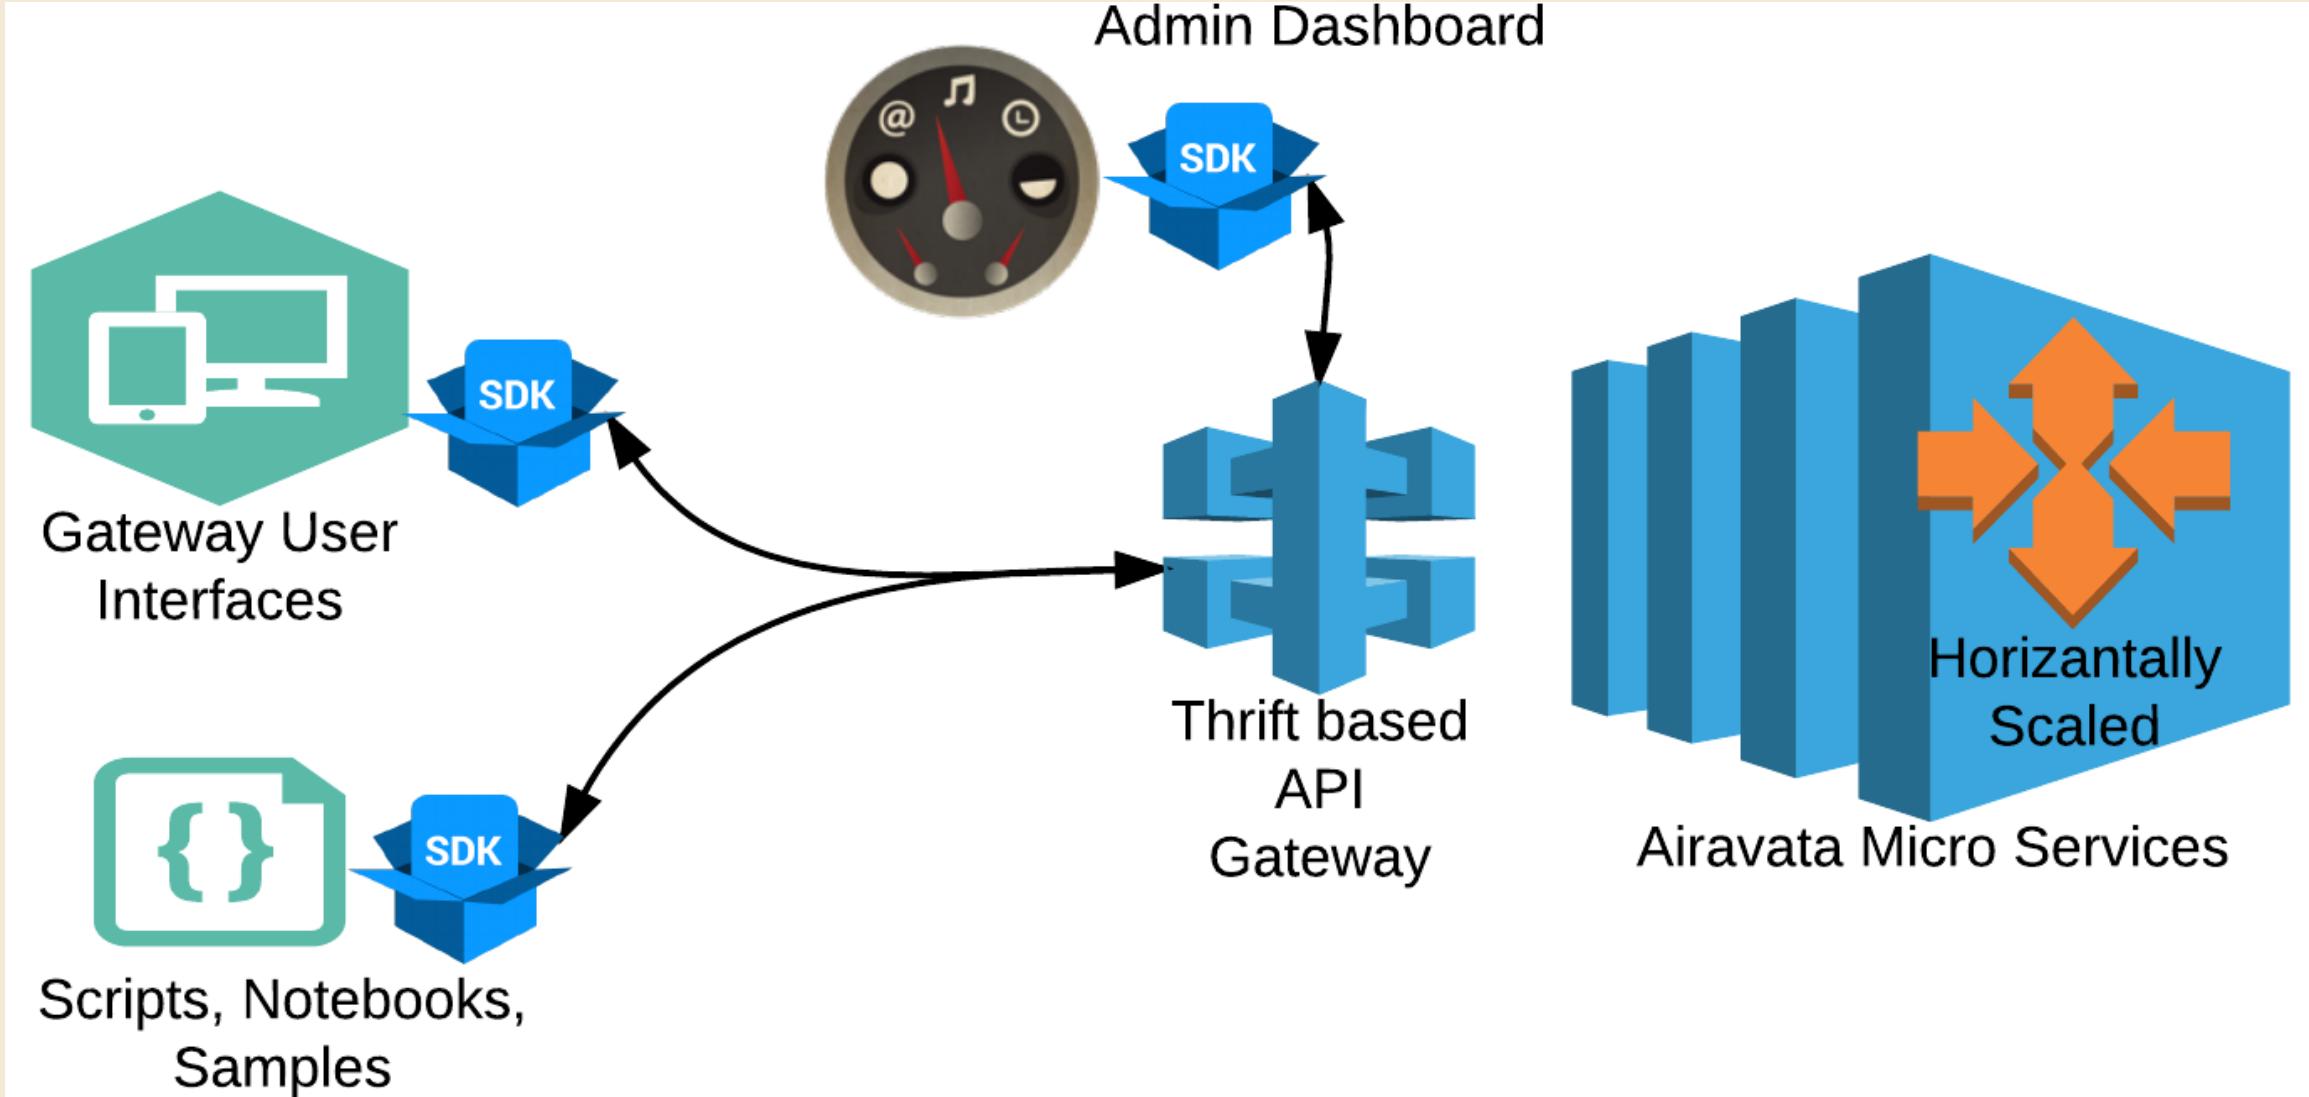
## First 2 week Projects

Will have two themes

January 13<sup>th</sup> 2017 – January 27<sup>th</sup> 2017



# Phase 1 Projects

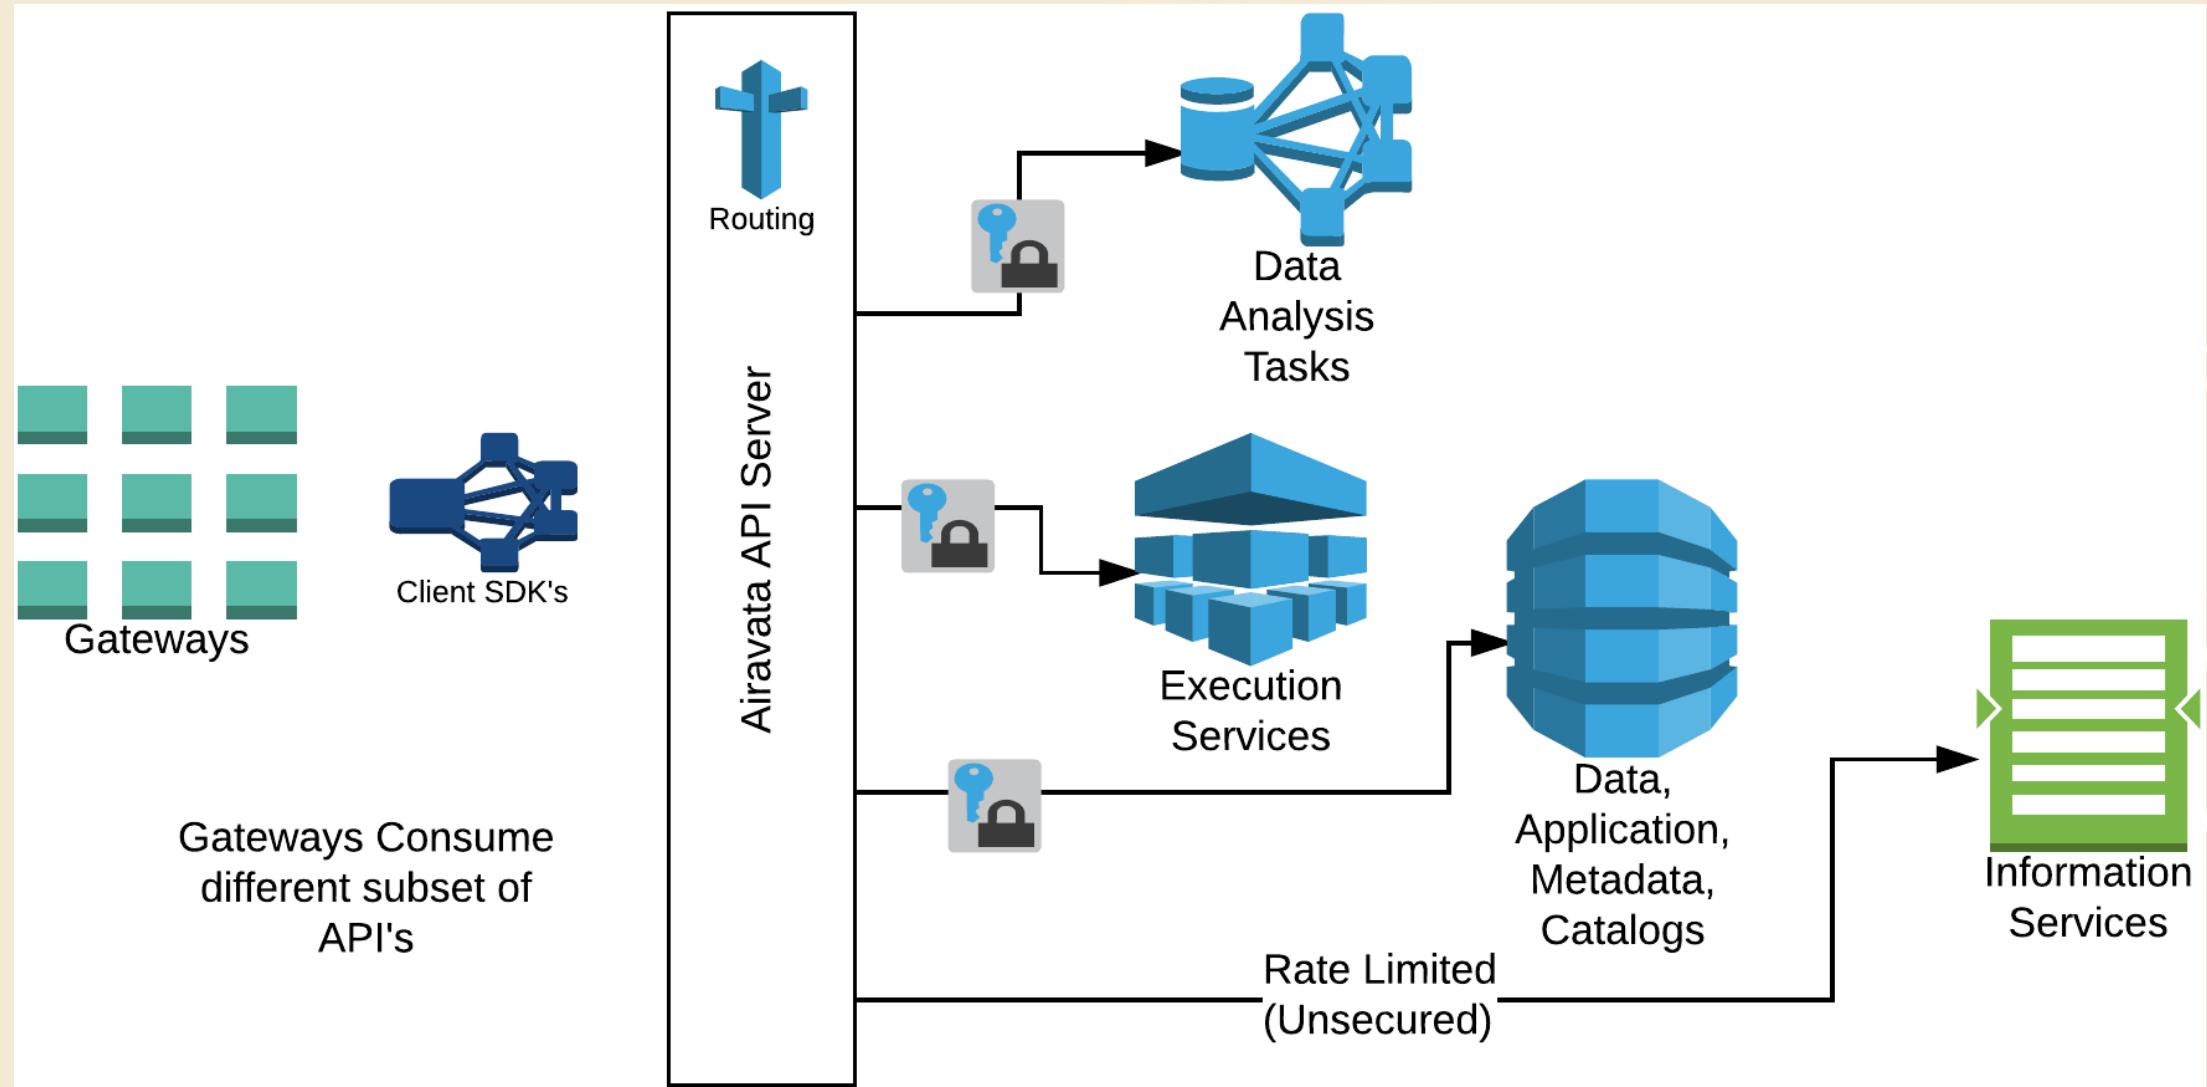


INDIANA UNIVERSITY BLOOMINGTON

**SCHOOL OF INFORMATICS AND COMPUTING**

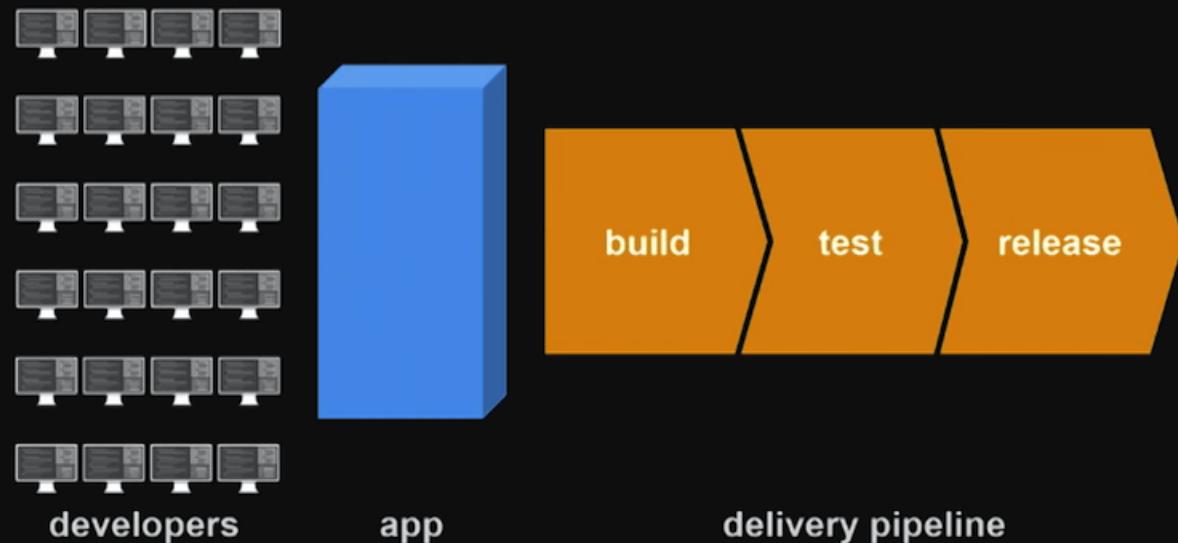
CSCI-B 649 Advanced Science Gateway Architectures

# Phase 2 Projects

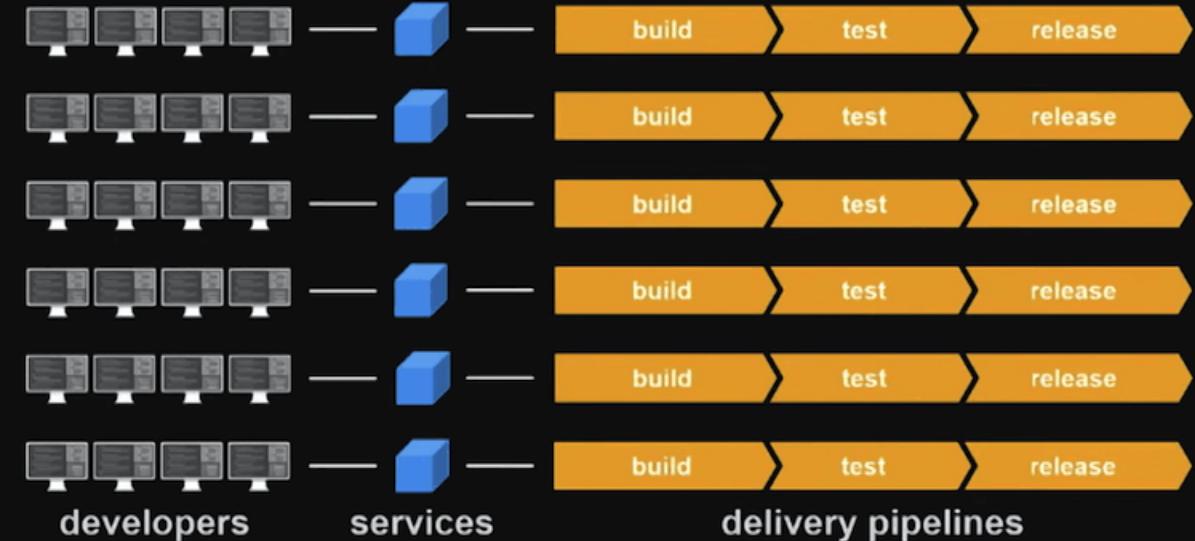


# Phased Releases

## Monolith development lifecycle



## Microservice development lifecycle

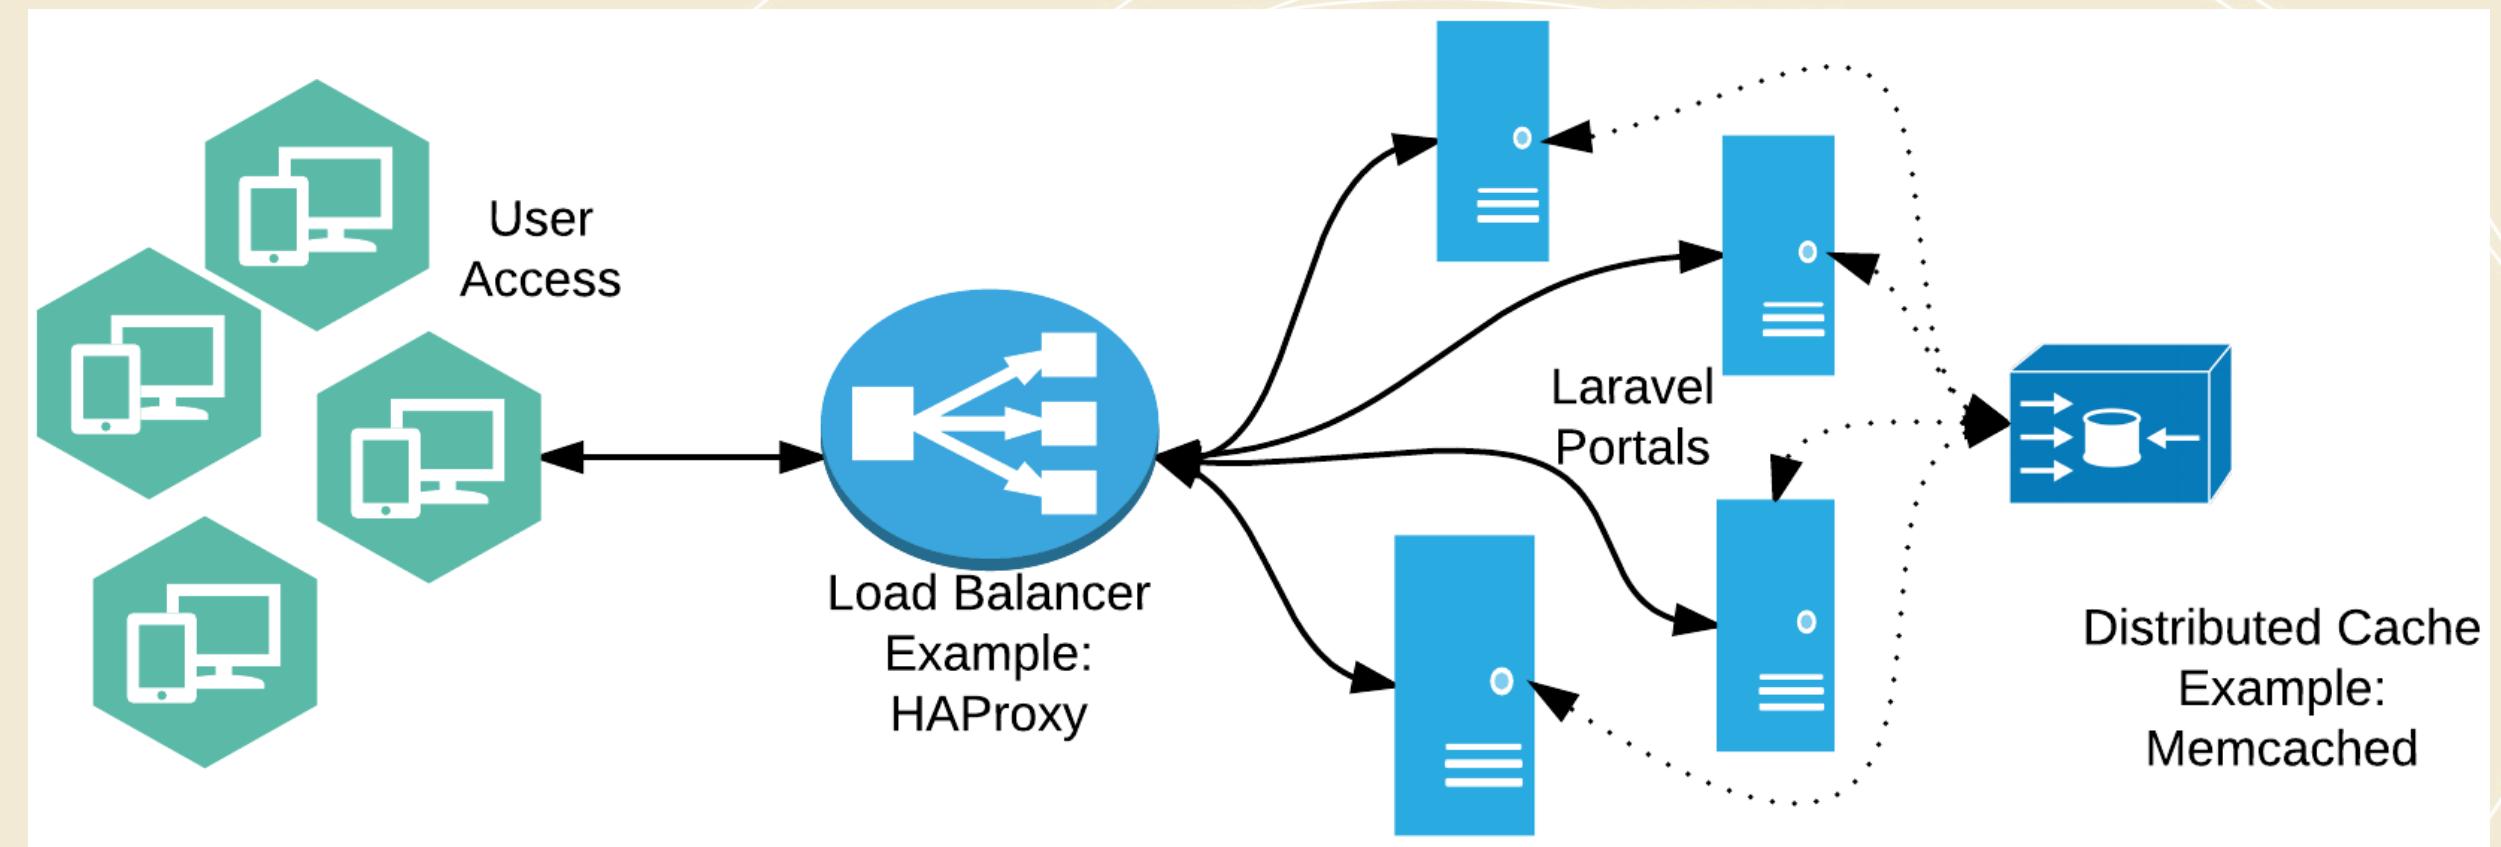


# Project Theme 1: Scalable Web Portal

- Load balance Laravel Portal.
- Fail Over to other regions in Amazon.
- Use distributed cache for session and security management.
- Easily Deployable.
- Have a Build and make it IDE and Developer friendly
- Facilitate micro-feature releases
- Connect laravel Portal with Thrift API Server



# First 2 weeks Goal 1:



# Example Test Case:

- Have multiple gateway instances:
  - sneha.scigap.org, marcus.scigap.org, mayank.scigap.org, eldho.scigap.org, ameya.scigap.org, anuj.scigap.org, sagar.scigap.org.....
  - Each of these portals have their own look and feel and functionality.
  - They all can be hosted from one web server.
- How can we scale the number of gateways.
- How can we not be nervous about hardware failures?
- How can we move from “pets to cattle’s”?



# “Pets vs Cattle” (Yes, again)



## Scale Up

- Servers are like pets.

Pets are given names, are unique, lovingly hand raised and cared for. When they get ill, you nurse them back to health



## Scale Out

- Servers are like cattle.

Cattle are given numbers and are almost identical to each other. When they get ill, you get another one.

“

“Future application architectures should use Cattle but Pets with strong configuration management are viable and still needed”

- Tim Bell, CERN

The above adapted from Tim Bell, CERN

<http://www.slideshare.net/noggin143/20121017-openstack-cern-accelerating-science>



# Project Theme 1: Load Balance API Gateway

- Thrift based multiplexed API Gateway
- Fault Tolerant
- Load Balanced
- Scalable
- Pluggable
- Facilitate Overlaid Features



# Timeline

- Next Friday Jan 20<sup>th</sup>
  - Present design choices
  - Discuss Hurdles
- Following Friday Jan 27<sup>th</sup>
  - Demonstrate Scalable, Reliable, Fault Tolerant Portal and API Server



# Thoughts?

Marlon Pierce, Suresh Marru  
[{marpierc, smarru}@iu.edu](mailto:{marpierc, smarru}@iu.edu)

