

Apache Thrift

Introduction & Tutorial

Marlon Pierce, Suresh Marru

CSCI-B 649 Science Gateway Architectures




INDIANA UNIVERSITY BLOOMINGTON
SCHOOL OF INFORMATICS AND COMPUTING

Q & A



INDIANA UNIVERSITY BLOOMINGTON
SCHOOL OF INFORMATICS AND COMPUTING

TIOBE Index



Programming Language “polyglotism”

- Modern distributed applications are rarely composed of modules written in a single language.
- Weaving together innovations made in a range of languages is a core competency of successful enterprises.
- Cross language communications are a necessity, not a luxury.
- In your projects you need to demonstrate this by using **three** or **more** languages.



Protocol Buffers

- “a language-neutral, platform-neutral, extensible way of serializing structured data for use in communications protocols, data storage, and more.”
- “Protocol buffers are a flexible, efficient, automated mechanism for serializing structured data – think XML, but smaller, faster, and simpler. ”
 - <https://developers.google.com/protocol-buffers/docs/overview>
- Started internally within Google in 2001 and Opened in 2008.



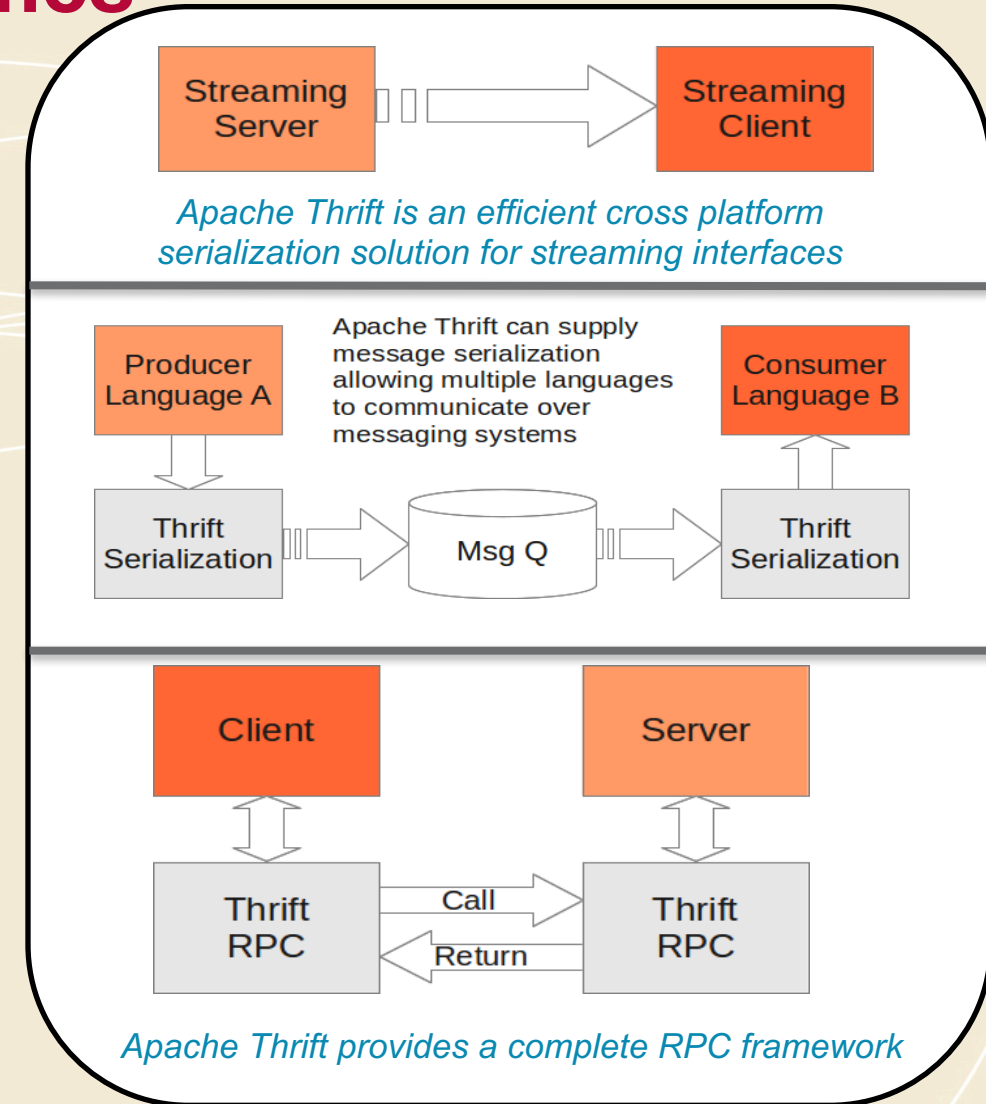
Apache Thrift

- Thrift is Facebook's implementation of Proto Buff open sourced under Apache.
- A high performance, scalable cross language serialization and RPC framework.
- Provides a full RPC Implementation with generated clients, servers, everything but the business logic.
- Thrift is fast and efficient, solutions for minimal parsing overhead and minimal size.



Thrift: Multiple Communication Schemes

- Streaming – Communications characterized by an ongoing flow of bytes from a server to one or more clients.
 - Example: An internet radio broadcast where the client receives bytes over time transmitted by the server in an ongoing sequence of small packets.
- Messaging – Message passing involves one way asynchronous, often queued, communications, producing loosely coupled systems.
 - Example: Sending an email message where you may get a response or you may not, and if you do get a response you don't know exactly when you will get it.
- RPC – Remote Procedure Call systems allow function calls to be made between processes on different computers.
 - Example: An iPhone app calling a service on the Internet which returns the weather forecast.



Source: Randy Abernethy. *The Programmer's Guide to Apache Thrift*, Manning Publications Co.



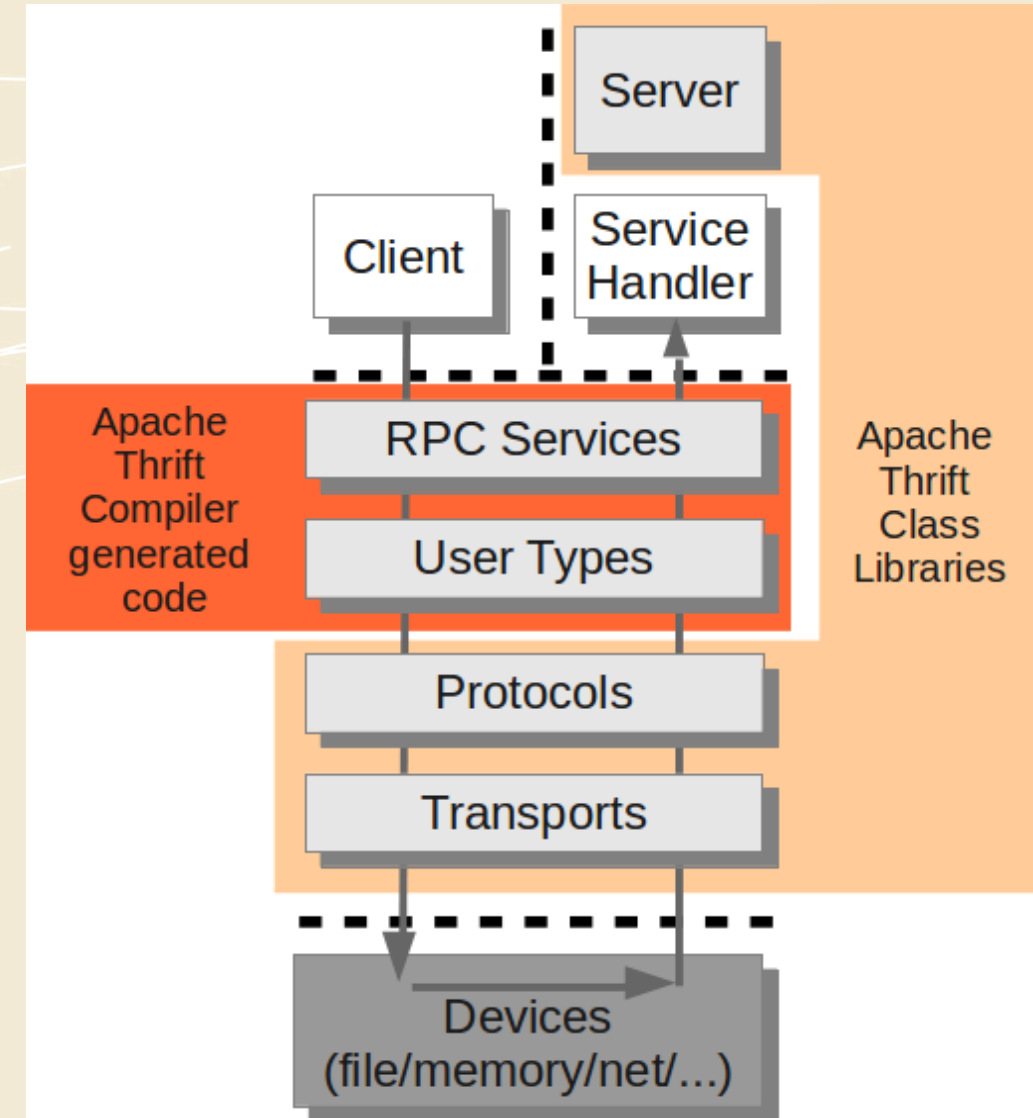
INDIANA UNIVERSITY BLOOMINGTON

SCHOOL OF INFORMATICS AND COMPUTING

CSCI-B 649 Science Gateway Architectures

Thrift for RPC Services

- User Code
 - client code calls RPC methods and/or [de]serializes objects
 - service handlers implement RPC service behavior
- Generated Code
 - RPC stubs supply client side proxies and server side processors
 - type serialization code provides serialization for IDL defined types
- Library Code
 - servers host user defined services, managing connections and concurrency
 - protocols perform serialization
 - transports move bytes from here to there



Source: Randy Abernethy. The Programmer's Guide to Apache Thrift, Manning Publications Co.



Thrift Resources

- Web
 - thrift.apache.org
 - github.com/apache/thrift
- Mail
 - Users: user-subscribe@thrift.apache.org
 - Developers: dev-subscribe@thrift.apache.org
- Chat
 - [#thrift](#)

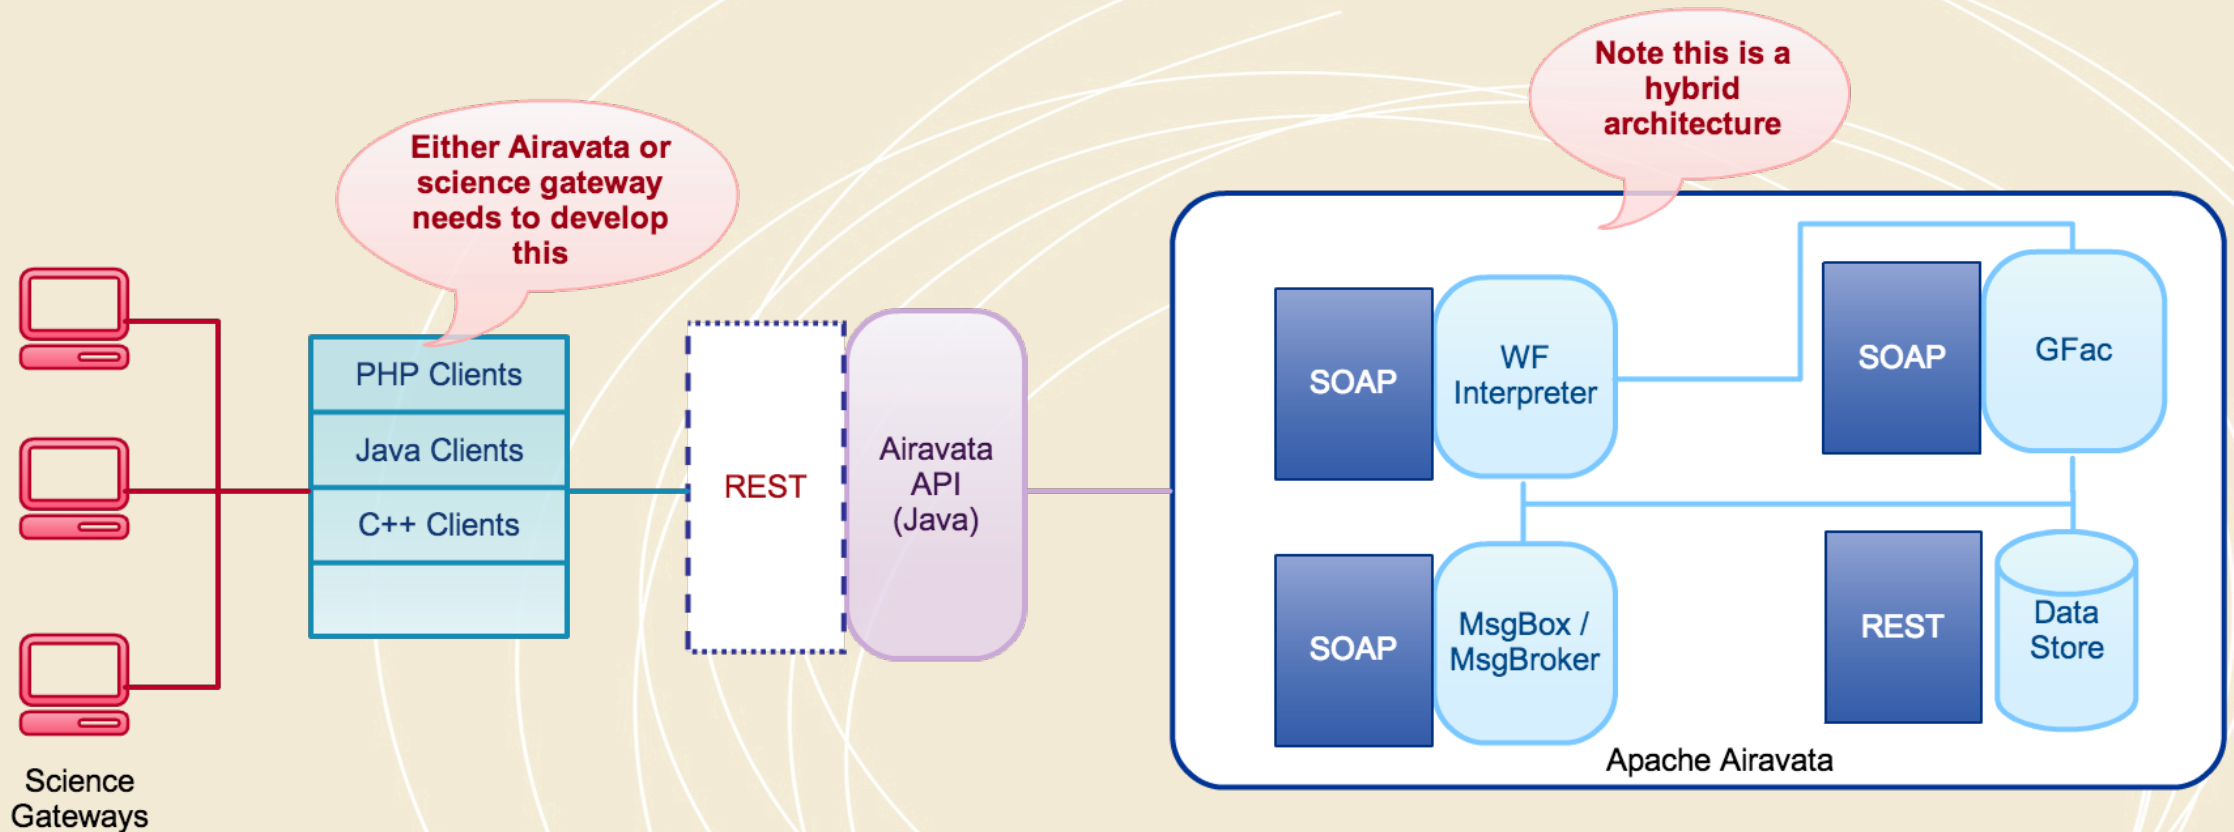


Thrift Experiences within Apache Airavata



INDIANA UNIVERSITY BLOOMINGTON
SCHOOL OF INFORMATICS AND COMPUTING

SOAP/REST based Airavata



(a simplified view - evolved over 10 years)



RPC using SOAP

- Pros
 - ability to separate out context and the payload
 - Already proven tools and techniques (XML, WSDL, WS-Security etc)
 - Clients can easily generate stubs using WSDLs
- Cons
 - Heavy weight
 - Data schema changes cannot be handled easily
 - Could not support broad range of clients



RPC using REST

- Pros
 - Flexible for data representation (JSON or XML).
 - Light weight.
 - Better performance.
- Cons
 - Multiple object layers.
 - No standard way to describe the service to the client.

Note: A detailed lecture on REST is forthcoming.



RPC using Thrift

- Integrating a language specific client library is easier than an open API like REST.
- Light framework
 - No multiple dependencies and servlet containers.
- Easy learning curve to get started
- If carefully crafted, the IDLs and framework support backward and forward compatibility

Note: These arguments will be similar for ProtoBuff and Avro



IDL's with Richer Data Structures

- Experiment data model is a **complex data model**
- Data structures : string, type-defs, integers, lists, sets
- Can refer other structs, enums

```
struct Experiment {  
  1: required string experimentID = DEFAULT_ID,  
  2: required string projectID = DEFAULT_PROJECT_NAME  
  3: optional i64 creationTime,  
  4: required string userName,  
  5: required string name,  
  6: optional string description,  
  7: optional string applicationId,  
  8: optional string applicationVersion,  
  9: optional string workflowTemplateId,  
  10: optional string workflowTemplateVersion,  
  11: optional UserConfigurationData userConfigurationData,  
  12: optional string workflowExecutionInstanceId,  
  13: optional list<DataObjectType> experimentInputs,  
  14: optional list<DataObjectType> experimentOutputs,  
  15: optional ExperimentStatus experimentStatus,  
  16: optional list<WorkflowNodeStatus> stateChangeList,  
  17: optional list<WorkflowNodeDetails> workflowNodeDetailsList,  
  18: optional list<ErrorDetails> errors  
}
```

Defining a struct

- Can send such **complex data model** over the wire
- Also note that Thrift support **exception handling** too.

```
string createExperiment(1: required experimentModel.Experiment experiment)  
  throws (1: airavataErrors.InvalidRequestException ire,  
         2: airavataErrors.AiravataClientException ace,  
         3: airavataErrors.AiravataSystemException ase)  
/**
```

Defining a service method



Thrift Experiences Summary

- Auto generated PolyGot clients.
- Clean design
- Multiple servers to choose from:
 - TSimpleServer - Simple single threaded server
 - TThreadPoolServer - Uses Java's built in ThreadPool management
 - TNonblockingServer - non-blocking TServer implementation
 - THsHaServer - extension of the TNonblockingServer to a Half-Sync/Half-Async server



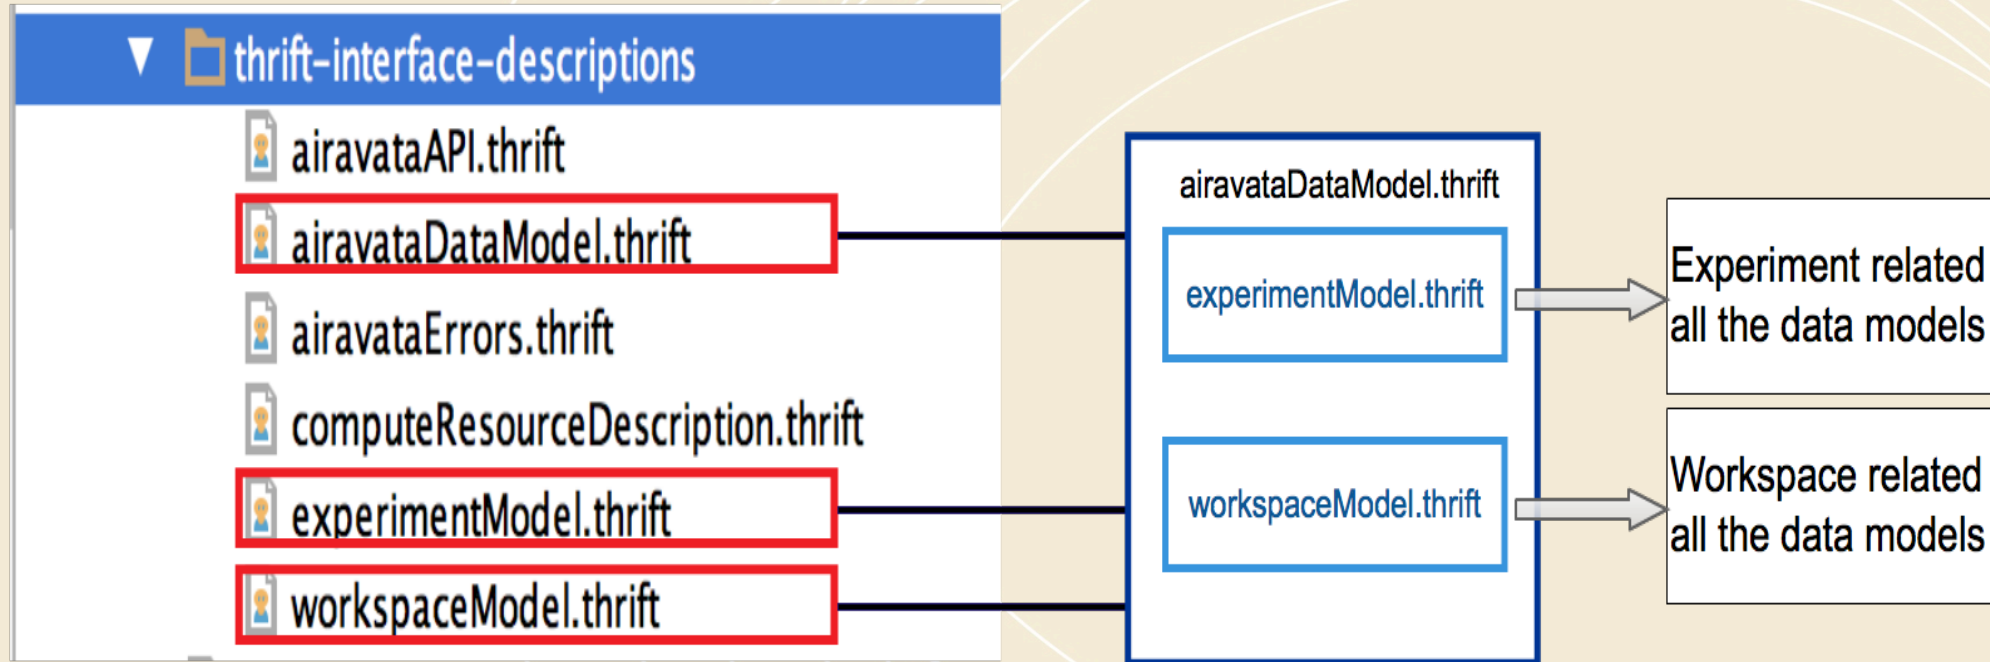
Experiences Contd..

- No explicit need for marshalling / unmarshalling.
- Data models can also be used internally.
- Servers are very robust - able to handle large number of concurrent requests.
- Easy to do modifications to the models, since code is auto-generated.
- Convenient way to achieve backward compatibility.



Lessons Learned

Modularize the data models in order to ease the maintenance



Lessons Learned Contd..

```
/**
 * A structure holding the experiment metadata and its child models.
 *
 * *
 * * userName:
 * * The user name of the targeted gateway end user on whose behalf the experiment is being created.
 * * the associated gateway identity can only be inferred from the security hand-shake so as to avoid
 * * authorized Airavata Clients mimicking an unauthorized request. If a gateway is not registered with
 * * Airavata, an authorization exception is thrown.
 * *
 * * experimentName:
 * * The name of the experiment as defined by the user. The name need not be unique as uniqueness is enforced
 * * by the generated experiment id.
 * *
 * * experimentDescription:
 * * The verbose description of the experiment. This is an optional parameter.
 */
```

```
struct Experiment {
1: required string experimentID = DEFAULT_ID,
2: required string projectID = DEFAULT_PROJECT_NAME
3: optional i64 creationTime,
4: required string userName,
5: required string name,
6: optional string description,
7: optional string applicationId,
8: optional string applicationVersion,
9: optional string workflowTemplateId,
10: optional string workflowTemplateVersion,
11: optional UserConfigurationData userConfigurationData,
12: optional string workflowExecutionInstanceId,
13: optional list<DataObjectType> experimentInputs,
14: optional list<DataObjectType> experimentOutputs,
15: optional ExperimentStatus experimentStatus,
16: optional list<WorkflowNodeStatus> stateChangeList,
17: optional list<WorkflowNodeDetails> workflowNodeDetailsList,
18: optional list<ErrorDetails> errors
}
```

Use "set" if you want to avoid duplicates

Add proper documentation to structs and services

```
/**
 * Terminate a running experiment.
 *
 * *
 * * @param airavataExperimentId
 * * The identifier for the requested experiment. This is returned during the create experiment step.
 * *
 * * @return
 * * This method call does not have a return value.
 * *
 * * @throws org.apache.airavata.api.error.InvalidRequestException
 * * For any incorrect forming of the request itself.
 * *
 * * @throws org.apache.airavata.api.error.ExperimentNotFoundException
 * * If the specified experiment is not previously created, then an Experiment Not Found Exception is thrown.
 * *
 * * @throws org.apache.airavata.api.error.AiravataClientException
 * * The following list of exceptions are thrown which Airavata Client can take corrective actions to resolve:
 * *
 * * UNKNOWN_GATEWAY_ID - If a Gateway is not registered with Airavata as a one time administrative
 * * step, then Airavata Registry will not have a provenance area setup. The client has to follow
 * * gateway registration steps and retry this request.
 * *
 * * AUTHENTICATION_FAILURE - How Authentication will be implemented is yet to be determined.
 * * For now this is a place holder.
 * *
 * * INVALID_AUTHORIZATION - This will throw an authorization exception. When a more robust security hand-shake
 * * is implemented, the authorization will be more substantial.
 * *
 * * @throws org.apache.airavata.api.error.AiravataSystemException
 * * This exception will be thrown for any Airavata Server side issues and if the problem cannot be corrected by the
 * * rather an Airavata Administrator will be notified to take corrective action.
 */
void terminateExperiment(1: string airavataExperimentId)
throws (1: airavataErrors.InvalidRequestException ire,
2: airavataErrors.ExperimentNotFoundException enf,
3: airavataErrors.AiravataClientException ace,
4: airavataErrors.AiravataSystemException ase)
}
```



Lessons Learned Contd..

- Thrift has limited support to handle null values.
 - Models can be complex with lot of other structs inside.
 - Enums cannot be passed as null over the wire even they are specified as optional in the struct
- Thrift has limited documentation and samples
 - airavata can be used as a reference
 - <https://github.com/apache/airavata>



Apache Thrift Tutorial



INDIANA UNIVERSITY BLOOMINGTON
SCHOOL OF INFORMATICS AND COMPUTING

4 Simple Steps to Create a RPC microservice

1. Define the service in a language neutral “Interface Description Language”.
2. Compile the IDL to generate Server and Client “stubs” in desired programming languages.
3. Plug the server implementation in the pre-generated server stub.
4. Call the remote services as if they are making local method calls.



Hands-on Tutorial

<https://github.com/airavata-courses/tutorials>

<https://github.com/apache/airavata-sandbox/tree/master/airavata-mock-multiplexed-api/>



Thank You!

Marlon Pierce, Suresh Marru
{marpierc, smarru}@iu.edu

