Cloud Facility for Scientific Gateways

J. Mike Lowe, Lead Systems Programmer
Pervasive Technologies Institute, UITS/RT, Indiana University
George Turner, Chief Systems Architect
Pervasive Technologies Institute, UITS/RT, Indiana University

CSCI-B 649, Topics in Systems
Science Gateway Architectures,
School of Informatics and Computing, Indiana University 24-Oct-2017

funded by the National Science Foundation
Award #ACI-1445604

Terms & Definitions

Jetstream
http://jetstream-cloud.org/

NSF funded by the National Science Foundation
Award #ACI-1445604

What is Jetstream

funded by the National Science Foundation
Award #ACI-1445604

What is Jetstream?

- **User-friendly**, widely accessible cloud environment
  - **User-selectable library** of preconfigured virtual machines
    - Interactive computing
    - Software maintained by domain specialist
    - No need for system administration skills
    - The “Atmosphere” side
  - **Programmable cyberinfrastructure**
    - Go beyond batch computing
    - Implement modern cloud computing techniques
    - Common modality for science gateways
    - The “API” side

Jetstream
http://jetstream-cloud.org/

NSF funded by the National Science Foundation
Award #ACI-1445604

“Long tail” of the Science

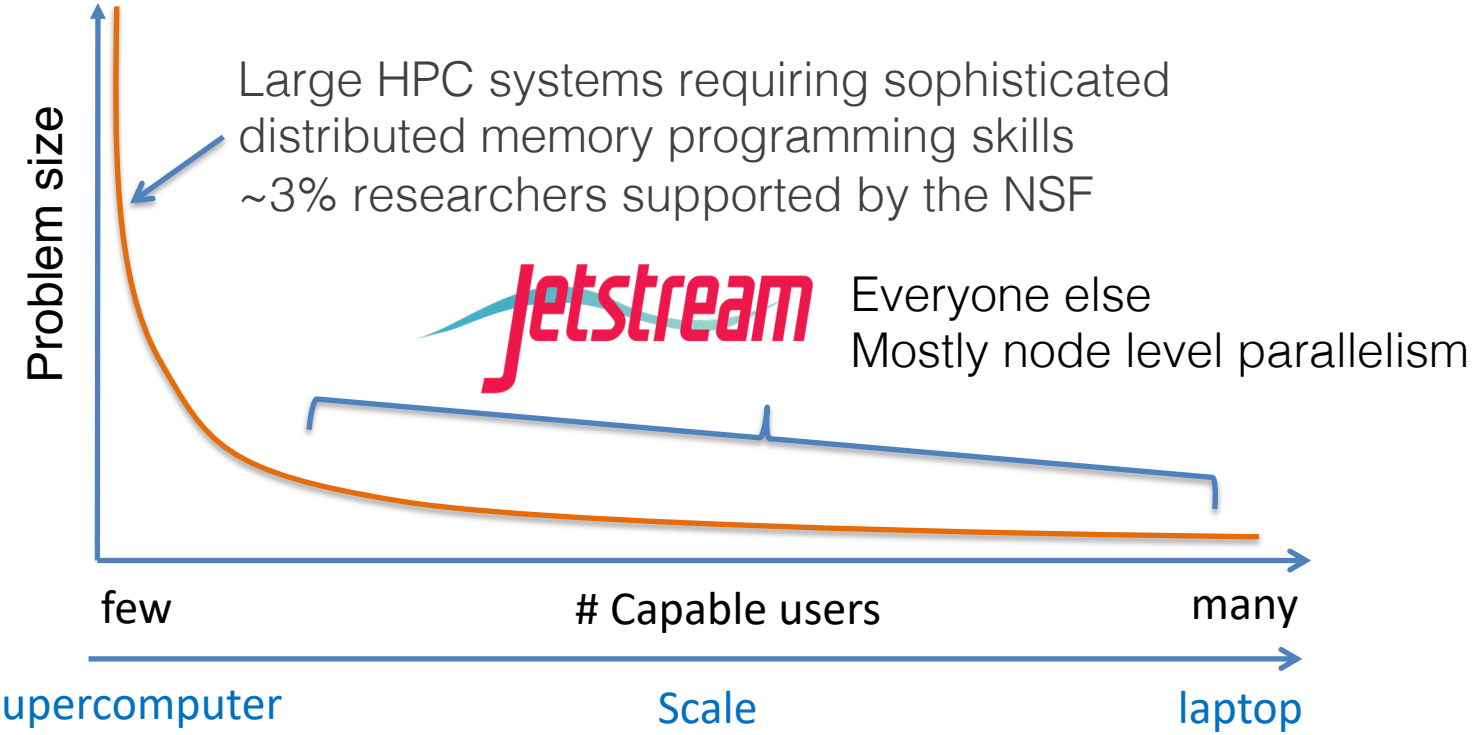

Jetstream
http://jetstream-cloud.org/

NSF
funded by the National Science Foundation
Award #ACI-1445604

What is Jetstream?

- **Primary goal** is to **expand the user base** of NSF's e**X**treme **D**igital (**XD**) program resources beyond the current community of users.
- **Lowering the hurdle** to onboard to XSEDE resources
  - Working to **ease** the **allocation request** process
  - **Easy-Button**; quick access but limited ability. (Beta)
- **Making Science Easy** for domain researchers, engineers, & educators
  - Domain software installed & maintained by the professionals
  - No sys-admin skills necessary

Jetstream
http://jetstream-cloud.org/

NSF funded by the National Science Foundation
Award #ACI-1445604

What is Jetstream? (Cont.)

- **Creating communities**
  - Domain developers **create, install, and maintain** the software
  - **Encourage collaboration** within the domains
  - **Operating system** level software is professionally **patched and maintained**
- **Repeatibility**: store & publish images via IU Scholarworks & create a DOI
- **Science Gateways**:

NSF funded by the National Science Foundation
Award #ACI-1445604

What is Jetstream?

- **Cloudy Technologies**: clouds are more the just virtual machines (VM)
  - **Old way**: robust (expensive) infrastructure, weak (cheap) software
    - You expect the hardware to not fail
    - State in maintained in volatile data structures
  - **Cloudy way**: commodity infrastructure, robust software
    - Expect & plan for infrastructure to fail
    - Put intelligence into the software to handle infrastructure failure
  - **Cows, not pets**:
    - pets have **state**, you name them, you get attached to them, you put forth great amount of care and effort
    - cows **do not** have **state**, you expect to have high turnover, you do not get attached to them, you give them numbers instead of names

Jetstream
http://jetstream-cloud.org/

NSF funded by the National Science Foundation
Award #ACI-1445604

What is Jetstream?

- **Software layers**
  - **Atmosphere** web interface
    - library of images, genertic, domain specific
    - simplify VM administration
  - **OpenStack:** software tools for building and managing cloud computing platforms for public and private clouds.
  - **KVM** hypervisor: what the VMs run on
  - **Ceph**: storage platform that stores data on a single distributed computer cluster, and provides interfaces for **object**-, **block**- and *file-level* storage.
  - **Operating systems**: CentOS, Ubuntu, Windows?
  - **Applications**; e.g. software developed by the domain specialist, gateways, etc.

Jetstream
http://jetstream-cloud.org/

NSF funded by the National Science Foundation
Award #ACI-1445604

API Access to Jetstream

- What was **unexpected**
  - **Demand** for **programmable cyberinfrastructure**
  - Great platform for learning **system administration skills**
  - Great platform for **teaching & learning cloudy technologies**
- **Command line clients**
- **Horizon dashboard** very popular; but, incomplete
- **Programmatic control**; python is popular
- **Slack channel** for collaboration API users of Jetstream

Jetstream
http://jetstream-cloud.org/

funded by the National Science Foundation
Award #ACI-1445604

OpenStack

funded by the National Science Foundation
Award #ACI-1445604

OpenStack Organization

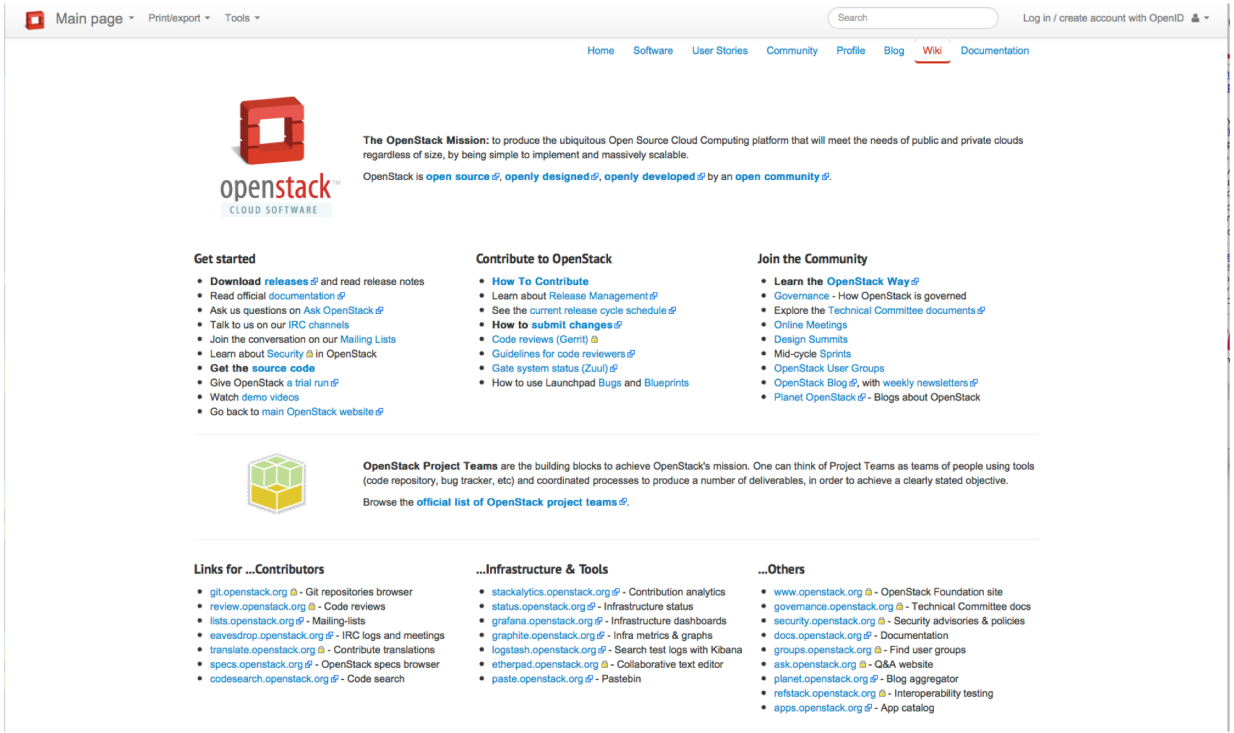

Jetstream
http://jetstream-cloud.org/

NSF
funded by the National Science Foundation
Award #ACI-1445604

OpenStack : the Project Navigator

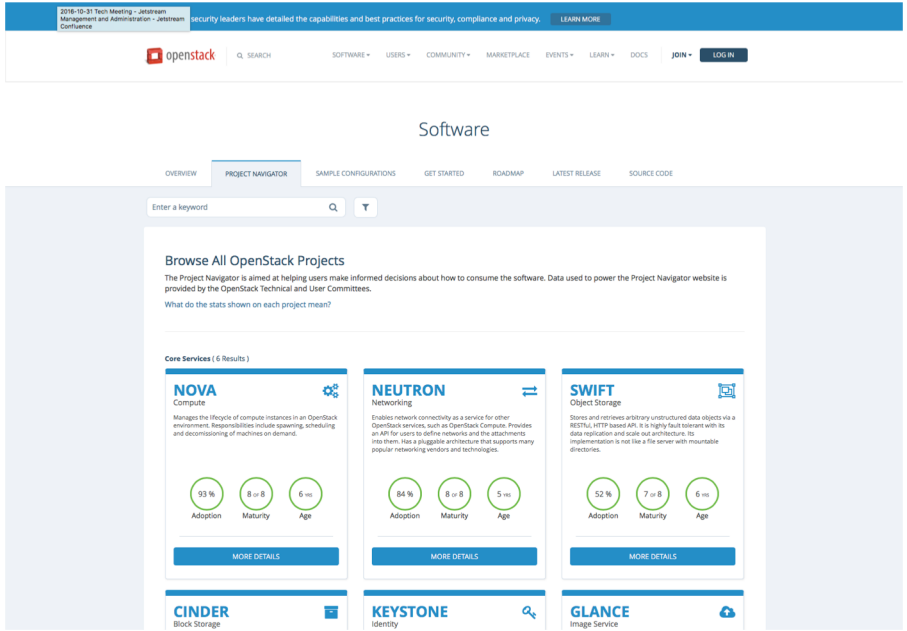

http://www.openstack.org/software/project-navigator/

Jetstream
http://jetstream-cloud.org/

funded by the National Science Foundation
Award #ACI-1445604

Openstack Projects …the core services

Service	Name	Adoption	Maturity	Age
Identity	Keystone	96%	7/8	5 yrs
Images	Glance	95%	6/8	7 yrs
Block device	Cinder	88%	7/8	5 yrs
Networking	Neutron	93%	7/8	5 yrs
Compute	Nova	95%	8/8	7 yrs
Object device	Swift	52%	7/8	7 yrs

https://www.openstack.org/software/project-navigator/

Jetstream
http://jetstream-cloud.org/

funded by the National Science Foundation
Award #ACI-1445604

Openstack Projects …some other services

Service	Name	Adoption	Maturity	Age
Dashboard	Horizon	87%	6/8	5 yrs
Telemetry	Ceilometer	55%	1/8	4 yrs
Orchestration	Heat	67%	6/8	4 yrs
Containers	Magnum	11%	2/8	2 yrs
Map/Reduce	Sahara	10%	3/8	3 yrs

https://www.openstack.org/software/project-navigator/

NSF funded by the National Science Foundation
Award #ACI-1445604

Openstack Projects …some other services

Service	Name	Adoption	Maturity	Age
Shared Filesystems	Manila	14%	5/8	3 yrs
Workflow	Mistral	5%	1/7	1 yr
Load Balancing as a Service	Octavia	>0%	1/7	1 yr

https://www.openstack.org/software/project-navigator/

Jetstream
http://jetstream-cloud.org/

NSF funded by the National Science Foundation
Award #ACI-1445604

Hardware & Infrastructure

Jetstream
http://jetstream-cloud.org/

funded by the National Science Foundation
Award #ACI-1445604

Production Cloud Hardware (per site)

Hardware	Number	Specifications	Function (IU)
Dell PowerEdge M630 blades	320	2X Intel E5-2680v3 "Haswell" 24 cores @ 2.5 GHz 128 GB RAM 2 TB local disk	Compute hosts OpenStack services
Dell PowerEdge R630 1U server	7	2X Intel E5-2680v3 "Haswell" 24 cores @ 2.5 GHz 128 GB RAM 2 TB local disk	Cluster management, High Availability, Databases, RabbitMQ
Dell PowerEdge R730xd 2U servers	20	2X Intel E5-2680v3 "Haswell" 24 cores @ 2.5 GHz 64 GB RAM 48 TB storage for Ceph pool	~1 PB Ceph storage
Dell S6000-ON network switches	9	32+2 40 Gb/s ports	Top of Rack & Spine 2 to 1 Fat Tree topology

Jetstream
http://jetstream-cloud.org/

funded by the National Science Foundation
Award #ACI-1445604

Jetstream Production Hardware

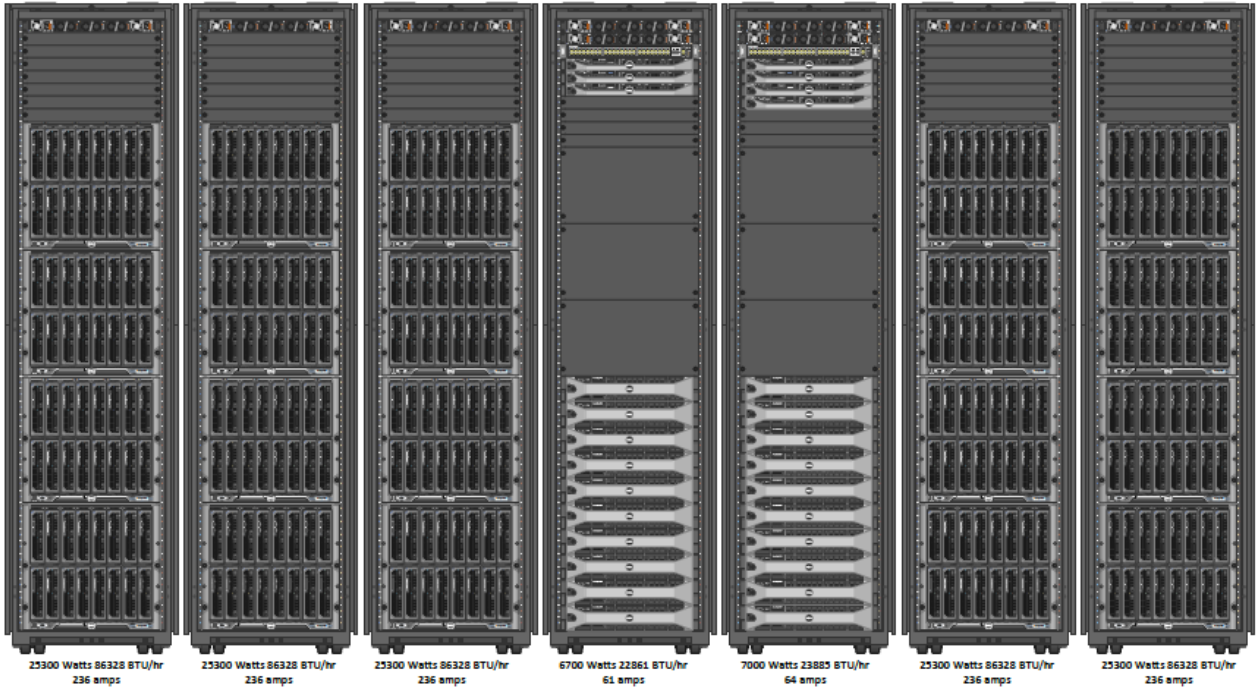

Jetstream
http://jetstream-cloud.org/

NSF
funded by the National Science Foundation
Award #ACI-1445604

Jetstream's Cooling Doors

ChilledDoor Rack Cooling System by motivair

½ the heat generated via Jetstream is removed via the cooling doors and ½ is removed via traditional CRACs

Colour related to health of the cooling system

Jetstream
http://jetstream-cloud.org/

NSF funded by the National Science Foundation
Award #ACI-1445604

Infrared image of Jetstream

Note yellow doors on left side: this is hot air blowing from the Wrangler Storage cluster onto the cooling doors

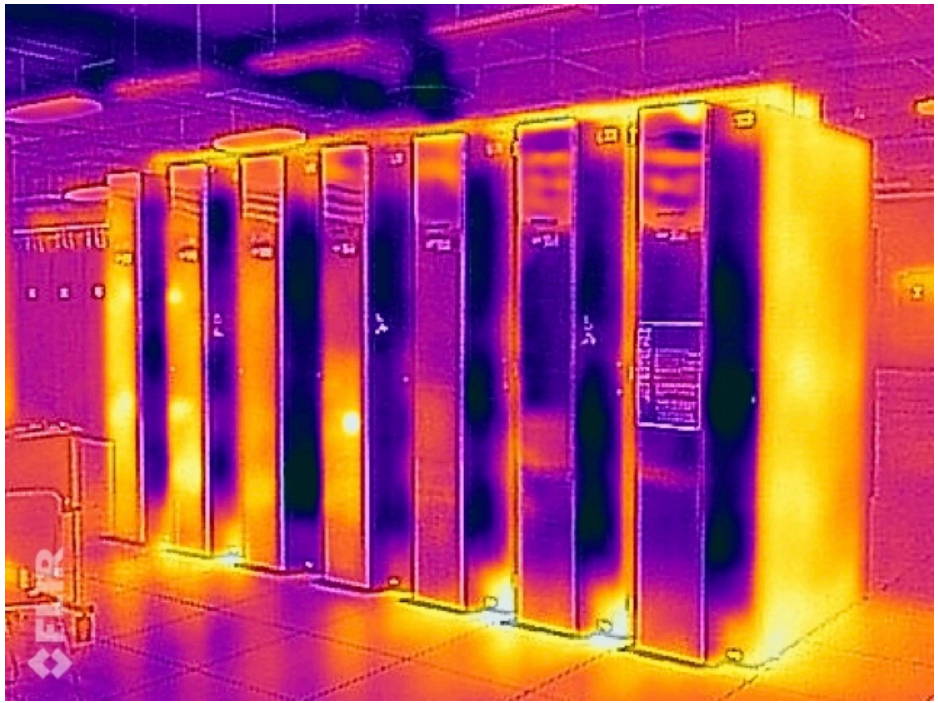

ChilledDoor Rack Cooling System by motivair

Yellow: hot
Blue: cold

Jetstream
http://jetstream-cloud.org/

funded by the National Science Foundation
Award #ACI-1445604

Benchmarks – single node

Jetstream
http://jetstream-cloud.org/

funded by the National Science Foundation
Award #ACI-1445604

HPCC results : VM vs BareMetal Comparison

Benchmark	VM/ BareMetal	Units	What's tested
HPL	97%	FLOPS	floating point execution rate for solving a system of linear equations
DGEMM	98%	FLOPS	floating point execution rate for double precision real matrix-matrix multiplication
Bandwidth	88%	B/s	bytes/unit_time it takes to transmit a 2MB message from one node to another
Latency	97%	s	time required to send an 8-byte message from one node to another

Jetstream
http://jetstream-cloud.org/

NSF funded by the National Science Foundation
Award #ACI-1445604

HPCC results : VM vs BareMetal Comparison (Cont.)

Benchmark	VM/ BareMetal	Units	What's tested
Random	80%	up/s	rate of random updates of memory
Stream	77%	B/s	sustained memory bandwidth
MPI-FFT	67%	FLOPS	floating point rate of execution of double precision complex one-dimensional Discrete Fourier Tranform
Ptrans	64%	B/s	rate of transfer for large arrays of data from multiprocessor's memory

Jetstream
http://jetstream-cloud.org/

NSF funded by the National Science Foundation
Award #ACI-1445604

Jetstream System Details

funded by the National Science Foundation
Award #ACI-1445604

VM Instance Sizes (Flavors)

Instance Type	vCPUs	RAM(GB)	Storage(GB)	Instances/Node
Tiny	1	2	8	46
Small	2	4	20	23
Medium	6	16	60	7
Large	10	30	120/60*	4
X-Large	22	60	240/60*	2
XX-Large	44	120	480/60*	1

Node config: dual Intel E-2680v3 “Haswell”, 24 physical cores/node @ 2.5 GHz, 128 GB RAM, dual 1 TB local disks.

* Effective 29-Mar-2017

Jetstream
http://jetstream-cloud.org/

funded by the National Science Foundation
Award #ACI-1445604

Jetstream Systems

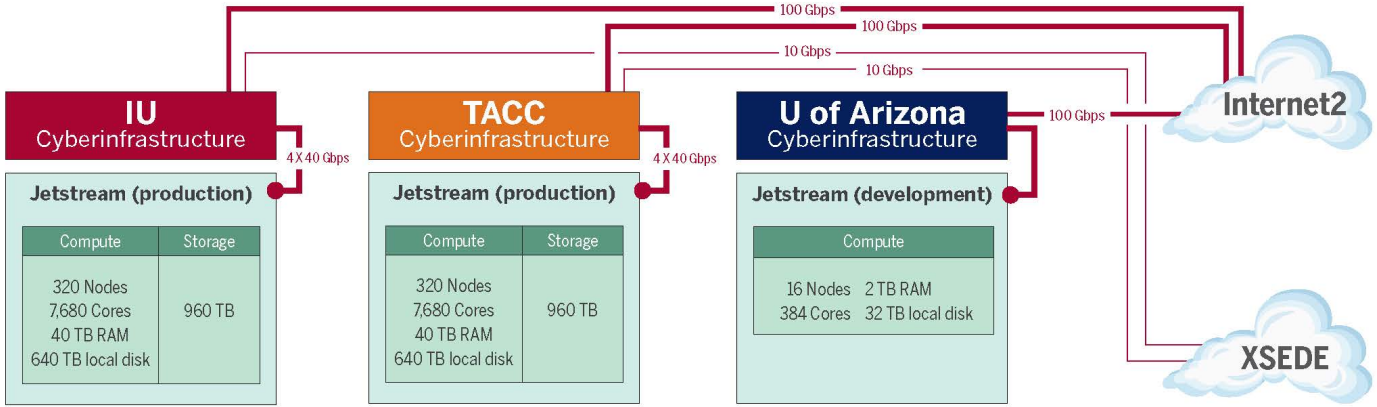

Jetstream
http://jetstream-cloud.org/

NSF
funded by the National Science Foundation
Award #ACI-1445604

Jetstream Systems (cont.)

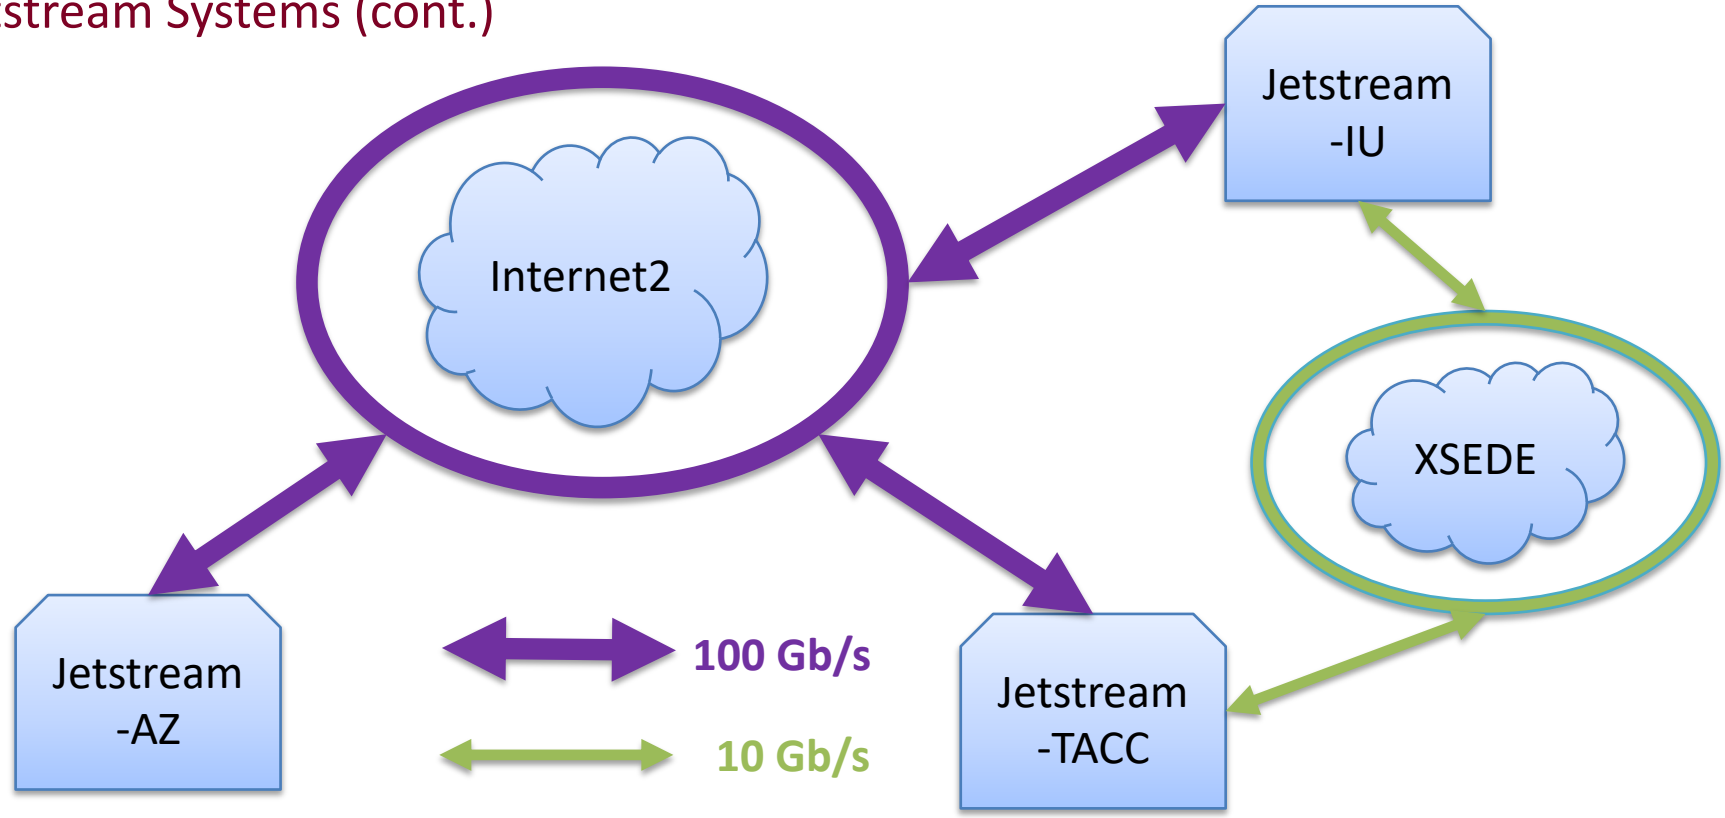

Jetstream
http://jetstream-cloud.org/

funded by the National Science Foundation
Award #ACI-1445604

Platform Overview

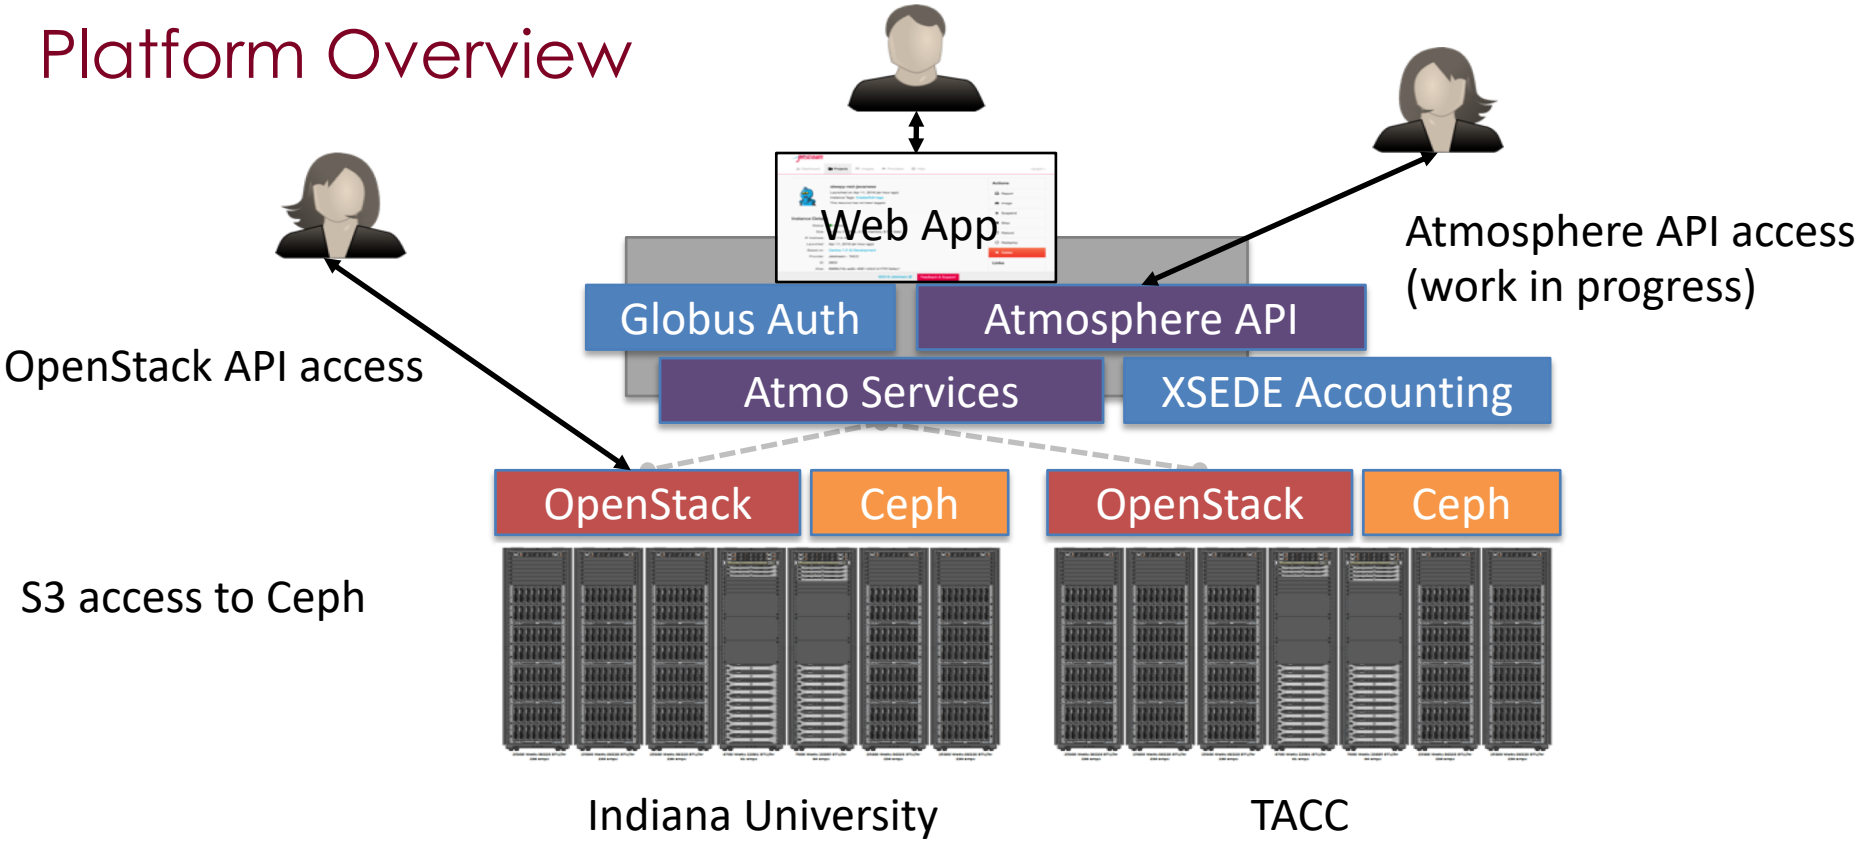

Jetstream
http://jetstream-cloud.org/

NSF
funded by the National Science Foundation
Award #ACI-1445604

Platform Overview

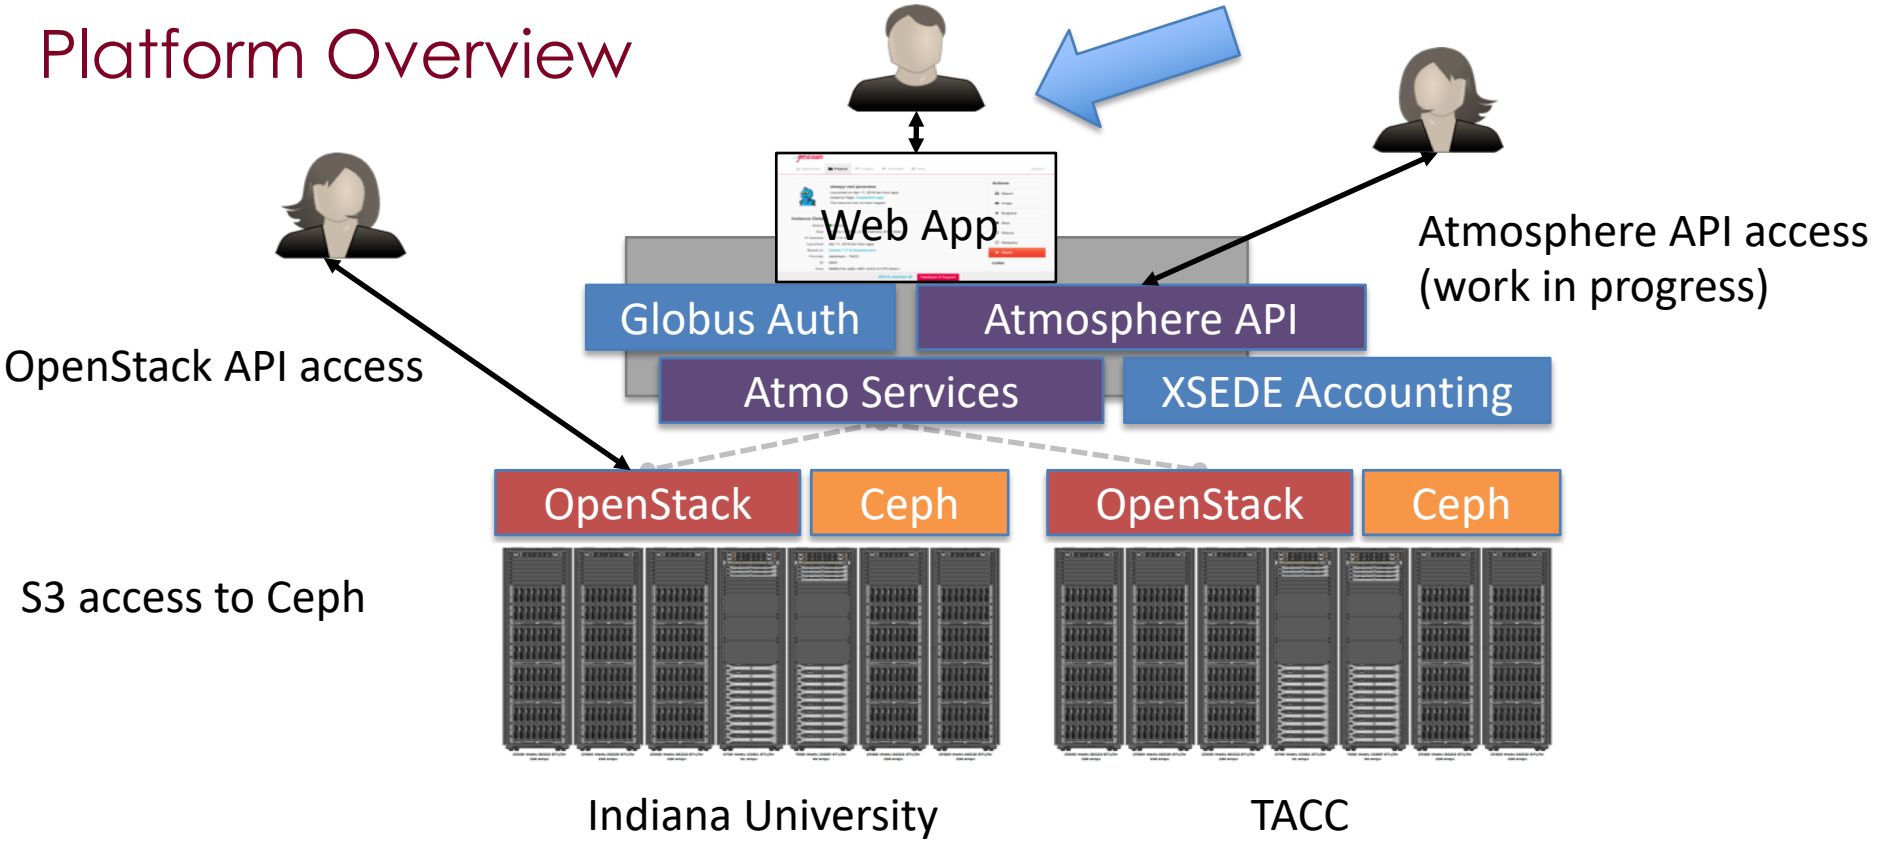

Jetstream
http://jetstream-cloud.org/

funded by the National Science Foundation
Award #ACI-1445604

Platform Overview

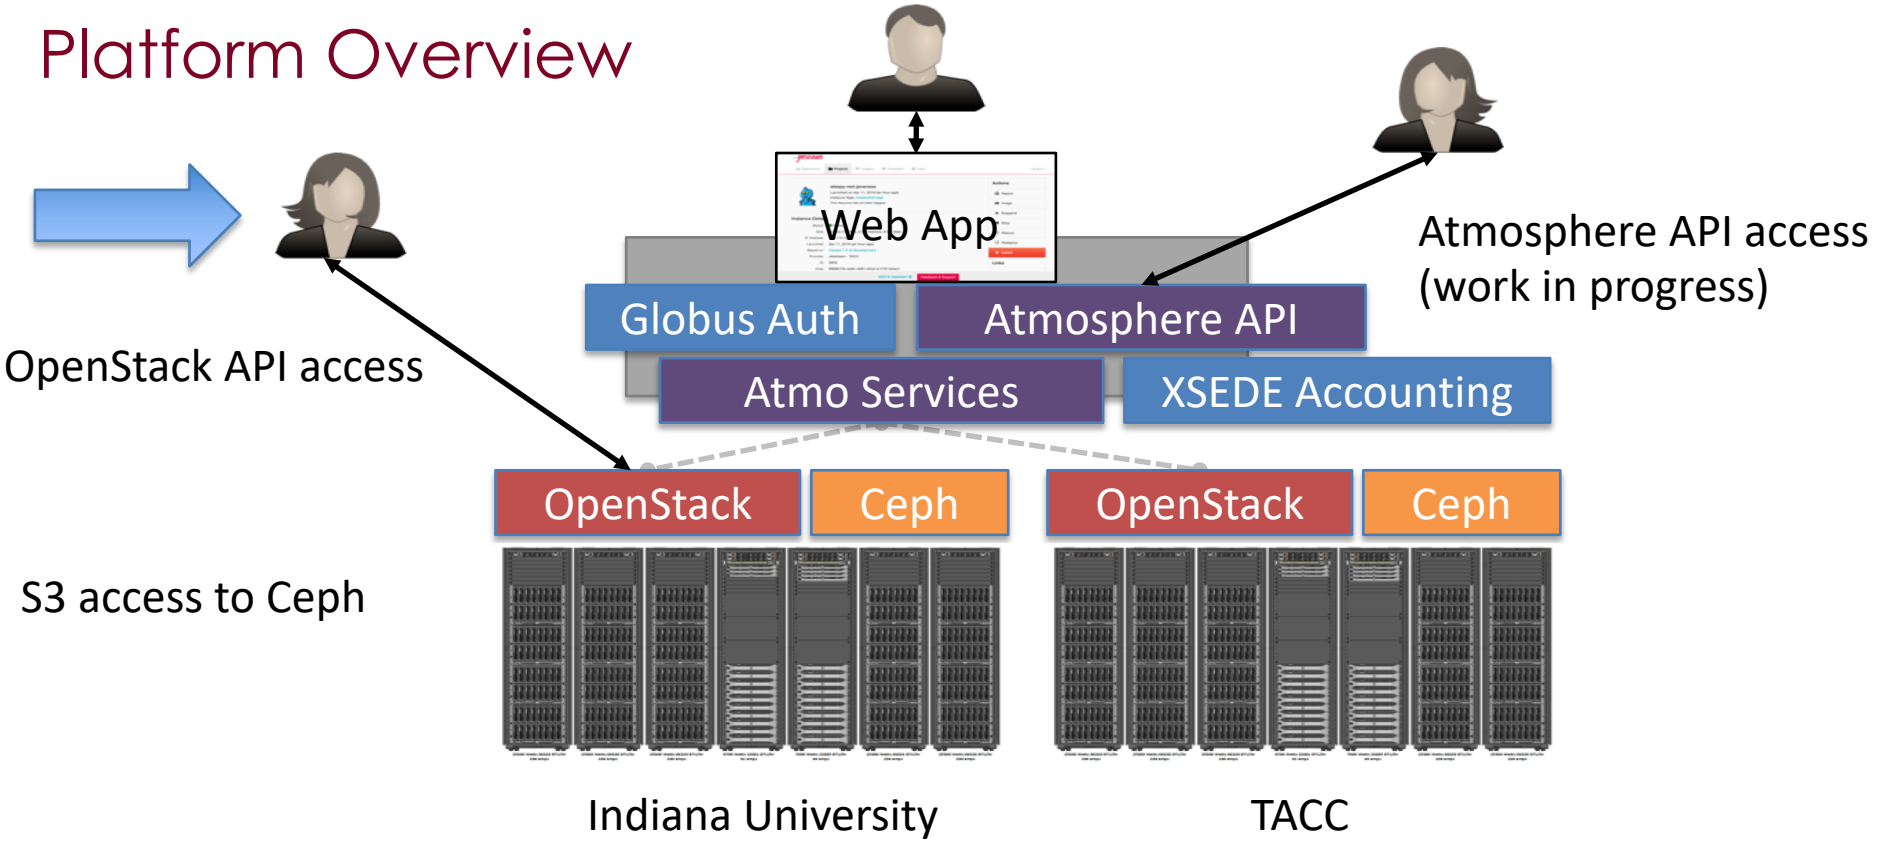

Jetstream
http://jetstream-cloud.org/

funded by the National Science Foundation
Award #ACI-1445604

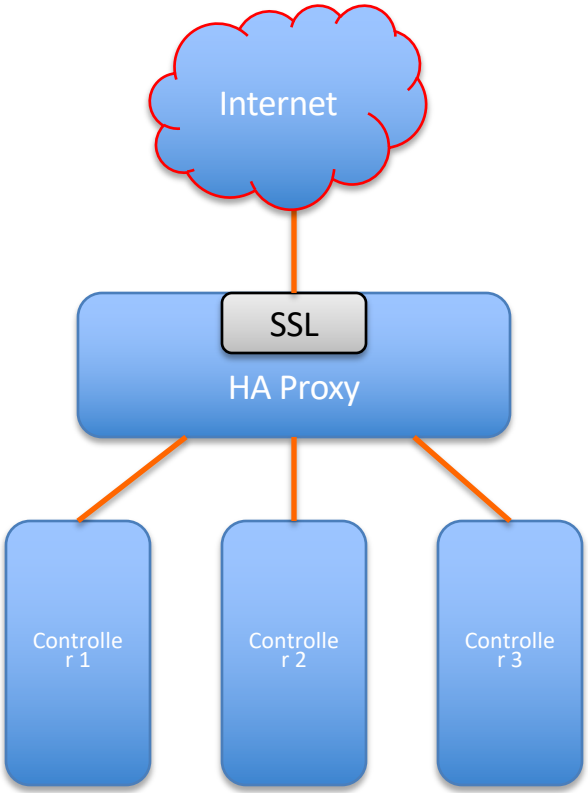

Jetstream
http://jetstream-cloud.org/

NSF
funded by the National Science Foundation
Award #ACI-1445604

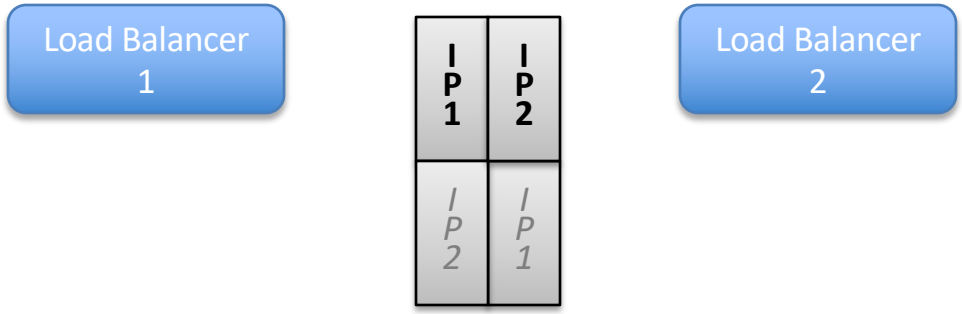

Jetstream
http://jetstream-cloud.org/
funded by the National Science Foundation
Award #ACI-1445604

High Availability layout for the databases

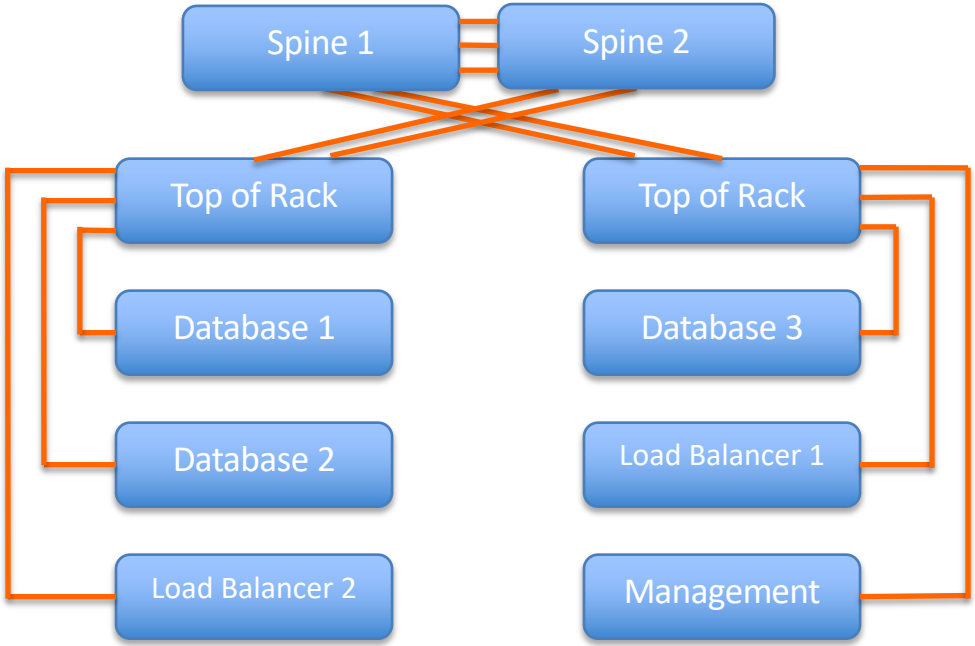

Jetstream
http://jetstream-cloud.org/

funded by the National Science Foundation
Award #ACI-1445604

Glance - Cinder - Ceph

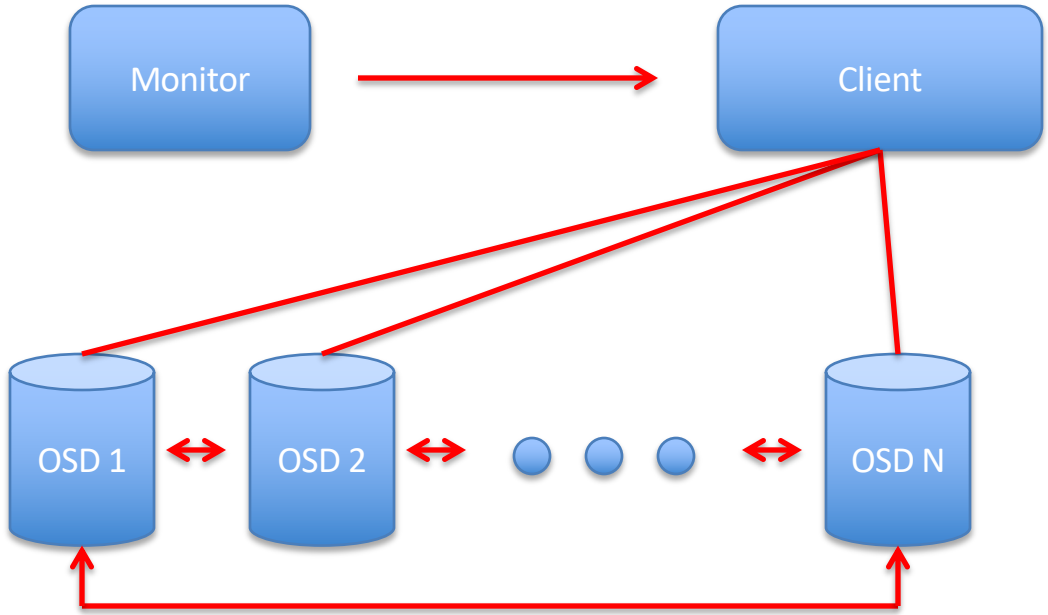

Jetstream
http://jetstream-cloud.org/

NSF
funded by the National Science Foundation
Award #ACI-1445604

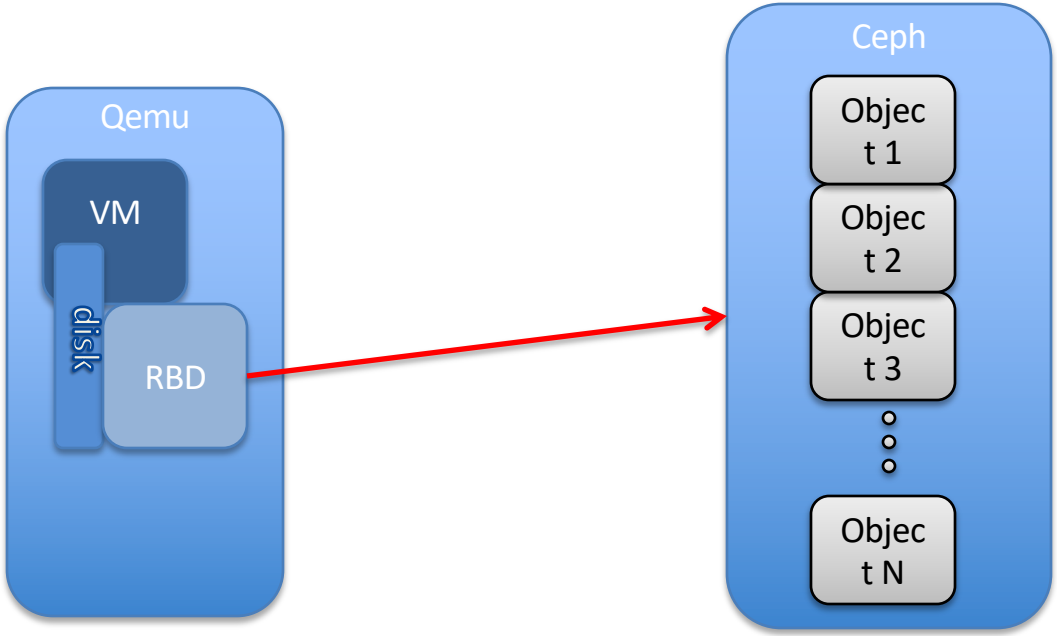

Jetstream
http://jetstream-cloud.org/

NSF
funded by the National Science Foundation
Award #ACI-1445604

OpenStack Overview

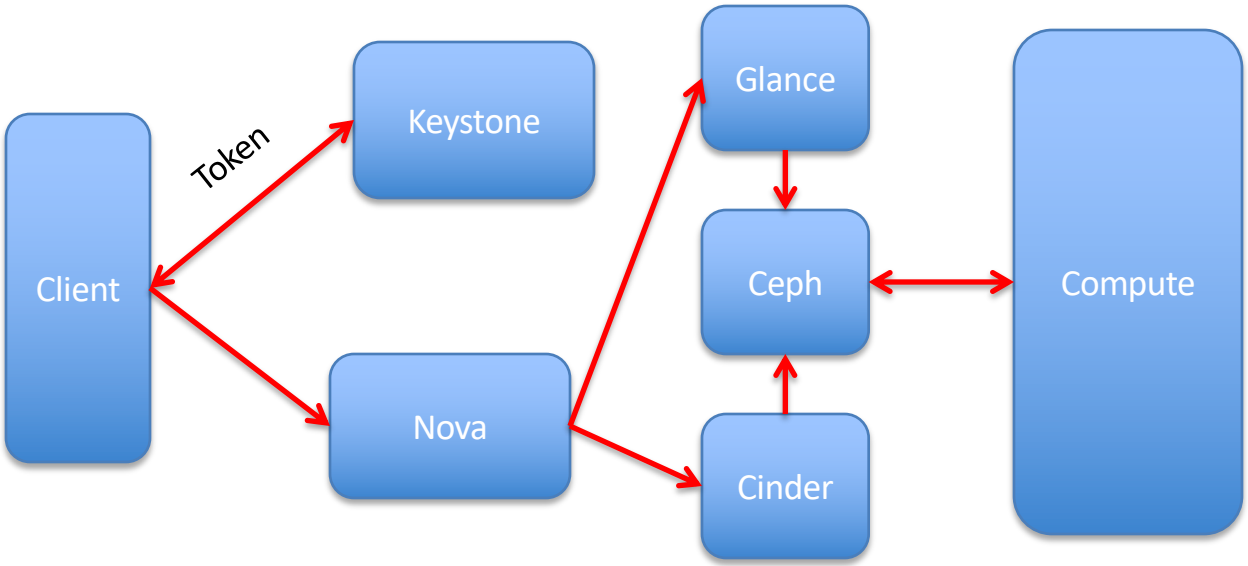

Jetstream
http://jetstream-cloud.org/

NSF
funded by the National Science Foundation
Award #ACI-1445604

VXLAN Packet

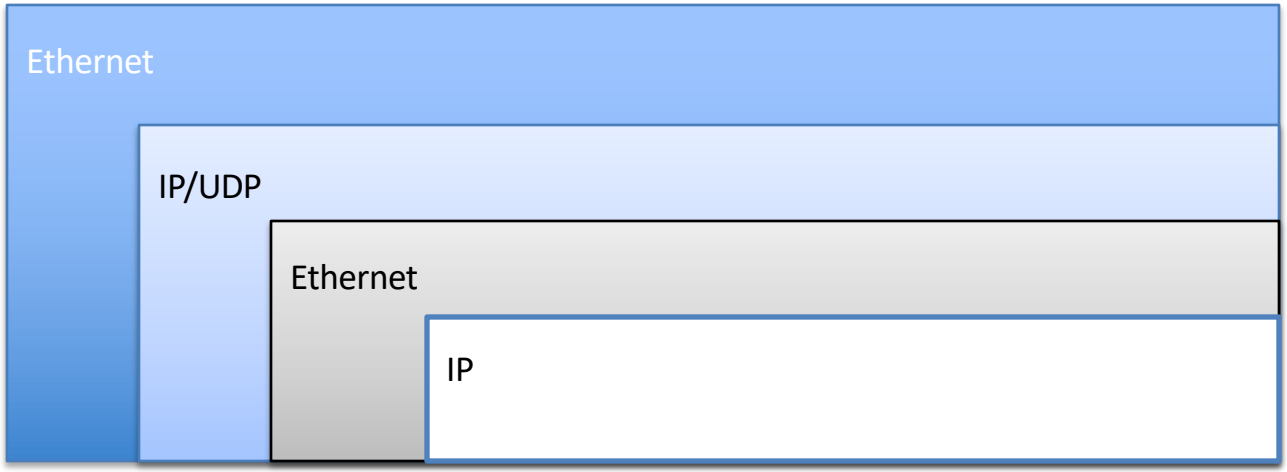

Jetstream
http://jetstream-cloud.org/

funded by the National Science Foundation
Award #ACI-1445604

Neutron Networking

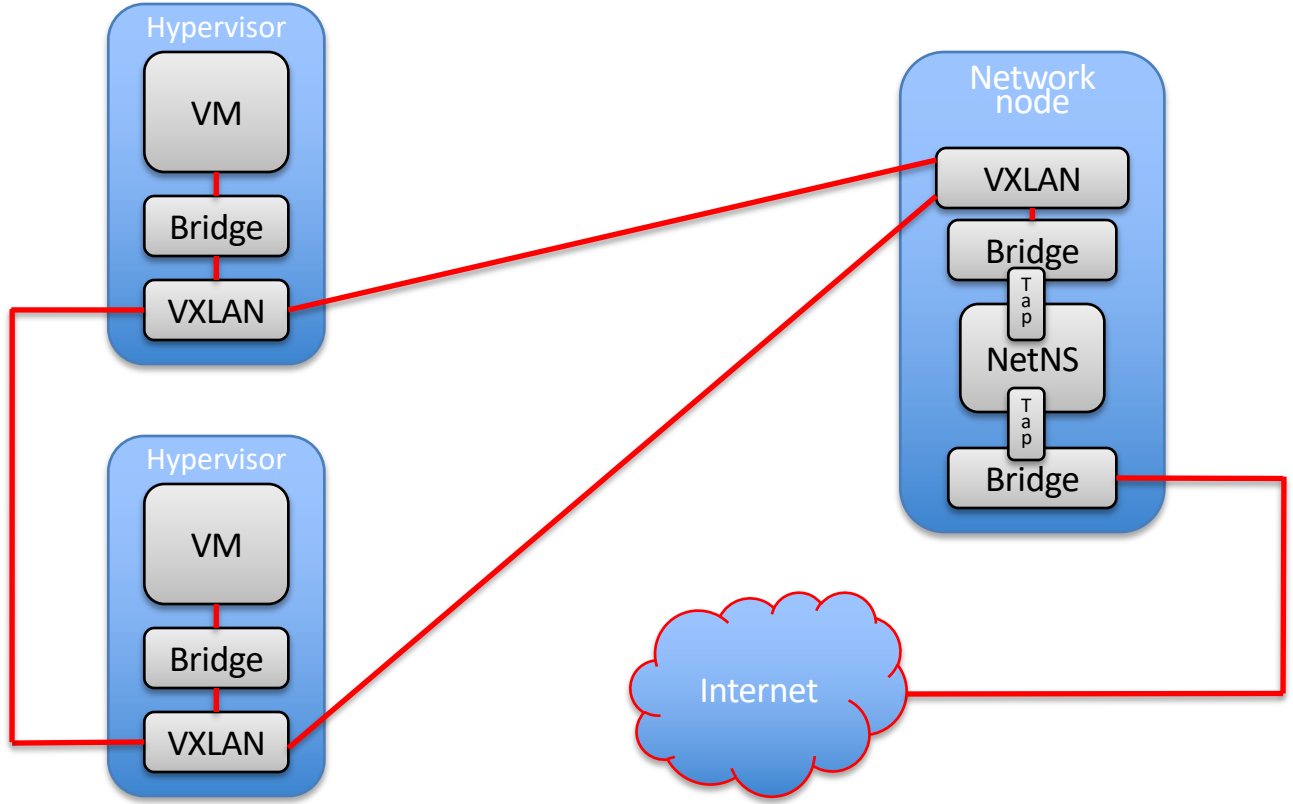

Jetstream
http://jetstream-cloud.org/

funded by the National Science Foundation
Award #ACI-1445604

Network Topology (cont.)

Sixteen blades per chasses
Two switches per chassis

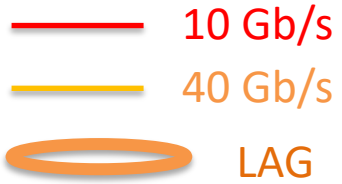

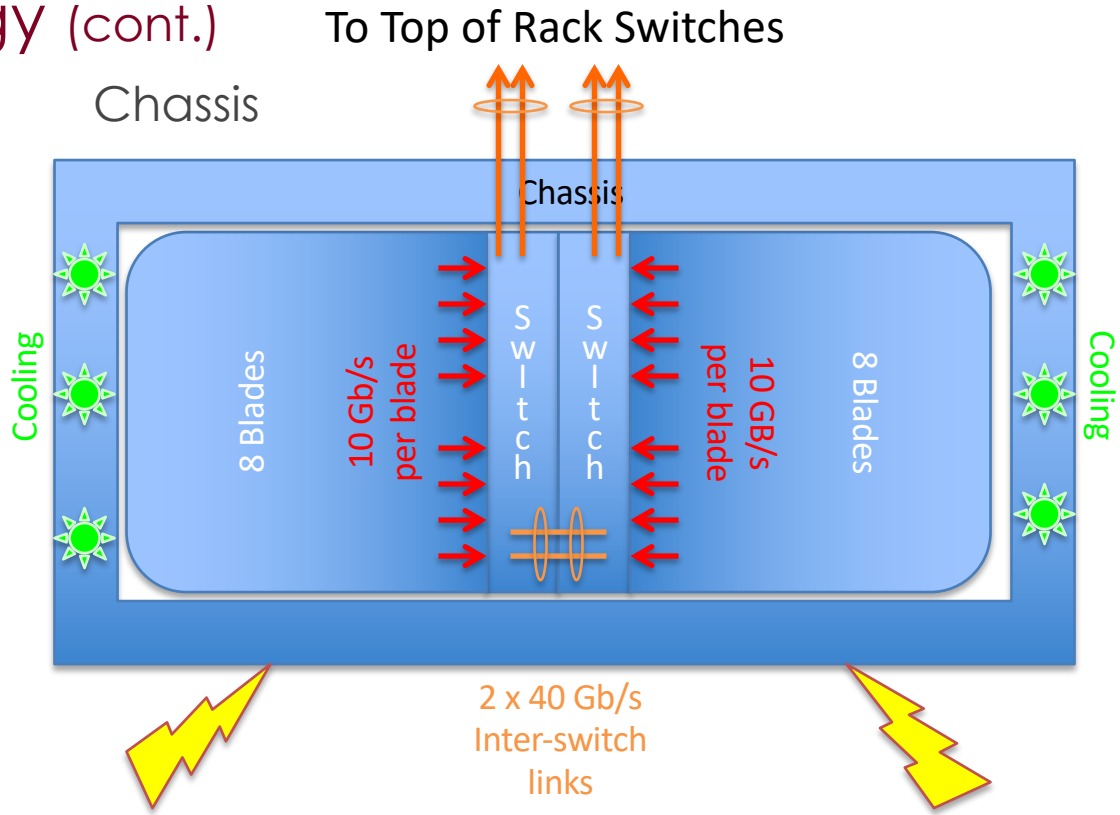

Jetstream
http://jetstream-cloud.org/

funded by the National Science Foundation
Award #ACI-1445604

Network Topology (cont.)

Chassis to Top of Rack

40 Gb/s

LAG

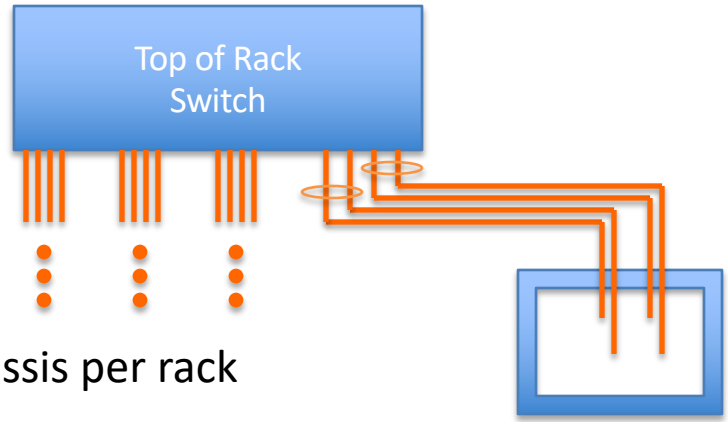

Four chassis per rack

Two switches per chassis

Jetstream
http://jetstream-cloud.org/

funded by the National Science Foundation
Award #ACI-1445604

Network Topology (cont.)

100 Gb/s uplink to Internet2

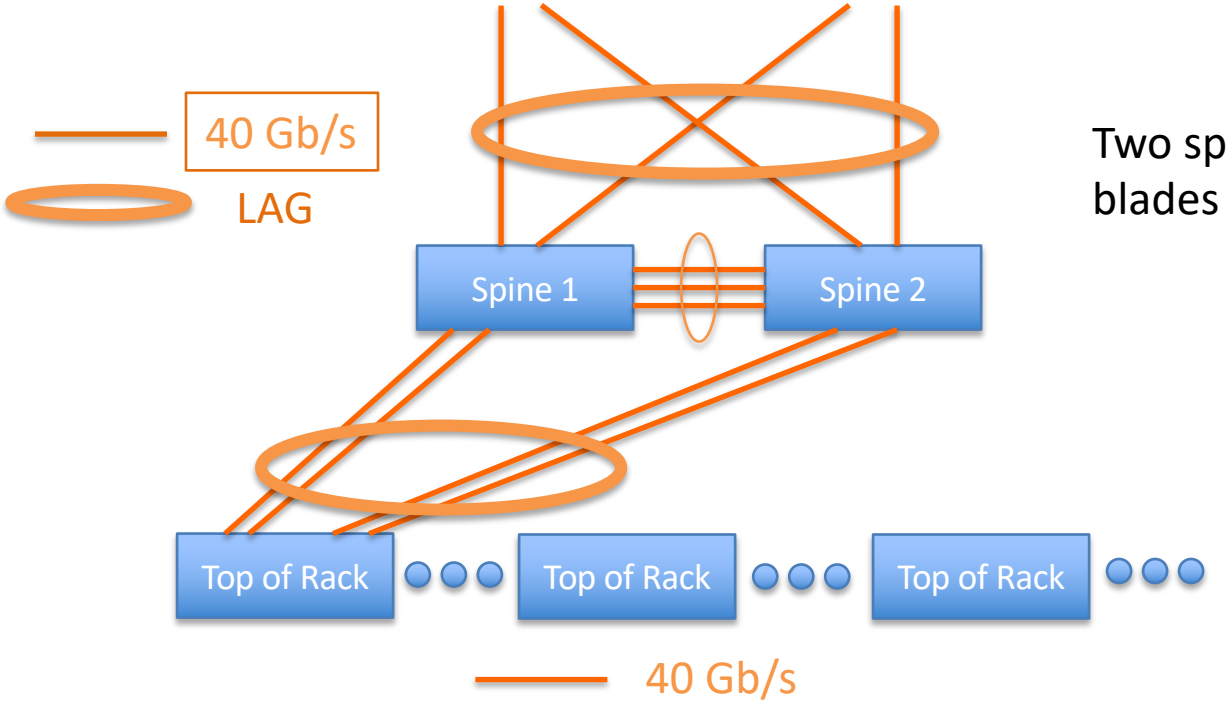

Two spines tie into two network blades in the datacenter switch

Seven racks tie into the two spine switches

Jetstream
http://jetstream-cloud.org/

funded by the National Science Foundation
Award #ACI-1445604

How do we onboard users onto Jetstream?

- An **XSEDE User Portal (XUP) account** is required. They are free! Get one at https://portal.xsede.org
- **Work with your XSEDE Campus Champion.**
- Submit an **allocations request**
  - Read the Allocations Overview - https://portal.xsede.org/allocations-overview
  - Writeup an allocation request – ***start with a Startup or Education request*** - https://portal.xsede.org/successful-requests
- **Easy Button**: instant access to small, limited instances while the allocation request is processed and the user is vetted.

Jetstream
http://jetstream-cloud.org/

NSF funded by the National Science Foundation
Award #ACI-1445604

Jetstream Information Sources

- Twitter: @jetstream-cloud
- Jetstream's web interface: https://use.jetstream-cloud.org/
  No login required to browse image library
- XSEDE User Portal account is required to actually login:
  https://portal.xsede.org Create account in seconds.
- Jetstream Home page: https://jetstream-cloud.org/
- Jetstream's public documentation: https://wiki.jetstream-cloud.org

Jetstream
http://jetstream-cloud.org/

NSF funded by the National Science Foundation
Award #ACI-1445604

Jetstream Information Sources (Cont.)

- Paper describing Jetstream Jetstream: A self-provisioned, scalable scince and engineering cloud environment
- Configuration management: https://github.com/jetstream-cloud/Jetstream-Salt-States

For questions, comments, etc. of any manner

help@jetstream-cloud.org

Jetstream
http://jetstream-cloud.org/

funded by the National Science Foundation
Award #ACI-1445604

Jetstream Partners

INDIANA UNIVERSITY
PERVASIVE TECHNOLOGY INSTITUTE

JOHNS HOPKINS
UNIVERSITY

THE UNIVERSITY OF ARIZONA

TACC

THE UNIVERSITY OF CHICAGO

UNC
THE ODUM INSTITUTE

MathWorks
Accelerating the pace of engineering and science

Jetstream
http://jetstream-cloud.org/

funded by the National Science Foundation
Award #ACI-1445604

Questions?

Project website: http://jetstream-cloud.org/
Project email: jethelp@iu.edu
Direct email: jomlowe@iu.edu, turnerg@iu.edu

funded by the National Science Foundation
Award #ACI-1445604