

Advanced Science Gateway Architectures

Marlon Pierce, Suresh Marru
Science Gateways Research Center

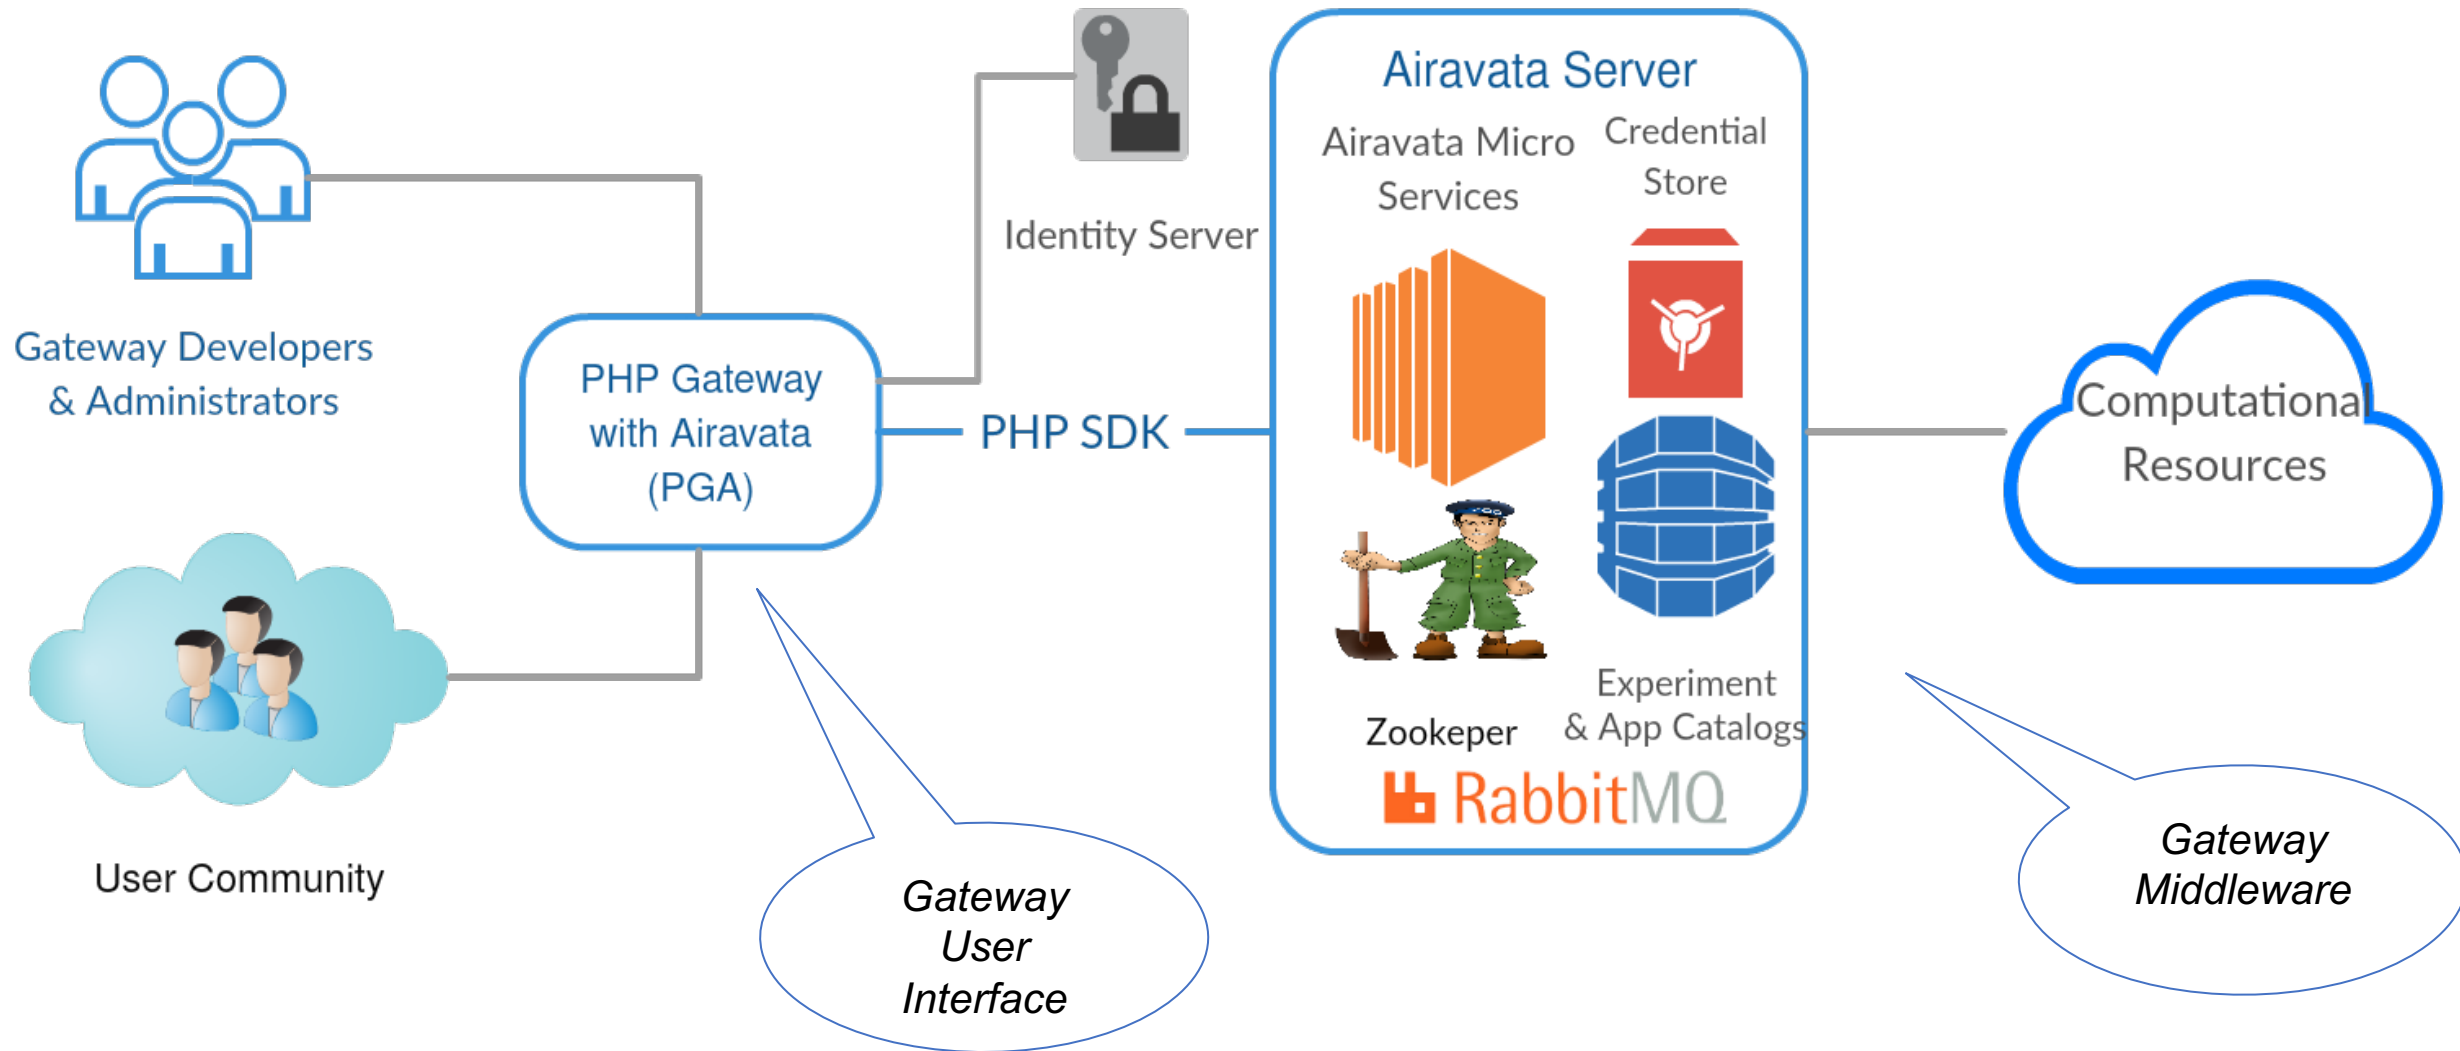
Office Hours (To Be Finalized)

- Wednesdays 10:00 am – 12:00 noon in CIB Multipurpose Room, as available.
 - This is Science Gateway Research Center's regular team hackathon time.
- Fridays 3:30-5:00 pm
- Other times by appointment
- Course HipChat: <https://www.hipchat.com/gMdyarVlz>

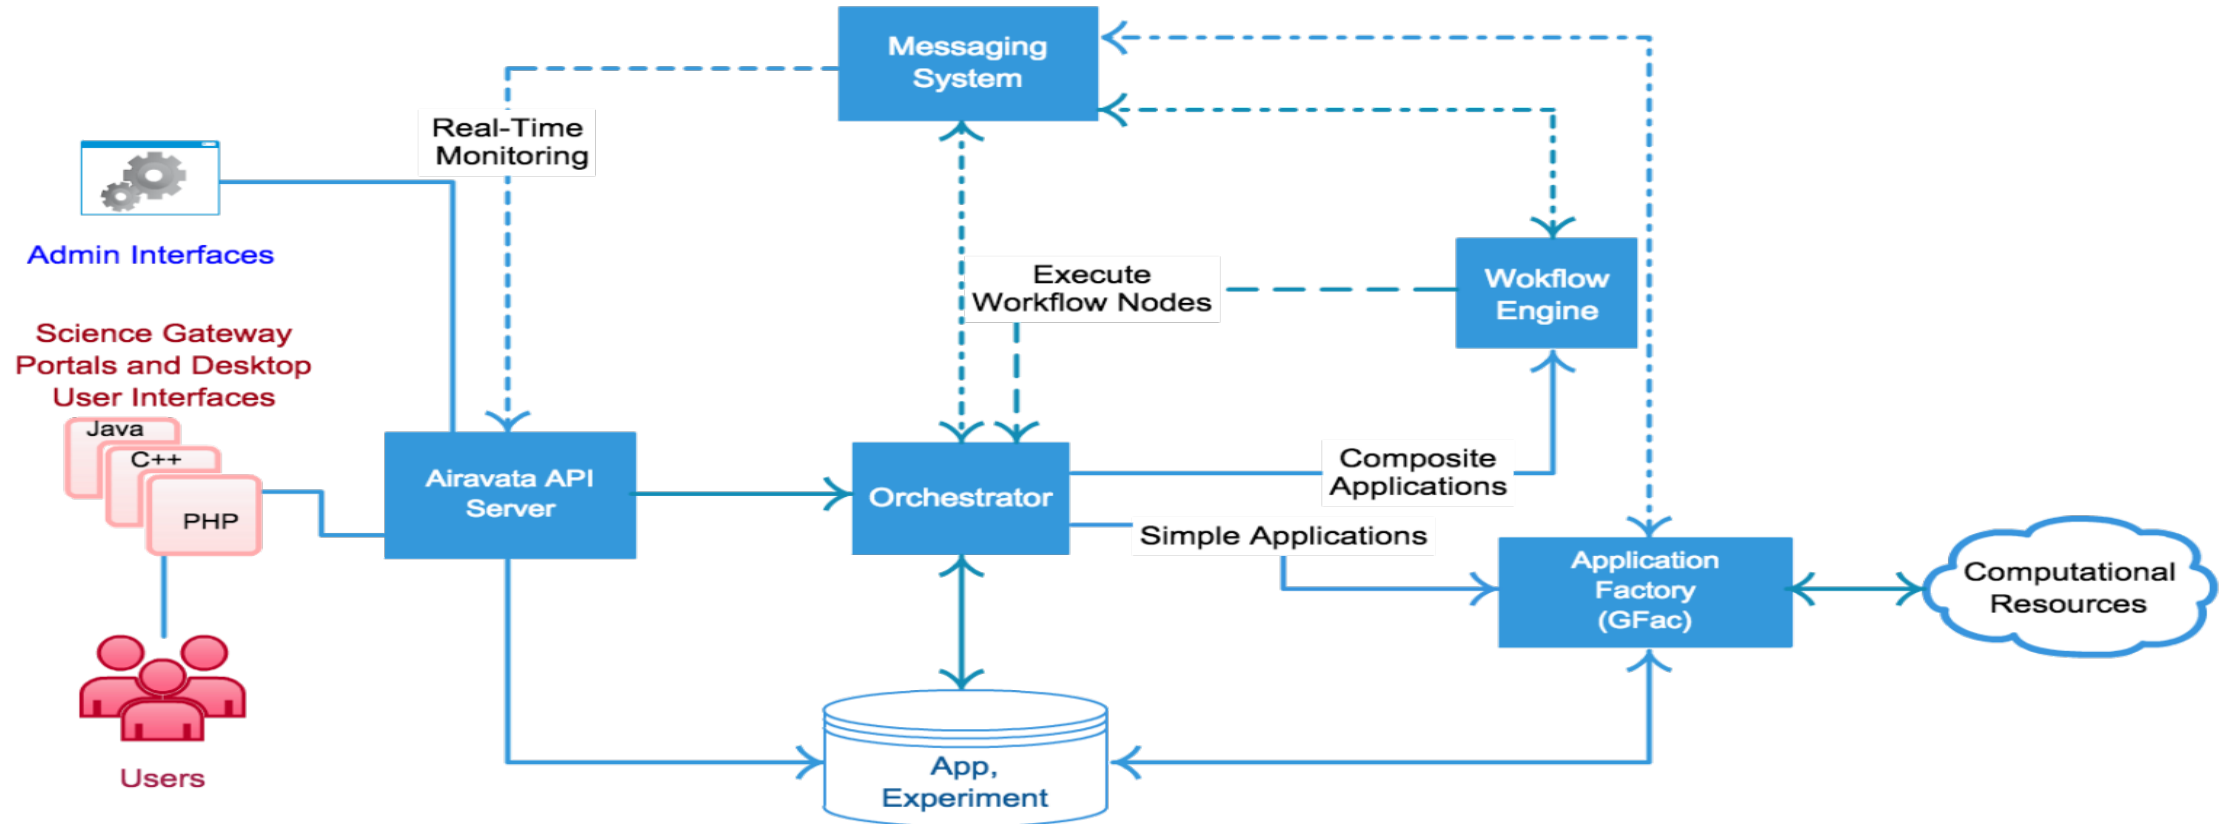
What is a Science Gateway?

Demo of PGA

Gateway Anatomy mapped to Airavata

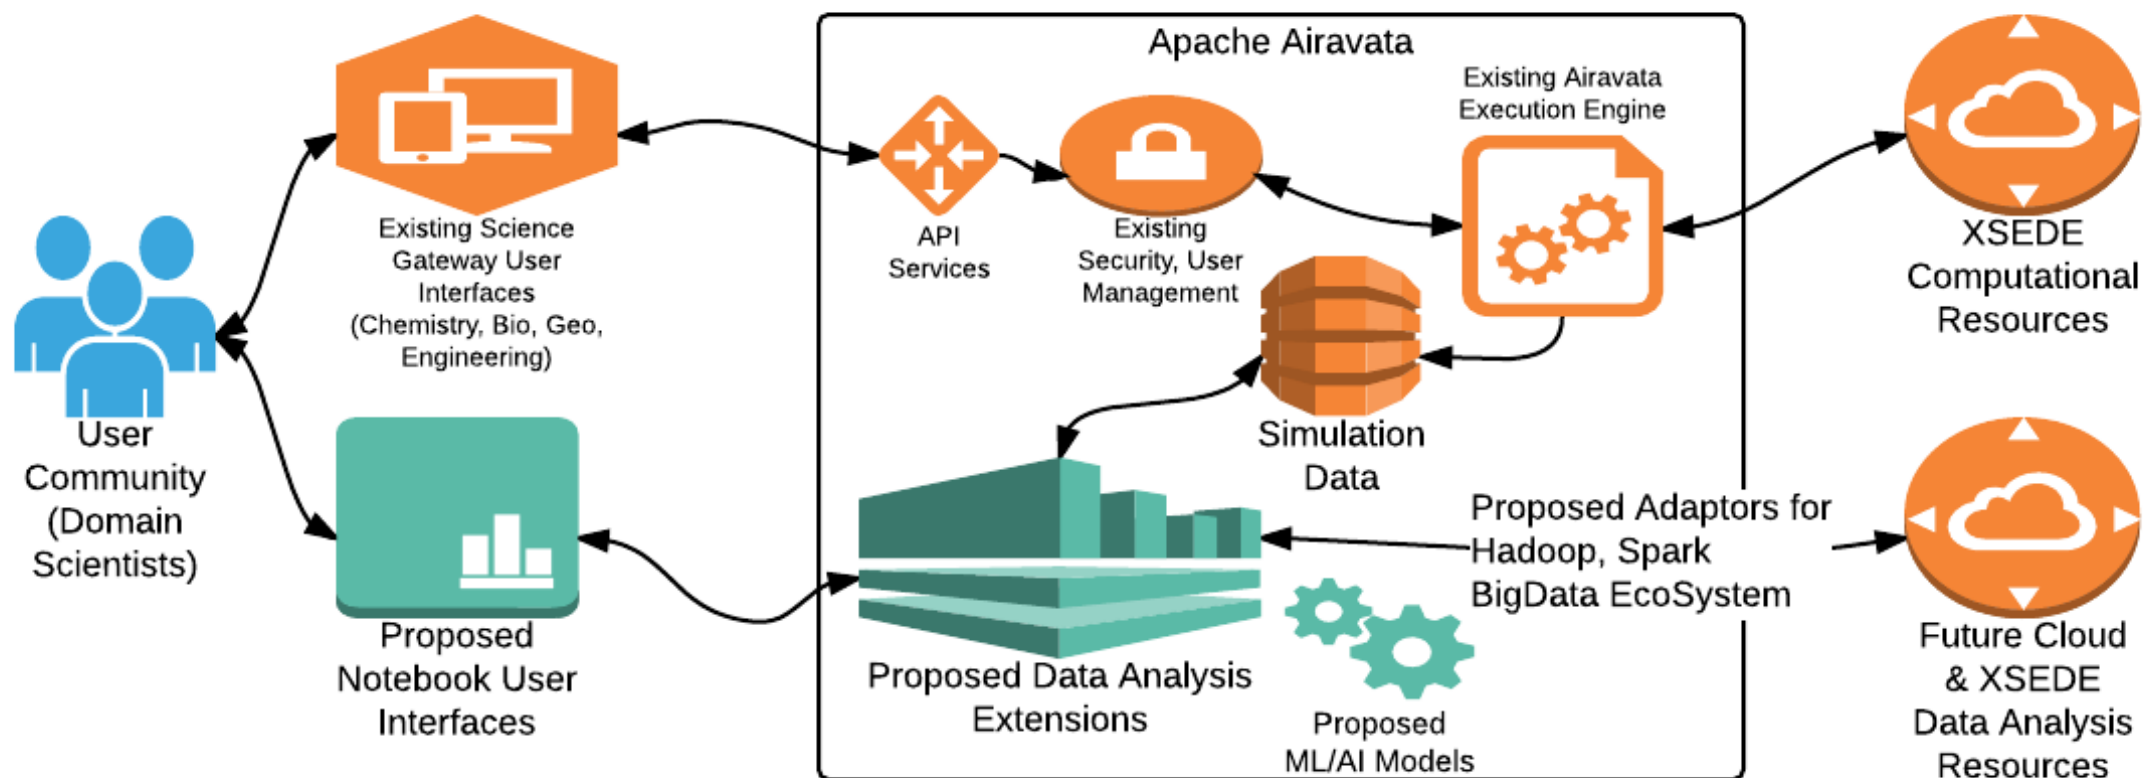


Airavata: Keep it simple, yet flexible



- External clients interact with Airavata API (based on Apache Thrift).
- Internally, components interact with each other through Component Programming Interfaces (thrift based CPIs).

Managing Computations → Data Analysis



Fall 2016 Class Summary

- Science gateways are distributed systems on top of distributed systems.
 - Gateways talk to supercomputers, clusters, clouds
 - Gateways themselves have multiple services.
- The microservice architecture is an important (but not the only) way to build gateways.
- Microservices depend on continuous integration and delivery techniques
 - Testing, updates, etc

Fine words, but how do you really implement this?

Some Unanswered Microservice Questions

- How should microservice communications be coordinated?
- Are all microservices equal?
 - Or do we have core services?
 - How do non-core services interact with core services?
 - Are core services read-only, or can non-core services write?
- For new requirements, should I modify existing services or should I create a new set for that specific use case?
- Where is the state of an experiment?
 - Key to load balancing, fault tolerance, etc.

Class themes are based on fundamental Airavata architectural issues.

Goal of the Class

- Students should learn core distributed systems concepts by contributing to Apache Airavata.
- We are all self-interested
 - You should use this course to expand your skills, demonstrate problem defining as well as problem solving skills, learn open source governance, and make concrete contributions to a top level Apache Software Foundation project.
 - We want to make Apache Airavata better. This may mean exploring some options that don't work out or only provide information.
- **“Contribute” means more than code.**
 - We want you to active discuss your work on the Apache Airavata developers' and architects' email lists.


A Note on Contributing Code

- This project will not initially focus on Apache Airavata's current code base.
 - We don't want to bias the class with some historical implementation decisions.
- But we will provide real Airavata problems for you to investigate.
- You will work in GitHub
- You can make your code contributions to Apache Airavata.


Expected Effort Level

- We expect you to put a lot of effort into the class.
- We expect that you will do this because you enjoy the work.

How effectively did out-of-class work (assignments, readings, practice, etc.) help you learn?



In a typical week, about how much time did you devote to this course? (Do not count scheduled class time, labs, etc.)



Course Structure

- The course will be based on projects.
 - Phase 1 and Phase 2
- Phase 1: The first two (maybe three) projects will be structured in order to create a mockup of Airavata 2.0
- Phase 1 projects are fully working codes, complete with version control, automated installation, integration, and deployment.
- The Phase 1 code base will be used to investigate real questions we have about Airavata's architecture and deployment.

Phase 2 Projects

- Phase 2 Projects are based on “themes”
- Themes will be used to investigate one or more real Airavata problems. These will be the subject of biweekly projects.
- You can work on different themes throughout the semester.

Grades

- Projects: 80%
 - There will be 8 projects, based on themes
 - These will use GitHub, CI/CD, and Amazon as many of you know.
 - Projects will be due every 2 weeks.
 - Each project will conclude with presentation and report
 - Grading will follow Fall 2016 guidance sheet
 - Grading will be based on individual effort
- Attendance, participation in class, and out of class interactions (via GitHub): 5%
- Final Report: 15%

Theme	Description
APIs, Data Model Representation and Storage	API Comparisons: REST, Swagger, Thrift, ProtoBuff etc.; Optimal serialization, deserialization, storage and search strategies; Internal (CPI) versus externally exposed (API) data models
Distributed Workload Distribution	Load Balancing Stateful vs Stateless Services; Message based vs hybrid message-RPC vs pure RPC; Work-Queue vs Master-Worker load balancing
DevOps Strategies	Fault tolerance testing and validation; Load balancing testing and validation; Continuous upgrades without downtime; Capacity testing; Comparison of hosting alternatives: OpenStack (Jetstream), Amazon AWS, Google, Microsoft Azure
Cyberinfrastructure Integration	Interacting with remote computing resources and data; Metascheduling; Gateway metadata and provenance; Alternative resource utilization
Cybersecurity	PHI, HIPAA, etc alignment for gateways; DevSecOps practices; Auditing, monitoring, and event detection; Client assurance
Data-Centric Computing and Gateways	Airavata and data stream processing; Internet of Things and Airavata; Searching and data mining structured data sets; Event-driven computations
Scientific User Environments	Gateways and interactive computing; Application monitoring; Visualization; Modern user interfaces and environments for science; Notebook integration
Application Toolboxes for Airavata Applications	Defining toolboxes: application and resource collections; Toolbox templates; Packaging, integrating, and sharing tool boxes

Mid-Course Corrections

N.B. This is our first time to do this type of course. We may adjust things as we go.